

- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11
- 12
- 13
- 14
- 15

3  
4  
5  
6  
7  
8  
9  
10  
11

4  
5  
6  
7  
8  
9  
10  
11

12

13

14

15

16 **Samenstelling van de promotiecommissie**

17

18

Promotoren: Prof. dr. H.J. van den Herik,  
Prof. dr. ir. H. X. Lin

Copromotor: Dr. J.L.A. Dubbeldam

Promotiecommissie: Prof. dr. A. Plaat,  
Prof. dr. J.N. Kok,  
Prof. dr. ir. A.W. Heemink (TU Delft) ,  
... (Buitenlands),  
Prof. dr. ir. B.J.A. Kröse  
Prof. dr. C.M. Jonker (SIKS)



SIKS Dissertation Series No. ...

20 The research reported in this thesis has been carried out under the auspices of SIKS, the  
21 Dutch Research School for Information and Knowledge Systems.

22

23 ISBN ....

24 Copyright © 2018, A.C. van Rossum

25 *All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or*  
26 *transmitted, in any form or by any means, electronically, mechanically, photocopying, recording*  
27 *or otherwise, without prior permission of the author.*

28 “ *The study of mental objects with reproducible properties is called mathematics.* ”  
29

30 The Mathematical Experience (Davis and Hersch, 1981)

31 “ *The study of physical objects with reproducible properties is called science.* ”  
32

33 The dawning of the age of stochasticity, Mathematics: frontiers and perspectives  
34 (Mumford, 2000)

# CONTENTS

36	<b>1 Introduction</b>	<b>1</b>
37	1.1 Scope of the Thesis . . . . .	1
38	1.2 Bayesian Nonparametrics . . . . .	2
39	1.3 Problem Statement and Research Questions . . . . .	2
40	1.4 Research Methodology . . . . .	3
41	1.5 Main Contribution . . . . .	4
42	1.6 Organization of the Thesis . . . . .	5
43	<b>2 Related Work</b>	<b>8</b>
44	2.1 Probability Theory . . . . .	8
45	2.1.1 Measure Theory . . . . .	9
46	2.1.2 Bayesian Inference . . . . .	16
47	2.2 Nonparametric Bayesian Models . . . . .	21
48	2.2.1 Beta Process . . . . .	27
49	2.2.2 Gamma Process . . . . .	28
50	2.2.3 Dirichlet Process . . . . .	29
51	2.2.4 Pitman-Yor Process . . . . .	31
52	2.2.5 Hierarchical Dirichlet Process . . . . .	32
53	2.3 Inference . . . . .	33
54	2.3.1 Inverse Transform Sampling . . . . .	33
55	2.3.2 Rejection Sampling . . . . .	34
56	2.3.3 Approximate Bayesian Computation . . . . .	35
57	2.3.4 Gibbs Sampling . . . . .	36
58	2.3.5 Metropolis-Hastings Sampling . . . . .	38
59	2.3.6 Split-Merge MCMC Sampling . . . . .	38
60	2.4 Chapter Conclusions . . . . .	39
61	<b>3 Nonparametric Bayesian Line Detection</b>	<b>42</b>
62	3.1 Infinite Line Model . . . . .	43
63	3.1.1 Posterior Predictive for a Line given Other Lines . . . . .	43
64	3.1.2 Likelihood of Data given a Line . . . . .	45
65	3.1.3 Conjugate Prior for a Line . . . . .	46
66	3.1.4 Posterior Predictive for a Line given Data . . . . .	47
67	3.2 Inference for the Infinite Line Model . . . . .	47

68	3.3	Accelerating Inference for the Infinite Line Model . . . . .	49
69	3.4	Performance of the Infinite Line Model . . . . .	50
70	3.4.1	Clustering Performance . . . . .	50
71	3.4.2	Two Examples . . . . .	51
72	3.5	Chapter Conclusions . . . . .	52
73	<b>4</b>	<b>Nonparametric Bayesian Segment Estimation</b>	<b>56</b>
74	4.1	Pareto Pairs . . . . .	56
75	4.2	Generative Process to Create a Line Segment . . . . .	59
76	4.3	Inference over a Line Segment . . . . .	60
77	4.4	Results . . . . .	63
78	4.5	Chapter Conclusions . . . . .	64
79	<b>5</b>	<b>Triadic Split-Merge Sampler</b>	<b>66</b>
80	5.1	The Class of Split-Merge Samplers . . . . .	66
81	5.2	Conventional Split-Merge Sampler . . . . .	67
82	5.2.1	Acceptance for the Split Step . . . . .	68
83	5.2.2	Acceptance for the Merge Step . . . . .	69
84	5.2.3	Sequentially-Allocated Merge-Split Sampler . . . . .	70
85	5.3	Triadic split-merge sampler . . . . .	70
86	5.3.1	Acceptance for the Split Step . . . . .	71
87	5.3.2	Acceptance for the Merge Step . . . . .	73
88	5.4	Results . . . . .	74
89	5.4.1	Implementation . . . . .	75
90	5.4.2	Comparison . . . . .	76
91	5.5	Chapter Conclusions . . . . .	77
92	<b>6</b>	<b>Adversarially Trained MCMC Kernels</b>	<b>79</b>
93	6.1	Data-Driven Inference . . . . .	79
94	6.2	Learning the Transition Operator . . . . .	80
95	6.2.1	Adversarial Training . . . . .	80
96	6.2.2	Variational Autoencoders . . . . .	81
97	6.2.3	Infusion Training . . . . .	81
98	6.3	Volumetric Models . . . . .	82
99	<b>7</b>	<b>Recommender Engine</b>	<b>85</b>
100	7.1	Application . . . . .	85
101	7.2	Model of Individuals . . . . .	86
102	7.2.1	Multi-modal Normal-Uniform Distribution Model . . . . .	86
103	7.2.2	Multi-modal Von-Mises-Uniform Distribution Model . . . . .	87
104	7.2.3	Hyperparameters . . . . .	88
105	7.3	Model of Groups . . . . .	89
106	7.4	Inference . . . . .	89
107	7.5	Results . . . . .	90
108	7.5.1	Artificial Dataset . . . . .	90
109	7.5.2	Real-world Dataset . . . . .	91
110	7.6	Discussion . . . . .	92

111	<b>8 Discussion and Conclusions</b>	<b>94</b>
112	<b>References</b>	<b>96</b>
113	<b>Appendices</b>	
114	<b>A Probabilistic Concepts</b>	<b>104</b>
115	A.1 Common Inequalities . . . . .	105
116	A.1.1 Markov's Inequality . . . . .	105
117	A.1.2 Chebyshev's Inequality . . . . .	106
118	A.1.3 Weak Law of Large Numbers . . . . .	106
119	A.1.4 Strong Law of Large Numbers . . . . .	107
120	A.1.5 Common Distributions . . . . .	107
121	<b>B Dirichlet-multinomial</b>	<b>109</b>
122	B.1 Categorical Distribution . . . . .	109
123	B.2 Multinomial Distribution . . . . .	110
124	B.3 N Categorical Distributions . . . . .	111
125	<b>C Gibbs Sampling</b>	<b>112</b>
126	C.1 Gibbs Sampling of Parameters . . . . .	114
127	<b>D Glossary</b>	<b>117</b>
128	<b>List of Figures</b>	<b>118</b>
129	<b>List of Tables</b>	<b>121</b>
130	<b>List of Abbreviations</b>	<b>122</b>
131	<b>Summary</b>	<b>124</b>
132	<b>Samenvatting</b>	<b>126</b>
133	<b>Acknowledgments</b>	<b>128</b>
134	<b>Curriculum Vitae</b>	<b>129</b>
135	<b>Publications</b>	<b>130</b>
136	<b>SIKS Dissertation Series</b>	<b>131</b>

# INTRODUCTION

137

138

## Contents

The thesis addresses nonparametric Bayesian methods in robotic vision. Nonparametric Bayesian models can be simultaneously employed to perform inference over the number of entities observed and over the shape or nature of these entities. This chapter introduces nonparametric Bayesian models, the research methodology based on the Bayesian methodology, the main contribution towards robotic vision, and the general organization of this thesis.

## Outline

The scope of thesis is to apply nonparametric Bayesian methods to robotic vision (Sect. 1.1). Bayesian nonparametric models define entities together with noise in such a way that inference can be performed in an optimal manner (Sect. 1.2). Particular problems in robotic vision that can benefit from Bayesian nonparametric methods are formulated and detailed (Sect. 1.3). The research methodology is described (Sect. 1.4). Our main contribution is to introduce nonparametric Bayesian models in robotic vision (Sect. 1.5). At the end of this chapter the organization of the thesis is given (Sect. 1.6).

## 1.1 Scope of the Thesis

In this thesis, modern Bayesian nonparametric methods are used to answer long-standing questions within computer vision and robotics. The following three challenging questions are typical examples. Is there a Bayesian form of line detection rather than applying the traditional Hough transform? Which of the nonparametric Bayesian priors can be used to detect multiple features simultaneously? What are efficient inference methods for these priors?

The scope of the thesis is the transfer of knowledge on Bayesian nonparametrics to well-described application domains. It will not establish a new body of work around a new family

of stochastic processes. The detailed application of complex models towards robotic vision is expected to help and encourage people in entirely different application domains, such as collaborative filtering, search engine optimization, and audio processing. All these different applications do not always need dedicated algorithms, but do deserve and can exploit the same optimal general inference techniques from Bayesian nonparametrics.

## 1.2 Bayesian Nonparametrics

In robotic vision (computer vision and depth perception) traditionally custom-made algorithms have been developed for a given task. There are specific methods to detect corners (e.g., Förstner and Gülch, 1987; Harris and Stephens, 1988; Shi and Tomasi, 1994), to detect edges (e.g., Sobel, 1970; Canny, 1986), to detect features (e.g., Hough, 1962), and to describe features (e.g., Lowe, 1999; Dalal and Triggs, 2005; Bay et al., 2006).

On the one hand, it is desirable that such sophisticated methods are generalizable to other application domains. On the other hand, it is important to take particular information about an application domain into account. The methods described in the previous paragraph are limited to their specific task. An example of limited generalizability can be found in the Hough transform. The Hough transform can be used to detect lines, but the way inference is performed in the algorithm does limit its application to basic forms of object detection. An example of limited specificity can be found in linear regression. Linear regression does not take into account real-world statistics.

Both generalization and specificity are formalized by a Bayesian model. A Bayesian model is general because it can be solved with general inference methods. One of such general inference methods is a Markov-Chain Monte Carlo method. It does not know anything about real-world statistics. A Bayesian model is also specific in that it can incorporate application-specific know-how by the definition of priors.

Typical problems in robotic vision will be about the recognition of several objects, multiple shapes, or objects that have multiple parts. Models that represent such objects do not have knowledge about the number of such objects, shapes, or parts. To incorporate application-specific know-how on the number of objects it is possible to define a prior that assigns a probability to this quantity. The number of objects can even be potentially infinite. The Bayesian models that define a prior on the number of objects, shapes, or parts are called nonparametric Bayesian models. This means that in contrast with conventional methods such as  $k$ -means clustering (Forgy, 1965; Lloyd, 1982) the number of objects does not need to be predefined.

## 1.3 Problem Statement and Research Questions

Many methods in robotics - and in particular in robotic vision - have been developed in times where computational resources were limited. Then, highly optimized algorithms have been developed, leveraging peculiarities of the application domain. Recent advances in Bayesian



201 methods, both with respect to concept development, as well as computational efficient so-  
 202 lution strategies, now open up new ways to solve old problems. However, extending only  
 203 the old methods themselves would lead to ad-hoc solution strategies that will miss benefits  
 204 from potential optimal and more widely applicable algorithms.

205 This observation leads us to the formulation of our problem statement (PS).

206 **PS:** *How can robotic problems effectively be generalized and their structure  
 exploited in a wider Bayesian framework?*

207 The problem statement is rather general. In our research, we focus on robotic vision, in the  
 208 form of point cloud recognition and depth perception. In particular, we look at objects, lines,  
 209 line segments, and more complex shapes.

210 The problem statement is divided into three research questions (RQs).

**RQ 1** How can we estimate the number of objects simultaneously with the  
 fitting of these objects?

211 **RQ 2** How can we estimate the number of lines simultaneously with line fitting  
 in computer vision?

**RQ 3** How can we recognize more general 3D objects?

## 212 1.4 Research Methodology

213 The research methodology advocated in this thesis follows the Bayesian methodology (cf.  
 214 Savage, 1972; Jaynes, 2003). So, our research methodology exists out of two phases. In the  
 215 first phase a Bayesian model is defined. This model exists of (1) a definition of parameters  
 216 and relations between these parameters, (2) a definition of the noise, and (3) the data. In  
 217 the second phase, the Bayesian method dictates all remaining unknowns, from the number  
 218 of parameters to the values of the parameters. To perform Bayesian inference efficiently new  
 219 methods are required if the model is complex (as is in the case of robotic vision).

220 The Bayesian methodology aims to establish the rationale for practical questions. The fol-  
 221 lowing two questions are clear examples.

- 222 ◦ If we observe a single point in an image, can we expect it to be part of a line?
- 223 ◦ If we have two lines and we live in a world with squares, what are we able to infer?

224 The two questions tap into our capabilities to define models that makes our prior knowledge  
 225 explicit. Moreover, if we are able to quickly assign (1) points to segments, (2) segments to  
 226 lines, (3) objects to categories, we can enrich it with all corresponding group properties  
 227 without the need to have them observed for this individual.

228 In robotic vision we take as an example the task of line detection. Both the Hough transform  
229 (Hough, 1962) and the RANSAC method (Bolles and Fischler, 1981) do detect lines, but they  
230 do not explicitly take noise into account. By applying Bayesian methodology to these tasks,  
231 the inference method becomes optimal in an information-theoretic sense. Also frequentist  
232 statisticians agree that nonparametric Bayesian models are consistent in the sense that they  
233 approach the underlying true distribution (Wasserman, 1998). There is no need to search  
234 for another method to infer lines in a line detection task. If someone would find a method  
235 that outperforms a Bayesian method it is either (1) because the signal or noise has not been  
236 correctly modeled after all, or (2) because the method overfits with respect to the available  
237 data. If approximations are used with respect to optimal Bayesian inference (either varia-  
238 tional approximations or Markov-Chain Monte Carlo), there are theoretical guarantees on  
239 convergence.

240 A well known problem with nonparametric Bayesian models is the curse of dimensionality.  
241 Compared to maximum likelihood methods or other non-probabilistic methods that do not  
242 take noise into account at all, the nonparametric Bayesian models require significant com-  
243 putational resources. Our research methodology first establishes the correct models, even  
244 if solving them seems computationally infeasible. Our approach is to develop subsequently  
245 approximations using more sophisticated samplers, so that the theoretical guarantees on  
246 convergence are preserved.

247 Due to the fact that the models are optimal by construction, there are no experiments re-  
248 quired to address the optimality in particular. However, experiments are still required to  
249 establish whether the models make sense. Yet, the methodology does also have limitations.  
250 For instance, we will not search over different noise models and limit priors to a particular  
251 hierarchical level.

## 252 1.5 Main Contribution

253 Our contribution to robotic vision can be subdivided into three parts that correspond with  
254 the three research questions.

255 The first part addresses the problem of inference about objects from a nonparametric Bayesian  
256 perspective. Contemporary methods in robotic vision do not allow for astute statements  
257 about their performance. In practice, this means that when using computer vision to detect  
258 cells under a microscope, someone cannot be confident about the number of detected cells.  
259 An autonomous cleaning robot in a supermarket cannot be confident about the aisle it is driv-  
260 ing into. To be able to properly take into account models and uncertainty simultaneously,  
261 Bayesian models have found mainstream adoption. State-of-the-art Bayesian methods that  
262 reason about the number of objects alongside object models are a recent object of study (cf.  
263 Ferguson, 1973; Hjort, 1990; Lijoi and Prünster, 2010; Joho et al., 2011). This thesis applies  
264 such nonparametric Bayesian models towards the applications of robotic vision and depth  
265 perception. Models such as the infinite line model and the infinite line segment model are  
266 introduced.

267 The second part addresses the problem of high-dimensional data, typified by the applications  
268 at hand. To efficiently sample more complex geometric structures, new MCMC (Markov-  
269 Chain Monte Carlo, Sect. 2.3.4) methods are required. This thesis introduces such an MCMC  
270 sampler, namely a new Split-Merge sampler, and applies it to this challenging application.

271 The third part addresses more complex robotic vision problems, in the form of object recog-  
272 nition of point clouds in 3D. It combines nonparametric Bayesian inference with models  
273 from deep learning.

## 274 1.6 Organization of the Thesis

275 **Chapter 1** (this chapter) introduces the problem of contemporary methods in computer  
276 vision and depth perception. Due to the fact that these methods are not optimal  
277 by construction, it is hard to articulate how they perform. The need for a Bayesian  
278 methodology is sketched briefly. The problem statement and the research questions  
279 are formulated. Moreover, the research methodology is described and the organization  
280 of the thesis is outlined.

281 **Chapter 2** describes (1) probability theory using measure theory, (2) random measures  
282 known as random processes of which five are described as nonparametric Bayesian  
283 models, and (3) six inference methods that infer model parameters of such nonpara-  
284 metric Bayesian models given the data. It is followed by a discussion that indicates  
285 which parts will be most useful for chapters 3 and 4.

286 **Chapter 3** examines a first nonparametric Bayesian model, i.e. the infinite line model. The  
287 infinite line model represents a countably infinite set of lines. Gibbs sampling is used  
288 to perform simultaneous inference over (1) the number of lines and (2) line parameter  
289 values such as slope and intercept.

290 **Chapter 4** examines a second nonparametric Bayesian model, i.e. the infinite line segment  
291 model. The infinite line segment model represents a countably infinite set of line  
292 segments. A split-merge MCMC sampling method is used to perform simultaneous  
293 inference over (1) the number of line segments and (2) line segment parameter val-  
294 ues such as slope, intercept, and segment size. Chapter 2-4 answer the first research  
295 question.

296 **Chapter 5** investigates a new MCMC method, the Triadic Split-Merge sampler. It is tailored  
297 to clustering problems and accelerates inference of the models in chapter 3 and 4. This  
298 chapter answers the second research question.

299 **Chapter 6** examines a third nonparametric Bayesian model, particular aimed at volumetric  
300 inference. This chapter answers the third research question.

301 **Chapter 7** applies the hierarchical sampling method to the domain of recommender en-  
302 gines. It estimates simultaneously the number of user types with a fitting procedure  
303 for the individual user. I want to change this chapter to something relevant to point  
304 clouds.

305 **Chapter 8** discusses the relevance of the developed models and inference methods. The  
306 answers to the research questions are discussed. Then the problem statement is an-  
307 swered and conclusions are formulated. Finally, recommendations are given and fu-  
308 ture research is envisaged.



## RELATED WORK

### Contents

In robotics depth sensors generate point clouds. The tasks of robotic object recognition, positioning, and navigation require models that represent such point clouds. It is unclear whether the current methods that perform inference over point clouds are appropriate for these tasks. The current models do not model uncertainty explicitly. This chapter presents models that can be used for point cloud modeling and that represent uncertainty. This (partially) answers research question RQ1. The chapter concludes with recommendations for the development of point cloud inference models. They will be implemented in a new model for line inference in Chapter 3 and line segment inference in Chapter 4.

### Outline

This chapter describes probability theory, and in particular, measure theory underlying random processes (Sect. 2.1). Five random processes are described, the Beta Process, the Gamma Process, the Dirichlet Process, the Pitman-Yor Process, and the Hierarchical Dirichlet Process. The random processes are presented as a Poisson process with a Lévy measure, a stick-breaking construction, and a sequential presentation (Sect. 2.2). These representations give rise to different inference methods. Six inference methods are described. Inverse transform sampling, rejection sampling, approximate Bayesian computation, Gibbs sampling, Metropolis-Hastings, and Split-Merge Markov chain Monte Carlo (Sect. 2.3). Inference about point clouds in the chapters to follow will use adaptations of the described models and inference methods for which some recommendations are given (Sect. 2.4).

## 2.1 Probability Theory

Modern probability is based on measure theory (Sect. 2.1.1). Measure theory will provide the means to formally describe random variables, random processes, and most generally,

random measures. A model represented by random measures can be fitted to the data using Bayesian inference (Sect 2.1.2).

### 2.1.1 Measure Theory

A random variable is a *function* that assigns values to a *set* of possible outcomes. The formal definition requires concepts such as “measurable function” and “probability space” from *measure theory* (Feller, 1950). Measure theory is used to generalize the notion of a random variable to that of a “random process”.

Informally, a measure generalizes the notion of size of Euclidean objects to sets and subsets. To definition of a measure is based on the definition of a  $\sigma$ -algebra. A  $\sigma$ -algebra ascribes a value to a sum of individual disjoint sets, even if they are infinite in number.

#### ▼ Definition 2.1 — $\sigma$ -algebra

A  $\sigma$ -algebra is a subset  $\Sigma \in 2^X$ , with  $X$  a set and  $2^X$  its powerset, with three requirements:

- $\Sigma$  is non-empty: at least one  $A \in X$  is in  $\Sigma$ ;
- $\Sigma$  is closed under complementation: if  $A$  in  $\Sigma$ , so is its complement  $A^c$ ;
- $\Sigma$  is closed under countable unions: if  $A_1, A_2, \dots$  in  $\Sigma$ , so is  $A = A_1 \cup A_2 \cup \dots$

The members of a  $\sigma$ -algebra are called *measurable sets*. Let  $X = \{1, 2, 3, 4\}$  and let us define a  $\sigma$ -algebra  $\Sigma = \{\emptyset, \{1\}, \{4\}, \{2, 3\}, \{1, 4\}, \{1, 2, 3\}, \{2, 3, 4\}, \{1, 2, 3, 4\}\}$ . Here  $\emptyset$  denotes the empty set. The complement of  $A$  is defined with respect to  $X$ :  $A \cup A^c = X$ . An example of closure under complementation: let  $A_1 = \{1\}$ , then  $A_1^c = \{2, 3, 4\}$  and  $A_1^c$  is indeed a member of  $\Sigma$ :  $A_1^c \in \Sigma$ . An example of closure under countable unions: let  $A_1 = \{1\}$  and  $A_2 = \{2, 3\}$ , then  $A_1 \cup A_2 = \{1, 2, 3\}$  and  $A_1 \cup A_2 \in \Sigma$ .

#### ▼ Definition 2.2 — *generated* $\sigma$ -algebra

A **generated  $\sigma$ -algebra**, with  $X$  a set and  $B \in 2^X$ , is the smallest  $\sigma$ -algebra  $\sigma(B)$  that contains all sets of  $B$ .

Let  $X = \{1, 2, 3, 4\}$  and  $B = \{\{1\}, \{2, 3\}\}$ , then the generated  $\sigma$ -algebra is the set  $\sigma(B) = \{\emptyset, \{1\}, \{2, 3\}, \{1, 4\}, \{2, 3, 4\}, \{1, 2, 3, 4\}\}$ . Here the sets in  $B$  are completed to the sets in  $\sigma(B)$  by obeying the requirements of a  $\sigma$ -algebra of closure under complementation and countable unions by addition (e.g.  $\{1, 4\}$  is added due to closure under completion with respect to  $\{2, 3\}$ ).

371

372

The notion of a  $\sigma$ -algebra (Fremlin, 2000) can be applied to solve the so-called Banach-Tarski paradox (Banach and Tarski, 1924). This paradox describes how a unit-ball in  $\mathbf{R}^3$  can be partitioned into a finite number of disjoint infinite sets (scattering of points) and then can be reassembled into two unit-balls again. This violates the intuitive notion of preservation of volume. If the measure  $\mu$  of the union of two disjoint sets is equal to the sum of the measures of the two sets, this is called *finite additivity*:  $\mu(\bigcup_{i=1}^N A_i) = \sum_{i=1}^N \mu(A_i)$ . In probability theory  $\sigma$ -*additivity* extends this to infinite disjoint sets:  $\mu(\bigcup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} \mu(A_i)$ . Measure theory solves the Banach-Tarski paradox by only assigning a measure to subsets that are measurable sets (Tao, 2011).

373

A *measure* assigns values to measurable sets (as stated before, measurable sets are members or subsets of  $\Sigma$ ).

375

376

### ▼ Definition 2.3 — *measure*

377

A **measure**  $\mu$  is a function from  $\Sigma$  to  $[-\infty, +\infty]$ , with three requirements:

378

- $\mu$  is non-negative:  $\mu(A) \geq 0$  for  $\forall A \in \Sigma$

379

- $\mu$  has a null empty set:  $\mu(\emptyset) = 0$

380

- $\mu$  is  $\sigma$ -additive:  $\mu(\bigcup_{i \in I_{\Sigma}} A_i) = \sum_{i \in I_{\Sigma}} \mu(A_i)$  for  $A_i$  disjoint

381

382

The first statement defines that a measure  $\mu$  only assigns non-negative values to sets in  $\Sigma$ .

383

The second statement equals the measure of the empty set  $\emptyset$  to 0. The third statement

384

defines that  $\sigma$ -additivity is required. For any two sets in  $\Sigma$  the measure of the union of the

385

sets equals the sum of the measures of the individual sets. Here  $I_{\Sigma}$  defines an index over

386

sets in  $\Sigma$ .

387

388

Informally, a measure relates the concepts of *sets* and *subsets* to notions of size. A measure can be seen as a *monotonically* increasing function. Let the set  $A$  in  $X$  be the interval  $[0, 1)$ , an uncountable (infinite) set of real numbers. Define the  $\sigma$ -algebra  $\{\emptyset, A\}$ . The empty set has measure 0, the set  $A$  has measure 1. Let us define the  $\sigma$ -algebra  $\{\emptyset, A_{0,0.5}, A_{0.5,1}, A\}$ . The set  $A_{0,0.5}$  corresponds to the interval  $[0, 0.5)$  and  $A_{0.5,1}$  to  $[0.5, 1)$ . Both sets are assigned measure 0.5 and their union has measure 1. This examples shows that with  $\sigma$ -additive unions, measures can be assigned to sets that are uncountable.

389

A *measurable space*  $(X, \Sigma)$  is defined as a pair.

390

391

### ▼ Definition 2.4 — *measurable space*

392

A **measurable space**  $(X, \Sigma)$  is a pair with:

393

- $X$  a set;

394

- $\Sigma$  a  $\sigma$ -algebra over  $X$ .

395



396 A *measure space*  $(X, \Sigma, \mu)$  is defined as a triple.

397 ▼ **Definition 2.5 — *measure space***

398  
399 A **measure space**  $(X, \Sigma, \mu)$  is a triple with:

- 400 ◦  $X$  a set;
- 401 ◦  $\Sigma$  a  $\sigma$ -algebra over  $X$ ;
- 402 ◦  $\mu$  a measure from  $\Sigma$  to  $[-\infty, \infty]$ .

404 A finite measure  $\mu$  assigns a finite real number to all  $A$ .

405 ▼ **Definition 2.6 — *finite measure***

406  
407 A **finite measure**  $\mu$  is a measure from  $\Sigma$  to  $[0, \infty)$ :

- 408 ◦  $\mu$  is non-negative:  $\mu(A) \geq 0$  for  $\forall A \in \Sigma$ ;
- 409 ◦  $\mu$  has a null empty set:  $\mu(\emptyset) = 0$ ;
- 410 ◦  $\mu$  is  $\sigma$ -additive:  $\mu(\bigcup_{i \in I_\Sigma} A_i) = \sum_{i \in I_\Sigma} \mu(A_i)$  for  $A_i$  disjoint;
- 411 ◦  $\mu$  for the whole sample space,  $X$ , is finite:  $\mu(X) = N$ .

413 A  $\sigma$ -finite measure allows  $A$  to be a countable union of sets with finite measure.

414 ▼ **Definition 2.7 —  $\sigma$ -finite measure**

415  
416 A  **$\sigma$ -finite measure**  $\mu$  is a finite measure with:

- 417 ◦  $X$  is a countable union of sets with finite measures

419 We will now define three measures: the *probability measure*, the *counting measure*, and the  
420 *Lebesgue measure*.

421 A *probability measure*,  $\mathbb{P}$ , is a finite measure which satisfies the following requirements:

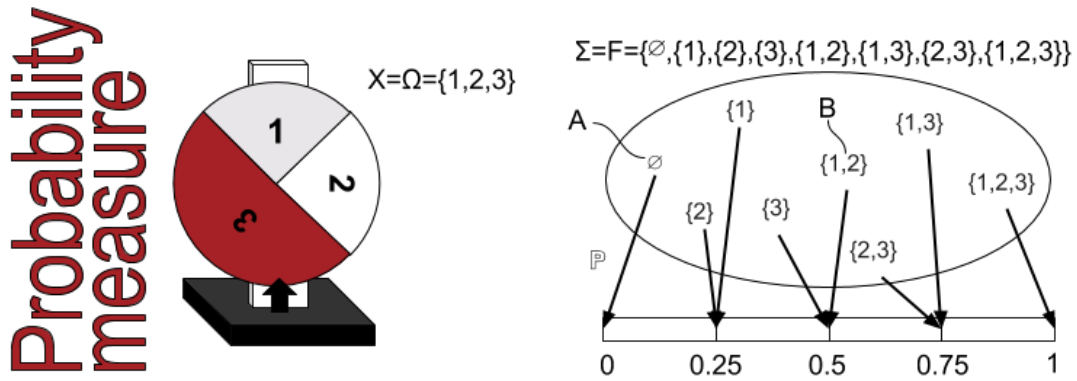
422 ▼ **Definition 2.8 — *probability measure***

423  
424 A **probability measure**  $\mathbb{P}$  is a measure  $\mu$  with:

- 425 ◦  $\mathbb{P}$  is non-negative:  $\mathbb{P}(A) \geq 0$  for  $\forall A \in \Sigma$
- 426 ◦  $\mathbb{P}$  has a null empty set:  $\mathbb{P}(\emptyset) = 0$
- 427 ◦  $\mathbb{P}$  is  $\sigma$ -additive:  $\mathbb{P}(\bigcup_{i \in I_\Sigma} A_i) = \sum_{i \in I_\Sigma} \mu(A_i)$  for  $A_i$  disjoint
- 428 ◦  $\mathbb{P}$  for the whole sample space,  $X$ , is unity:  $\mathbb{P}(X) = 1$

These requirements are called the Kolmogorov axioms (Kolmogoroff, 1933). The probability measure is a measure and therefore obeys the same three requirements of a non-negative measure for each set, the existence of a null empty set, and  $\sigma$ -additivity. A fourth requirement restricts the measure for the whole space  $X$  to 1.

In Fig. 2.1 the probability measure is visualized as a mapping from the probability space to the unit interval  $[0, 1]$ .



**Figure 2.1:** A probability measure  $\mathbb{P}$  mapping the probability space for 3 events to the unit interval. Left: a turning wheel representing three possible outcomes of which the third is twice as likely as the other two outcomes. Right: a probability measure  $\mathbb{P}$  assigned to each outcome. The empty set,  $A = \emptyset$ , has probability measure 0. The set of encountering either 1 or 2,  $B = \{1, 2\}$ , has probability measure 0.5. Adapted from Wikipedia.

A probability space  $(X, \Sigma, \mathbb{P})$  is a measure space  $(X, \Sigma, \mu)$  with the probability measure  $\mathbb{P}$  as its measure  $\mu$ .

### ▼ Definition 2.9

A probability space  $(X, \Sigma, \mathbb{P})$  is a triple with:

- $X$  a set
- $\Sigma$  a  $\sigma$ -algebra over  $X$ .
- $\mathbb{P}$  a probability measure from  $\Sigma$  to  $[0, 1]$ .

The triple for the probability space  $(X, \Sigma, \mathbb{P})$  is also written as  $(\Omega, \mathbb{F}, \mathbb{P})$ . The space  $X$  is the event space  $\Omega$ , the set of *elementary outcomes*. The  $\sigma$ -algebra over subsets of  $\Omega$  is denoted by  $\mathbb{F}$ . The probability measure  $\mathbb{P}$  assigns a value on the unit interval  $[0, 1]$  to every event in  $\mathbb{F}$ .

The *counting measure* forms the basis for the definition of discrete probabilities (Schilling, 2005).

---

**▼ Definition 2.10**


---

A **counting measure**  $\nu$  on a space  $X$  is a measure  $\mu$  with:

- $\nu$  is non-negative and integer-valued for  $\forall A \in \Sigma$
  - $\nu < \infty$  for  $\forall A \in \Sigma$  if  $A$  bounded (of finite size)
  - $\nu = \infty$  if  $\exists A \in \Sigma$  with  $A$  unbounded (infinite)
- 

A counting measure is a measure that is integer-valued. Every set  $A$  has a measure that is a positive integer or zero. If and only if the set  $A$  is unbounded is its counting measure infinite.

The *Borel  $\sigma$ -algebra* defines a  $\sigma$ -algebra for the real line  $\mathbb{R}$ .

---

**▼ Definition 2.11**


---

A **Borel  $\sigma$ -algebra**  $\mathbb{B}_\sigma$  on  $\mathbb{R}$  is the smallest  $\sigma$ -algebra that contains all open subsets of  $\mathbb{R}$ :

- $\mathbb{B} = \Sigma(U)$  with  $U = U \subseteq \mathbb{R}$ :  $U$  is open
- 

The Borel  $\sigma$ -algebra contains all open subsets of  $\mathbb{R}$ . The property of closure under complementation of a  $\sigma$ -algebra means that it also contains the closed subsets of  $\mathbb{R}$ . If  $A = (0, 1)$ , then  $A^c = \{[-\infty, 0], [1, \infty]\}$ .

A *Borel measure* assigns values to subsets of  $\mathbb{B}_\sigma$ .

---

**▼ Definition 2.12**


---

A **Borel measure**  $\mu$  is a function from  $\Sigma = \mathbb{B}_\sigma$  to  $[-\infty, +\infty]$ , with the three measure requirements:

- $\mu$  is non-negative:  $\mu(A) \geq 0$  for  $\forall A \in \Sigma$
  - $\mu$  has a null empty set:  $\mu(\emptyset) = 0$
  - $\mu$  is  $\sigma$ -additive:  $\mu(\bigcup_{i \in I_\Sigma} A_i) = \sum_{i \in I_\Sigma} \mu(A_i)$  for  $A_i$  disjoint
- 

The *Borel space* is a measurable space with a Borel  $\sigma$ -algebra rather than a general  $\sigma$ -algebra.

---

**▼ Definition 2.13**


---

A **Borel space**  $(X, \mathbb{B}_\sigma)$  is a pair with:

- $X$  a set
  - $\mathbb{B}_\sigma$  a Borel  $\sigma$ -algebra over  $X$ .
-

487 A *complete measure space* is a measure space in which every subset of every null set is mea-  
 488 surable.

489

#### ▼ Definition 2.14

490

491 A **complete measure space**  $(X, \Sigma, \mu)$ :

492

- $S \subseteq N \in \Sigma$  and  $\mu(N) = 0 \Rightarrow S \in \Sigma$

493

494 The Borel space is not a complete measure space. There are sets in the Borel  $\sigma$ -algebra that  
 495 are of measure zero and that have subsets that are undefined.

496 The *Lebesgue measure* defines a size to subsets of  $\mathbb{R}^n$  that completes the Borel measure  
 497 (Lebesgue, 1902). It makes use of the notion of an *outer measure*.

498

#### ▼ Definition 2.15

499

500 An **outer measure**  $\phi$  on a space  $\mathbb{R}$  is a measure  $\mu$  with:

501

- $\phi$  is non-negative and real-valued for  $\forall A \in \Sigma$

502

- $\phi$  has a null empty set:  $\phi(\emptyset) = 0$

503

- $\phi$  is  $\sigma$ -subadditive:  $\phi(\bigcup_{i \in I_\Sigma} A_i) \leq \sum_{i \in I_\Sigma} \phi(A_i)$  for  $\forall A_i$

504

- $\phi$  is monotone:  $A \subseteq B$  implies  $\phi(A) \leq \phi(B)$

505

- $\phi$  is translation-invariant:  $\phi(A + x) = \phi(A)$  for  $\forall A \in \Sigma$  and  $\forall x \in \mathbb{R}$

506

507 An outer measure relaxes  $\sigma$ -additivity of disjoint sets of  $X$  to  $\sigma$ -subadditivity for any se-  
 508 quence of sets. Intuitively, the outer measure of a set is an upper bound on the size of a  
 509 set.

510

#### ▼ Definition 2.16

511

512 A **Lebesgue outer measure**  $\lambda$  on a space  $\mathbb{R}^n$  is an outer measure  $\phi$  with:

513

- $\lambda(A) = \inf \left\{ \sum_{k=1}^{\infty} l(I_k) : (I_k)_{k \in \mathbb{N}} \text{ is a sequence of open intervals with } A \subseteq \bigcup_{k=1}^{\infty} I_k \right\}$

514

515 Here  $A \subseteq \mathbb{R}$  is a subset of the real line. The Lebesgue outer measure  $\lambda$  is the infimum (greatest  
 516 lower bound) of the sum of the lengths  $l(I) = b - a$  of the intervals  $I = [a, b]$ .

517 The *Lebesgue measure* is defined through the Lebesgue outer measure.

518

#### ▼ Definition 2.17

519

520 A **Lebesgue measure**  $m$  on a space  $\mathbb{R}^n$  is a Lebesgue outer measure  $\lambda$  with:

521

- $m(B) = \lambda(B \cup A) + \lambda(B \cup A^c)$

522

523 A measurable function is defined between two measurable spaces.

---

**▼ Definition 2.18**


---

A **measurable function**  $f : X \rightarrow Y$  fulfills:

$$\circ f^{-1}(E) \in \Sigma \quad \text{for} \quad \forall E \in T$$

with both  $(X, \Sigma)$  and  $(Y, T)$  measurable spaces.

---

A measurable function *preserves the structure* of the corresponding measurable spaces (captured through the  $\sigma$ -algebras).

A random variable is a measurable function between two measurable space, with as domain a measurable space that is a probability space.

---

**▼ Definition 2.19**


---

A  $(X, \Sigma)$ -valued **random variable**  $X$  is a measurable function from probability space  $(\Omega, \mathbb{F}, \mathbb{P})$  to measurable space  $(X, \Sigma)$ .

---

A random variable is a  $(X, \Sigma)$ -valued random variable with the following choice for the codomain and  $\sigma$ -algebra:

---

**▼ Definition 2.20**


---

A **random variable**  $X$  is a measurable function from probability space  $(\Omega, \mathbb{F}, \mathbb{P})$  to the real line with the Borel  $\sigma$ -algebra  $(\mathbb{R}, \mathbb{B}_{\mathbb{R}})$ .

---

The codomain is the real line  $\mathbb{R}$  and the Borel  $\sigma$ -algebra.

Random variables can be generalized to complex random variables or random elements of any type. A *complex random variable* is a measurable function from  $\Omega$  to  $\mathbb{C}$ . A *random elephant* is a measurable function from  $\Omega$  to a suitable space of elephants (Kingman, 1993).

---

**▼ Definition 2.21**


---

A **random measure** is a function  $\xi : \Omega \times X \rightarrow [0, +\infty]$  from probability space  $(\Omega, \mathbb{F}, \mathbb{P})$  to measurable space  $(X, \Sigma)$  such that  $\xi(\cdot, X)$  is a random variable on  $(\Omega, \mathbb{F}, \mathbb{P})$  and  $\xi(\omega, \cdot)$  is a measure on  $\Sigma$ .

---

We have encountered a random variable, and a probability measure  $\mathbb{P}$  on the original probability space. Now, one might wonder, are probabilities not logically assigned to elements on the measurable space that is the codomain of this random variable? Why does it map

to a measurable space and not a measure space actually? This is because (through the  $\sigma$ -algebras of both spaces, or more precisely the random variable itself) the probability measure is *induced* on the target space. This is known as a *probability distribution*:

561

### ▼ Definition 2.22

562

Given a random variable  $X$  from  $(\Omega, \mathbb{F}, \mathbb{P})$  to  $(\mathbb{R}, \mathbb{B}_\sigma)$ , the **probability distribution**  $\mu$  is the induced probability measure:  $\mu(B) = \mathbb{P}(X^{-1}(B))$  for all Borel sets  $B \in \mathbb{B}_\sigma$ .

565

The measurable *function*  $X$  is inverted:  $X^{-1}(\cdot)$ . The measure  $\mu$  exists on  $(\mathbb{R}, \mathbb{B}_\sigma)$  just as  $\mathbb{P}$  exists on  $(\Omega, \mathbb{F})$ . The notation for the measure  $\mu$  does not include the original probability space or  $\sigma$ -algebra. The complete notation for the probability distribution  $\mu$  can be written as a function  $f$  of  $X$ :

$$f_X(x) = f_{X,(\Omega,\mathbb{F},\mathbb{P}),(\mathbb{R},\mathbb{B}_\sigma)}(x) \quad (2.1)$$

At the left  $X$  denotes the random variable,  $x \in \Omega$  are the (elementary) outcomes on the sample space  $\Omega$ . At the right the complete notation adds  $\mathbb{F}, \mathbb{P}$  and  $\mathbb{R}, \mathbb{B}_\sigma$ . The shorthand notation at the left will be used to indicate the real line with a Borel  $\sigma$ -algebra as codomain.

A random variable  $X$  is *distributed as*  $f_X(x)$ , notation:

$$X \sim f_X(x) \quad (2.2)$$

A *random process* is an *ordered* set of random variables. The set can be a sequence of random variables in a time series. It can be a series of steps in the spatial domain, called a random field.

577

### ▼ Definition 2.23

578

A **random process**  $X$  is a collection  $\{X_t : t \in T\}$  with  $X_t$  an  $(S, \Sigma)$ -valued random variable on  $\Omega$  and  $(\Omega, \mathbb{F}, \mathbb{P})$  a probability space,  $(S, \Sigma)$  a measurable space, and  $T$  a totally *ordered* set.

582

A random process is a probability distribution with a domain that is a set of probability distributions. A random process is a distribution over distributions, a hierarchy over distribution.

## 2.1.2 Bayesian Inference

Let  $x$  be a  $(S, \Sigma_S, \mu_S)$ -valued random variable<sup>1</sup>,  $y$  a  $(T, \Sigma_T, \mu_T)$ -valued random variable, then we can construct  $z$ , a  $(C, \Sigma_C, \mu_C)$ -valued random variable with the latter being a subset of the product set of  $x$  and  $y$ :  $C \in S \otimes T$ .

<sup>1</sup>The lowercase  $x$  is used instead of  $X$  in the context of probability distributions as in Eq. 2.1.

---

**▼ Definition 2.24**


---

A **product space**  $(S \otimes T, \Sigma_{S \otimes T})$  and **product measure**  $\mu_{S \otimes T}$  has  $\sigma$ -algebra  $\Sigma_{S \otimes T} = \sigma(F \otimes G : F \in \Sigma_S, G \in \Sigma_T)$  and  $\mu_{S \otimes T}(F \otimes G) = \mu_S(F) \otimes \mu_T(G)$  with  $(S, \Sigma_S, \mu_S)$  and  $(T, \Sigma_T, \mu_T)$  two  $\sigma$ -finite measure spaces.

---

The **joint probability distribution**  $P_C$  is a probability measure on the product  $\sigma$ -algebra  $\Sigma_C$ . As function of the random variables  $x$  and  $y$  the joint probability distribution is written as  $x, y(x, y)$ ,  $f(x, y)$ , or  $p(x, y)$ .

A  $\sigma$ -algebra is *independent* in the following sense.

---

**▼ Definition 2.25**


---

Let  $(\Omega, \mathbb{F}, P)$  be a probability space and  $\mathbb{A}$  and  $\mathbb{B}$  be a sub- $\sigma$ -algebras of  $\mathbb{F}$ .  $\mathbb{A}$  and  $\mathbb{B}$  are **independent  $\sigma$ -algebras** if:

$$\circ P(A \cap B) = P(A)P(B) \quad \forall A \in \mathbb{A} \text{ and } B \in \mathbb{B}$$


---

Two random variables  $x$  and  $y$  are independent if and only if the  $\sigma$ -algebras that they generate are independent.

---

**▼ Definition 2.26**


---

Let  $(\Omega, \mathbb{F}, P)$  be a probability space,  $\mathbb{G} \subseteq \mathbb{F}$  a sub- $\sigma$ -algebra of  $\mathbb{F}$ , and  $X : \Omega \rightarrow \mathbb{R}$  a real-valued random variable ( $\mathbb{F}$ -measurable with respect to the Borel  $\sigma$ -algebra  $\mathbb{B}_\sigma$  on  $\mathbb{R}$ ). There exists a function  $\mu : \mathbb{B}_\sigma \times \Omega \rightarrow \mathbb{R}$  such that  $\mu(\cdot, \omega)$  is a probability measure on  $\mathbb{B}_\sigma$  for each  $\omega \in \Omega$  and  $\mu(H, \cdot) = P(X \in H | \mathbb{G})$  (almost surely) for every  $H \in \mathbb{B}_\sigma$ . For any  $\omega \in \Omega$ , the function  $\mu(\cdot, \omega) : \mathbb{B}_\sigma \rightarrow \mathbb{R}$  is called a **conditional probability distribution** of  $X$  given  $\mathbb{G}$ .

---

Informally, a conditional probability is described with a sub- $\sigma$ -algebra which only presents part of the structure of the full  $\sigma$ -algebra. As function of the random variables  $x$  and  $y$  the conditional probability distribution of  $y$  given  $x$  is written as  $f_{Y|X}(y|x)$ ,  $f(y|x)$ , or  $p(y|x)$ .

The random variables  $x$  and  $\theta$  define a Bayesian model with observations  $x$  and parameters  $\theta$ .

---

**▼ Definition 2.27**


---

A **Bayesian model**  $f(x, \theta)$  defines a function between observations  $x$  and state or parameters  $\theta$  with both  $x$  and  $\theta$  random variables.

---

In a **supervised learning** task both  $\theta$  and  $x$  are known. In an **unsupervised learning** task  $x$  is known, but  $\theta$  is unknown. The random variable  $\theta$  is called a hidden or latent variable. The random variable  $\theta$  can be any random element: a random vector, a random matrix.

Let the observations  $x$  be a sequence  $x_0, x_1, \dots$ , then the observations  $x_i$  can be distributed *independent and identically*:

---

▼ **Definition 2.28**

A collection of random variables  $x = \{x_0, x_1, \dots\}$  is **independent and identically distributed (i.i.d.)** if:

- the probability distribution  $p(x_i)$  is the same for  $\forall x_i \in x$
  - each  $x_i$  is independent with respect to  $x_j$  with  $i \neq j$
- 

The observations  $x_i$  can be distributed in an *exchangeable* sequence in which any order is equally likely:

---

▼ **Definition 2.29**

A sequence of random variables  $x = \{x_0, x_1, \dots\}$  is **exchangeable** if for any finite permutation  $\rho$  of the indices  $0, 1, \dots$ :

- the joint probability distribution of the permuted sequence  $p(x_{\rho(0)}, x_{\rho(1)}, \dots)$  equals that of the original sequence  $p(x_0, x_1, \dots)$
- 

The joint probability distribution of i.i.d. observations given parameters can be written as a product:

$$p(x_0, \dots, x_{k-1} | \theta) = \prod_{i=0}^{k-1} p(x_i | \theta) \quad (2.3)$$

---

▼ **Definition 2.30**

The **likelihood function** is defined as:

$$\mathcal{L}(\theta | x) = p(x | \theta) \quad (2.4)$$


---

The likelihood of the parameters  $\theta$  given observations  $x$  is the probability of these observations given the parameter values.

The likelihood function definition allows us to find an optimal set of parameter values given the observations. The probability  $p(x | \theta)$  can be maximized (Aldrich and Others, 1997):

---

▼ **Definition 2.31**

---



660 **Maximum Likelihood** is defined as:

$$\theta^* \in \operatorname{argmax}_{\theta} \prod_{i=0}^{k-1} p(x_i|\theta) \quad (2.5)$$

661

662

---

663 The Maximum Likelihood method finds the maximum of  $\prod_{i=0}^{k-1} p(x|\theta)$  for all possible pa-  
 664 rameter values  $\theta$ . The maximum in Maximum Likelihood does not need to be unique (Steel,  
 665 1994). The notation makes this explicit by writing  $\theta^*$  as a member (denoted by the  $\in$  sym-  
 666 bol) of the outcomes of the argmax operation (and does not use the equal sign).

667 A function  $f(\cdot)$  and the logarithm of a function  $\log f(\cdot)$  have the same maxima. This is due  
 668 to the fact that the logarithm is a monotonic function (a monotonically increasing function).  
 669 The log of a product of logarithms is equal to the sum of the individual logarithms.

670

#### ▼ Definition 2.32

671

672 **Maximum log-Likelihood** is defined as:

$$\theta^* \in \operatorname{argmax}_{\theta} \sum_{i=0}^{k-1} \log p(x_i|\theta) \quad (2.6)$$

673

674

---

675 In the case we have information about the parameters  $\theta$  we can model this with a probability  
 676 distribution.

677

#### ▼ Definition 2.33

678

679 A **prior probability distribution** defines a probability distribution  $p(\theta)$  to parameters  
 680  $\theta$  without a dependency on the observations  $x$ .

681

---

682 Given the definition of a prior probability distribution, we can define *Maximum A Posteriori*  
 683 estimation.

684

#### ▼ Definition 2.34

685

686 **Maximum A Posteriori** estimation:

$$\theta^* \in \operatorname{argmax}_{\theta} \sum_{i=0}^{k-1} \log p(x_i|\theta) + \log p(\theta) \quad (2.7)$$

687

688

---

689 If we are not only interested in the parameter  $\theta^*$  that maximizes  $p(x|\theta)$  and  $p(\theta)$ , but in the  
 690 complete distribution for  $p(\theta)$  we need Bayes' theorem described by Laplace (marquis de  
 691 Laplace, 1820).

692

### ▼ Definition 2.35

693

**Bayesian inference** uses Bayes' theorem:

694

$$f(\theta|x) = p(\theta|x) = \frac{\overbrace{p(x|\theta)}^{\text{likelihood}} \overbrace{p(\theta)}^{\text{prior}}}{\underbrace{p(x)}_{\text{normalization constant}}} = \frac{p(x|\theta)p(\theta)}{\int p(x|\theta)p(\theta)d\theta} \quad (2.8)$$

695

696

697 Bayes' theorem describes the posterior probability  $p(\theta|x)$  as the likelihood times the prior  
 698 probability distribution divided by a normalization constant, also called the evidence. The  
 699 normalization constant is not a function of the parameters  $\theta$ . If a function is known except  
 700 for the normalization constant this is indicated by the "proportional to" symbol  $\propto$ .

$$f(\theta|x) \propto p(x|\theta)p(\theta) \quad (2.9)$$

701 In Bayesian inference  $p(\theta|x)$  is calculated. In contrast, in Maximum Likelihood and Maxi-  
 702 mum A Posteriori only parts of Eq. 2.8 are calculated, respectively  $p(x|\theta)$  and  $p(x|\theta)p(\theta)$ .  
 703 In Section 2.3 inference methods will be described that approximate Bayesian inference.  
 704 Approximation is required in the case closed-form expressions are not available. If the in-  
 705 ference task only requires Maximum A Posteriori, approximation methods are also available  
 706 (Daume, 2007), but this is outside of the scope of the current thesis.

707 There are two supervised learning models, a generative model and a discriminative model.

708

### ▼ Definition 2.36

709

A **generative** model defines the joint probability distribution  $p(x, \theta)$  and uses Bayes  
 710 rule to define  $p(x|\theta)$ .

711

712

713

### ▼ Definition 2.37

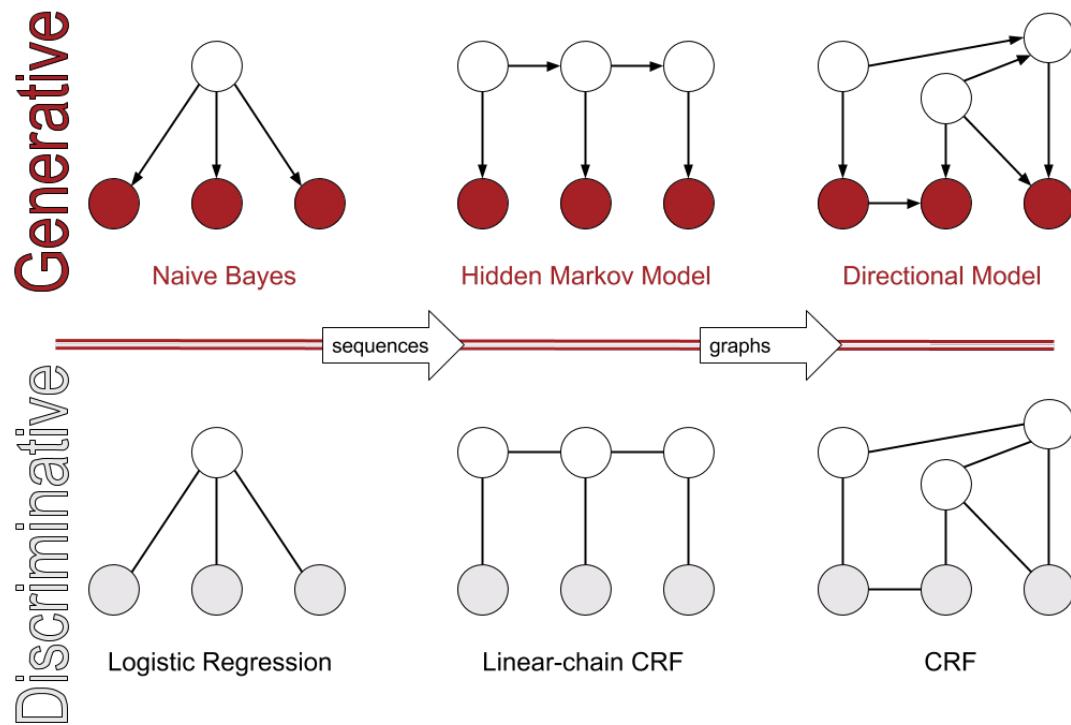
714

A **discriminative** model defines the conditional probability distribution  $p(x|\theta)$  directly.

715

716

717



**Figure 2.2:** Generative models: Naive Bayes (Russell et al., 1995), Hidden Markov Models (Baum and Petrie, 1966), and Directional Models (Koller and Friedman, 2009). Discriminative models: Logistic Regression, Linear-chain Conditional Random Fields, and (general) Conditional Random Fields. Figure adapted from Sutton and McCallum (2011).

Fig. 2.2 shows generative and discriminative models. There is no definitive reason to use a generative model above a discriminative model or vice-versa. First, a discriminative model has lower asymptotic error, but a generative model approaches its asymptotic error faster in the case of a Naive Bayes classifier versus Logistic Regression (Jordan, 2002), but see also (Xue and Titterton, 2008) which doubts the existence of such precisely defined regimes. Asymptotic error denotes the error with an increasing number of samples. Second, the prior  $p(\theta)$  in the generative model provides a principled way to handle missing information, while the direct modeling of decision boundaries in a discriminative model often leads to better performance in a classification task (Jaakkola et al., 1999). Apart from generative models and discriminative models, there are hybrid models (Bouchard and Triggs, 2004; Raina et al., 2003; Bosch et al., 2008). In this thesis we will limit ourselves to generative models.

## 2.2 Nonparametric Bayesian Models

A model can be composed out of a finite set of probability distributions. An example of such a model is the Naive Bayes model.

---

**▼ Definition 2.38**


---

The **Naive Bayes model** is a product over a finite number  $k \neq \infty$  of probability distributions  $p(x_i|\theta)$  multiplied by the prior distribution  $p(\theta)$ :

$$p(\theta|x) \propto p(\theta) \prod_{i=0}^{k-1} p(x_i|\theta) \quad (2.10)$$


---

A finite mixture model is a sum over a finite number of probability distributions:

---

**▼ Definition 2.39**


---

A **finite mixture model** is a sum over a finite number  $k \neq \infty$  of probability distributions  $p(x_i)$ , with each distribution weighted by a factor  $w_i$  with  $\sum_i w_i = 1$ .

$$p(x) = \sum_{i=0}^{k-1} w_i p(x_i) \quad (2.11)$$


---

The mixture model is finite in the sense that there are only  $k \neq \infty$  distributions summed up. The weights of the individual distributions  $p(x_i)$  are normalized (sum up to one) such that the weighted sum of probability distributions is itself a probability distribution.

An infinite mixture model is a sum over an infinite number of probability distributions:

---

**▼ Definition 2.40**


---

A **infinite mixture model** is a sum over an infinite number of probability distributions  $p(x_i)$ , with each distribution weighted by a factor  $w_i$  with  $\sum_i w_i = 1$ .

$$p(x) = \sum_{i=0}^{\infty} w_i p(x_i) \quad (2.12)$$


---

The infinite mixture model is a sum over an infinite number of probability distributions with weights that sum up to one. In this way it assigns a finite value to a countably infinite set of functions.

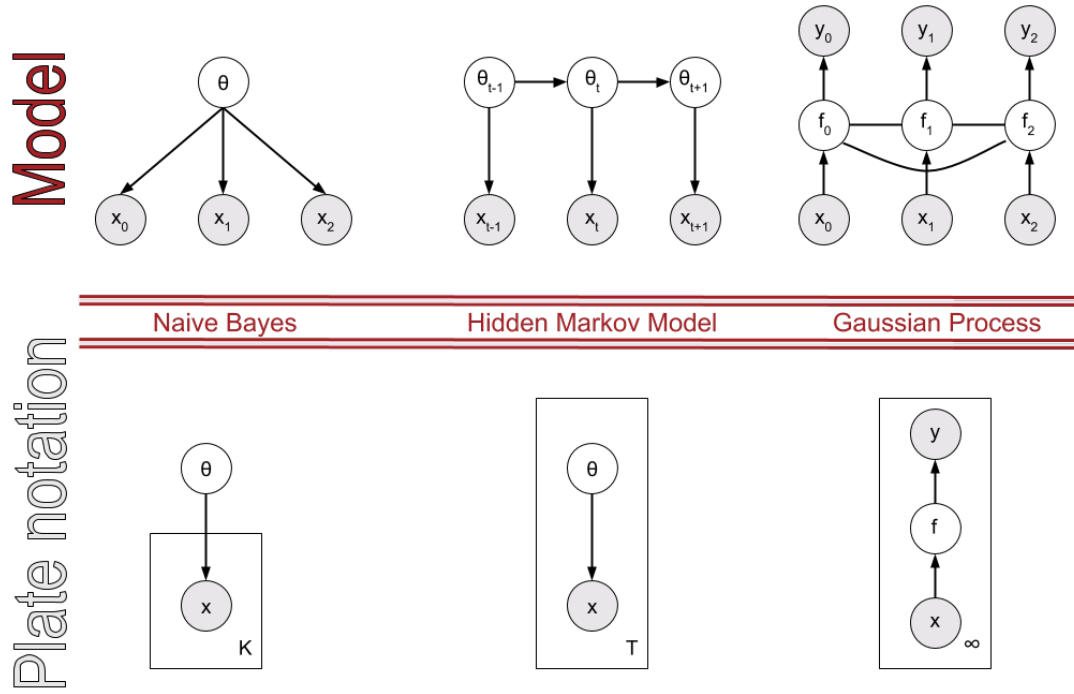
In section 2.1.1 random elements are described in general. Random elements can vary from random vectors, random distributions, random clusters (partitions), to random trees. Table 2.1 describes the random elements and the corresponding examples of random processes in the literature.

**Table 2.1:** A list of mathematical structures and random processes that can be used to generate them.

Structure	Example
Distribution on functions	Gaussian Process
Distribution on distributions	Beta Process
	Gamma Process
	Dirichlet Process
	Polya Tree
Distribution on partitions	Chinese Restaurant Process
	Pitman-Yor Process
Distribution on hierarchical partitions	Dirichlet Diffusion Tree
	Kingman's coalescence
Distribution on sparse binary matrices	Indian Buffet Process
Distribution on integer-valued matrices	Gamma-Poisson Process
Distribution on kd-trees	Mondrian Process

762 The Gaussian Process (Rasmussen and Williams, 2006) describes a distribution on functions.  
 763 The Beta Process (Hjort, 1990), the Gamma Process (Ferguson, 1974), the Dirichlet Process  
 764 and the Polya Tree (Ferguson, 1973) describe a distribution on distributions. The Chinese  
 765 Restaurant Process (Aldous, 1985) and Pitman-Yor Process (Pitman and Yor, 1997) describe  
 766 a distribution on partitions. The Dirichlet Diffusion Tree (Neal, 2001) and Kingman's coales-  
 767 cence (Kingman, 1965) describe a distribution on hierarchical partitions. The Indian Buffet  
 768 Process (Ghahramani and Griffiths, 2005) describes a distribution over sparse binary matri-  
 769 ces. The Gamma-Poisson Process (Titsias, 2008) describes a distribution over integer-valued  
 770 matrices. The Mondrian Process (Roy and Teh, 2009) describes a distribution over kd-trees.

771 Random processes and mixture models are visually represented by a method called *plate*  
 772 *notation* (Koller and Friedman, 2009). Sets of variables are represented in a plate, a rectan-  
 773 gular region (see Fig. 2.3).



**Figure 2.3:** Top: graphical model of a Naive Bayes, Hidden Markov Model, and Gaussian Process. Bottom: corresponding plate notation of the Naive Bayes, Hidden Markov Model, and Gaussian Process. Observed variables are denoted by a circle that is shaded. The plate notation is a representation that does not capture all dependencies. The dependencies between the states  $\theta_0$  and  $\theta_1$  in the Hidden Markov Model are not represented for example.

774 Plate notation is a representation that does not preserve all dependencies between variables.  
 775 The dependencies between the states in the Hidden Markov Model are not represented for  
 776 example.

777 Some random process are mathematically represented by a completely random measure  
 778 (Kingman, 1967), which is defined as follows:

779 **▼ Definition 2.41**

780 A **completely random measure** is a random measure  $\mu : \Omega \times X \rightarrow [0, +\infty]$  from  
 781 probability space  $(\Omega, \mathbb{F}, \mathbb{P})$  to measurable space  $(X, \Sigma)$  with

- 782
- 783 ◦ for any collection of disjoint sets  $A_1, \dots, A_k \in \Sigma$  and  $A_i \cap A_j = \emptyset$  for  $i \neq j$  a mutual  
 784 independency between  $\mu(A_1), \dots, \mu(A_k)$ .
- 785

786 It is shown (Kingman, 1967) that a completely random measure can be decomposed into  
 787 three components:

- 788 1. a deterministic function
- 789 2. a countable set of non-negative random masses at deterministic locations

790 3. a countable set of non-negative random masses at random locations

791 The first component is a deterministic function. The second component has non-negative  
 792 random masses, also called atoms, on deterministic locations. The third component is the  
 793 one of interest. It has a set of random masses (atoms) that can be represented as a Poisson  
 794 random measure on  $\mathbb{R}^+ \otimes X$  with mean measure  $\nu$  which is known as the Lévy intensity  
 795 measure (Favaro et al., 2013).

**Table 2.2:** Lévy measure of the Beta Process (Wang and Carin, 2012), Gamma Process (Knowles et al., 2014), the Dirichlet Process (Lijoi and Prünster, 2010) (indirectly through  $F = 1 - e^{-\nu}$ ).

Random Process	Lévy measure
Beta Process	$\nu(da, dw) = H(da)\alpha w^{-1}(1-w)^{\alpha-1}dw$
Gamma Process	$\nu(da, dw) = H(da)w^{-1}e^{-\alpha w}dw$
Dirichlet Process	$\nu(da, dw) = H(da)\frac{e^{-w\alpha(x, \infty)}}{1-e^{-w}}dw$

796 For Lévy measure decompositions of other processes such as the Indian Buffet Process, see  
 797 Wang and Carin (2012).

798 Recall definition 2.29 for exchangeable sequences. There is a theorem (?) that states that  
 799 there is parameter  $\theta$  such that the data  $x_i$  is conditionally independent given this parameter  
 800 for exchangeable sequences.

#### 801 ▼ Definition 2.42

802 **de Finetti's theorem.** A sequence  $\{x_0, x_1, \dots\}$  of  $(X, \Sigma_X)$ -valued random variables is an  
 803 infinitely exchangeable sequence if and only if there exist a measure  $\mu(d\theta)$  on  $\theta$  such  
 804 that  
 805

$$p(x_0, \dots, x_{k-1}) = \int_{\Sigma_X(X)} \prod_{i=0}^{k-1} p(x_i|\theta) \mu(d\theta) \quad \forall k \geq 1 \quad (2.13)$$

806  
 807  
 808 De Finetti's theorem states that if we have exchangeable data, we have a parameter  $\theta$ , a  
 809 likelihood  $p(x|\theta)$  and some measure  $\mu$  on  $\theta$ , such that the data  $(x_0, \dots, x_{k-1})$  is conditionally  
 810 independent.

811 De Finetti's theorem is not limited to exchangeable sequences. There are corresponding  
 812 theorem's for other exchangeable objects (Orbanz and Roy, 2015):

813 The *stick-breaking presentation* (Freedman and Diaconis, 1983) or residual allocation model  
 814 (Hoppe, 1986; Sawyer and Hartl, 1985)

#### ▼ Definition 2.43

**Table 2.3:** Several exchangeable structures have a theorem that describes an underlying measure that can be sampled i.i.d. Exchangeable sequences (de Finetti, 1930), increments (Bühlmann, 1960), partitions (Kingman, 1978), arrays (Aldous, 1981), and Markov chains (Diaconis and Freedman, 1980).

Mathematical Objects	Theorem
Exchangeable Sequences	de Finetti
Exchangeable Increments	Bühlmann
Exchangeable Partitions	Kingman
Exchangeable Arrays	Aldous-Hoover
Exchangeable Markov Chains	Diaconis-Freedman

An infinite sequence of random variables  $\phi = \{\phi_0, \phi_1, \dots\}$  has a **stick-breaking presentation** with parameters  $\alpha$  and  $\beta$  denoted by  $\phi \sim GEM(\alpha, \beta)$ .

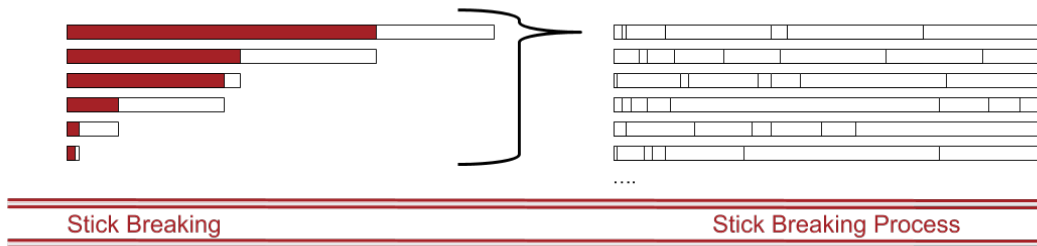
$$w_k \stackrel{i.i.d.}{\sim} Beta(1 - \beta, \alpha + k\beta) \quad k = 1, \dots, K \quad (2.14)$$

$$\phi_k = w_k \prod_{i=1}^{k-1} (1 - w_i) \quad (2.15)$$

815

816

817 The stick-breaking presentation samples repeatedly from a  $Beta(1 - \beta, \alpha + k\beta)$  distribution.  
 818 The result of the process is a vector of  $k$  weights  $\phi_k$ . The abbreviation *GEM* stands for  
 819 Griffiths, Engen, and McCloskey (Ewens, 1990; Ethier, 1990). There is also a variant of  
 820 GEM with a single parameter  $\alpha$  which can be obtained by setting  $\beta = 0$ . In that case  $w_k$  are  
 821 drawn from a  $Beta(1, \alpha)$  distribution.



**Figure 2.4:** The stick-breaking presentation. Left: at the first row, the stick is broken at  $x_0$ , at the next rows the remaining part of the stick is broken  $x_i$  with  $i > 0$ . Only six iterations are shown. Right: samples of a stick-breaking process. The first row shows the stick ratios from the stick-breaking presentation at the left. The next rows show other samples from the same process.

822 Fig. 2.4 visualized the stick-breaking process. A stick of fixed length 1 gets broken at a  
 823 position sampled from a Beta distribution. An infinite number of times it is broken. A stick-  
 824 breaking process generates many of these broken sticks. The stick-breaking presentation is  
 825 convenient for extensions that allow the atoms to vary according to some other stochastic



process (Dunson et al., 2012) and for approximations by truncating the stick-breaking to a limited number of times (Kurihara et al., 2007).

### 2.2.1 Beta Process

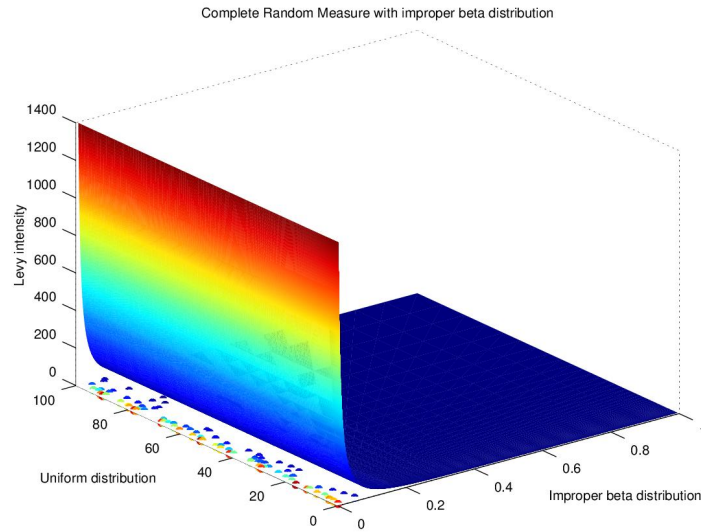
A Beta Process (Hjort, 1990) is a random process with a countably infinite collection of weighted atoms in a space  $(X, \mathbb{B})$  with weights that are in between  $[0, 1]$  (Hjort, 1990).

#### ▼ Definition 2.44

Let  $(X, \mathbb{B})$  be a Borel space,  $\nu$  a finite measure, and  $\alpha > 0$  a scale parameter, then a **Beta Process** is a Lévy process on  $(X, \mathbb{B})$  with its Lévy measure  $\nu$  corresponding to the density:

$$\nu(dw) = \alpha w^{-1}(1-w)^{\alpha-1}dw \quad (2.16)$$

with  $w > 0$ .



**Figure 2.5:** A Completely Random Measure with a Lévy intensity defined on the product space  $\Omega \otimes (0, 1)$ . Here  $\Omega$  is a bounded interval on which the base measure  $B_0 = U(0, 100)$  is defined. On  $(0, 1)$  we define an improper beta distribution  $\alpha w^{-1}(1-w)^{\alpha-1}$ . In this example  $\alpha = 10$ . This is how a Beta Process can be generated from a nonhomogeneous spatial Poisson point process. This has been visualized before (Jordan, 2010). The image is produced by rejection sampling using a homogeneous Poisson point process at  $\max(\nu)$  over  $w = [0.01, 0.9]$ . For  $w \rightarrow 0$  this maximum would go to  $\infty$  and all points would be rejected. Hence, the points should be denser for  $w$  around 0 and should be seen as an approximation of the actual process.

In Fig. 2.5 the Beta Process is generated from a Completely Random Measure with a Lévy intensity defined on  $\Omega \otimes (0, 1)$  (Thibaux and Jordan, 2007). In this case  $\Omega$  is the so-called base measure  $B_0$  and is assumed uniform over a bounded region. The  $(0, 1)$  space is equipped with an improper Beta distribution. It is called improper or degenerate because the scale parameter of the standard Beta distribution is set to zero. This has the consequence that the

integral is infinite:  $\nu(\Omega \otimes (0, 1)) = \infty$ . It is due to the fact that the density  $w^{-1}(1-w)^{\alpha-1}$  goes to infinity for  $w \rightarrow 0$ . That means that a countable infinite number of points can be obtained from the Poisson process.

The Beta Process has a sequential representation in the form of the Indian Buffet Process:

#### ▼ Definition 2.45

An **Indian Buffet Process** is a sequential process that is an exchangeable distribution over matrices:

$$p(z_{i,j} = k | z_{0,0}, \dots, z_{i-1,K_+}) = \begin{cases} \frac{n_{-i,k}}{i} & \text{if } k \leq K_+ \\ \frac{\lambda^{k_{new}} e^{-\lambda}}{k_{new}!} & \text{if } k > K_+ \end{cases} \quad (2.17)$$

Here  $\lambda = \alpha/i$ ,  $k_{new} = K_+ - k$ . The  $i$ 'th data item samples an existing column with a probability of the number of times it has been sampled before divided by its index,  $n_{-i,k}/i$ . It samples a new column with a probability according to a Poisson distribution,  $\lambda^{k_{new}} e^{-\lambda} / k_{new}!$ . The conditional form of the sequential presentation describes a closed-form solution for Gibbs sampling, section 2.3.4 (?).

The Beta Process (BP) is used in linguistics (He et al., 2013; Vanhainen and Salvi, 2012), computer vision (Zhou et al., 2011; Gao and Sun, 2013),

### 2.2.2 Gamma Process

A Gamma process is a random process with independent gamma distributed increments (Ferguson, 1974), formally:

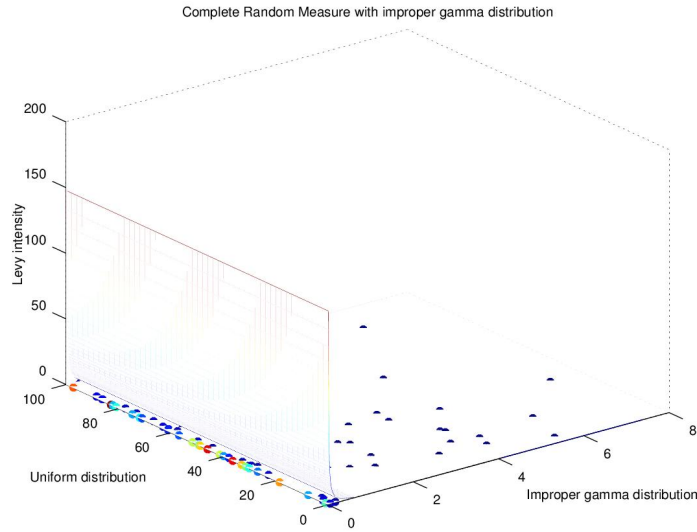
#### ▼ Definition 2.46

Let  $(X, \mathbb{B})$  be a Borel space,  $\nu$  a finite measure, and  $\alpha > 0$  a scale parameter, then a **Gamma Process** is a Lévy process on  $(X, \mathbb{B})$  with its Lévy measure  $\nu$  corresponding to the density:

$$\nu(dw) = w^{-1} e^{-\alpha w} dw \quad (2.18)$$

with  $w > 0$ .

The Gamma Process is used in risk theory (Dufresne et al., 1991), spatial statistics (Wolpert and Ickstadt, 1998; Rao and Teh, 2009), erosion (Singpurwalla, 1997; Abdel-Hameed, 2012), and finance (Madan and Seneta, 1990; K  chler and Tappe, 2008).



**Figure 2.6:** A Completely Random Measure with a Lévy intensity defined on the product space  $\Omega \otimes \mathbb{R}$ . Here  $\Omega$  is a bounded interval on which the base measure  $B_0 = U(0, 100)$  is defined. On  $\mathbb{R}$  we define an improper gamma distribution  $w^{-1}e^{-\alpha w}$ . In this example  $\alpha = 1$ . This is how a Gamma Process can be generated from a nonhomogeneous spatial Poisson point process. The image is produced by rejection sampling in the same way as Fig. 2.5.

### 873 2.2.3 Dirichlet Process

874 The Dirichlet Process (DP) is, just as the Beta Process and the Gamma Process, a distribution  
875 over distributions (Ferguson, 1973).

#### ▼ Definition 2.47

A **Dirichlet process**  $DP$  over a set  $S$  can be used to draw sample paths  $X$ :

$$X \sim DP(\alpha, H)$$

with  $\alpha$  the dispersion parameter and  $H$  a measure on  $S$  and for which any measurable partition  $\{B_0, \dots, B_{n-1}\} \in S$  is drawn from a Dirichlet distribution:

$$(X(B_0), \dots, X(B_{n-1})) \sim \text{Dirichlet}(\alpha H(B_0), \dots, \alpha H(B_{n-1}))$$

876

877

878 The Lévy intensity of the Dirichlet Process is complicated, because it is a so-called normalized  
879 process, see Regazzini et al. (2003).

880 De Finetti's theorem can be used to define the so-called Chinese Restaurant Process (CRP).  
881 This is a *distribution over partitions*. The Dirichlet Process has with the Chinese Restaurant  
882 Process a sequential representation:

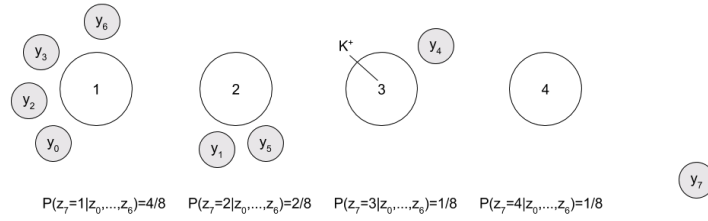
▼ Definition 2.48

A **Chinese Restaurant Process** is a sequential process that is an exchangeable distribution over partitions:

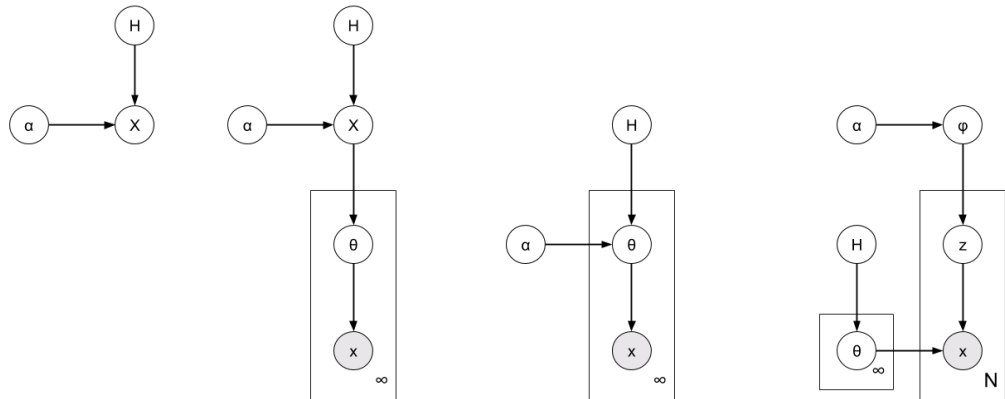
$$p(z_i = k | z_0, \dots, z_{i-1}) = \begin{cases} \frac{n_k}{\alpha + i} & \text{if } k \leq K_+ \\ \frac{\alpha}{\alpha + i} & \text{if } k > K_+ \end{cases} \quad (2.19)$$

The conditional probability of a cluster assignment  $z_i$  for data item  $y_i$  given the cluster assignments  $z_0, \dots, z_{i-1}$  is proportional to the number of data items  $n_k$  assigned to an existing cluster  $k$ , or proportional to  $\alpha$  for a new cluster.

The Chinese Restaurant Process is visualized in Fig. 2.8.



**Figure 2.7:** The Chinese Restaurant Process with  $i$  customers already sitting down. A new customer  $y_{i=7}$  arrives and gets assigned,  $z_i$ . This is an existing table  $\{1, 2, 3\}$  with a probability proportional to the number of customers  $n_i$  sitting at that table:  $n_i/(\alpha + i)$ , or a new, empty table 4 with probability  $1/(\alpha + i)$ . In the visualized Chinese Restaurant Process the dispersion factor  $\alpha = 1$ .



Dirichlet Process      Dirichlet Process Mixture      Chinese Restaurant Process      Stick Breaking Process

**Figure 2.8:** The Dirichlet Process  $X \sim DP(\alpha, H)$ . The Dirichlet Mixture Model where  $X$  is a prior for a sum  $\sum_i w_i p(x_i | \theta_i)$ , a Chinese Restaurant Process with  $X$  marginalized out, and the Stick Breaking Process with a distribution over partition sizes  $\pi$  and indicator variables  $z_i$ .

---

**▼ Definition 2.49**


---

The **stick-breaking presentation** of the Dirichlet Process states that if

$$\phi_k \sim GEM(\alpha, 0) \quad (2.20)$$

$$\theta_k \sim H \quad (2.21)$$

$$G = \sum_{k=1}^{\infty} \phi_k \delta(\theta, \theta_k) \quad (2.22)$$

893 then  $G \sim DP(\alpha, H)$ .

894

---

895 The weights  $\phi_k$  are sampled from the stick-breaking process  $GEM(\alpha, 0)$ . The parameter  
 896 values  $\theta_k$  are sampled from the base measure  $H$ . To sample from the Dirichlet Process we  
 897 have to sample these parameters with the given weights.

898 If the stick-breaking process is used as a prior for a mixture, then the cluster assignments  $z_i$   
 899 are sampled according to the mixing proportions  $\phi$ :

$$\phi \sim GEM(\alpha, 0) \quad (2.23)$$

$$\theta_k \sim H \quad (2.24)$$

$$G = \sum_{k=1}^{\infty} \phi_k \delta(\theta, \theta_k) \quad (2.25)$$

$$z_i \sim Mult(\phi) \quad (2.26)$$

$$x_i \sim F(\theta_{z_i}) \quad (2.27)$$

900 Here  $\theta_k = \theta_{z_i}$  for observation with index  $i$  and cluster assignment  $k$ :  $z_i = k$ .

## 901 2.2.4 Pitman-Yor Process

902 The Pitman-Yor Process introduces another parameter  $d$  with respect to the Dirichlet Process.  
 903 It has been developed by Pitman and Yor as the two-parameter Poisson-Dirichlet distribution  
 904 (Pitman and Yor, 1997). The Pitman-Yor Process has the following definition:

---

**▼ Definition 2.50**


---

A Pitman-Yor process  $PY$  over a set  $S$  can be used to draw sample paths  $X$ :

$$X \sim PY(d, \alpha, H)$$

905 with  $\alpha > -d$  a strength parameter,  $0 \leq d < 1$  a discount parameter, and  $H$  a measure  
 906 on  $S$ .

907

---

908 The Pitman-Yor Process (PYP) generalizes the DP. The Pitman-Yor Process has a stick-  
 909 breaking presentation in which sticks are drawn from  $GEM(\alpha, \beta)$ . The Dirichlet Process  
 910 has a stick-breaking presentation in which sticks are drawn from  $GEM(\alpha, 0)$ , see Def. 2.49.

### ▼ Definition 2.51

---

The **stick-breaking presentation** of the PYP states that if

$$\phi_k \sim GEM(\alpha, \beta) \quad (2.28)$$

$$\theta_k \sim H \quad (2.29)$$

$$G = \sum_{k=1}^{\infty} \phi_k \delta(\theta, \theta_k) \quad (2.30)$$

911 then  $G \sim PYP(\alpha, H)$ .

912

---

913 The Pitman-Yor Process is used in quite a few applications, such as language models (Teh  
 914 et al., 2006), scene segmentation (Sudderth and Jordan, 2009), speech induction (Blunsom  
 915 and Cohn, 2011), and time series (Bassetti et al., 2014).

## 916 2.2.5 Hierarchical Dirichlet Process

917 The Hierarchical Dirichlet Process (HDP) extends the Dirichlet Mixture Model with a hier-  
 918 archical structure (Teh et al., 2006).

### ▼ Definition 2.52

---

A Hierarchical Dirichlet process  $HDP$  over a set  $S$  can be used to draw sample paths  $X$ :

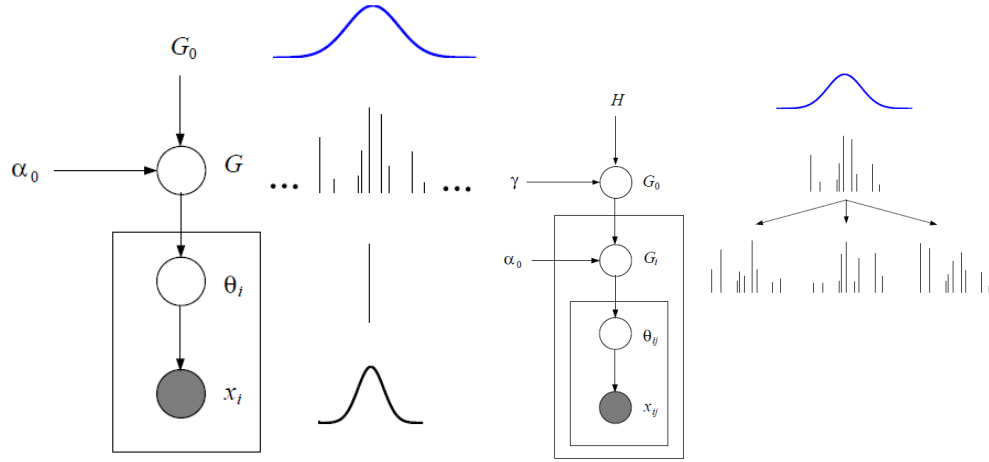
$$G_0 \sim DP(\gamma, H)$$

$$X_i \sim DP(\alpha_0, G_0) \text{ for each group } i$$

919 with a Dirichlet Process with a general  $\gamma$  dispersion parameter and base distribution  $H$   
 920 as a measure on  $S$  of which the generated distributions  $G_0$  are used as base distribution  
 921 for each group distribution  $X_i$ .

922

---



(a) Dirichlet Process Mixture Model. Each draw from the process corresponds to a parameter associated with a distribution (in this case a Gaussian). The only freedom left to express by  $G_i$  is in the weights of those atoms. This reflects a decomposition in a structural and non-structural component.

**Figure 2.9:** The difference visualized between a Dirichlet Process mixture and a Hierarchical Dirichlet Process. It illustrates also that the input of a Dirichlet process does not have to be a continuous function. If it is a continuous distribution it will become a discrete distributed almost surely. If it is a discrete distribution, it will have atoms at the locations where the discrete distribution had its probability mass concentrated.

923 The Hierarchical Dirichlet Process uses the outcome of a Dirichlet Process as a starting point  
 924 to define multiple distributions with atoms at the same locations, while they come equipped  
 925 with different weights. So, the Dirichlet Process on the lower level uses not a continuous  
 926 distribution as input, but a discrete one, generated by the DP at the top layer. Note, that  
 927 the Dirichlet Process will create an a.s. discrete distribution, but it can be fed a discrete  
 928 distribution as prior just fine.

## 929 2.3 Inference

930 There will be six inference methods described, all sampling methods. Section 2.3.1 describes  
 931 inverse transform sampling. Section 2.3.2 describes rejection sampling. Section 2.3.3 de-  
 932 scribes approximate Bayesian computation. Section 2.3.4 describes Gibbs sampling. Sec-  
 933 tion 2.3.5 describes Metropolis-Hastings. Section 2.3.6 describes Split-Merge MCMC.

### 934 2.3.1 Inverse Transform Sampling

935 Let  $p_f(x)$  and  $p_g(x)$  be two probability distributions with  $p_f(x) + p_g(x) = 1$ . Sample with  
 936 probability  $p_f(x)$  by drawing from a uniform distribution  $u \sim U(0, 1)$ , then if  $u < p_f(x)$ ,  
 937 accept  $x$ , if not reject  $x$ . This can be generalized to more than two probability distributions.

938 Let  $p_f(x)$ ,  $p_g(x)$ ,  $p_h(x)$  be multiple probability distributions. Draw  $u$  from the uniform distri-  
 939 bution  $U(0, 1)$ . If  $u < p_f(x)$  we select  $f$ , else if  $u < p_f(x) + p_g(x)$  we select  $g$ , and otherwise

we select  $g$ . The procedure is called inverse transform sampling. The term “inverse” stems from the fact that we return  $x$  (or  $f(x)$ ) given  $u$ . Inverse transform sampling is common to a lot of sampling methods. Often when there are samples from a uniform distribution this is an inverse transform sampling step.

### 2.3.2 Rejection Sampling

Let  $f(x)$  be a complicated function from which it is hard to take samples. Let  $g(x)$  be a simple function that is easy to sample from. Then we can sample from  $f(x)$  by making sure  $Mg(x) > f(x)$ :

$$X = S(f(x), g(x)) \quad (2.31)$$

The rejection sampling method (Halperin and Burrows, 1960) for  $f(x)$  is described in algorithm 1:

---

**Algorithm 1** Rejection sampling for  $f(x)$

---

```

1: procedure REJECTION SAMPLING( $f(x), g(x)$ )      ▷ Target and proposal distribution.
2:   for  $t = 1 \rightarrow T$  do
3:      $x^t \sim g(x)$                                 ▷ Generate  $x^t$  from  $g(x)$ 
4:      $u \sim U(0, 1)$                                 ▷ Inverse transform sampling
5:      $p_0 = f(x)/(Mg(x))$ 
6:     if  $u < p_0$  then
7:        $X = X \cup x^t$                                 ▷ Accept
8:     end if
9:   end for
10:  return  $X$                                           ▷  $X$  will have the distribution of  $f(x)$ 
11: end procedure

```

---

We can also use rejection sampling to sample from the posterior  $f(\theta|x)$  given that we know the form of the likelihood function and that we can sample from the prior. We know that we can sample from the posterior by sampling from  $p(\theta)p(x|\theta)$ . We also know that the prior  $p(\theta)$  necessarily has to be larger than  $p(\theta)p(x|\theta)$  for any observation, because  $p(x|\theta)$  is a probability, hence smaller than one.

$$\Theta = S(\sim p(\theta), p(x|\theta), x) \quad (2.32)$$

In our notation we make explicit that we need  $p(x|\theta)$  for each combination of observations and parameters, but that we only need to sample  $\sim$  from the prior  $p(\theta)$ .



**Algorithm 2** Rejection sampling for  $f(\theta | x)$ 


---

```


1: procedure REJECTION SAMPLING( $p(\theta), p(x|\theta), x$ )      ▷ Requires prior, likelihood and
   observations.
2:   for  $t = 1 \rightarrow T$  do
3:      $\theta^t \sim p(\theta)$                                 ▷ Generate  $\theta^t$  from prior
4:      $u \sim U(0, 1)$                                     ▷ Inverse transform sampling
5:      $p_0 = p(x | \theta)$ 
6:     if  $u < p_0$  then
7:        $\Theta = \Theta \cup \theta^t$                             ▷ Accept
8:     end if
9:   end for
10:  return  $\Theta$                                           ▷  $\Theta$  will have the distribution of  $f(\theta|x)$ 
11: end procedure

```

---

957 In algorithm 2 the envelope distribution  $p(\theta)$  and the target distribution  $p(\theta)p(x | \theta)$ , cancel  
 958 in such way that only  $p(x | \theta)$  remains.

959 Most examples illustrate rejection sampling by estimating the area of a circle, but let us  
 960 visualize the method in the context of sampling (Fig. 2.10).



**Figure 2.10:** A Gaussian is placed over the complex target probability density function. Subsequently the samples that fall in between these two ‘envelopes’ are rejected. This results in a sampling scheme that follows exactly the more complicated probability density function. Note that if the function is scaled by a factor, the sampling scheme stays the same. Such a scaling factor is only important if we want, for example, to know the area under the graph.

### 961 2.3.3 Approximate Bayesian Computation

962 In Approximate Bayesian Computation (ABC) (Rubin and Others, 1984) the likelihood func-  
 963 tion does not need to be calculated<sup>2</sup> (Sisson and Fan, 2011). In contrast, it is assumed that

<sup>2</sup>ABC is also called likelihood-free computation

964 there is a model available that allows to generate observations given the (searched for) pa-  
 965 rameters. In ABC for each configuration of parameters a set of observations is generated.

$$\Theta = S(\sim p(\theta), X, \sim M(\theta), d(X^t, X), \epsilon) \quad (2.33)$$

966 Approximate Bayesian computation uses many tuning parameters. Its most salient charac-  
 967 teristic though is that it generates pseudo-observations through  $\sim M(\theta)$ .

---

**Algorithm 3** Approximate Bayesian computation

---

```

1: procedure APPROXIMATE BAYESIAN COMPUTATION( $p(\theta), X, M, d, \epsilon$ )    ▷ Requires prior,
   observations, model, distance function, and threshold.
2:   for  $t = 1 \rightarrow T$  do
3:      $\theta^t \sim p(\theta)$                                               ▷ Generate  $\theta$  from prior
4:      $X^t \sim M(\theta)$                                               ▷ Simulate observations  $X^t$  from model  $M$ 
5:      $\rho = d(X^t, X)$     ▷ Calculate distance between simulated and actual observations
6:     if  $\rho \leq \epsilon$  then
7:        $\Theta = \Theta \cup \theta^t$     ▷ Accept  $\theta^t$  if distance falls under threshold  $\epsilon$ .
8:     end if
9:   end for
10:  return  $\Theta$     ▷  $\Theta$  will have the distribution of  $f(\theta|X)$ 
11: end procedure

```

---

968 The term Bayesian reflects the fact that a prior is involved. The weight of this prior can be  
 969 manipulated by the threshold  $\epsilon$ . If this threshold is set very low, the prior plays no role and  
 970 only observations are taken into account. If  $\epsilon$  is set extremely high, all  $\theta$  coming from the  
 971 prior will be accepted, and the actual observations are not used in the process. There are  
 972 several disadvantages to approximate Bayesian computation:

- 973     ◦ A set of simulated observations has to be compared with the actual observations. This  
 974       becomes unwieldly if there are many observations.
- 975     ◦ It is possible to use summary statistics rather than the observations themselves. If  
 976       these are sufficient statistics there will be no information lost. If not, there will be  
 977       information loss in practice
- 978     ◦ The distance function suffers from the curse of dimensionality. In the case that the  
 979       dimensionality of the individual observations becomes high, or the number of param-  
 980       eters becomes large, it gets increasingly difficult to come up with a distance function  
 981       which is efficient and accurate at the same time.

### 982 2.3.4 Gibbs Sampling

983 Gibbs sampling (Geman and Geman, 1984) is similar to the *coordinate descent* optimization  
 984 algorithm (Wright, 2015). In coordinate descent a local minimum of a function is found by  
 985 iteratively performing a line search along one coordinate direction at a time. Gibbs sampling

is named after the physicist Gibbs and belongs to the family of Markov chain Monte Carlo (MCMC) methods. Gibbs sampling optimizes over one variate in the multivariate probability distribution at a time. The update value is set and fixed. Then, the next variate is chosen in a round-robin like manner.

$$\Theta = S(X, \sim p(\theta_i | \theta_{-i}, X), \sim p(\theta), B) \quad (2.34)$$

The multiple parameters in the multivariate probability distribution are denoted by  $\theta$ . The parameters are denoted individually with  $\theta_i$ . The set of all parameters except for  $i$  is denoted by  $\theta_{-i}$ . If we sample a parameter we write  $\theta^t$  with  $t$  the iteration or sampling round. The set of parameter samples has capital letter  $\Theta$ .

---

**Algorithm 4** Gibbs sampling
 

---

```

1: procedure GIBBS SAMPLING( $p(\theta_i | \theta_{-i}, X), p(\theta), X, B$ )           ▷ Requires parameters,
   observations and burn-in.
2:    $\theta^0 \sim p(\theta)$                                            ▷ Set parameters to some initial value
3:   for  $t = 1 \rightarrow T$  do
4:     for  $i = 1 \rightarrow k$  do
5:        $\theta_i^t \sim p(\theta_i^{t-1} | \theta_{-i}^t, X)$            ▷ Generate  $\theta_i^t$  from the full conditional probability
6:     end for
7:      $\Theta = \Theta \cup \theta^t$ 
8:   end for
9:    $\Theta_{B:T} \in \Theta$                                            ▷ Get  $\Theta_T$  set, from burn-in  $B$  to end of run  $T$ 
10:   $\Theta \sim \Theta_{B:T}$                                            ▷ Sample  $\Theta$  from correlated  $\Theta_{B:T}$ 
11:  return  $\Theta$ 
12: end procedure

```

---

Gibbs samples are Markovian, the conditional probability only takes into account values at the previous time step  $t - 1$ . When running the Gibbs sampling algorithm long enough, it will visit all possible states eventually. The Markovian property makes subsequent steps correlated, hence when finally extracting the parameter probabilities, it is important not to consider subsequent steps. It is also important to run the algorithm for a while, so it does not suffer from a bad choice of initial parameter values, this is called burn-in. Gibbs sampling works properly because the time the algorithm spends in parts of the space is directly related to the probability of getting into that part of the space.

In the physics literature Gibbs sampling is known as Glauber dynamics or the heat bath algorithm. Observe that Gibbs sampling does not necessary require an actual calculation of the conditional probability in all cases. The obvious exception is for the observations, which are already known. Another, neat optimization procedure arises when conjugate priors are used. A conjugate prior leads to a posterior distribution that can be described analytically. In such a case it is computationally unnecessary to perform sampling. It is much faster to use the actual available analytic description. This is commonly called collapsed Gibbs sampling.

### 1009 2.3.5 Metropolis-Hastings Sampling

1010 Another MCMC algorithm is Metropolis-Hastings (Metropolis et al., 1953), likewise used for  
 1011 high-dimensional distributions. Metropolis-Hastings calculates an acceptance factor  $\alpha$  which  
 1012 takes into account if a step should be taken according to a predefined proposal distribution.  
 1013 In case this step is not accepted, the current sample is resampled.

$$\Theta = S(X, \theta^0, Q(\theta^{t+1}|\theta^t), f(\theta, X)) \quad (2.35)$$

1014 Here we need  $Q(\theta^{t+1}|\theta^t)$  explicitly as well as samples from it.

---

#### Algorithm 5 Metropolis-Hastings sampling

---

```

1: procedure METROPOLIS-HASTINGS SAMPLING( $\theta^0, X, Q, f$ )  $\triangleright$  Requires initial parameters,
   observations, proposal distribution, and function proportional to desired distribution
2:   for  $t = 1 \rightarrow T$  do
3:      $\theta^{t+1} \sim Q(\theta^{t+1}|\theta^t)$   $\triangleright$  Sample from proposal distribution  $Q$ 
4:      $\alpha = \frac{f(\theta^{t+1}, X^{t+1})Q(\theta^t|\theta^{t+1})}{f(\theta^t, X^t)Q(\theta^{t+1}|\theta^t)}$   $\triangleright$  Calculate acceptance
5:      $u \sim U(0, 1)$   $\triangleright$  Inverse transform sampling
6:     if  $\alpha > u$  then
7:        $\Theta = \Theta \cup \theta^{t+1}$   $\triangleright$  Accept  $\theta^{t+1}$ 
8:     else
9:        $\Theta = \Theta \cup \theta^t$   $\triangleright$  Reuse previous sample (note, different from rejection)
10:    end if
11:  end for
12:  return  $\Theta$   $\triangleright$   $\Theta$  will be samples from the distribution  $f(\theta|x)$ 
13: end procedure

```

---

1015 A particular choice of a Metropolis-Hastings step is that of a proposal distribution that does  
 1016 not depend on the state of the chain. This is already suggested by Hastings and is called the  
 1017 independence sampler.

### 1018 2.3.6 Split-Merge MCMC Sampling

1019 The discussed sampling methods do not assume much structure in the model. This means  
 1020 that in hierarchical models sampling either occurs through updating the to be estimated  
 1021 quantities observation by observation or cluster by cluster. Splitting a cluster in two requires  
 1022 moving data points one by one to this new cluster. Much more efficient sampling methods  
 1023 can be obtained if we would be able to reason about larger chunks of cluster assignments.

1024 Split-merge sampling can update cluster assignments for multiple observations at once. It is  
 1025 an adaptation of acceptance method in the Metropolis-Hastings algorithm.

**Algorithm 6** Split-Merge MCMC sampling

---

```

1: procedure SPLIT-MERGE MCMC SAMPLING( $\theta^0, X, Q, f$ )    ▷ Requires initial parameters,
   observations, proposal distribution, and function proportional to desired distribution
2:   for  $t = 1 \rightarrow T$  do
3:      $i \sim D(0, N - 1)$                                 ▷ Sample observation  $i$  discretely
4:      $j \sim D(0, N - 1)$                                 ▷ Sample observation  $j$  discretely
5:     if  $c_i == c_j$  then
6:        $c_{old} = c_i$ 
7:        $\theta_{c_{new}}^{t+1} \sim Q(\theta^{t+1} | \theta^t)$           ▷ Sample from proposal distribution  $Q$ 
8:       for  $k \in c_{old}$  do
9:          $c_k \sim C(c_{old}, c_{new})$                     ▷ Assign to new cluster categorically
10:      end for
11:     else
12:        $c_{merge} = c_i$ 
13:       for  $k \in c_j$  do
14:          $c_k = c_{merge}$                                 ▷ Assign all to first cluster
15:       end for
16:     end if
17:      $\alpha = \frac{f(\theta^{t+1}, X^{t+1})Q(\theta^{t+1} | \theta^t)}{f(\theta^t, X^t)Q(\theta^t | \theta^{t+1})}$           ▷ Calculate acceptance
18:      $u \sim U(0, 1)$                                     ▷ Inverse transform sampling
19:     if  $\alpha > u$  then
20:        $\Theta = \Theta \cup \theta^{t+1}$                         ▷ Accept  $\theta^{t+1}$ 
21:     else
22:        $\Theta = \Theta \cup \theta^t$                             ▷ Reuse previous sample (note, different from rejection)
23:     end if
24:   end for
25:   return  $\Theta$                                           ▷  $\Theta$  will be samples from the distribution  $f(\theta|x)$ 
26: end procedure

```

---

1026 The exact acceptance probability depends on the model at hand. For the mixture model with  
 1027 a Dirichlet Process as prior, its performance is further improved by adjusting the assignment  
 1028 process from random to observation-supported by introducing intermediate restricted Gibbs  
 1029 sampling steps (Jain and Neal, 2004, 2007). Similarly, there are other variants that incor-  
 1030 porate data fit to the splitting step. Labels can for example be calculated sequentially (Dahl,  
 1031 2003) or methods can be used that postulate subcluster structure within clusters to stream-  
 1032 line reasoning over split and merge sets (?).

## 1033 2.4 Chapter Conclusions

1034 Chapter 3 describes Gibbs sampling to perform inference over an infinite set of lines. Gibbs  
 1035 sampling requires conditional probabilities. These are given in closed-form because there  
 1036 is a conjugate description of the line parameters given the points that form the lines in this  
 1037 application.

1038 Chapter 4 describes Split-Merge MCMC sampling to perform inference over an infinite set  
1039 of line segments. The parameters for line segments do not have a conjugate description.  
1040 Metropolis-Hastings can be used to perform inference over the line segments, but the search  
1041 space is quite large. The Split-Merge MCMC method performs faster inference than Metropolis-  
1042 Hastings because it is able to split and merge line segments (with multiple points ascribed  
1043 to them) at once.



# NONPARAMETRIC BAYESIAN LINE DETECTION

## Contents

The nonparametric Bayesian models from the literature (Chapter 2) can be applied to perform inference over point clouds. An example of a point cloud are points distributed over lines in a two-dimensional space. Traditionally, RANSAC and the Hough transform have been used to perform inference over such lines. This chapter uses a nonparametric Bayesian model to perform inference over a countably infinite number of lines. Given a prior with respect to the noise and distribution of points over the lines, Bayesian inference describes the optimal procedure to perform line fitting.

## Published in

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. Nonparametric Bayesian Line Detection. *International Conference on Pattern Recognition and Methods*, ICPRAM 2016, Rome, Italy, February 24-26, 2016. Best paper award in theory and methods track.

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. *Fundamentals of Nonparametric Bayesian Line Detection*. Springer, 2017.

## Outline

The infinite line model describes a collection of lines with a Dirichlet Process as prior (Sect. 3.1). Inference in the infinite line model is performed through Gibbs sampling. (Sect. 3.2). Gibbs sampling over parameters converges slowly. It can be accelerated through sampling over clusters (Sect. 3.3). The inference method results are measured using clustering performance measures (Sect. 3.4). The chapter summarizes the findings (Sect. 3.5) and introduces extensions which will be handled in the next chapters.

In computer vision and particularly in robotics, traditionally the task of line detection has been performed through sophisticated, but ad-hoc methods. We will give two examples



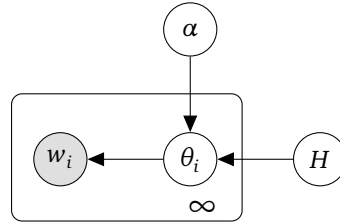
of such methods. RANSAC (Bolles and Fischler, 1981) is a method that iteratively tests a hypothesis. A line is fitted through a subset of points. Then other points that are in consensus with this line (according to a certain loss function) are added to the subset. This procedure is repeated till a certain performance level is obtained. The Hough transform (Hough, 1962) is a deterministic approach which maps points in the image space to curves in the so-called Hough space of slopes and intercepts. A line is extracted by getting the maximum in the Hough space.

There are four main problems with these methods. First, the extension of RANSAC or Hough to the detection of multiple lines is nontrivial (Zhang and Křsecká, 2007; Gallo et al., 2011; Chen et al., 2001). Second, the noise level is hardcoded into model parameters and it is not possible to incorporate knowledge about the nature of the noise. Third, it is hard to extend the model to hierarchical forms, for example, to lines that form more complicated structures such as squares or volumetric forms. Fourth, there are no results known with respect to any form of optimality of the mentioned algorithms.

In this chapter we postulate a method to perform inference over the number of lines and over the fitting of points on that line using the nonparametric Bayesian methods from chapter 2.

### 3.1 Infinite Line Model

The Dirichlet Process described as prior for a mixture distribution (Fig. 2.8 in section 2.2.3) can be used in this particular case as a prior for the distribution of points over a countably infinite set of lines.



**Figure 3.1:** The Infinite Line Model in the Chinese Restaurant Process representation (compare with Fig. 2.8). Top:  $\alpha$ , the concentration parameter of the Dirichlet Process. Bottom, left to right:  $w_i$ , the observation, an individual point in a 2D space;  $\theta_i$ , the parameters (intercept, slope) of the line belonging to observation  $w_i$ ;  $H$ , the base distribution from which line parameter values are sampled.

The infinite line model is visualized (Fig. 3.2) through plate notation (Buntine, 1994). In section 3.1.1 it is described how  $\theta_i$  is sampled from  $H$  and  $\alpha$ . In section 3.1.2 it is described how  $w_i$  is sampled from  $\theta_i$ . In section 3.1.3 the prior  $H$  for  $\theta_i$  is described.

#### 3.1.1 Posterior Predictive for a Line given Other Lines

Let us reiterate the definition of the Dirichlet Process. Let  $H$  be a distribution over  $\Theta$ , let  $\alpha$  be scalar. The Dirichlet Process generates a distribution  $G \sim DP(\alpha, H)$ :

$$G(\theta_1, \dots, \theta_\infty) \sim DP(\alpha, H(\theta_1, \dots, \theta_\infty)) \quad (3.1)$$

1098 A Dirichlet Process assigns a Dirichlet distribution to every parameter partition  $\Theta_1, \dots, \Theta_r$ :

$$(G(\Theta_1), \dots, G(\Theta_r)) \sim Dir(\alpha H(\Theta_1), \dots, \alpha H(\Theta_r)) \quad (3.2)$$

1099 The Dirichlet is conjugate to the categorical:

$$(G(\Theta_1), \dots, G(\Theta_r)) \mid \theta_1, \dots, \theta_n \sim Dir(\alpha H(\Theta_1) + n_1, \dots, \alpha H(\Theta_r) + n_r) \quad (3.3)$$

$$n_k = \sum_{j=1}^n \delta_{\theta_j}(\Theta_k)$$

1100 In the above notation,  $\delta_{\theta_j}(\Theta_k)$  is a Dirac measure (a generalization of the Dirac delta func-  
 1101 tion), also known as an indicator function. Given a set  $\Theta_k$  with a  $\sigma$ -algebra over subsets of  
 1102  $\Theta$ :

$$\delta_{\theta_j}(\Theta_k) = 1_{\Theta_k}(\theta_j) = \begin{cases} 1 & \text{if } \theta_j \in \Theta_k \\ 0 & \text{if } \theta_j \notin \Theta_k \end{cases} \quad (3.4)$$

1103 The posterior for the Dirichlet Process base distribution and dispersion parameter is a Dirich-  
 1104 let Process with adjusted parameters:

$$G(\cdot) \mid \theta_1, \dots, \theta_n \sim DP\left(\alpha + n, \frac{\alpha}{\alpha + n} H(\cdot) + \frac{n}{\alpha + n} \frac{\sum_{j=1}^n \delta_{\theta_j}(\cdot)}{n}\right) \quad (3.5)$$

1105 The posterior base distribution  $G$  is a weighted average between the prior base distribution  
 1106  $H$  and the empirical distribution  $n^{-1} \sum_{j=1}^n \delta_{\theta_j}$  with the weights respectively  $\alpha$  and  $n$  (nor-  
 1107 malized). The dispersion parameter  $\alpha$  is updated to  $\alpha + n$ . Note that  $\delta_{\theta_j}(\cdot)$  is a distribution,  
 1108 the Dirac measure Eq. 3.4.

1109 The posterior predictive for a new parameter  $\theta_{n+1}$  has the form:

$$P(\theta_{n+1} \in \Theta_k \mid \theta_1, \dots, \theta_n) = \frac{1}{\alpha + n} \left( \alpha H(\Theta_k) + \sum_{j=1}^n \delta_{\theta_j}(\Theta_k) \right) \quad (3.6)$$

1110 In other words, the posterior predictive of  $\theta_{n+1}$  given the parameters  $\theta_1, \dots, \theta_n$  in Eq. 3.6 has  
 1111 exactly the same form as the posterior base distribution  $G$  given the parameters  $\theta_1, \dots, \theta_n$   
 1112 (Blackwell and MacQueen, 1973) in Eq. 3.5, namely:

$$\theta_{n+1} \mid \theta_1, \dots, \theta_n \sim \frac{1}{\alpha + n} \left( \alpha H(\theta_{n+1}) + \sum_{j=1}^n \delta(\theta_j - \theta_{n+1}) \right) \quad (3.7)$$

1113 A normal Dirac delta function  $\delta(\theta_j - \theta_{n+1})$  can be used here, which is only non-zero when  
 1114  $\theta_j$  is equal to  $\theta_{n+1}$ .

1115 Equivalently, if we describe  $\theta_n$  conditioned on  $\theta_1, \dots, \theta_{n-1}$  we have to run over  $n - 1$  rather  
 1116 than  $n$  parameters:

$$\theta_n \mid \theta_1, \dots, \theta_{n-1} \sim \frac{1}{\alpha + n - 1} \left( \alpha H(\theta_n) + \sum_{j=1}^{n-1} \delta(\theta_j - \theta_n) \right) \quad (3.8)$$

1117 Due to the exchangeability property we can also consider any other parameter update (Neal,  
 1118 2000):

$$\theta_i \mid \theta_{-i} \sim \frac{1}{\alpha + n - 1} \left( \alpha H(\theta_i) + \sum_{j \neq i} \delta(\theta_j - \theta_i) \right) \quad (3.9)$$

1119 The notation  $\theta_{-i}$  means every parameter  $\theta$  except for the one equal to  $\theta_i$ .

### 1120 3.1.2 Likelihood of Data given a Line

1121 The likelihood of data given line parameters is defined to be according to the **Bayesian linear**  
 1122 **regression** model. The Bayesian linear regression model for a single line (Box and Tiao,  
 1123 2011) assumes a linear relationship between the independent  $x_i$  and dependent variables  
 1124  $y_i$  with Gaussian noise added in the  $y$ -direction. The individual points  $i$  are drawn from a  
 1125 Normal distribution:

$$y_i \sim \mathcal{N}(x_i \beta, \sigma^2) \quad (3.10)$$

1126 The (column) vector  $\beta$  maps the (row) vector with independent variables  $x_i$  to the depen-  
 1127 dent variable  $y_i$ . The noise is normally distributed with standard deviation  $\sigma$  along the  
 1128 dimension of the dependent variable.

1129 In a 2D point cloud the point  $p$  is represented by  $(x_p, y_p)$ . The points are mapped into an  
 1130 intercept-slope representation through defining  $X_i = [1, x_p]$  and  $y_i = y_p$ . The vector  $\beta$  will  
 1131 then contain the  $y$ -intercept as the first value, the slope as the second value.

1132 All observations that belong to the same single line lead to a likelihood function that corre-  
 1133 sponds to a normally distributed random variable with  $y$  and  $X$  as parameters:

$$p(y \mid X, \beta, \sigma^2) \propto \sigma^{-n} \exp \left( -\frac{1}{2\sigma^2} (y - X\beta)^T (y - X\beta) \right) \quad (3.11)$$

1134 The dependent variable is now a column vector of values  $y$  and each observation has a row  
 1135 of independent variables in  $X$ . The vector  $\beta$  and the standard deviation  $\sigma$  are shared across  
 1136 all observations. The term  $y - X\beta$  is written out like this:

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} - \begin{pmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix} \quad (3.12)$$

1137 Note that Eq. 3.11 has exactly the same form for a single point or multiple points that belong  
 1138 to the same line. Hence, we have the probability of a point  $w_i$  given the line parameters  
 1139  $\theta_k = (\beta_k, \sigma_k)$ :

$$F(w_i, \theta_k) = p(w_i | \theta_k) = p(w_i | \beta_k, \sigma_k^2) = p(y_i | X_i, \beta_k, \sigma_k^2) \quad (3.13)$$

1140 To get the full distribution  $p(w_i, \beta, \sigma^2)$  we will need also  $p(\beta, \sigma^2)$ .

### 1141 3.1.3 Conjugate Prior for a Line

1142 The conjugate prior for the likelihood in Eq. 3.11 is a product of a prior for the standard  
 1143 deviation  $p(\sigma)$  and the conditional probability of the line coefficients given the standard  
 1144 deviation  $p(\beta | \sigma^2)$ .

$$p(\sigma^2, \beta) = p(\sigma^2)p(\beta | \sigma^2) \quad (3.14)$$

1145 The standard deviation  $\sigma$  is sampled from an Inverse-Gamma (IG) distribution:

$$p(\sigma) \propto (\sigma^2)^{-(\nu_0/2+1)} \exp\left(-\frac{1}{2\sigma^2} \nu_0 s_0^2\right) \quad (3.15)$$

1146 This is an  $IG(a_0, b_0)$  with  $a_0 = \nu_0/2$  and  $b_0 = 1/2 \nu_0 s_0^2$ . The conditional with respect to the  
 1147 line coefficients has a normal distribution as prior:

$$p(\beta | \sigma^2) \propto \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2} (\beta - \mu_0)^T \Lambda_0 (\beta - \mu_0)\right) \quad (3.16)$$

1148 Let us collect  $\Lambda_0, \mu_0, a_0, b_0$  into  $\lambda_0$ , we have now a description of our base distribution  $H$ :

$$H(\theta_k) = NIG(\theta_k; \lambda_0) \quad (3.17)$$

1149 NIG is an abbreviation of the Normal-Inverse-Gamma distribution. The standard deviation  
 1150 is sampled from the Gamma distribution with  $a_0$  and  $b_0$  as hyperparameters and the line  
 1151 coefficients from a Normal distribution:

$$\begin{aligned} \sigma_k &= \tau_k^{-1/2} & \tau_k &\sim \mathcal{G}(a_0, b_0) \\ \mu_k &\sim \mathcal{N}(\mu_0, \sigma^2 \Lambda_0^{-1}) \end{aligned} \quad (3.18)$$

### 1152 3.1.4 Posterior Predictive for a Line given Data

1153 Due to the fact that it is a conjugate distribution we have a simplified description for updating  
 1154 the parameters at once, given a set of observations. The sufficient statistics are updated  
 1155 (Minka, 2000) according to:

$$\begin{aligned}\Lambda_n &= (X^T X + \Lambda_0) \\ \mu_n &= \Lambda_n^{-1}(\Lambda_0 \mu_0 + X^T y) \\ a_n &= a_0 + n/2 \\ b_n &= b_0 + 1/2(y^T y + \mu_0^T \Lambda_0 \mu_0 - \mu_n^T \Lambda_n \mu_n)\end{aligned}\tag{3.19}$$

1156 Let us collect  $\Lambda_0, \mu_0, a_0, b_0$  into  $\lambda_0$  and  $\Lambda_n, \mu_n, a_n, b_n$  into  $\lambda_n$ . Let us collect a set of our obser-  
 1157 vations and  $(X, y)_k$  into  $w_k$ . The update for the sufficient statistics can then be summarized  
 1158 as:

$$\lambda_n = U_{ss}(\lambda_0, w)\tag{3.20}$$

1159 If we combine this update with sampling  $\theta_k$  from  $\lambda_n$  according to Eq. 3.17, then we obtain:

$$p(\theta_k | \lambda_0, w_k) \propto F(w_k, \theta_k) H(\theta_k; \lambda_0) = p(\theta_k | \lambda_n) = NIG(\theta_k; \lambda_n)\tag{3.21}$$

1160 Sampling of  $NIG(\theta_k; \lambda_n)$  is as in Eq. 3.18, but with  $\lambda_n$  rather than  $\lambda_0$ .

1161 Let us integrate over  $\theta$  (through the function  $H$ ):

$$Q(w_k, \lambda_0) = \int_{\Theta} F(w_k, \theta) dH(\theta; \lambda_0)\tag{3.22}$$

## 1162 3.2 Inference for the Infinite Line Model

1163 The posterior predictive for parameters (see Eq. 3.9) combined with observations  $w_i$  is de-  
 1164 scribed by:

$$p(\theta_i | \theta_{-i}, w_i) \propto r_i H_i(\theta_i) + \sum_{j \neq i} L_{i,j} \delta(\theta_j - \theta_i)\tag{3.23}$$

1165 The posterior is proportional (indicated by  $\propto$ ) to three terms. First, the  $\alpha$  weighted posterior  
 1166 predictive  $r_i$  for a new cluster. Second, the posterior to sample from  $H_i(\theta_i)$  with probability  
 1167  $r_i$ . Third, the likelihood of an observation given a line  $L_{i,j}$ :

$$r_i = \alpha Q(w_i, \lambda_0) = \alpha \int_{\Theta} F(w_i, \theta) dH(\theta)\tag{3.24}$$

1168 The posterior  $H_i(\theta)$  is the normalized product of the prior distribution  $H(\theta_i)$  with the like-  
 1169 lihood  $F(w_i, \theta_i)$  for a single observation  $w_i$ .

$$H_i(\theta_i) \propto H(\theta_i)F(w_i, \theta_i) \quad (3.25)$$

$$L_{i,j} = F(w_i, \theta_j) \quad (3.26)$$

1170 Sampling a new cluster parameter from  $H_i(\theta_i)$  is done with probability:

$$p(\theta_{new}) = \frac{r_i}{r_i + \sum_{j \neq i} L_{i,j}} \quad (3.27)$$

1171 We can use this to derive the parameters  $\theta_i$ .

---

**Algorithm 7** Gibbs sampling over parameters  $\theta_i$

---


```

1: procedure GIBBS ALGORITHM 1( $w, \lambda_0, \alpha$ )           ▷ Accepts points  $w$ , hyperparameters  $\lambda_0, \alpha$  and
   returns  $k$  line coordinates
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:        $r_i = \alpha Q(w_i, \lambda_0)$            ▷ Posterior predictive of  $w_i$  given hyper parameters (Eq. 3.24)
5:       for all  $j = 1 : N, j \neq i$  do
6:          $L_{i,j} = F(w_i, \theta_j)$            ▷ Likelihood for a line given observation (Eq. 3.26)
7:       end for
8:        $p(\theta_{new}) = \frac{r_i}{r_i + \sum_{j \neq i} L_{i,j}}$    ▷ Probability of sampling a new parameter (Eq. 3.27)
9:        $u \sim U(0, 1)$ 
10:      if  $p(\theta_{new}) > u$  then                 ▷ Sample with probability  $p(\theta_{new})$ 
11:         $\lambda_n = U_{ss}(w_i, \lambda_0)$            ▷ Update sufficient statistics with  $w_i$  (Eq. 3.20)
12:         $\theta_i \sim NIG(\theta_i; \lambda_n)$          ▷ Sample  $\theta_i$  from NIG (Eq. ??)
13:      else
14:         $\theta_i$  sampled from existing clusters   ▷ Sample old cluster
15:      end if
16:    end for
17:  end for
18:  return summary on  $\theta_k$  for  $k$  lines
19: end procedure

```

---

1172 This Gibbs algorithm is described in its general form before (Neal, 2000) (Algorithm 1). We  
 1173 perform a loop in which for  $T$  iterations each  $\theta_i$  belonging to observation  $w_i$  is updated  
 1174 in sequence. First, the posterior predictive for  $w_i$  given the hyperparameters  $p(w_i | \lambda_0)$  is  
 1175 calculated. Second, the likelihood  $L_{i,j}$  for all  $\theta_j$  given  $w_i$  (with  $j \neq i$ ) is calculated. Third,  
 1176 the fraction with  $r_i$  defines if  $\theta_i$  will be sampled from a new cluster or if one of the existing  
 1177 clusters will be sampled. Fourth, a new cluster is sampled, the sufficient statistics are up-  
 1178 dated with information on  $w_i$  and thereafter  $\theta$  is sampled from a Normal-Inverse-Gamma  
 1179 distribution with the updated hyperparameters. That, or an existing cluster will be sampled.



**Figure 3.2:** The Infinite Line Model in the stick-breaking representation (compare with Fig. 2.8). From left to right:  $\alpha$ , the concentration parameter of the Dirichlet Process;  $(\phi_1, \dots, \phi_k)$ , the partition of points over lines;  $z_i$ , the assignment parameters that link observation  $w_i$  with line  $k$ ;  $w_i$ , the observation, an individual point with  $x$  and  $y$  coordinates;  $\theta_k$ , the parameters of line  $k$ ;  $\lambda_0$ , the base measure from which the line parameter values are sampled.

### 3.3 Accelerating Inference for the Infinite Line Model

It is also possible to iterate only over the clusters. The derivation takes a few steps (Neal, 2000) but leads to a simple update for the component indices that only depends on the number of data items per cluster, the parameter  $\alpha$ , and the data at hand.

---

#### Algorithm 8 Gibbs sampling over clusters $c_k$

---

```

1: procedure GIBBS ALGORITHM 2( $w, \lambda_0, \alpha$ )    ▷ Accepts points  $w$  and hyperparameters  $\lambda_0$  and  $\alpha$ ,
   returns  $k$  line coordinates
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:        $c = \text{cluster}(w_i)$                     ▷ Get cluster  $c$  currently assigned to observation  $w_i$ 
5:        $\lambda_c = \text{downdate}(w_i, \lambda_c)$         ▷ Adjust sufficient statistics for cluster  $c$  by removing
   observation  $w_i$ 
6:        $m_c = m_c - 1$     ▷ Adjust cluster size  $m_c$  (observation  $w_i$  removed reduces it with one)
7:       for all  $k = 1 : K$  do
8:          $L_k = m_k F(w_i, \theta_k)$             ▷ Update likelihood for cluster  $k$  given observation  $w_i$ 
9:       end for
10:       $r_i = Q(w_i, \lambda_0)$                     ▷ Posterior predictive of  $w_i$  given hyper parameters
11:       $p(\text{new}) = \frac{r_i}{r_i + \sum_k L_k}$           ▷ Sample new or old?
12:      if  $p(\text{new})$  then
13:         $\lambda_k = U_{ss}(w_i, \lambda_0)$             ▷ Update sufficient statistics with observation  $w_i$ 
14:         $\theta_i \sim \text{NIG}(\lambda)$                   ▷ Sample  $\theta_i$  from NIG
15:      else
16:         $k$  sampled from existing clusters
17:         $\lambda_k = \text{update}(w_i, \lambda_k)$           ▷ Restore sufficient statistics with observation  $w_i$ 
18:      end if
19:       $m_k = m_k + 1$                             ▷ Increment cluster size  $m_k$ 
20:    end for
21:    for all  $k = 1 : K$  do
22:       $\theta_k \sim \text{NIG}(\lambda_k)$                 ▷ Sample  $\theta_k$  from NIG
23:    end for
24:  end for
25:  return summary on  $\theta_k$  for  $k$  lines
26: end procedure

```

---

The probability to sample from a cluster depends on the number of items in that cluster (except the data item at hand). This is expressed in equation 3.28.

$$p(c_i = c \text{ and } c_i = c_j \text{ and } i \neq j \mid c_{-i}, w_i, \alpha, \theta) \propto \frac{n_{c,-i}}{\alpha + n - 1} F(w_i \mid \theta_i) \quad (3.28)$$

1186 The probability to sample a new cluster only depends on  $\alpha$  and the total number of data  
1187 items. This is described in equation 3.29.

$$p(c_i \in \Omega(c) \text{ and } c_i \neq c_j \text{ and } i \neq j \mid c_{-i}, \alpha) \propto \frac{\alpha}{\alpha + n - 1} \int F(w_i \mid \theta_i) dH(\theta) \quad (3.29)$$

1188 Here  $\Omega(c)$  denotes all admitted values for  $c_i$ .

1189 The importance of conjugacy is obvious from Eq. 3.29, it will lead to an analytic form of the  
1190 integral. The inference method using equations 3.28 and 3.29 is described in section ??.

1191 Directly sampling over the clusters is described in its general form (Neal, 2000) (Algo-  
1192 rithm 2).

1193 Rather than updating each  $\theta_i$  per observation  $w_i$ , an entire cluster  $\theta_k$  is updated. In Algo-  
1194 rithm 7 the update of a cluster would require a first observation to generate a new cluster  
1195 at  $\theta_j$  and then moving all observations of the old cluster  $\theta_i$  to  $\theta_j$ .

1196 Algorithm 8 follows the same procedure in excluding  $w_i$  from calculating the likelihood. This  
1197 requires the previously mentioned “downdate” from the corresponding sufficient statistics.  
1198 In Algorithm 8 after all observations have been iterated over and assigned the corresponding  
1199 cluster  $k$ , an outer loop iterates over all clusters to obtain new parameters  $\theta$  from the NIG  
1200 prior.

## 1201 3.4 Performance of the Infinite Line Model

1202 The Infinite Line Mixture Model (see section ??) is able to fit an infinite number of lines  
1203 through a point cloud in two dimensions. These lines are no line segments, but infinite  
1204 lines. However, to test the model a variable number of lines are generated of a length that is  
1205 considerably larger compared to the spread caused by the standard deviation of points from  
1206 that line.

1207 As described before, Gibbs sampling leads to correlated samples. We choose to get the  
1208 Maximum A Posterior estimates for our clusters by picking the median values for all the  
1209 parameters involved.

### 1210 3.4.1 Clustering Performance

1211 The results are measured using conventional metrics for clustering performance. For exam-  
1212 ple the Rand Index describes the accuracy of cluster assignments (Rand, 1971):

$$R = \frac{a + b}{a + b + c + d} \quad (3.30)$$



Here  $a$  numbers the pair of points that belong to the same cluster, both at ground truth as well as after the inference procedure. Likewise  $b$  numbers the pair of points that belong to different clusters in both sets. The values  $c$  and  $d$  describe discrepancies between the ground truth and the results after inference. A Rand Index of one means that there have been no mistakes.

The clustering performance is separate from the line estimation performance. If the points are not properly assigned, the line will not be estimated correctly. Due to the fact that line estimation has this secondary effect, this performance is not taken into account. Moreover, from lines that generated only a single, or very few points, we can extract point assignments, but line coefficients are impossible to derive. This would lead to introducing a threshold for the number of points per cluster. Moreover, the performance would then need to be measured by weighting the fitting versus the assignment.

The performance of Algorithm 1 can be seen in Fig. 3.3 and is rather disappointing. On average the inference procedure agrees upon the ground truth for 75% of the cases considering the Rand Index. Moreover, if we adjust for chance as with the Adjusted Rand Index, the performance drops to only having 25% correct!

Algorithm 2 leads to stellar performance measures (Fig. 3.4). Apparently updating entire clusters at once with respect to their parameter values leads at times to perfect clustering, bringing the performance metrics close to their optimal values.

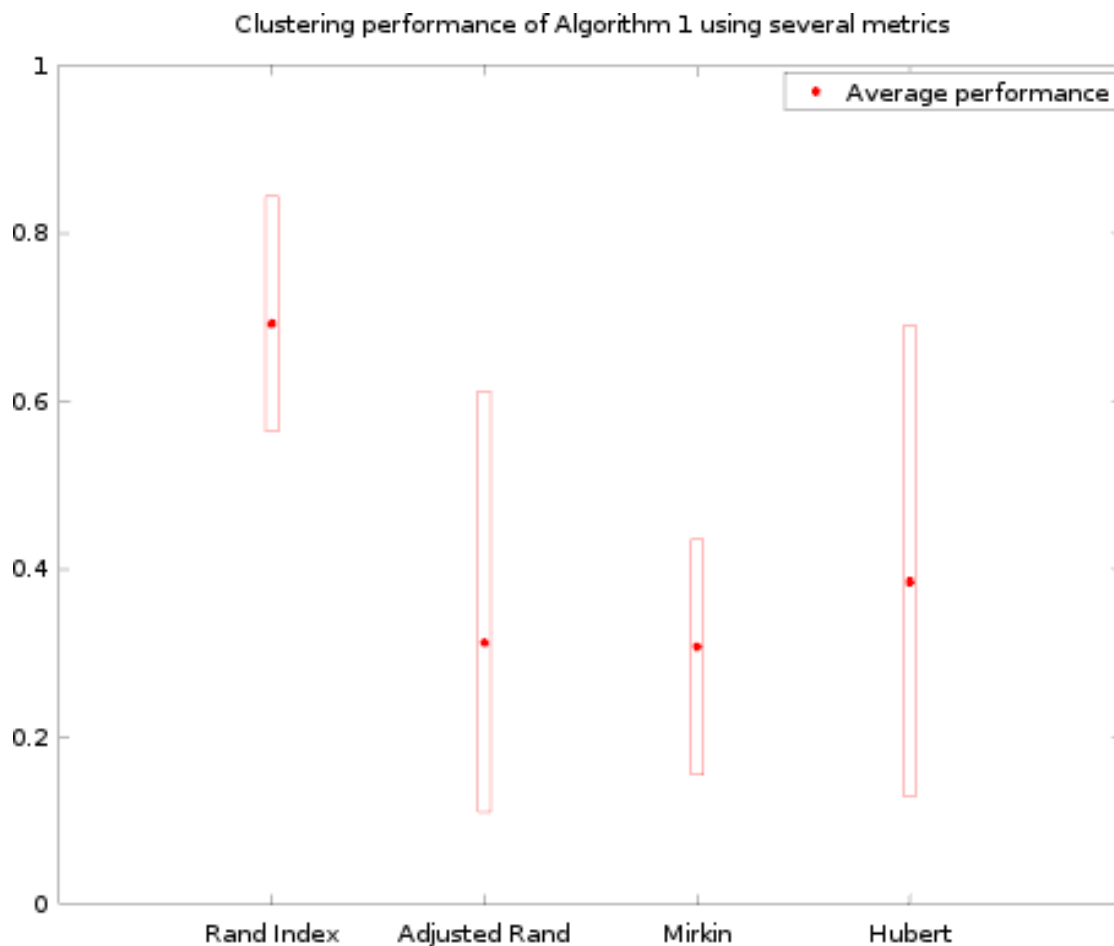
The lack of performance of Algorithm 1 is not only caused by slower mixing (time required to reach the steady state distribution). Also when allowing it ten times the number of iterations of Algorithm 2, it still does not reach the same performance levels. A line seems to form local regions of high probability making it difficult for points to postulate slightly changed line coordinates.

### 3.4.2 Two Examples

In the following we show two examples to understand the inference process better. Figure 3.5 shows the assignment after a single Gibbs step in Algorithm 1. There is a single line that is represented by two clusters. Algorithm 1 does not have merge or split steps to group these clusters at once, it thus has to move each data point one by one. By the way, there are split-merge algorithms that take these more sophisticated Gibbs steps into account (Jain and Neal, 2004).

The example in Fig. 3.6 shows that a single point as an outlier is not a problem for our method. A single point might throw off Bayesian linear regression, but because there are multiple lines to be estimated in our Infinite Line Mixture Model, this single point is assigned its own line.

The extension to more points as outliers would of course require us to postulate a distribution for these outlier points as well. A uniform distribution might for example be used in tandem with the proposed model. This however would lead to a non-conjugate model and hence different inference methods.

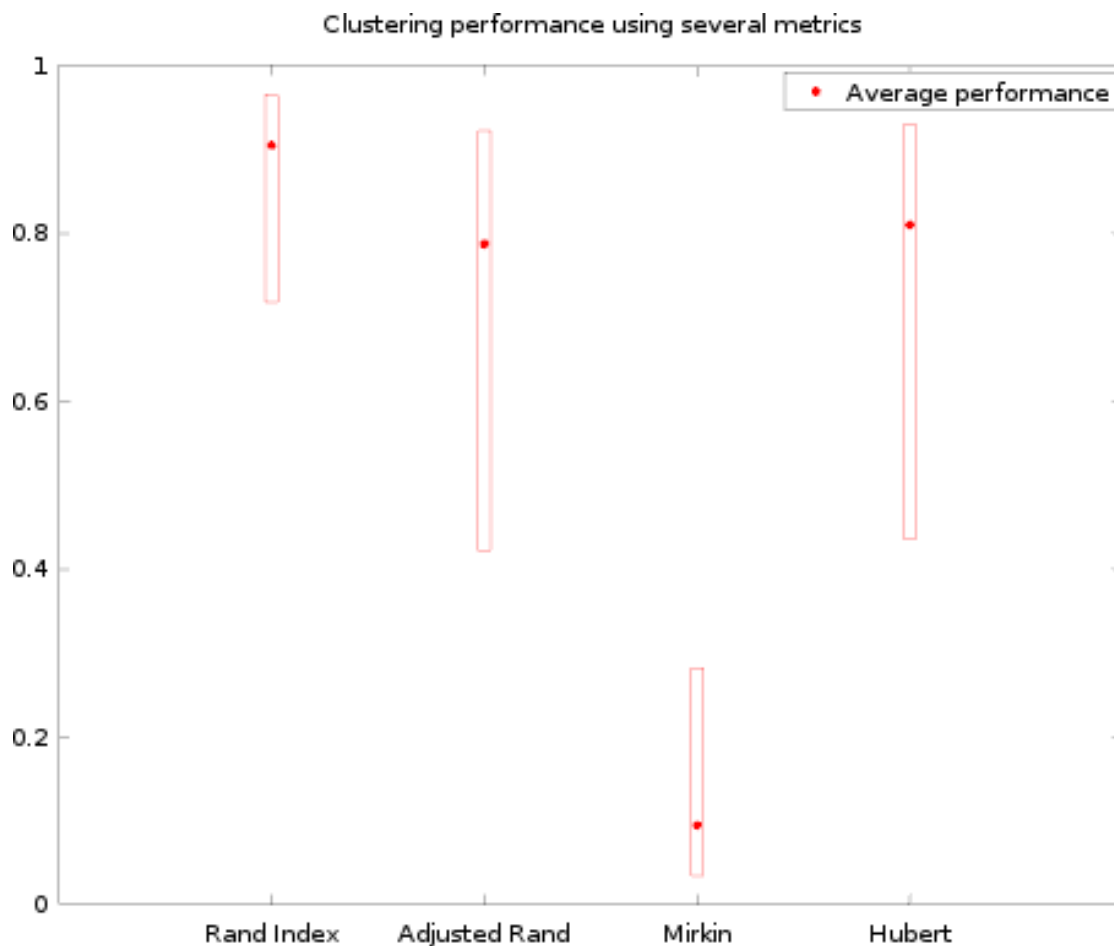


**Figure 3.3:** The performance of Algorithm 1 with respect to clustering is measured using the Rand Index, the Adjusted Rand Index, the Mirkin metric, and the Hubert metric. A figure of 1 means perfect clustering for all metrics, except Mirkin's where 0 denotes perfect clustering.

## 1252 3.5 Chapter Conclusions

1253 The Infinite Line Mixture Model that is proposed extends the familiar Bayesian linear regres-  
 1254 sion model to an infinite number of lines using a Dirichlet Process as prior. The model is a  
 1255 full Bayesian method to detect multiple lines. A full Bayesian method, in contrast to ad-hoc  
 1256 methods such as the Hough transform or RANSAC, means optimal inference (Zellner, 1988)  
 1257 given the model and noise definition.

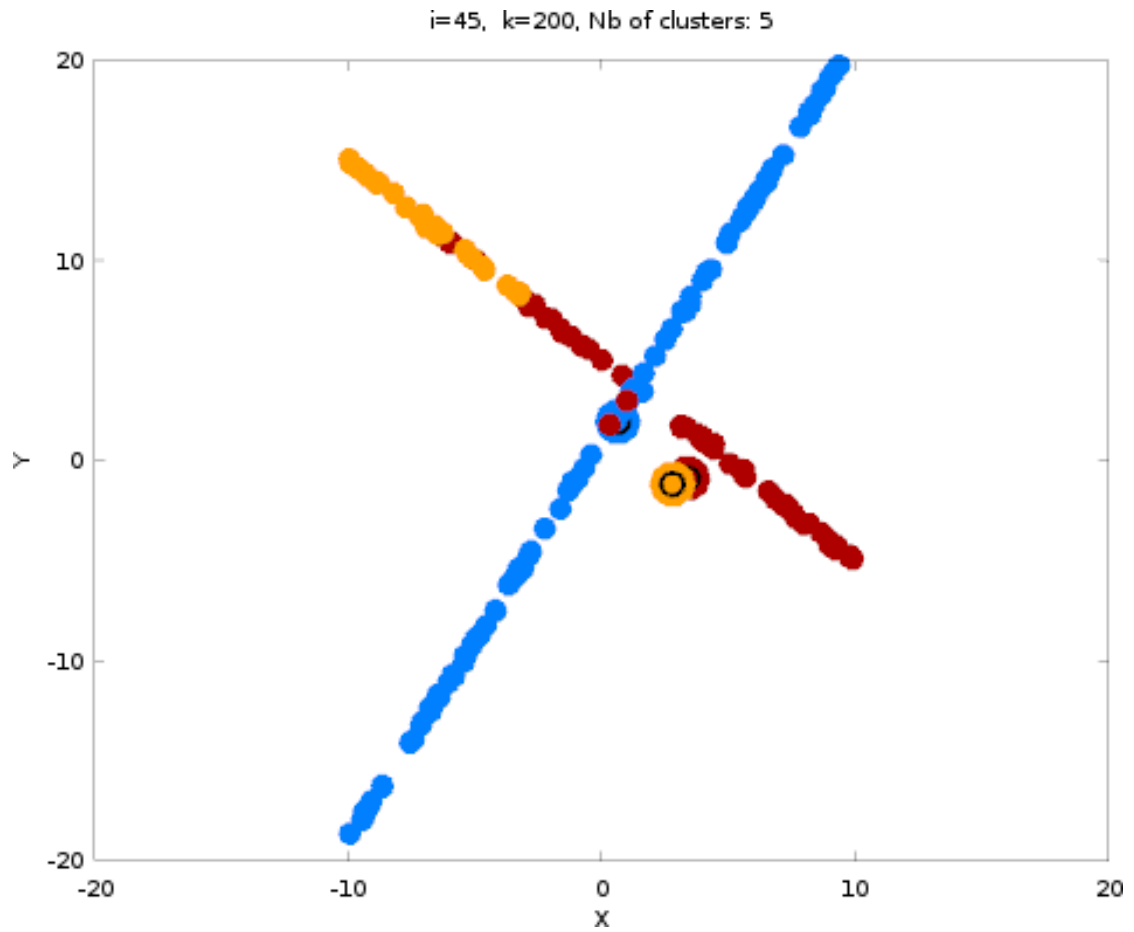
1258 Results in section ?? show high values for difference performance metrics for clustering,  
 1259 such as the Rand Index, the Adjusted Rand Index, and other metrics. The Bayesian model  
 1260 is solved through two types of algorithms. Algorithm 7 iterates over all observations and  
 1261 suffers from slow mixing. The individual updates makes it hard to reassign large number of  
 1262 points at the same time. Algorithm 8 iterates over entire clusters. This allows updates for  
 1263 groups of points leading to much faster mixing. Note, that even optimal inference results  
 1264 in occasional misclassifications. The dataset is generated by a random process. Hence,



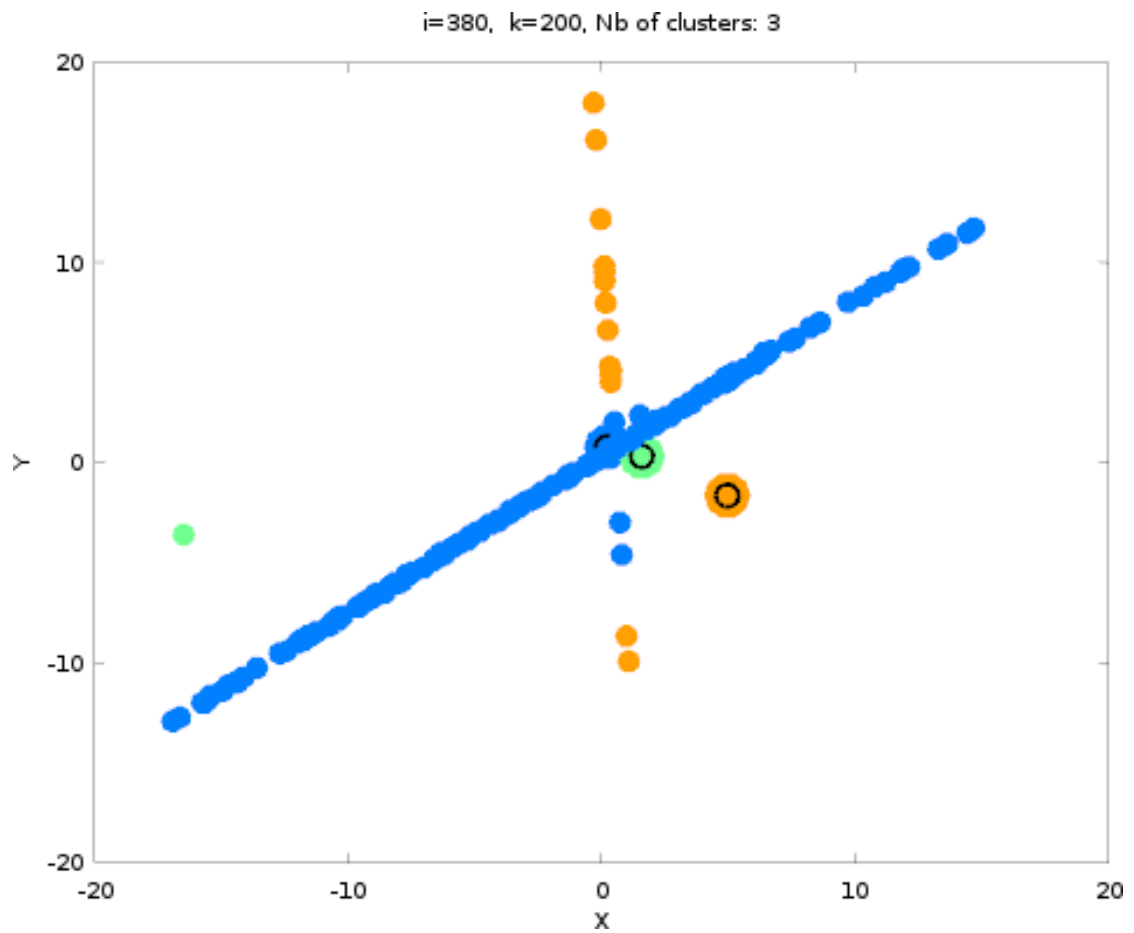
**Figure 3.4:** The performance of Algorithm 2 with respect to clustering is measured using the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert metric. A figure of 1 means perfect clustering for all metrics, except Mirvin's where 0 denotes perfect clustering.

occasionally two lines are generated with almost the same slope and intercept. Points on these lines are impossible to assign to the proper line.

The essential contribution in this chapter is the introduction of a fully Bayesian method to infer lines and there are two ways in which the postulated model can to be extended for full-fledged inference in computer vision as required in robotics. First, the extension of lines in 2D to planes in 3D. This is quite a trivial extension that does not change anything of the model except for the dimension of the data points. Second, somehow a prior needs to be incorporated to limit the lines of infinite length, to line segments. To restrict points on the lines to a uniform distribution of points over a line segment, a symmetric Pareto distribution can be used as prior (for the end points). This would subsequently allow for a hierarchical model in which these end points are in their turn part of more complicated objects. Hence, the Infinite Line Mixture Model is an essential step towards the use of Bayesian methods (and thus properly formulated priors) for robotic computer vision.



**Figure 3.5:** One of the Gibbs steps in the inference of two particular lines. The points are more or less distributed according to the lines, but one line exists out of two large clusters. The line coordinates are visualized by a double circle. The x-coordinate is the y-intercept of the line, the y-coordinate is the slope.



**Figure 3.6:** The assignment of a line to a single point. There are three clusters found, rather than only the obvious two.

1278

1279

1280

# NONPARAMETRIC BAYESIAN SEGMENT ESTIMATION

1281

## Contents

1282

1283

1284

1285

1286

1287

The nonparametric Bayesian model for line estimation (Chapter 3) does not take into account lines that are of finite length. In this chapter, we introduce a Bayesian method to perform inference over such line segments. In this model our prior for the extend of the line segment is a symmetric Pareto distribution. Due to the fact that the prior and likelihood is not a conjugate pair a more general inference method is used, namely Gibbs sampling with auxiliary variables.

1288

## Published in

1289

1290

1291

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. Non-parametric Segment Detection. *Proceedings of the Eighth European Starting AI Researcher Symposium, STAIRS 2016*, The Hague, Netherlands, August 26-September 2, 2016.

1292

## Outline

1293

1294

1295

1296

1297

1298

1299

The model is using a Normal-Inverse-Gamma and a Pareto prior for an individual line segment (Sect. 4.1). These line segments are generated using a Dirichlet Process (Sect. 4.2). This generative process is used to perform inference using Gibbs sampling over auxiliary variables (Sect. 4.3). The results for inference over line segments are compared with these for lines. (Sect. 4.4). Finally, weak aspects of the current MCMC method are established (Sect. 4.5) which will form the basis for new inference methods in the next chapters.

1300

## 4.1 Pareto Pairs

1301

1302

1303

Lines in a two-dimensional space are mathematical objects that can be described by two parameters. To limit a line to a line segment, four parameters are required. There are

two parametrizations that come to mind. First, a center-point parametrization, in which parameters describe the center of a line segment, the slope of the line through the center, and the size of the line segment. Second, an endpoint parametrization, in which parameters describe the locations of the two endpoints. These parametrizations are equivalent, but generalizations can be intuitive or cumbersome. The generalization to a line segment from a two-dimensional space to a three-dimensional space, requires the endpoints to be positions in a 3D space. The center-point parametrization would require a nonintuitive description of the angles in particular directions. The generalization to squares and rectangles or shapes with many endpoints, might benefit from the center-point parametrization.

There seems to be no statistical description of data points distributed over a line segment that has a conjugate prior form. A line segment itself, however, has a conjugate form! Suppose that we have a prior for the location of endpoints on the x-axis. Given the data we want to update the location of the endpoints. By leaving out the spread of the data over the segment, we can do this using a conjugate Bayesian construction.

The data is distributed according to a symmetric uniform distribution. Hence the likelihood is given by:

$$p(x | a) \sim \mathcal{U}(-a, a) = \begin{cases} \frac{1}{2a} & \text{for } x \leq |a| \\ 0 & \text{otherwise} \end{cases} \quad (4.1)$$

Here the uniform distribution is centered around 0 and extends with size  $a$  in both directions. It is possible to shift the entire distribution with  $b$ . For now, let's continue with one endpoint at  $a$  and one endpoint at  $-a$ .

A prior for the (endpoints of a) symmetric uniform distribution is a symmetric Pareto distribution:

$$p(a) \sim \mathcal{P}_s(\lambda, k) = \begin{cases} \frac{1}{2} k \lambda^k |a|^{-k-1} & |a| \geq \lambda \\ 0 & \text{otherwise} \end{cases} \quad (4.2)$$

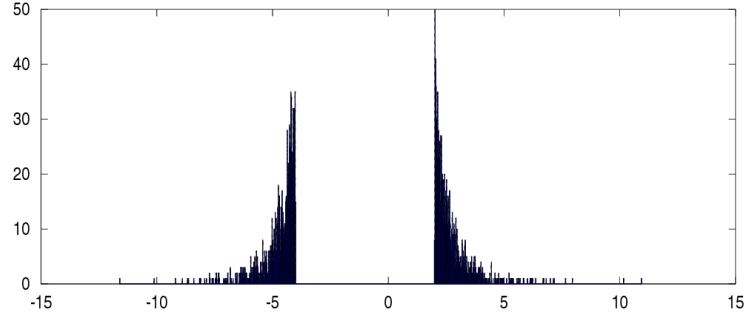
The factor  $\frac{1}{2}$  stems from the fact that the symmetric Pareto distribution is now mirrored across the y-axis. Hence, the probability density is half of that of the normal Pareto distribution for the positive x-axis.

If we would just sample from a symmetric Pareto distribution, we can sample multiple times from the positive x-axis. To actually sample endpoints of segments we have to sample pairs of points.

$$p(a, b) \sim \mathcal{P}_p(\lambda_m, \lambda_n, k) \quad (4.3)$$

1331 We can describe this process as first sampling  $a$  and  $b$  from a categorical distribution to  
 1332 decide which one will be the left endpoint and which one the right endpoint. Then we  
 1333 sample the right endpoint from a normal Pareto distribution and the left endpoint from a  
 1334 mirrored Pareto distribution.

1335 The sampling of Pareto pairs is visualized in Fig. 4.1.



**Figure 4.1:** Sampling of Pareto pairs. The parameters are  $\lambda_m = 2$ ,  $\lambda_n = -2$ ,  $k = 5$ , and we have sampled  $N = 1000$  pairs.

1336 The Pareto distribution is a conjugate prior for the uniform distribution, with updated hy-  
 1337 perparameters:

$$p(a \mid D) = \mathcal{P}(c, N + k) \quad (4.4)$$

1338 The data is denoted by  $D = \{x_0, \dots, x_{N-1}\}$ , the parameter  $k$  is adjusted with the number of  
 1339 data points  $N$ , and the parameter  $c$  is the maximum of  $\{m, \lambda\}$  with  $m$  the maximum value  
 1340 in  $D$ .

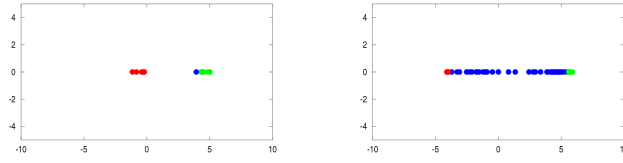
1341 Given this description for the posterior for a single point, the posterior for a Pareto pair  
 1342 can be found by sampling in parallel from a Pareto distribution  $\mathcal{P}(c_n, N + k_n)$  with  $c_n$  the  
 1343 maximum of the data points  $D$  and  $\lambda_n$ ,  $N$  the number of Pareto pairs, and  $k_n$  the hyperprior  
 1344 for the endpoint at the right. Plus sampling from a Pareto distribution  $\mathcal{P}(c_m, N + k_m)$  with  
 1345  $c_m$  the minimum of the data points  $D$  and  $\lambda_m$ ,  $N$  the same, and  $k_m$  the hyperprior for the  
 1346 endpoint at the left.

1347 If  $k_n \neq -k_m$  the distribution is shifted such that  $k'_n = -k'_m$ . This makes the form of the  
 1348 probability distribution symmetric with respect to the  $y$  axis. In the end the results are  
 1349 shifted back. This transformation makes sense for pairs of points. We do not want the  
 1350 two scale parameters of the Pareto distribution to influence the symmetry of the overall  
 1351 distribution.

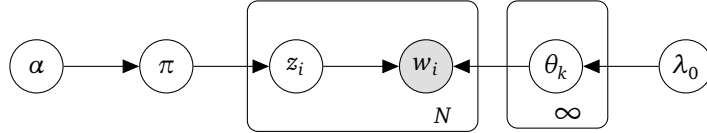
1352 Sampling from the Pareto distribution is through inverse transform sampling. By sampling  
 1353 from  $U(0, 1)$  with 1 included, we transform according to  $k/U^{1/a}$ .

1354 Fig. 4.2 shows how the endpoints are updated given the data. An uninformative prior is  
 1355 used. In this case the hyperparameters  $k_n$  and  $k_m$  are set close to 0, thus the data will wash  
 1356 out the prior immediately. Naturally, it is possible to set them quite large. In that case it





**Figure 4.2:** Consider the data uniformly distributed on a line segment and a symmetric Pareto prior for both the endpoints, then we can update the estimate for the endpoints given the data as visualized. Each subfigure shows an adjustment of the endpoints given more data points (1, 3, 10, and 100 data points).



**Figure 4.3:** The Bayesian linear regression model for multiple line segments in plate notation is the same as for the Infinite Line Model. The Dirichlet process is defined at the left with concentration parameter  $\alpha$ . It generates the partitions  $(\pi_1, \dots, \pi_k)$  with assignment parameters  $z_i$  that denote which observation  $w_i$  belongs to which cluster  $k$ . The cluster is summarized through the parameter set  $\theta_k$  and has  $\lambda_0$  as its hyperparameter. The parameter set  $\theta_k$  includes parameters that signify the line itself such as slope and y-intercept, plus the parameters that denote the extend of the segment.

1357 must be noted that the data will never be able to “correct for” this prior. Note also that the  
 1358 maximum and minimum operators are quite sensitive to outliers as well.

## 1359 4.2 Generative Process to Create a Line Segment

1360 To be able to perform inference over a line segment in a two-dimensional space, we’ll have  
 1361 to map somehow these points to a one-dimensional space.

1362 In the case of a line we can sample  $\theta_i$  from a Normal-Inverse-Gamma with hyperparameter  
 1363  $\lambda_{temp}$ . The latter we have in closed form given observations through a single update.

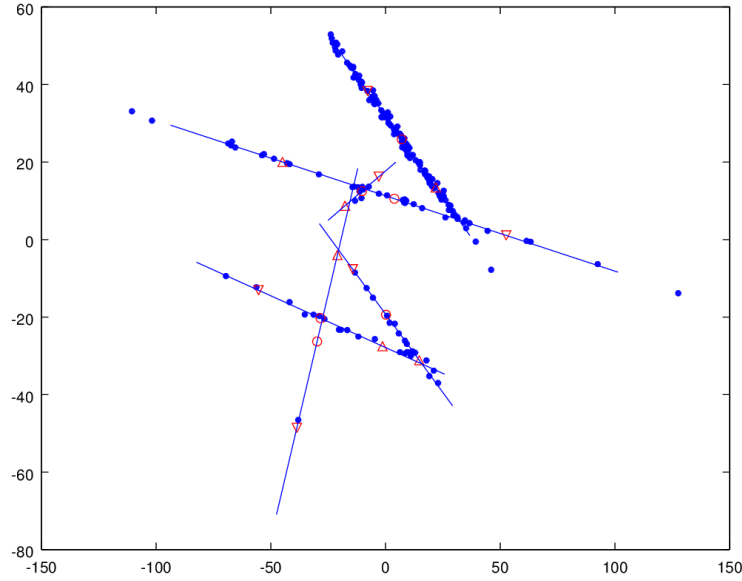
1364 In the case of a line segment there is no known conjugate prior available. Let’s reiterate the  
 1365 Dirichlet Process basis for our nonparametric model:

$$\begin{aligned} G &\sim DP(\alpha, H) \\ \theta_i &| G \sim G \\ w_i &| \theta_i \sim F(\theta_i) \end{aligned} \tag{4.5}$$

1366 Again  $F$  describes the mapping from parameters  $\theta_i$  to observations  $w_i$ . As described, for line  
 1367 segments this mapping is different from that of lines.

$$\begin{aligned} F(\theta_i) &= \mathcal{N}(\mu_i + H(v_i), \Sigma_i) \\ H(v_i) &= \mathcal{N}(v_i, 1)\gamma_i \\ \gamma_i &\sim \mathcal{N}(0, 1) \end{aligned} \tag{4.6}$$

1368 The probability density  $F$  is a Gaussian with a mean that is additively distributed according  
 1369 to another distribution  $H$ . The latter distribution originates from the product of a normal  
 1370 distribution with a value sampled from a normal distribution. Fig. 4.4 shows how points are  
 1371 generated from the described distribution.



**Figure 4.4:** Line segments generated through a Dirichlet Process. The Dirichlet Process itself is again the same. But now four parameters are generated. A normal-inverse-Wishart distribution is used to generate the center of the line segment, and an inverse-Wishart distribution to generate one of the endpoints of the line segments (the other end point is mirrored through its center). Points are generated normally over the line segments, with an additional Gaussian component to indicate the deviation from the line segment from the normal-inverse-Wishart.

1372 To generate lines uniformly, only  $\gamma_i$  needs adjustment:

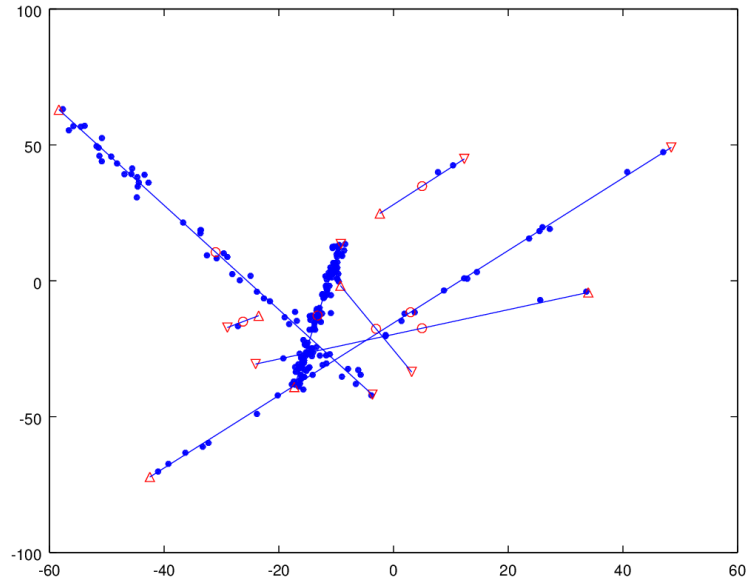
$$\begin{aligned}
 F(z_i) &= \mathcal{N}(\mu_i + H_i(v_i), \Sigma_i) \\
 H(v_i) &= \mathcal{N}(v_i, 1)\gamma_i \\
 \gamma_i &\sim \mathcal{U}(-1, 1)
 \end{aligned}
 \tag{4.7}$$

1373 Fig. 4.5 displays the adjustment with points generated uniformly over the line segment.

1374 The descriptions in Eq. 4.6 and 4.7 are clearly not conjugate setups. This means that infer-  
 1375 ence over line segments requires more complicated sampling strategies.

### 1376 4.3 Inference over a Line Segment

1377 To perform inference over a line segment our model is not conjugate anymore. This re-  
 1378 quires a sampling algorithm that does not make use of conjugacy. An algorithm that does



**Figure 4.5:** Line segments generated through a Dirichlet Process. Compared to Fig. 4.4 the points are generated uniformly over the line segments: points are not generated outside of the line segments.

1379 not assume conjugacy is described in its general form before (Neal, 2000) (Algorithm 8). The  
 1380 sampling process proposes  $m$  new values for the parameters directly from the hyperparam-  
 1381 eters. These are called auxiliary parameters. Now, to establish to which cluster a certain  
 1382 observation  $w_i$  need to be assigned, the likelihood of each existing and new clusters alike  
 1383 are compared. The weight of an old cluster is defined through the number of data points  
 1384 assigned to it. The weight of a new cluster is defined through  $\alpha/m$ .

1385 After every data item is assigned a cluster, the cluster parameters themselves are updated  
 1386 given the assigned data items. In a conjugate model the sufficient statistics can be updated  
 1387 at once, given such observations. In a nonconjugate model we will need to update  $\theta_j$  by  
 1388 sampling from  $p(\theta_j | y)$ .

**Algorithm 9** Gibbs sampling over auxiliary variables (a  $\theta_i$ )

---

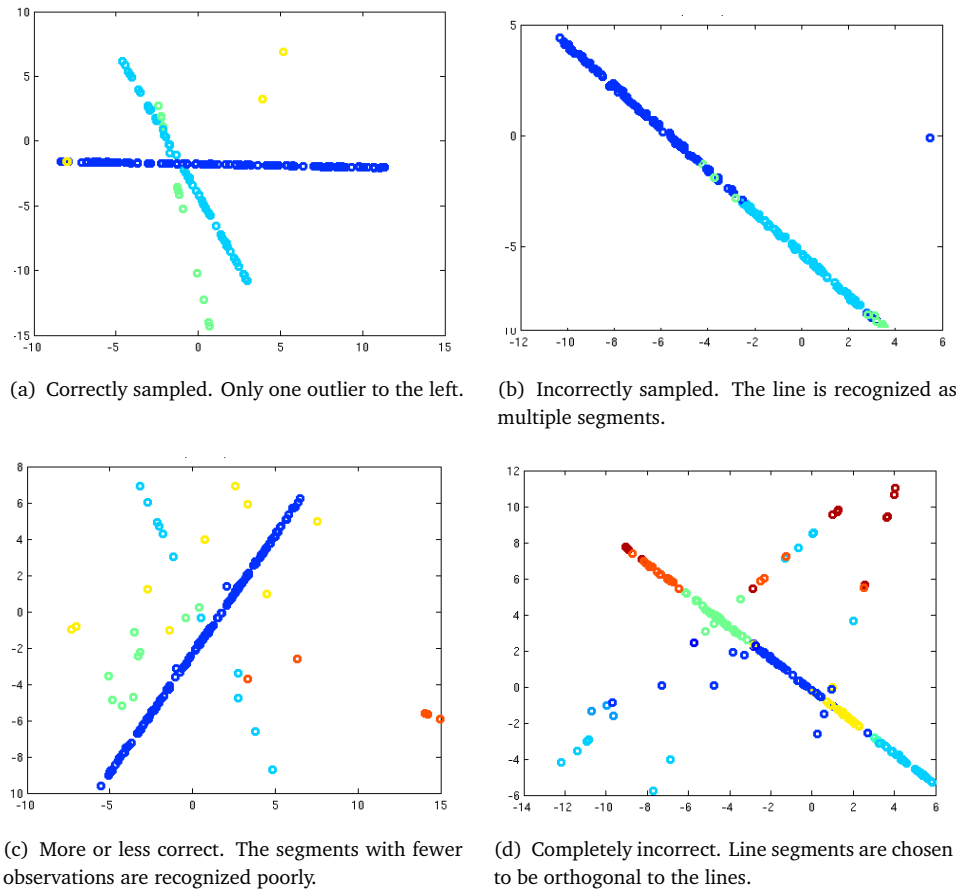
```

1: procedure GIBBS ALGORITHM WITH AUXILIARY VARIABLES( $w, \lambda_0, \alpha$ )  $\triangleright$  Accepts points  $w$ ,
   hyperparameters  $\lambda_0, \alpha$ , number of auxiliary variables  $m$ , and returns  $k$  line coordinates
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:       for all  $j = 1 : m$  do
5:          $\theta_j \sim NIG(\lambda_0)$   $\triangleright$  Sample  $\theta_j$  from NIG
6:       end for
7:       for all  $j = 1 : K + m, j \neq i$  do
8:          $L_j = \text{likelihood}(w_i, \theta_j)$   $\triangleright$  Update likelihood for all theta (except  $\theta_i$ ) given
           observation  $w_i$ 
9:       end for
10:       $P_{-i=1:K} = b \sum_{-i} L_{-i}$   $\triangleright$  Calculate probability of existing cluster
11:       $P_{-i=K:K+m} = b\alpha/mL_m L_{-i}$   $\triangleright$  Calculate probability of new cluster
12:       $\theta_i = \theta_j$  according to above  $P_{-i}$   $\triangleright$  Sample  $\theta_i$  accord. to above prob
13:      Remove unused clusters
14:    end for
15:    for all  $j = 1 : K$  do
16:       $\theta_j \sim p(\theta_j | y)$   $\triangleright$  Update  $\theta_j$ 
17:    end for
18:  end for
19:  return summary on  $\theta_k$  for  $k$  line segments
20: end procedure

```

---

## 4.4 Results



**Figure 4.6:** Bayesian point estimates of the sampling process with varying outcomes.

There is one phenomenon that is very noticeable in Fig. 4.6. Line segments that form a larger line segment are not recognized as such by the inference method.

The results over a larger dataset can be measured with clustering metrics as visualized in Fig. 4.7. The Rand Index, Adjusted Rand Index, and Hubert metrics show all reduced performance compared to line detection where there are no constraints on segment size.



**Figure 4.7:** Segment detection performs much worse than line detection across all three clustering performance indicators. Perfect clustering is indicated by 1.0 for Rand Index, Adjusted Rand Index, and Hubert.

## 4.5 Chapter Conclusions

Segment estimation is a much harder problem than line estimation (Chap. 3). In this chapter we used an advanced method, namely MCMC sampling with auxiliary variables to perform inference over an infinite set of line segments. The parameters for line segments do not have a conjugate description. Metropolis-Hastings has been used to perform inference over the line segments, but the search space is quite large. The auxiliary variable MCMC method is indeed faster than ordinary Metropolis-Hastings thanks to postulating multiple new lines than only one.

However, the segment estimation problem is a challenge for the current inference methods. The target probability density has a lot of modes that each needs to be found and are separated by very low probability regions. In Chapter 5 we will introduce new sampling methods that will cope with these challenges.



# TRIADIC SPLIT-MERGE SAMPLER

1409	<b>Contents</b>	The nonparametric Bayesian model for line estimation, the infinite line model (Chapter 3) thanks to its conjugate properties has been solved with moderately straightforward sampling methods. The additional constraints that limit lines to line segments (Chapter 4) reduced convergence of the underlying MCMC sampling method (a Gibbs method with auxiliary variables) to sub-par results.
1410		
1411		
1412		
1413		
1414		
1415		This chapter introduces a new sampling method called the triadic split-merge sampler.
1416		
1417	<b>Published in</b>	A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. Triadic Split-Merge Sampler. <i>The 10th International Conference on Machine Vision, ICMV 2017, Vienna, Austria, November 13-November 15, 2017.</i>
1418		
1419		
1420	<b>Outline</b>	The class of split-merge samplers, part of MCMC samplers, are introduced (Sect. 5.1). A conventional split-merge sampler, labeled the dyadic split-merge sampler is detailed (Sect. 5.2). The new split-merge sampler, the triadic split-merge sampler is introduced (Sect. 5.3). The results for inference over lines is compared between the conventional and the new sampler (Sect. 5.4). Finally, although this sampler already improves on the state-of-the-art we see in the chapter conclusions (Sect. 5.5) how we further improve the inference procedure, which will be the basis of the next chapter.
1421		
1422		
1423		
1424		
1425		
1426		
1427		
1428		

## 5.1 The Class of Split-Merge Samplers

We will consider a Dirichlet process as a prior on the distribution over parameters  $G$ . The form of this model is:



$$\begin{aligned}
y_i | \theta_i &\sim F(\theta_i) \\
\theta_i | G &\sim G \\
G &\sim DP(H, \alpha)
\end{aligned} \tag{5.1}$$

## 1432 5.2 Conventional Split-Merge Sampler

1433 The conventional split-merge sampler Jain and Neal (2004) splits a single cluster into two  
 1434 clusters, and merges two clusters into a single cluster. Hence, this split-merge sampler oper-  
 1435 ates on two clusters at each time step, for which reason we will call it a dyadic split-merge  
 1436 sampler in constrast with our approach.

---

### Algorithm 10 Dyadic split-merge sampler

---

```

1: procedure DYADIC SPLIT-MERGE SAMPLER( $c$ )                                ▷ Accepts cluster assignments
    $c$  of length  $N$  (besides Metropolis-Hastings acceptance factors  $a(c', c)$  and a split procedure e.g.
   SIMPLERANDOMSPLIT) and returns a (potentially) updated cluster assignment vector  $c'$ .
2:    $i \sim U(1, N)$                                                          ▷ Sample  $i$  random uniformly over cluster assignments.
3:    $j \sim U(1, N) \cap i$                                                  ▷ Sample  $j$  also random uniformly, but with  $j \neq i$ .
4:    $S_R = \{c_i, c_j\}$                                                      ▷ Sampled clusters  $c_i, c_j$ .
5:    $S_I = \{c_x\}$  with  $c_x \in S_R$  for  $x \in \{1, \dots, N\}$                 ▷ All data in clusters  $c_i, c_j$ .
6:    $S_E = S \cap S_R$                                                      ▷ All data in clusters  $c_i, c_j$  excluding  $S_R$ .
7:    $N_S = \text{unique}(S_R)$ 
8:   if  $N_S = 1$  then                                                     ▷ Case:  $i, j$  belong to the same cluster.
9:      $c_i^{(2)} = c_k$  with  $c_k \notin \{c_1, \dots, c_N\}$                     ▷ Sample new cluster for  $c_i^{(2)}$ .
10:     $c_j^{(2)} = c_j^{(1)}$                                                   ▷ Keep  $c_j$  the same.
11:     $c_e^{(2)} = \text{SPLITPROCEDURE}(S_E, c_i^{(2)}, c_j^{(2)})$                 ▷ After  $c_i^{(2)}, c_j^{(2)}$  assign  $S_E$ .
12:    for all  $m \notin S_I$  do
13:       $c_m^{(2)} = c_m^{(1)}$                                               ▷ Data points in clusters other than  $c_i, c_j$  are not adjusted.
14:    end for
15:     $c' = \{c_i^{(2)}, c_j^{(2)}, c_e^{(2)}, c_m^{(2)}\}$ 
16:     $a = a_{\text{split}}(c', c)$  according to Eq. 5.3                        ▷ MH acceptance for a split.
17:  else                                                                 ▷ Case:  $i, j$  belong to different clusters  $c_i \neq c_j$  ( $N_S = 2$ ).
18:    for all  $q \in S_I$  do
19:       $c_q^{(1)} = c_j^{(2)}$                                               ▷ Assign all data points in  $c_i$  and  $c_j$  to  $c_j$ .
20:    end for
21:    for all  $m \notin S_I$  do
22:       $c_m^{(1)} = c_m^{(2)}$                                               ▷ Data points in clusters other than  $c_i, c_j$  are not adjusted.
23:    end for
24:     $c' = \{c_q^{(1)}, c_m^{(1)}\}$ 
25:     $a = a_{\text{merge}}(c', c)$  according to Eq. 5.10                      ▷ MH acceptance for a merge.
26:  end if
27:   $u \sim U(0, 1)$                                                          ▷ Sample  $u$  between 0 or 1 uniformly.
28:  if  $a < u$  then
29:     $c' = c$                                                              ▷ Reject  $c'$  by setting it to  $c$ 
30:  end if
31:  return  $c'$ , the (updated) cluster assignment vector:  $c \rightarrow c'$ .
32: end procedure

```

---

1437 In algorithm 10 the notation  $c_i^{(2)}$  is used to signify that the cluster assignment  $c_i$  has 2 clusters  
 1438 under consideration. In the dyadic algorithm we could have used  $c_i^{merge}$  and  $c_i^{split}$ , however  
 1439 in the triadic algorithm (see algorithm 13) with multiple split and merge operations the  
 1440 latter notation would become confusing.

---

**Algorithm 11** Simple random split
 

---

```

1: procedure SIMPLERANDOMSPLIT( $S, c_0, c_1$ )  $\triangleright$  Accepts unassigned set  $S$  and cluster indices  $c_0, c_1$ ,
   returns cluster assignment  $c'_m$ .
2:   for all  $m \in S$  do
3:      $c'_m \sim Cat(c_0, c_1)$  with equiprobable  $p(c_0) = p(c_1) = \frac{1}{2}$ .
4:   end for
5:   return  $c'_m$ , the cluster assignment for  $S$ .
6: end procedure
  
```

---

1441 The dyadic split-merge sampler in Algorithm 10 samples two distinct data items. If the data  
 1442 items belong to the same cluster a split step is attempted. If the data items belong to different  
 1443 clusters a merge step is attempted. The split procedure itself is the so-called simple random  
 1444 split (Algorithm 11) that assigns data items with the same probability to one of the parts of  
 1445 the splitted cluster without consideration for data fit.

### 1446 5.2.1 Acceptance for the Split Step

1447 The acceptance ratio contains the Metropolis ratio to step from  $c$  to  $c'$ :

$$\frac{P(c')L(c'|y)}{P(c)L(c|y)} \quad (5.2)$$

1448 Additionally, the Hastings correction is applied because of the asymmetry of the proposal  
 1449 distribution in the form of  $q(c|c')/q(c'|c)$ :

$$a_{split}(c^{(2)}, c^{(1)}) = \min \left[ 1, \frac{q(c^{(1)}|c^{(2)})}{q(c^{(2)}|c^{(1)})} \frac{P(c^{(2)})}{P(c^{(1)})} \frac{L(c^{(2)}|y)}{L(c^{(1)}|y)} \right] \quad (5.3)$$

1450 The notation  $c^{(2)}$  is used to indicate that the cluster index vector is referencing 2 unique  
 1451 clusters (in this case after the split step).

1452 The prior distribution is represented by a Chinese Restaurant Process with concentration  
 1453 parameter  $\alpha$  and no discount factor. Data not yet assigned is assigned with probability  
 1454  $\alpha/(n + \alpha)$  to a new cluster and with probability  $n_c/(n + \alpha)$  to an existing cluster  $c$ . Here  $n$   
 1455 are the total number of assigned data points,  $n_c$  are the number of data points assigned to  
 1456 cluster  $c$ . There are  $D$  clusters. Hence, the prior over clusters:

$$P(c) = \frac{\Gamma(\alpha)}{\Gamma(\alpha + n)} \alpha^D \prod_{c_l} \Gamma(n_{c_l}) = \alpha^D \frac{\prod_{c_l} (n_{c_l} - 1)!}{\prod_{k=1}^n (\alpha + k - 1)} \quad (5.4)$$

1457 In the prior distribution ratio before and after the split step many of the factors drop out.  
 1458 There is one factor  $\alpha$  remaining and the number of data points in the splitted cluster is part

of the equation. There is no dependency on other clusters or the total number of data points and we can simplify the formula using the beta function  $B(a, b)$ :

$$\frac{P(c^{(2)})}{P(c^{(1)})} = \alpha \frac{(n_{c_i^{(2)}} - 1)!(n_{c_j^{(2)}} - 1)!}{(n_{c_i^{(1)}} - 1)!} = \alpha B(n_{c_i^{(2)}}, n_{c_j^{(2)}}) \quad (5.5)$$

The likelihood can be written as:

$$L(c|y) = \prod_{c=1}^D \prod_{k:c_k=c} p(y_k|\phi) \quad (5.6)$$

Here we assume no conjugacy between  $F(y_k, \phi)$  and prior distribution  $H(\phi)$  and hence write  $p(y_k|\phi)$  rather than the conjugate construction  $\int F(y_k, \phi) dH_{k,c}(\phi)$  (see Dahl (2005)). The likelihood ratio becomes:

$$\frac{L(c^{(2)}|y)}{L(c^{(1)}|y)} = \frac{\prod_{k:c_k^{(2)}=c_i^{(2)}} p(y_k|\phi) \prod_{k:c_k^{(2)}=c_j^{(2)}} p(y_k|\phi)}{\prod_{k:c_k^{(1)}=c_i^{(1)}} p(y_k|\phi)} \quad (5.7)$$

The split step determines the probability of a particular split. Given that already two data points are assigned to distinct clusters, only the remaining ones have to be assigned with equal probability to  $c_i^{(2)}$  and  $c_j^{(2)}$ :

$$q(c^{(2)}|c^{(1)}) = \left(\frac{1}{2}\right)^{-2+n_{c_i^{(2)}}+n_{c_j^{(2)}}} = \left(\frac{1}{2}\right)^{-2+n_{c_i^{(1)}}} \quad (5.8)$$

The probability of the reverse of the split operation is exactly 1. There is only one way in which a single cluster could have risen from a split cluster, hence:

$$\frac{q(c^{(1)}|c^{(2)})}{q(c^{(2)}|c^{(1)})} = \frac{1}{\left(\frac{1}{2}\right)^{n_{c_i^{(2)}}+n_{c_j^{(2)}}-2}} = 2^{-2+n_{c_i^{(1)}}} \quad (5.9)$$

### 5.2.2 Acceptance for the Merge Step

Acceptance of a merge step consists of the same components as that of the split step.

$$a_{merge}(c^{(1)}, c^{(2)}) = \min \left[ 1, \frac{q(c^{(2)}|c^{(1)})}{q(c^{(1)}|c^{(2)})} \frac{P(c^{(1)})}{P(c^{(2)})} \frac{L(c^{(1)}|y)}{L(c^{(2)}|y)} \right] \quad (5.10)$$

$$\frac{P(c^{(1)})}{P(c^{(2)})} = \alpha^{-1} \frac{(n_{c_i^{(1)}} - 1)!}{(n_{c_i^{(2)}} - 1)!(n_{c_j^{(2)}} - 1)!} = \frac{1}{\alpha B(n_{c_i^{(2)}}, n_{c_j^{(2)}})} \quad (5.11)$$

$$\frac{L(c^{(1)}|y)}{L(c^{(2)}|y)} = \frac{\prod_{k:c_k^{(1)}=c_i^{(1)}} p(y_k|\phi)}{\prod_{k:c_k^{(2)}=c_i^{(2)}} p(y_k|\phi) \prod_{k:c_k^{(2)}=c_j^{(2)}} p(y_k|\phi)} \quad (5.12)$$

1474

$$\frac{q(c^{(2)}|c^{(1)})}{q(c^{(1)}|c^{(2)})} = \left(\frac{1}{2}\right)^{-2+n_{c_i^{(1)}}} = 2^{2-n_{c_i^{(1)}}} \quad (5.13)$$

1475 The ratios of the merge step are the inverse of the ratios of the split step.

1476 **5.2.3 Sequentially-Allocated Merge-Split Sampler**

1477 A variant on the conventional split-merge sampler is the Sequentially Allocated Merge-Split<sup>1</sup>  
 1478 (SAMS) sampler Dahl (2003). The simple random split procedure of Algorithm 11 is re-  
 1479 placed by a procedure that sequentially assigns observations to clusters rather than splitting  
 1480 the data random uniformly over the splitted clusters.

---

**Algorithm 12** Sequentially Allocated Merge-Split
 

---

```

1: procedure SAMS( $S, c_0, c_1$ )    ▷ Accepts unassigned set  $S$ , cluster indices  $c_i$ , and  $p(y_k|\theta_{c_i})$  with
    $i = 0, 1$ , returns cluster assignment  $c'_m$ .
2:    $T = \text{random\_shuffle}(S)$ 
3:   for all  $m \in T$  do
4:      $p(c_m = c_0 | c_0, c_1, \theta_{c_0}, \theta_{c_1}) = \frac{N_0 p(y_k | \theta_0)}{N_0 p(y_k | \theta_0) + N_1 p(y_k | \theta_1)}$ 
5:      $p(c_m = c_1 | c_0, c_1, \theta_{c_0}, \theta_{c_1}) = 1 - p(c_m = c_0 | c_0, c_1, \theta_{c_0}, \theta_{c_1})$ 
6:      $c'_m \sim p(c_m | c_0, c_1, \theta_{c_0}, \theta_{c_1})$ 
7:   end for
8:   return  $c'_m$ , the cluster assignment for  $S$ .
9: end procedure

```

---

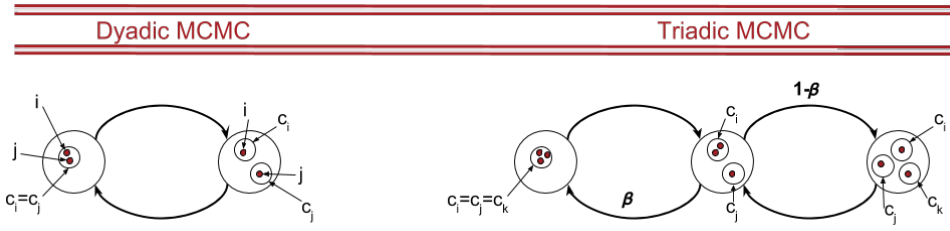
1481 In contrast to the simple random split, observations  $y_k$  are used in the SAMS to obtain cluster  
 1482 assignments that correspond with the data rather than cluster assignments independent of  
 1483 the data.

1484 **5.3 Triadic split-merge sampler**

1485 The triadic split-merge sampler uses up to three clusters for a split or merge step (Fig. 5.1).

---

<sup>1</sup>In the naming of split-merge or merge-split samplers, the order of merge split does not bear any significance.



**Figure 5.1:** Right: dyadic MCMC picks two data items  $i, j$  random uniformly. If both are in the same cluster a split towards two clusters is attempted. If both are in distinct clusters a merge towards one cluster is attempted. Left: triadic MCMC picks three data items  $i, j, k$  random uniformly. If all three are in the same cluster a split towards two clusters is attempted. If the three items are in two clusters either a split into three (with probability  $1 - \beta$ ) or a merge into a single cluster (with probability  $\beta$ ) is attempted. If the three data items are in three distinct clusters a merge is attempted. There are no direct transitions from a single cluster to three clusters or the other way around.

1486 The intuition behind the triadic split-merge sampler is twofold:

- 1487 ○ In the dyadic sampler there is a large asymmetry between split and merge steps. There  
 1488 is only one way in which two clusters can be merged into one single cluster, while there  
 1489 are many ways in which one single cluster can be split into two clusters. This asym-  
 1490 metry is reduced by transitioning between two and three clusters. This is a straight-  
 1491 forward improvement in balancing split and merge steps (for alternatives, see Wang  
 1492 and Russell (2015)).
- 1493 ○ In practical optimization problems it might be useful to form a third cluster out of  
 1494 subsets of two other clusters. The dyadic MCMC sampler requires immediate steps in  
 1495 which (1) one of these clusters is split into two, (2) the other is split into two, and  
 1496 (3) the two new clusters are merged. This means that (a) mixing and hence conver-  
 1497 gence will be slow and (b) the intermediate steps might have very low probability and  
 1498 function as an unnecessary barrier between high probable states.

1499 Sampling random uniformly for three unique items is implemented through a random shuf-  
 1500 fle algorithm, in particular the modern version of the Fisher-Yates shuffle introduced by  
 1501 Durstenfeld Durstenfeld (1964) and picking the first three items.

### 1502 5.3.1 Acceptance for the Split Step

1503 In the triadic split-merge sampler there are two splitting steps. It is possible to split according  
 1504 to the dyadic split-merge sampler. However, given two clusters there are (split) jumps to  
 1505 three states as well as (merge) jumps to single states again. To account for this asymmetry  
 1506 another Hastings correction is applied to establish detailed balance.

$$a_{split}(c^{(2)}, c^{(1)}) = \min \left[ 1, \frac{r(c^{(1)}|c^{(2)}) q(c^{(1)}|c^{(2)}) P(c^{(2)}) L(c^{(2)}|y)}{r(c^{(2)}|c^{(1)}) q(c^{(2)}|c^{(1)}) P(c^{(1)}) L(c^{(1)}|y)} \right] \quad (5.14)$$

**Algorithm 13** Triadic split-merge sampler

---

```

1: procedure TRIADIC SPLIT-MERGE SAMPLER( $c$ )                                ▷ Accepts
   cluster assignments  $c$  of length  $N$  (besides Metropolis-Hastings acceptance factors  $a(c', c)$  and a
   split procedure) and returns a (potentially) updated cluster assignment vector  $c'$ .
2:    $i \sim U(1, N)$                                 ▷ Sample  $i$  random uniformly over cluster assignments.
3:    $j \sim U(1, N) \cap i$                             ▷ Sample  $j$  also random uniformly, but with  $j \neq i$ .
4:    $k \sim U(1, N) \cap \{i, j\}$                         ▷ Sample  $k$  random uniformly, but with  $k \neq j, k \neq i$ .
5:    $S_R = \{c_i, c_j, c_k\}$                                 ▷ Sampled clusters  $c_i, c_j, c_k$ .
6:    $S_I = \{c_x\}$  with  $c_x \in S_R$  for  $x \in \{1, \dots, N\}$         ▷ All data in clusters  $c_i, c_j, c_k$ .
7:    $S_E = S_I \cap S_R$                                 ▷ All data in clusters  $c_i, c_j, c_k$  excluding  $S_R$ .
8:    $N_S = \text{unique}(S_R)$ 
9:    $u \sim U(0, 1)$                                 ▷ Sample  $u$  between 0 or 1 uniformly.
10:  if  $N_S = 1$  then                                ▷ Case:  $i, j, k$  belong to the same cluster.
11:    return  $c' = \text{DYADIC SPLIT-MERGE SAMPLER}(c)$ 
12:  else if  $N_S = 2$  and  $u < \beta$  then                ▷ Case: a cluster with one item and one with two items and
    $u < \beta$ .
13:    return  $c' = \text{DYADIC SPLIT-MERGE SAMPLER}(c)$ 
14:  else if  $N_S = 2$  and  $u \geq \beta$  then                ▷ Case: a cluster with one item and one with two items and
    $u \geq \beta$ .
15:     $c_i^{(3)} = c_k$  with  $c_k \notin \{c_1, \dots, c_N\}$         ▷ Sample new cluster for  $c_i^{(3)}$ .
16:     $c_j^{(3)} = c_j^{(2)}$                                 ▷ Keep  $c_j$  the same.
17:     $c_e^{(3)} = \text{SPLITPROCEDURE}(S_E, c_i^{(3)}, c_j^{(3)})$         ▷ After  $c_i^{(3)}, c_j^{(3)}$  assign  $S_E$ .
18:    for all  $m \notin S_I$  do
19:       $c_m^{(3)} = c_m^{(2)}$                                 ▷ Data points in clusters other than  $c_i, c_j$  are not adjusted.
20:    end for
21:     $c' = \{c_i^{(3)}, c_j^{(3)}, c_e^{(3)}, c_m^{(3)}\}$ 
22:     $a = a_{\text{split}}(c', c)$  according to Eq. 5.14        ▷ MH acceptance for a split.
23:  else                                ▷ Case:  $i, j, k$  belong to three different clusters  $c_i \neq c_j \neq c_k$  ( $N_S = 3$ ).
24:     $S_L = S_I \cap \{c_i^{(3)}, c_j^{(3)}\}$                 ▷ Data in clusters  $c_i, c_j, c_k$  except for  $i$  and  $j$  itself.
25:     $\{c_i^{(2)}, c_j^{(2)}\} = \text{SAMS}(S_L, c_i^{(3)}, c_j^{(3)})$         ▷ Assign data points in  $c_i, c_j, c_k$  to  $c_i, c_j$ .
26:    for all  $m \notin S_L$  do
27:       $c_m^{(2)} = c_m^{(3)}$                                 ▷ Data points in clusters other than  $S_L$  are not adjusted.
28:    end for
29:     $c' = \{c_i^{(2)}, c_j^{(2)}, c_m^{(2)}\}$ 
30:     $a = a_{\text{merge}}(c', c)$  according to Eq. 5.21        ▷ MH acceptance for a merge.
31:  end if
32:   $u \sim U(0, 1)$                                 ▷ Sample  $u$  between 0 or 1 uniformly.
33:  if  $a < u$  then
34:     $c' = c$                                 ▷ Reject  $c'$  by setting it to  $c$ 
35:  end if
36:  return  $c'$ , the (updated) cluster assignment vector:  $c \rightarrow c'$ .
37: end procedure

```

---

1507 Here we have one additional term compared to the split step from one cluster to two clusters:

$$\frac{r(c^{(1)}|c^{(2)})}{r(c^{(2)}|c^{(1)})} = \frac{\beta}{1} \quad (5.15)$$

1508 The parameter  $\beta$  is free to control, as long as  $0 < \beta < 1$  (to maintain ergodicity). The  
1509 transition from two states to three states is another split step:

$$a_{split}(c^{(3)}, c^{(2)}) = \min \left[ 1, \frac{r(c^{(2)}|c^{(3)})}{r(c^{(3)}|c^{(2)})} \frac{q(c^{(2)}|c^{(3)})}{q(c^{(3)}|c^{(2)})} \frac{P(c^{(3)})}{P(c^{(2)})} \frac{L(c^{(3)}|y)}{L(c^{(2)}|y)} \right] \quad (5.16)$$

1510 The fraction with  $r$ :

$$\frac{r(c^{(2)}|c^{(3)})}{r(c^{(3)}|c^{(2)})} = \frac{1}{1 - \beta} \quad (5.17)$$

1511 The fraction with  $q$  uses the total number of data points  $n_c$  in the clusters:

$$\frac{q(c^{(2)}|c^{(3)})}{q(c^{(3)}|c^{(2)})} = \frac{\left(\frac{1}{2}\right)^{n_c-2}}{\left(\frac{1}{3}\right)^{n_c-3}} = (3^{n_c-3})(2^{2-n_c}) = \left(\frac{3}{2}\right)^{n_c} \frac{2^2}{3^3} \quad (5.18)$$

1512 To move from 2 clusters to 3 clusters the probability is a  $1/3$  for each cluster index in vector  
1513  $c$  (except for the three data items already selected randomly, hence  $n_c - 3$ ). To move back,  
1514 the probability is a  $1/2$  and there are only two data items randomly assigned beforehand.  
1515 The fraction with  $P$  uses the number of data points in each of the clusters before and after  
1516 the step:

$$\frac{P(c^{(3)})}{P(c^{(2)})} = \alpha \frac{(n_{c_i^{(3)}} - 1)!(n_{c_j^{(3)}} - 1)!(n_{c_k^{(3)}} - 1)!}{(n_{c_i^{(2)}} - 1)!(n_{c_j^{(2)}} - 1)!} = \alpha \frac{B(n_{c_i^{(3)}}, n_{c_j^{(3)}}, n_{c_k^{(3)}})}{B(n_{c_i^{(2)}}, n_{c_j^{(2)}})} \quad (5.19)$$

1517 Here we introduced a generalized Beta function  $B(a, b, c) = \Gamma(a)\Gamma(b)\Gamma(c)/\Gamma(a + b + c)$  with  
1518  $\Gamma(x) = (x - 1)!$  the Gamma function. The likelihood ratio becomes:

$$\frac{L(c^{(3)}|y)}{L(c^{(2)}|y)} = \frac{\prod_{m:c_m^{(3)}=c_i^{(3)}} P(y_m|\phi) \prod_{m:c_m^{(3)}=c_j^{(3)}} P(y_m|\phi) \prod_{m:c_m^{(3)}=c_k^{(3)}} P(y_m|\phi)}{\prod_{m:c_m^{(2)}=c_i^{(2)}} P(y_m|\phi) \prod_{m:c_m^{(2)}=c_j^{(2)}} P(y_m|\phi)} \quad (5.20)$$

### 1519 5.3.2 Acceptance for the Merge Step

1520 The merge step from two to one cluster is analogous to the split step:

$$a_{merge}(c^{(1)}, c^{(2)}) = \min \left[ 1, \frac{r(c^{(2)}|c^{(1)})}{r(c^{(1)}|c^{(2)})} \frac{q(c^{(2)}|c^{(1)})}{q(c^{(1)}|c^{(2)})} \frac{P(c^{(1)})}{P(c^{(2)})} \frac{L(c^{(1)}|y)}{L(c^{(2)}|y)} \right] \quad (5.21)$$

1521 The merge step from three clusters to two clusters is:

$$a_{merge}(c^{(2)}, c^{(3)}) = \min \left[ 1, \frac{r(c^{(3)}|c^{(2)})}{r(c^{(2)}|c^{(3)})} \frac{q(c^{(3)}|c^{(2)})}{q(c^{(2)}|c^{(3)})} \frac{P(c^{(2)})}{P(c^{(3)})} \frac{L(c^{(2)}|y)}{L(c^{(3)}|y)} \right] \quad (5.22)$$

1522 Note that all the fractions in Eq. 5.22 are the reverse of the fractions in Eq. 5.16. Inverting  
1523 Eq. 5.17–5.20 will be left to the reader.

One additional issue we have to consider. When merging three clusters into two we can (1) distribute the data over all three clusters or (2) alternatively, keep the data in two clusters assigned to these clusters and only distribute the data in the third cluster over the other two clusters. The second and alternative option however would introduce unnecessary asymmetry with the merge step. In other words, Eq. 5.23 is not the inverse of Eq. 5.18. In contrast, the equation is similar to splitting one cluster across two as in Eq. 5.9:

$$\frac{q_{alt}(c^{(3)}|c^{(2)})}{q_{alt}(c^{(2)}|c^{(3)})} = 2^{-2+n_c} \quad (5.23)$$

Hence the first option is entertained and the  $q$ -fraction is exactly the inverse of Eq. 5.18.

Another choice has been made, namely to exclude direct operations between a single cluster and three clusters. This is because factors like:

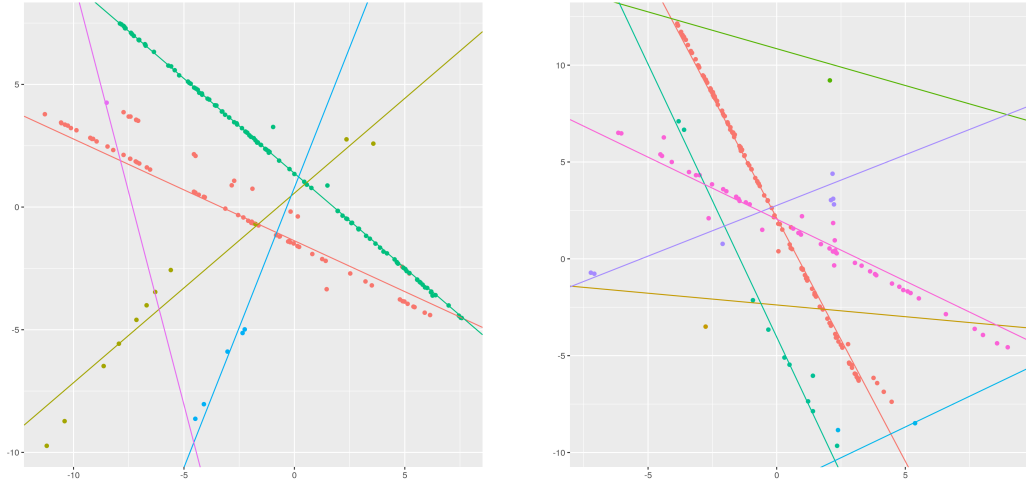
$$\frac{P(c^{(3)})}{P(c^{(1)})} = \alpha^2 \frac{(n_{c_i^{(3)}} - 1)!(n_{c_j^{(3)}} - 1)!(n_{c_k^{(3)}} - 1)!}{(n_{c_i^{(1)}} - 1)!} \quad (5.24)$$

become very small and although compensated by a large  $q$  fraction, remain further away from an acceptance factor of 1. Note that by the ability to split a single cluster into two and then into three, there is no ergodic argument to introduce also the immediate step.

## 5.4 Results

The problem we use to test our sampler is a well-known problem in computer vision, namely that of the inference of line parameters (slope and intercept) given data points. Rather than ordinary linear regression, in computer vision there is a mixture of lines that have to be estimated. Moreover, the number of lines is not known beforehand. To solve this problem we use the Dirichlet process mixture (Eq. 5.1) with a normal distribution  $N(0, \sigma_0)$  to generate the line parameters and a likelihood function that defines points to be uniformly distributed across a line of length 20 and deviating from the line according to a normal distribution  $N(0, \sigma_1)$ .





**Figure 5.2:** Two examples of fitting a mixture of lines to data items scattered over a two-dimensional space. The lines drawn are inferred using one of the methods in this paper. The lines are not the ground truth, but are meant to demonstrate the typical errors made by fitting methods. Note for example that there are mistakes in both the assignment of points to lines as well as the line parameters (slope and intercept).

#### 1545 5.4.1 Implementation

1546 The sampler is open-source<sup>2</sup> implemented in C++ which means that (a) it is computationally  
 1547 fast, (b) it can be run on embedded devices if a cross-compiler is available and the Eigen3  
 1548 library is ported. Note, that due to the fact that the simulator uses a lot of random numbers  
 1549 the system should use a modern compiler (g++-6 or newer) and should have enough entropy  
 1550 available<sup>3</sup>. Rather than a random scan, the implementation uses a fixed scan as advocated  
 1551 in the literature MacEachern (2007).

1552 To speed up the sampler most calculations are done in log-space. Consider  $v = u + 1$ . The  
 1553 ratio with probabilities (Eq. 5.5 and 5.19) becomes:

$$\log \frac{P(c^{(v)})}{P(c^{(u)})} = \log(\alpha) + \sum_i \log \Gamma(n_{c_i^{(v)}}) - \sum_i \log \Gamma(n_{c_i^{(u)}}) \quad (5.25)$$

1554 The fraction with  $q(\cdot)$  (Eq. 5.9 and 5.18) becomes:

$$\log \frac{q(c^{(v-1)}|c^{(v)})}{q(c^{(v)}|c^{(v-1)})} = (v - n_c - 1) \log(v - 1) - (v - n_c) \log(v) \quad (5.26)$$

1555 The fraction with  $r$  becomes for example (Eq. 5.17):

$$\log \frac{r(c^{(2)}|c^{(3)})}{r(c^{(3)}|c^{(2)})} = -\log(1 - \beta) \quad (5.27)$$

<sup>2</sup>Code can be found at <https://code.annevanrossum.nl/noparama>.

<sup>3</sup>On Linux this can be checked in `/proc/sys/kernel/random/entropy_avail`.

1556 The log-probability to calculate the likelihood given by a multivariate Normal distribution is  
 1557 well-known.

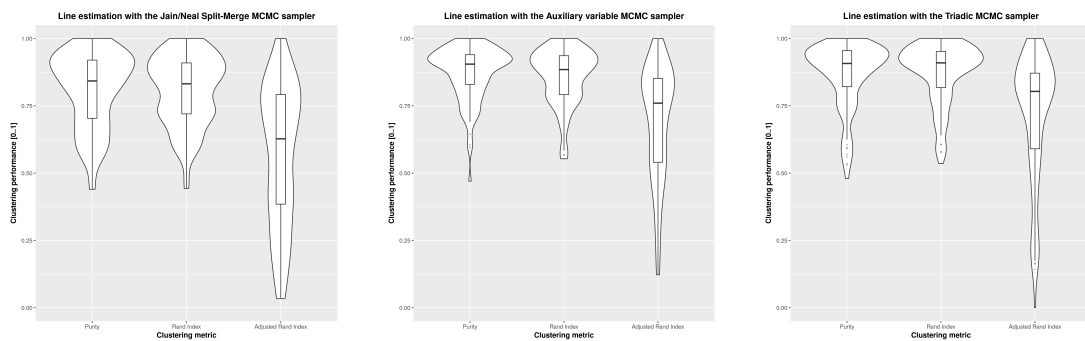
## 1558 5.4.2 Comparison

1559 The Triadic sampler using SAMS is compared with the Jain-Neal Dyadic sampler using SAMS  
 1560 and an auxiliary variable sampler with  $m = 3$  (see algorithm 8 in Neal (2000)).

Method	Purity	Rand Index	Adjusted Rand Index
Dyadic sampler	0.80960	0.80580	0.56382
Auxiliary variables	0.87235	0.85879	0.68224
Triadic sampler	0.86405	0.87188	0.71067

**Table 5.1:** The purity, rand index, and adjusted rand index establishing the quality of the clustering method. The closer the values to one, the better the method performed. The purity metric assigns high values to clusters that do not have data points from other clusters (but does not penalize the number of clusters). The rand index computes similarity between clusters taking false negatives and false positives into account. The adjusted rand index accounts for chance. The adjusted rand index is most useful in our comparison.

1561 In Table 5.1 the line estimation problem is compared for the dyadic sampler, an auxiliary  
 1562 variables sampler, and the proposed triadic sampler. The simulation is run with  $\beta = 0.1$  so  
 1563 that a significant number of steps are tried between two and three clusters (rather than only  
 1564 between one and two clusters).



**Figure 5.3:** The same results as in Table 5.1, but visualized in a violin plot. The distribution over metric values are displayed in a vertical fashion. From left to right the distribution shifts to one, signifying better clustering performance.

1565 In Fig. 5.3 the different metrics are visualized in the form of violin plots. The improvement  
 1566 in clustering is especially visible with the adjusted rand index.

## 1567 **5.5 Chapter Conclusions**

1568 A new split-merge sampler has been introduced, implemented, and applied to the computer  
1569 vision problem of line estimation. The sampler outperforms existing samplers, such as the  
1570 ordinary (dyadic) split-merge sampler Jain and Neal (2004) and auxiliary variable sampler  
1571 Neal (2000).

1572 Although the proposed split-merge sampler is able to mix considerably faster through a mix-  
1573 ture model, it does not use global jumps directly based on the data at hand. It is reasonable  
1574 to suggest that MCMC methods benefit from combining the local jumps with global jumps,  
1575 for example by a mixture of the local Metropolis-Hastings sampler with a Metropolized in-  
1576 dependence sampler Jampani et al. (2015). We will introduce such a sampler in chapter  
1577 6.



# ADVERSARIALLY TRAINED MCMC KERNELS

**Contents** To use MCMC for volumetric inference it is necessary to be able to accelerate the algorithms even further. Volumetric objects exhibit more structure, which is reflected by symmetry.

**Outline** We describe MCMC methods that cope with symmetric objects.

## 6.1 Data-Driven Inference

There are three aspects we would like to address in our inference engine.

The first aspect aims to have structure within our inference engine. The proposal distribution in a Markov chain, although moderately complex in the previous chapter, does not have much knowledge about the model at hand. An artificial border is maintained that does not allow the inference engine to have knowledge about the model. The purpose of this is never articulated in particular. However, it is logical from a separation of concern. Such an inference engine (1) does not need to receive any information about the model and (2) is guaranteed to be general in the sense that it is not tailored to a particular model. This is nicely articulated by Tran et al. (2017) from which we quote.

“Many existing probabilistic programming languages treat the inference engine as a black box, abstracted away from the model. These cannot capture probabilistic inferences that reuse the model’s representation - a key idea in recent advances in variational inference, generative adversarial networks, and also in more classic inference.”

The second aspect concerns the data at hand. In MCMC the position for the chain is driven by (1) the prior, (2) the prior and the likelihood, (3) a sequence of priors and likelihood, (4)

a sequence of priors, likelihood and proposal distributions, basically anything, except for the data itself. Data-driven approaches would namely destroy the convergence of the Markov chain. To start an MCMC sampler in a data-driven manner and continue in a data-oblivious manner is a possible solution (Zhang and Perez-Cruz, 2017). Even better, it is possible to use a Metropolized independence sampler Jampani et al. (2015). Such a sampler samples independently from the previous state and uses global information. However, to work well its proposal distribution needs to match the target distribution quite well. Although, when combined with a local sampler, it might be sufficient to just be able to match the modes of the target distribution well.

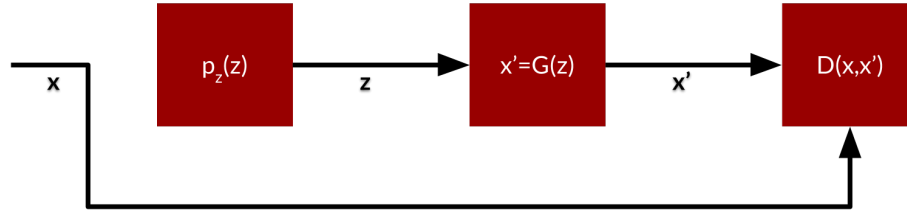
The third aspect concerns the way we build our MCMC engine. The split-merge sampler of the previous chapter has been meticulously designed. If we admit a data-driven approach, we might as well adjust our MCMC engine using training samples. Note, that this training will be across a set of line, box, or scenes mixtures. The MCMC engine will not be able to learn just the parameters of a particular visual object. It will learn how to jump around (optionally, adaptively) from one visual object to the next or from one cluster configuration to the next. In other words, it will be able to teach itself to become a Triadic Split-Merge sampler if that happens to be a good engine. Is it possible to constrain the search through MCMC kernels such that its result is always converging in an MCMC sense? If we aim to learn the transition operator of our Markov chains, there is new literature at hand that makes use of deep nets.

## 6.2 Learning the Transition Operator

There are multiple methods that can be used in a generative setting. We will discuss the three most prominent ones: (1) generative adversarial networks, (2) variational autoencoders, and (3) infusion training. This is far from an extensive categorization, worth studying are variational walkback (Goyal et al., 2017), stacked generative adversarial networks (Huang et al., 2016), generative latent optimization (Bojanowski et al., 2017), deep learning through the use of non-equilibrium thermodynamics (Sohl-Dickstein et al., 2015), denoising autoencoders, or generative stochastic networks, to name just a few.

### 6.2.1 Adversarial Training

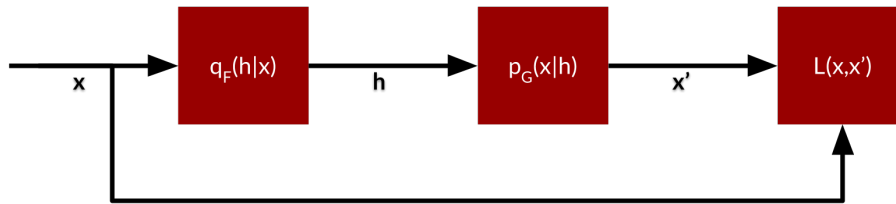
Adversarial training has been extensively studied since the article on generative adversarial networks by Goodfellow et al. (2014). A particular adversarial setup for training an MCMC has been suggested as well (Song et al., 2017). The generator samples from a Markov chain. A discriminator subsequently needs to judge if its incoming data comes from the generator or if it is sampled from the actual data set. To start the process, the generator can run the chain from the model as well as from the data.



**Figure 6.1:** Left:  $p_z(z)$  is a (prior) random distribution that generates random variables  $z$ . Middle: the generator  $G$  maps the random variables  $z$  to simulated data points  $x'$ . Right: the discriminator  $D(x, x')$  compares the simulated data  $x'$  with the real data  $x$ . The generator tries to generate samples in such way that the discriminator has difficulties distinguishing the simulated from the real data.

## 6.2.2 Variational Autoencoders

Variational autoencoders (Kingma and Welling, 2013; Rezende et al., 2014) are ordinary autoencoders with additional constraints on the latent variables. The latent variables in autoencoder parlance are called the code. In a variational autoencoder the latent variables are forced to approximately describe a unit Gaussian distribution. The autoencoder is trained using a loss function that is composed out of (1) a generative loss, a mean squared error that measures how accurately the network reconstructs its input, and (2) a latent loss, a KL-divergence that measures how closely the latent variables match a unit Gaussian. To optimize the KL divergence a reparameterization trick is applied. The encoder does not generate a vector with real values, but generates a vector with means and standard deviations instead.



**Figure 6.2:** Left:  $q_F(h|x)$  maps the data  $x$  to (hidden) random variables  $h$ . Middle:  $p_G(x|h)$  maps the hidden random variables to reconstructed data  $x'$ . Right:  $L(x, x')$  measures the similarity between  $x$  and  $x'$ .

## 6.2.3 Infusion Training

The transition operator can also be learned directly through infusion training. In infusion training we gradually adjust totally unstructured noise to a target distribution as well. In this method a particular data point is ‘infused’ into the Markov chain to bias the model sampling to move towards this particular data point (and not another). In contrast to a generative loss that is a mean squared error, this promises to have less blurry reconstructions.

### 6.3 Volumetric Models

In image processing autoencoders have been used for 2D shape recognition. To apply the same type of models to 3D point clouds, these point clouds are represented through voxels. The application of data-driven deep learning techniques, be it autoencoders, generative adversarial networks, or adversarial autoencoders promises similar good results in these 3D settings than in the current computer vision tasks.

The 3D ShapeNet model (Wu et al., 2015) exists of 3D voxel input that is piped through several stages with an increasing number of filters. The used voxel representation is a binary tensor. It assigns a value of 1 to each voxel that is inside the 3D object mesh and a value 0 to each voxel outside the mesh (empty space). The voxel sizes are fixed as well as the grid size (in this particular model the grid exists of 30x30x30 voxels). The inference model is a Deep Belief Network (DBN). Convolution operators, in the form of filters over small neighbourhoods, are used to reduce the number of model parameters (30x30x30 fully connected would be really many weights). The DBN is used in a supervised setting where shapes are trained with object labels. The model subsequently learns to generate shapes given an object label.

An unsupervised method in the form of a convolutional (volumetric) autoencoder (Sharma et al., 2016) has been applied to the same type of data. This (denoising) autoencoder, coined VConv-DAE maps from an entire voxel grid to another voxel grid. This work uses a combination of standard techniques, a dropout layer, a deconvolution layer, ReLu as well as sigmoid activation, but it is not in particular tailored to 3D point clouds.

Other representations than voxels are used. For example collections of 2D views and transformation parameters (Dosovitskiy et al., 2017). The most interesting are methods that work with raw data, the point cloud themselves. This alleviate the need to process the data and does not inadvertently increase the data dimensions, for example by artificially introduce voxels where there is no object present.

PointNet (?) directly operates on point clouds. To handle the input as a set of points (unordered), it uses a symmetric function over  $n$  input vectors and outputs a vector that is invariant to the input order. Typically sum and multiplication operators are such symmetric functions. After input and feature transforms by multi-layer perceptrons, a max pooling operator is used to map the input to a global feature. In the ModelNet40 shape classification benchmark there are more than 12000 CAD models from 40 object categories. PointNet achieves state of the art results compared to volumetric methods for a fraction of the computational costs.

Point clouds are also directly used in so-called deep kd-networks (Klokov and Lempitsky, 2017). A kd-tree is constructed by recursively picking the coordinate axis with the largest range of point coordinates and splitting the set of points into two subsets of equal size. These subsets are recursed into successively. The recursion stops at a particular level, depth  $D$ . The kd-networks are purported to outperform for example PointNet amongst other model architectures.



1694 A deep permutation equivariant (for semisupervised learning) and permutation invariant  
1695 (for supervised learning) network has also been directly applied to point clouds (Ravanbakhsh  
1696 et al., 2016). It does not reach the ModelNet40 accuracy levels from PointNet or the kd-  
1697 networks though.

1698 PointNet++ Qi et al. (2017) introduces hierarchical structure to PointNet. This fits better  
1699 non-uniform point distributions and seems to surpass kd-nets again on the ModelNet40 task.



1700

1701

## RECOMMENDER ENGINE

1702	<b>Contents</b>	The described nonparametric Bayesian models (Chapter 3, 4, 5, and 6)
1703		are not limited to computer vision tasks. This chapter describes a rec-
1704		ommender engine in which groups of runners are extracted from data
1705		collected from social media.
1706	<b>Outline</b>	We (1) introduce the form of the data at hand, (2) describe a multi-modal
1707		Von Mises-Uniform distribution to model the individual runners, (3) use
1708		a Dirichlet Process prior to group people, (4) use the previously described
1709		MCMC methods to perform inference, (5) show the results on an artificial
1710		and real-world data set, and (6) discuss ways with which the model can
1711		be expanded.

### 7.1 Application

1713 The data of people exercising can be considered binary (someone is either exercising or not  
 1714 in a particular timeslot). We do have however more information available. We know how  
 1715 often people have been exercising in a timeslot. This data has the form as visualized in  
 1716 Table 7.1.

**Table 7.1:** Example of the type of data about the timing of exercising. A person is represented by row, her preferences by column. There is not a predefined number of users or groups of users.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
0	0	0	0	4	23	38	9	12	6	2	7	2	3	2	7	5	3	2	0	3	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	4	11	4	4	2	1	2	0	0	0
0	0	0	0	0	1	0	0	1	2	0	0	0	1	0	2	3	1	1	3	2	0	0	0
0	0	0	0	0	0	5	4	1	3	11	3	7	3	3	4	3	10	23	3	0	0	0	0
0	0	0	0	0	0	8	35	23	12	4	41	14	11	8	7	14	38	36	6	5	1	2	0
0	0	0	0	0	0	0	1	9	2	0	3	0	0	0	1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	29	9	8	7	5	3	0	2	0	1	8	6	1	4	1	2	0	0
0	0	0	0	0	2	14	12	7	1	0	2	0	2	0	3	4	6	4	3	9	2	0	1
0	0	0	0	0	0	2	0	0	0	0	0	0	0	0	0	3	3	2	0	1	0	0	0
0	0	0	0	0	0	0	1	1	0	1	1	0	0	0	2	1	2	3	0	1	0	0	0
0	0	0	0	0	0	1	0	0	0	0	0	0	2	0	4	5	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	3	2	1	0	1	1	0	0	0	0	3	11	12	0	0	0
0	0	0	0	0	0	0	0	0	0	2	1	0	0	0	1	0	3	5	4	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	18	4	0	0	0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	2	0	0	1	9	14	14	2	1	4	15	4	0	0
0	0	0	0	0	0	0	0	2	4	2	1	3	5	2	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	0	1	1	0	0	0	1	1	3	3	1	3	2	3	0	0
0	0	0	0	0	0	0	1	2	1	3	2	0	1	2	0	0	0	0	1	2	0	0	0
0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	8	4	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	4	6	2	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	0	1	1	6	2	1	0	0	0	0	1	1	0	0	0	0
0	0	0	0	0	0	2	2	3	12	18	12	12	2	4	3	5	6	43	37	16	1	2	0

1717 The first row defines the time of day, starting at midnight till the last timeslot from 23.00 till  
 1718 00.00. In this particular case it shows that nobody is running from midnight to four o'clock  
 1719 in the morning, understandably so.

## 1720 7.2 Model of Individuals

### 1721 7.2.1 Multi-modal Normal-Uniform Distribution Model

1722 We first postulate the likelihood function for the moments at which people exercise through  
 1723 the day as defined in Eq. 7.1.

$$f(x|\theta) = w_0 \mathcal{U}(a, b) + \sum_{i=1}^2 w_i \mathcal{N}(\mu_i, \sigma_i) \quad (7.1)$$

1724 The likelihood (Eq. 7.1) is built up out of three probability density functions: one Uniform dis-  
 1725 tribution  $\mathcal{U}(a, b)$  with  $a$  and  $b$  as parameters and two Normal distributions  $\mathcal{N}(\mu_i, \sigma_i)$  with  
 1726 mean  $\mu_i$  and  $\sigma_i$ . The distributions are weighted by the factors  $w_0, w_1, w_2$ . The collection of  
 1727 parameters for the likelihood function is referred to by  $\theta = \{a, b, \mu_1, \sigma_1, \mu_2, \sigma_2, w_0, w_1, w_2\}$ .

1728 This probability density function  $f(x|\theta)$  will have the form as in Fig. 7.1. The uniform  
 1729 distribution generates values here between 00:00 and 24:00. There are on top of that the  
 1730 two Normal distributions that form peaks at certain moments during the day.



**Figure 7.1:** The likelihood function for the moments at which people decide to exercise during the day. On the horizontal axis time, on the vertical axis the frequency of exercising.

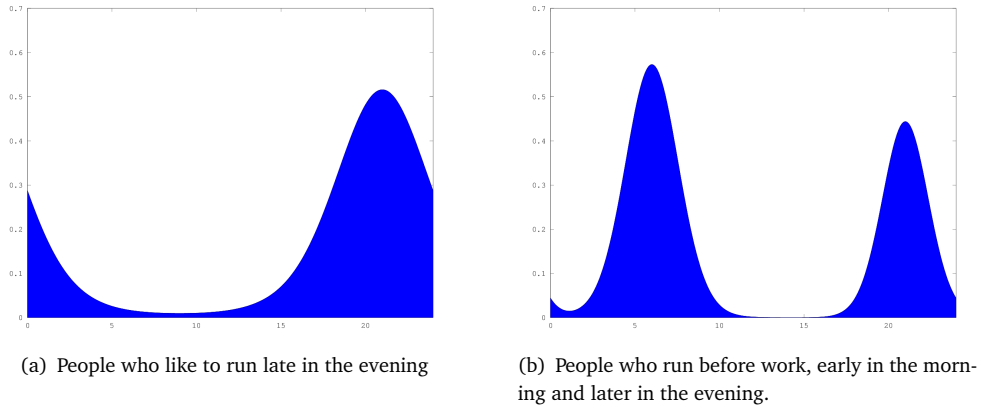
1731 There is something noticable in Fig. 7.1, namely that the 24 hours of a day cause the Nor-  
 1732 mal distribution to be cut off. Especially in Fig. 7.1 (a) there should be some considerable  
 1733 likelihood of running in the wee hours of the morning between 00:00 and 01:00.

## 1734 7.2.2 Multi-modal Von-Mises-Uniform Distribution Model

1735 There are several options to define a distribution over a limited range  $T$ . A so-called wrapped  
 1736 distribution is a distribution defined over the unity circle. By just multiplying it with  $T/(2\pi i)$   
 1737 it can be used to define a probability density function over a day ( $T = 24$ ).

$$f(x|\theta) = w_0 \mathcal{U}(a, b) + \sum_{i=1}^2 w_i \mathcal{VM}(\mu_i, \kappa_i) \quad (7.2)$$

1738 The likelihood (Eq. 7.2) is again built up out of three probability density functions: one  
 1739 Uniform distribution  $\mathcal{U}(a, b)$  with  $a$  and  $b$  as parameters and two Von Mises distributions  
 1740  $\mathcal{VM}(\mu_i, \kappa_i)$  with mean  $\mu_i$  and  $\kappa_i$ . The parameters  $\mu_i$  will be scaled and shifted with  $[a, b]$  so  
 1741 all variables within this range fall on the unity circle. The parameter  $\kappa_i$  plays the same role as  
 1742  $\sigma_i$  for the Normal distribution. The distributions are weighted by the factors  $w_0, w_1, w_2$ . The  
 1743 collection of parameters for the likelihood function is referred to by  $\theta = \{a, b, \mu_1, \kappa_1, \mu_2, \kappa_2, w_0, w_1, w_2\}$ .



**Figure 7.2:** The improved likelihood function for the moments at which people decide to exercise during the day using the Von Mises distribution. On the horizontal axis time, on the vertical axis the frequency of exercising.

1744 The likelihood with Von Mises distributions rather than Normal distributions is visualized  
 1745 in Fig. 7.2. The Von Mises distributions capture behavior in deviations from a standard  
 1746 exercise time. It does not take into account structured deviations, for example running late  
 1747 from work every Thursday, or weekends, or a person either running in the morning or in the  
 1748 evening, but never both at the same day.

### 1749 7.2.3 Hyperparameters

1750 The parameters  $\theta$  for each user  $j$  are either fixed or generated from prior distributions. The  
 1751 hyperparameters  $a$  and  $b$  for the Uniform distribution are set to 0 and 24. People can run  
 1752 potentially any time of the day. The hyperparameters  $\mu_i, \kappa_i$  for the Von Mises distributions  
 1753 are generated from a Uniform-Exponential distribution (Eq. 7.3). The Uniform distribution  
 1754 reflects the fact that if people run on regular times, this time can be any time of the day. The  
 1755 Exponential distribution defines a prior on how much people deviate from such a regular  
 1756 time to exercise.

$$\begin{aligned} f(\mu_i|a, b) &= \mathcal{U}(a, b) \\ f(\kappa_i|\lambda) &= \mathcal{E}(\lambda) \end{aligned} \tag{7.3}$$

1757 The Uniform distribution generates  $\mu_i$  between  $a$  and  $b$ . The parameter  $\kappa_i$  is generated from  
 1758 an Exponential distribution with hyperparameter  $\lambda$ . If  $\lambda$  is set to be small ( $< 0.5$ ) we have  
 1759 a high likelihood that  $\kappa$  can be large and we have pronounced peaks. In contrary, if  $\lambda$  is set  
 1760 to be large, the Von Mises distribution likely approaches the Uniform distribution due to a  
 1761 higher chance of sampling a small value for  $\kappa$ .

1762 The weights we sample from a normalized product of a zero-deflated Bernoulli distribution  
 1763 and a Dirichlet distribution (Eq. 7.4).

$$f(w_i|p, \alpha_i) = \mathcal{B}(p)\mathcal{D}(\alpha_i)/Z \quad (7.4)$$

1764 The Bernoulli distribution samples zeros and ones with probability of  $p = 0.5$ , leading to  
 1765 3-vectors like 001, 101, etc. The distribution is corrected in such way that the chance to  
 1766 sample 000 is zero. The Bernoulli distribution only gives weights  $w_i = 0$  or  $w_i = 1$ , hence  
 1767 it is multiplied with a Dirichlet distribution. The Dirichlet samples a 3-vector with weights  
 1768 between zero and one where the weights  $\sum_i w_i = 1$ . A symmetric Dirichlet distribution  
 1769 with  $\alpha = 1$  is similar to the Uniform distribution over the simplex. It is set slightly more  
 1770 towards favoring particular distributions with  $\alpha_i = 1/3$  (we assume that people sample one  
 1771 or two distributions, and rarely from all three). The product with the zero-deflated Bernoulli  
 1772 distribution is made up to sum up to one again by normalizing the result.

1773 We can combine Eq. 7.3 and Eq. 7.4 in Eq. 7.5:

$$\theta \sim \mathcal{U}(a, b)\mathcal{E}(\lambda)\mathcal{B}(p)\mathcal{D}(\alpha_i)/Z \quad (7.5)$$

1774 To sample the parameters  $\theta$  this is the base distribution we will encounter in the next sec-  
 1775 tion 7.3.

## 1776 7.3 Model of Groups

1777 Each person's exercise schedule is represented by a Von-Mises-Uniform distribution. People  
 1778 that are similar do have exercise schedules that can be represented by the same Von-Mises-  
 1779 Uniform distribution. To group similar schedules we define a nonparametric discrete distri-  
 1780 bution over a potentially infinite number of groups with each person assigned to a group.

1781 A Dirichlet Process (Eq. 7.6) is a distribution over distributions that can be used as a prior  
 1782 for such a nonparametric discrete distribution.

$$DP(\alpha, H) \quad (7.6)$$

1783 The Dirichlet Process has (1) a hyperparameter  $\alpha$ , which defines the likelihood that there are  
 1784 many clusters versus few clusters (although it doesn't say anything about its actual count),  
 1785 and (2) a base distribution  $H$ , the distribution that generates  $\theta$  (Eq.7.5).

## 1786 7.4 Inference

1787 The implementation of the model makes use of Gibbs sampling with auxiliary variables Jain  
 1788 and Neal (2007).

1789 Details on this algorithm can be found in Jain and Neal (2007) and previous work of the  
 1790 authors van Rossum et al. (2016a,b).

**Algorithm 14** Gibbs sampling over auxiliary variables

---

```

1: procedure GIBBS ALGORITHM WITH AUXILIARY VARIABLES( $w, \lambda_0, \alpha$ )      ▷ Accepts schedule  $w$ ,
   hyperparameters  $\lambda_0, \alpha$ , number of auxiliary variables  $m$ , and returns  $k$  groups
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:       for all  $j = 1 : m$  do
5:          $\theta_j \sim H(\lambda_0)$                                           ▷ Sample  $\theta_j$  from base distribution  $H$  in Eq. 7.5
6:       end for
7:       for all  $j = 1 : K + m, j \neq i$  do
8:          $L_j = \text{likelihood}(w_i, \theta_j)$       ▷ Update likelihood for all theta (except  $\theta_i$ ) given  $w_i$ 
9:       end for
10:       $P_{-i=1:K} = b \sum_{-i} L_{-i}$                                           ▷ Calculate probability of existing cluster
11:       $P_{-i=K:K+m} = b\alpha/mL_mL_{-i}$       ▷ Calculate probability of new cluster
12:       $\theta_i = \theta_j$  according to above  $P_{-i}$       ▷ Sample  $\theta_i$  accord. to above prob
13:      Remove unused clusters
14:    end for
15:    for all  $j = 1 : K$  do
16:       $\theta_j \sim p(\theta_j | y)$                                           ▷ Update  $\theta_j$ 
17:    end for
18:  end for
19:  return summary on  $\theta_k$  for  $k$  groups of runners
20: end procedure

```

---

## 7.5 Results

The algorithm is run first on an artificial dataset of which we know the ground truth (Sect. 7.5.1) and next on the real-world dataset from Twitter (Sect. 7.5.2).

### 7.5.1 Artificial Dataset

The algorithm has been used on a self-generated dataset (generated from the probability density function as in Fig. 7.2). In this case we have the ground truth that establishes which exercise schedule comes from which probability density function. Using this ground truth we can calculate how often our algorithm makes a mistake, grouping a person with people that belong to another group. The results with this dataset are perfect (Rand Index equal to one: perfect clustering). The results are of the form of Table 7.2 (except by a permutation of indices), hence are not shown (the indices have no intrinsic meaning).

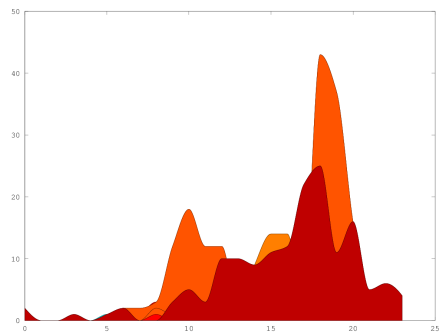
**Table 7.2:** Top row: A sequence of cluster indices indicates the ground truth. Each cluster index represents a multi-modal Von-Mises-Uniform distribution with different parameters  $\theta$ . Bottom row: A sequence of cluster indices that are the result of the described algorithm. Each cluster index represents again a multi-modal Von-Mises distribution. Errors would be represented by an inconsistent mapping from the top row to the bottom row.

1	4	4	5	2	2	2	1	3	2	1	5	1	1	4	4	3	1	3	1	2	...
4	5	5	1	2	2	2	4	3	2	4	1	4	4	5	5	3	4	3	4	2	...

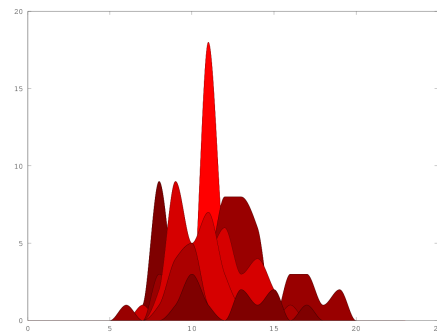


## 1802 7.5.2 Real-world Dataset

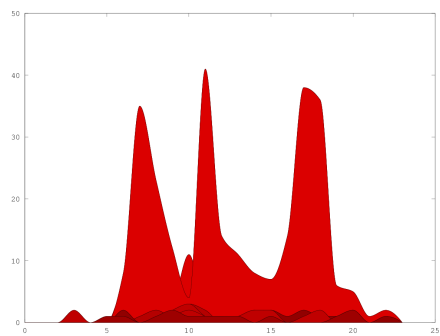
1803 The results when we actually use the collected Twitter dataset can be best visualized (Fig. 7.3).  
 1804 This dataset consists of around 4000 moments at which people decide to run. The dataset  
 1805 is subsequently filtered on regulars, people that at least have run a few times.



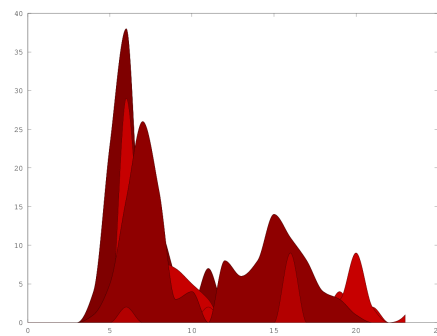
(a) People who predominally seem to like to run in the early evening.



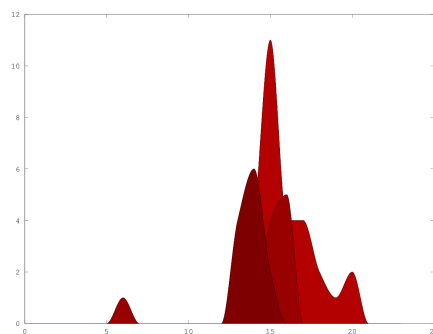
(b) People who run during the day, not later than at 20:00 o'clock.



(c) People who seem to run all the time (interesting collection though!).



(d) People who run early in the morning.



(e) People who run around 16:00 o'clock and almost never run at another time.

**Figure 7.3:** After running the algorithm, the above figures show the different categories of runners that have been found.

1806 The results in Fig. 7.3 show that there is a categorization of people indeed. After the al-  
 1807 gorithm has done its work it is possible to assign a certain label in a post-hoc manner, e.g.  
 1808 “People who run around 16:00 o'clock and almost never run at another time.”.

## 1809 7.6 Discussion

1810 The postulated model allows to reason about groups of people performing exercising, with-  
1811 out predefining what these groups constitute apart from very general characteristics such  
1812 as that there might be preferred times of day to exercise. A Dirichlet process is used as a  
1813 nonparametric Bayesian model in which both the number of groups and the assignment of  
1814 people exercising are learnt from the data.

1815 There are several directions in which this research can be extended. First, more data would  
1816 be very helpful. Only a limited number of people are posting consistently their training  
1817 data online. Tapping into the data of current fitness promoting companies would allow the  
1818 model to wash out the prior a bit more and adjust to the data. Second, we also collected  
1819 weather data over this time period and also expect the day of the week to have significant  
1820 influence. When there will be more data available these are logically dimensions to include  
1821 in the dataset. Third, it would be interesting to study if people relate to the group of people  
1822 they have been categorized with. Can this help or support their exercise regime?



1823



1824

## DISCUSSION AND CONCLUSIONS



## REFERENCES

- 1827 M Abdel-Hameed. Optimal replacement policies for devices subject to a gamma wear process. *The*  
 1828 *theory and applications of reliability*, pages 397–412, 2012.
- 1829 David J Aldous. Representations for partially exchangeable arrays of random variables. *Journal of*  
 1830 *Multivariate Analysis*, 11(4):581–598, 1981.
- 1831 David J Aldous. *Exchangeability and related topics*. Springer, 1985.
- 1832 John Aldrich and Others. R.A. Fisher and the making of Maximum Likelihood 1912-1922. *Statistical*  
 1833 *Science*, 12(3):162–176, 1997.
- 1834 Charles E Antoniak. Mixtures of Dirichlet processes with applications to Bayesian nonparametric  
 1835 problems. *The annals of statistics*, pages 1152–1174, 1974.
- 1836 Stefan Banach and Alfred Tarski. Sur la décomposition des ensembles de points en parties respec-  
 1837 tivement congruentes. *Fund. math*, 6(1):924, 1924.
- 1838 Federico Bassetti, Roberto Casarin, and Fabrizio Leisen. Beta-product dependent Pitman–Yor pro-  
 1839 cesses for Bayesian inference. *Journal of Econometrics*, 180(1):49–72, 2014.
- 1840 Leonard E Baum and Ted Petrie. Statistical inference for probabilistic functions of finite state Markov  
 1841 chains. *The annals of mathematical statistics*, 37(6):1554–1563, 1966.
- 1842 Herbert Bay, Tinne Tuytelaars, and Luc Van Gool. Surf: Speeded up robust features. In *European*  
 1843 *conference on computer vision*, pages 404–417. Springer, 2006.
- 1844 David Blackwell and James B MacQueen. Ferguson distributions via Pólya urn schemes. *The annals*  
 1845 *of statistics*, pages 353–355, 1973.
- 1846 Phil Blunsom and Trevor Cohn. A hierarchical Pitman-Yor process HMM for unsupervised part of  
 1847 speech induction. In *Proceedings of the 49th Annual Meeting of the Association for Computational*  
 1848 *Linguistics: Human Language Technologies-Volume 1*, pages 865–874. Association for Computa-  
 1849 tional Linguistics, 2011.
- 1850 Piotr Bojanowski, Armand Joulin, David Lopez-Paz, and Arthur Szlam. Optimizing the latent space  
 1851 of generative networks. *arXiv preprint arXiv:1707.05776*, 2017.
- 1852 Robert C Bolles and Martin A Fischler. A RANSAC-Based Approach to Model Fitting and Its Application  
 1853 to Finding Cylinders in Range Data. In *IJCAI*, volume 1981, pages 637–643, 1981.
- 1854 Anna Bosch, Andrew Zisserman, and Xavier Muñoz. Scene classification using a hybrid generative/dis-  
 1855 criminative approach. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 30(4):712–  
 1856 727, 2008.
- 1857 Guillaume Bouchard and Bill Triggs. The tradeoff between generative and discriminative classifiers.  
 1858 In *16th IASC International Symposium on Computational Statistics (COMPSTAT'04)*, pages 721–  
 1859 728, 2004.

- George E P Box and George C Tiao. *Bayesian inference in statistical analysis*, volume 40. John Wiley & Sons, 2011.
- H Bühlmann. *Austauschbare stochastische Variablen und ihre Grenzwertsätze*. PhD thesis, ETH Zürich, 1960.
- Wray L Buntine. Operations for learning with graphical models. *JAIR*, 2:159–225, 1994.
- John Canny. A computational approach to edge detection. *IEEE Transactions on pattern analysis and machine intelligence*, (6):679–698, 1986.
- Haifeng Chen, Peter Meer, and David E Tyler. Robust regression for data with multiple structures. In *Computer Vision and Pattern Recognition, 2001. CVPR 2001. Proceedings of the 2001 IEEE Computer Society Conference on*, volume 1, pages I—1069. IEEE, 2001.
- David B Dahl. An Improved Merge-Split Sampler for Conjugate Dirichlet Process Mixture Models. Technical report, University of Wisconsin–Madison, November 2003.
- David B Dahl. Sequentially-Allocated Merge-Split Sampler for Conjugate and Nonconjugate Dirichlet Process Mixture Models. Technical report, Texas A&M University, November 2005.
- Navneet Dalal and Bill Triggs. Histograms of oriented gradients for human detection. In *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR’05)*, volume 1, pages 886–893. IEEE, 2005.
- Hal Daume. Fast search for Dirichlet process mixture models. In *International Conference on Artificial Intelligence and Statistics*, pages 83–90, 2007.
- B de Finetti. Funzione caratteristica di un fenomeno aleatorio. *Atti Reale Accademia Nazionale dei Lincei*, VI:86–133, 1930.
- Persi Diaconis and David Freedman. de Finetti’s theorem for Markov chains. *The Annals of Probability*, pages 115–130, 1980.
- Alexey Dosovitskiy, Jost Tobias Springenberg, Maxim Tatarchenko, and Thomas Brox. Learning to generate chairs, tables and cars with convolutional networks. *IEEE transactions on pattern analysis and machine intelligence*, 39(4):692–705, 2017.
- François Dufresne, Hans U Gerber, and Elias S W Shiu. Risk theory with the gamma process. *Astin Bulletin*, 21(02):177–192, 1991.
- David B Dunson, Ya Xue, and Lawrence Carin. The matrix stick-breaking process. *Journal of the American Statistical Association*, 2012.
- Richard Durstenfeld. Algorithm 235: Random Permutation. *Communications of the ACM*, 7(7):420, July 1964. ISSN 0001-0782. doi: 10.1145/364520.364540. URL <http://doi.acm.org/10.1145/364520.364540>.
- Michael D Escobar and Mike West. Bayesian density estimation and inference using mixtures. *Journal of the american statistical association*, 90(430):577–588, 1995.
- Stewart N Ethier. The distribution of the frequencies of age-ordered alleles in a diffusion model. *Advances in Applied Probability*, pages 519–532, 1990.
- Warren John Ewens. Population genetics theory-the past and the future. In *Mathematical and statistical developments of evolutionary theory*, pages 177–227. Springer, 1990.
- Stefano Favaro, Yee Whye Teh, and Others. MCMC for normalized random measure mixture models. *Statistical Science*, 28(3):335–359, 2013.
- William Feller. An introduction to probability theory and its applications. Vol. I. 1950.
- Thomas S Ferguson. A Bayesian analysis of some nonparametric problems. *The annals of statistics*, pages 209–230, 1973.
- Thomas S Ferguson. Prior distributions on spaces of probability measures. *The annals of statistics*, pages 615–629, 1974.
- E.W. Forgy. Cluster analysis of multivariate data: Efficiency vs. interpretability of classifications. *Biometrics*, 21:768–769, 1965.

- Wolfgang Förstner and Eberhard Gülch. A fast operator for detection and precise location of distinct points, corners and centres of circular features. In *Proc. ISPRS intercommission conference on fast processing of photogrammetric data*, pages 281–305, 1987.
- David Freedman and Persi Diaconis. On inconsistent Bayes estimates in the discrete case. *The Annals of Statistics*, pages 1109–1118, 1983.
- David Heaven Fremlin. *Measure theory*, volume 4. Torres Fremlin, 2000.
- Orazio Gallo, Roberto Manduchi, and Abbas Rafii. CC-RANSAC: Fitting planes in the presence of multiple surfaces in range data. *Pattern Recognition Letters*, 32(3):403–410, 2011.
- Qing-Bin Gao and Shi-Liang Sun. Human activity recognition with beta process hidden Markov models. In *Machine Learning and Cybernetics (ICMLC), 2013 International Conference on*, volume 2, pages 549–554. IEEE, 2013.
- Stuart Geman and Donald Geman. Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, (6):721–741, 1984.
- Zoubin Ghahramani and Thomas L Griffiths. Infinite latent feature models and the Indian buffet process. In *Advances in neural information processing systems*, pages 475–482, 2005.
- Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in neural information processing systems*, pages 2672–2680, 2014.
- Anirudh Goyal, Nan Rosemary Ke, Surya Ganguli, and Yoshua Bengio. Variational walkback: Learning a transition operator as a stochastic recurrent net. *arXiv preprint arXiv:1711.02282*, 2017.
- Max Halperin and G L Burrows. The Effect of Sequential Batching for Acceptance–Rejection Sampling Upon Sample Assurance of Total Product Quality. *Technometrics*, 2(1):19–26, 1960.
- Chris Harris and Mike Stephens. A combined corner and edge detector. In *Alvey vision conference*, volume 15, page 50. Citeseer, 1988.
- Li He, Hairong Qi, and Russell Zaretzki. Beta process joint dictionary learning for coupled feature spaces with application to single image super-resolution. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 345–352, 2013.
- Nils Lid Hjort. Nonparametric Bayes estimators based on beta processes in models for life history data. *The Annals of Statistics*, pages 1259–1294, 1990.
- Fred M Hoppe. Size-biased filtering of Poisson-Dirichlet samples with an application to partition structures in genetics. *Journal of Applied Probability*, pages 1008–1012, 1986.
- Paul V.C. Hough. Method and Means for Recognizing Complex Patterns, Dec 1962. URL <https://www.google.com/patents/US3069654>. Patent US 3069654 A.
- Xun Huang, Yixuan Li, Omid Poursaeed, John Hopcroft, and Serge Belongie. Stacked generative adversarial networks. *arXiv preprint arXiv:1612.04357*, 2016.
- Tommi S Jaakkola, David Haussler, and Others. Exploiting generative models in discriminative classifiers. *Advances in neural information processing systems*, pages 487–493, 1999.
- Sonia Jain and Radford M. Neal. A Split-Merge Markov Chain Monte Carlo Procedure for the Dirichlet Process Mixture Model. *Journal of Computational and Graphical Statistics*, 13(1):158–182, 2004. ISSN 10618600. URL <http://www.jstor.org/stable/1391150>.
- Sonia Jain and Radford M Neal. Splitting and Merging Components of a Nonconjugate Dirichlet Process Mixture Model. *Bayesian Analysis*, 2(3):445–472, 2007.
- Varun Jampani, Sebastian Nowozin, Matthew Loper, and Peter V Gehler. The informed sampler: A discriminative approach to bayesian inference in generative computer vision models. *Computer Vision and Image Understanding*, 136:32–44, 2015.
- Edwin T Jaynes. *Probability theory: The logic of science*. Cambridge university press, 2003.



- 1955 Dominik Joho, Gian Diego Tipaldi, Nikolas Engelhard, Cyrill Stachniss, Wolfram Burgard, Mar-  
 1956 tin Senk, Felix Faber, Maren Bennewitz, Clemens Eppner, Attila Görög, and Others. Unsuper-  
 1957 vised Scene Analysis and Reconstruction Using Nonparametric Bayesian Models. *Robotics and*  
 1958 *Autonomous Systems (RAS)*, 59(5):319–328, 2011.
- 1959 A Jordan. On discriminative vs. generative classifiers: A comparison of logistic regression and naive  
 1960 Bayes. *Advances in neural information processing systems*, 14:841, 2002.
- 1961 Michael I Jordan. Hierarchical models, nested models and completely random measures. *Frontiers of*  
 1962 *Statistical Decision Making and Bayesian Analysis: in Honor of James O. Berger*. New York: Springer,  
 1963 2010.
- 1964 Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint*  
 1965 *arXiv:1312.6114*, 2013.
- 1966 J F C Kingman. *Poisson processes*, volume 3 of *Oxford Studies in Probability*. The Clarendon Press  
 1967 Oxford University Press, New York, 1993. ISBN 0-19-853693-3.
- 1968 J F C Kingman. Some further analytical results in the theory of regenerative events. *Journal of*  
 1969 *Mathematical Analysis and Applications*, 11:422–433, 1965.
- 1970 J F C Kingman. Random partitions in population genetics. In *Proceedings of the Royal Society of*  
 1971 *London A: Mathematical, Physical and Engineering Sciences*, volume 361, pages 1–20. The Royal  
 1972 Society, 1978.
- 1973 John Kingman. Completely random measures. *Pacific Journal of Mathematics*, 21(1):59–78, 1967.
- 1974 Roman Klovov and Victor Lempitsky. Escape from cells: Deep kd-networks for the recognition of  
 1975 3d point cloud models. In *2017 IEEE International Conference on Computer Vision (ICCV)*, pages  
 1976 863–872. IEEE, 2017.
- 1977 David Knowles, Zoubin Ghahramani, and Konstantina Palla. A reversible infinite HMM using nor-  
 1978 malised random measures. In *Proceedings of the 31st International Conference on Machine Learning*  
 1979 *(ICML-14)*, pages 1998–2006, 2014.
- 1980 Daphne Koller and Nir Friedman. *Probabilistic graphical models: principles and techniques*. MIT press,  
 1981 2009.
- 1982 A Kolmogoroff. *Grundbegriffe der wahrscheinlichkeitsrechnung*, volume 2 of *Ergebnisse der Mathe-*  
 1983 *matik und ihrer Grenzgebiete*. Springer-Verlag, 1933.
- 1984 Uwe Küchler and Stefan Tappe. Bilateral Gamma distributions and processes in financial mathemat-  
 1985 ics. *Stochastic Processes and their Applications*, 118(2):261–283, 2008.
- 1986 Kenichi Kurihara, Max Welling, and Yee Whye Teh. Collapsed Variational Dirichlet Process Mixture  
 1987 Models. In *IJCAI*, volume 7, pages 2796–2801, 2007.
- 1988 Henri Lebesgue. Intégrale, longueur, aire. *Annali di Matematica Pura ed Applicata (1898-1922)*, 7  
 1989 (1):231–359, 1902.
- 1990 Antonio Lijoi and Igor Prünster. Models beyond the Dirichlet process. *Bayesian nonparametrics*, 28:  
 1991 80, 2010.
- 1992 S. Lloyd. Least squares quantization in pcm. *IEEE Trans. Inform. Theory*, 28:129–137, 1982. Originally  
 1993 as an unpublished Bell laboratories Technical Note (1957).
- 1994 David G Lowe. Object recognition from local scale-invariant features. In *Computer vision, 1999.*  
 1995 *The proceedings of the seventh IEEE international conference on*, volume 2, pages 1150–1157. Ieee,  
 1996 1999.
- 1997 Steven N MacEachern. Comment on "Splitting and Merging Components of a Nonconjugate Dirichlet  
 1998 Process Mixture Model" by Jain and Neal. *Bayesian Analysis*, 2(3):483–494, 2007.
- 1999 Steven N MacEachern and Peter Müller. Estimating mixture of Dirichlet process models. *Journal of*  
 2000 *Computational and Graphical Statistics*, 7(2):223–238, 1998.
- 2001 Dilip B Madan and Eugene Seneta. The variance gamma (VG) model for share market returns. *Journal*  
 2002 *of business*, pages 511–524, 1990.
- 2003 Pierre-Simon marquis de Laplace. *Théorie analytique des probabilités*. V. Courcier, 1820.

- 2004 Nicholas Metropolis, Arianna W. Rosenbluth, Marshall N. Rosenbluth, Augusta H. Teller, and Edward  
2005 Teller. Equation of State Calculations by Fast Computing Machines. *The Journal of Chemical Physics*,  
2006 21(6):1087–1092, 1953. doi: 10.1063/1.1699114.
- 2007 Thomas Minka. Bayesian linear regression. Technical report, Citeseer, 2000.
- 2008 Radford M Neal. Defining Priors for Distributions Using Dirichlet Diffusion Trees. Technical report,  
2009 University of Toronto, Toronto, Ontario, Canada, 2001.
- 2010 Radford M Neal. Markov Chain Sampling Methods for Dirichlet Process Mixture Models. *Journal of*  
2011 *Computational and Graphical Statistics*, 9(2):249–265, 2000.
- 2012 Peter Orbanz and Daniel M Roy. Bayesian models of graphs, arrays and other exchangeable random  
2013 structures. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 37(2):437–461, 2015.
- 2014 Jim Pitman and Marc Yor. The two-parameter Poisson-Dirichlet distribution derived from a stable  
2015 subordinator. *The Annals of Probability*, pages 855–900, 1997.
- 2016 Charles Ruizhongtai Qi, Li Yi, Hao Su, and Leonidas J Guibas. Pointnet++: Deep hierarchical feature  
2017 learning on point sets in a metric space. In *Advances in Neural Information Processing Systems*,  
2018 pages 5105–5114, 2017.
- 2019 Rajat Raina, Yirong Shen, Andrew McCallum, and Andrew Y Ng. Classification with hybrid gen-  
2020 erative/discriminative models. In *Advances in neural information processing systems*, page None,  
2021 2003.
- 2022 William M Rand. Objective criteria for the evaluation of clustering methods. *Journal of the American*  
2023 *Statistical association*, 66(336):846–850, 1971.
- 2024 Vinayak Rao and Yee W Teh. Spatial normalized gamma processes. In *Advances in neural information*  
2025 *processing systems*, pages 1554–1562, 2009.
- 2026 Carl Edward Rasmussen and Christopher K I Williams. *Gaussian processes for machine learning*, vol-  
2027 ume 2. 2006.
- 2028 Siamak Ravanbakhsh, Jeff Schneider, and Barnabas Poczos. Deep learning with sets and point clouds.  
2029 *arXiv preprint arXiv:1611.04500*, 2016.
- 2030 Eugenio Regazzini, Antonio Lijoi, and Igor Prünster. Distributional results for means of normalized  
2031 random measures with independent increments. *Annals of Statistics*, pages 560–585, 2003.
- 2032 Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and  
2033 approximate inference in deep generative models. *arXiv preprint arXiv:1401.4082*, 2014.
- 2034 Daniel M Roy and Yee W Teh. The mondrian process. In *Advances in neural information processing*  
2035 *systems*, pages 1377–1384, 2009.
- 2036 Donald B Rubin and Others. Bayesianly justifiable and relevant frequency calculations for the applied  
2037 statistician. *The Annals of Statistics*, 12(4):1151–1172, 1984.
- 2038 Stuart Russell, Peter Norvig, and Artificial Intelligence. Artificial Intelligence: A modern approach.  
2039 *Artificial Intelligence. Prentice-Hall, Englewood Cliffs*, 25:27, 1995.
- 2040 Leonard J Savage. *The foundations of statistics*. Courier Corporation, 1972.
- 2041 Stanley Sawyer and Daniel Hartl. A sampling theory for local selection. *Journal of Genetics*, 64(1):  
2042 21–29, 1985.
- 2043 René L Schilling. *Measures, integrals and martingales*, volume 13. Cambridge University Press, 2005.
- 2044 Abhishek Sharma, Oliver Grau, and Mario Fritz. Vconv-dae: Deep volumetric shape learning without  
2045 object labels. In *European Conference on Computer Vision*, pages 236–250. Springer, 2016.
- 2046 Jianbo Shi and Carlo Tomasi. Good features to track. In *Computer Vision and Pattern Recognition,*  
2047 *1994. Proceedings CVPR’94., 1994 IEEE Computer Society Conference on*, pages 593–600. IEEE,  
2048 1994.
- 2049 Nozer Singpurwalla. Gamma processes and their generalizations: an overview. In *Engineering prob-*  
2050 *abilistic design and maintenance for flood protection*, pages 67–75. Springer, 1997.

- 2051 Scott A Sisson and Yanan Fan. Likelihood-free MCMC. *Handbook of Monte Carlo*, ed. Brooks, A.,  
2052 Gelman, A., Jones, GL, Meng XL, pages 313–335, 2011.
- 2053 I Sobel. *Camera Models and Perception*. PhD thesis, Stanford University, Stanford, CA, 1970.
- 2054 Jascha Sohl-Dickstein, Eric A Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised  
2055 learning using nonequilibrium thermodynamics. *arXiv preprint arXiv:1503.03585*, 2015.
- 2056 Jiaming Song, Shengjia Zhao, and Stefano Ermon. A-nice-mc: Adversarial training for mcmc. *arXiv*  
2057 *preprint arXiv:1706.07561*, 2017.
- 2058 Mike Steel. The maximum likelihood point for a phylogenetic tree is not unique. *Systematic Biology*,  
2059 43(4):560–564, 1994.
- 2060 Erik B Sudderth and Michael I Jordan. Shared segmentation of natural scenes using dependent  
2061 Pitman-Yor processes. In *Advances in Neural Information Processing Systems*, pages 1585–1592,  
2062 2009.
- 2063 Charles Sutton and Andrew McCallum. An introduction to conditional random fields. *Machine Learn-*  
2064 *ing*, 4(4):267–373, 2011.
- 2065 Terence Tao. *An introduction to measure theory*, volume 126. American Mathematical Soc., 2011.
- 2066 Yee Whye Teh, Michael I Jordan, Matthew J Beal, and David M Blei. Hierarchical Dirichlet processes.  
2067 *Journal of the American statistical association*, 101(476), 2006.
- 2068 Romain Thibaux and Michael I Jordan. Hierarchical beta processes and the Indian buffet process. In  
2069 *International Conference on Artificial Intelligence and Statistics*, pages 564–571, 2007.
- 2070 Michalis K Titsias. The infinite gamma-Poisson feature model. In *Advances in Neural Information*  
2071 *Processing Systems*, pages 1513–1520, 2008.
- 2072 Dustin Tran, Matthew D Hoffman, Rif A Saurous, Eugene Brevdo, Kevin Murphy, and David M Blei.  
2073 Deep probabilistic programming. *arXiv preprint arXiv:1701.03757*, 2017.
- 2074 Anne C. van Rossum, Hai Xiang Lin, Johan Dubbeldam, and H. Jaap van den Herik. Nonparametric  
2075 Bayesian Line Detection - Towards Proper Priors for Robotic Computer Vision. In *Proceedings of*  
2076 *the 5th International Conference on Pattern Recognition Applications and Methods*, pages 119–127,  
2077 Feb 2016a. ISBN 978-989-758-173-1. doi: 10.5220/0005673301190127.
- 2078 Anne C. van Rossum, Hai Xiang Lin, Johan Dubbeldam, and H. Jaap van den Herik. Nonparametric  
2079 Segment Detection. In *Proceedings of the 8th European Starting AI Researcher Symposium (STAIRS)*,  
2080 pages 203–208, Aug 2016b. ISBN 978-989-758-173-1. doi: 10.5220/0005673301190127.
- 2081 Niklas Vanhainen and Giampiero Salvi. Word Discovery with Beta Process Factor Analysis. In *INTER-*  
2082 *SPEECH*, pages 799–802, 2012.
- 2083 Wei Wang and Stuart Russell. A Smart-dumb/Dumb-smart Algorithm for Efficient Split-merge MCMC.  
2084 In *Proceedings of the Thirty-First Conference on Uncertainty in Artificial Intelligence*, UAI15, pages  
2085 902–911, Arlington, Virginia, United States, 2015. AUAI Press. ISBN 978-0-9966431-0-8. URL  
2086 <http://dl.acm.org/citation.cfm?id=3020847.3020940>.
- 2087 Yingjian Wang and Lawrence Carin. Levy measure decompositions for the beta and gamma processes.  
2088 *arXiv preprint arXiv:1206.4615*, 2012.
- 2089 Larry Wasserman. Asymptotic properties of nonparametric Bayesian procedures. In *Practical non-*  
2090 *parametric and semiparametric Bayesian statistics*, pages 293–304. Springer, 1998.
- 2091 Robert L Wolpert and Katja Ickstadt. Poisson/gamma random field models for spatial statistics.  
2092 *Biometrika*, 85(2):251–267, 1998.
- 2093 Stephen J Wright. Coordinate descent algorithms. *Mathematical Programming*, 151(1):3–34, 2015.
- 2094 Zhirong Wu, Shuran Song, Aditya Khosla, Fisher Yu, Linguang Zhang, Xiaoou Tang, and Jianxiong  
2095 Xiao. 3d shapenets: A deep representation for volumetric shapes. In *Proceedings of the IEEE*  
2096 *conference on computer vision and pattern recognition*, pages 1912–1920, 2015.
- 2097 Jing-Hao Xue and D Michael Titterton. Comment on “AJOn discriminative vs. generative clas-  
2098 sifiers: A comparison of logistic regression and naive Bayes”. *Neural processing letters*, 28(3):  
2099 169–187, 2008.

- 2100 Arnold Zellner. Optimal information processing and Bayes's theorem. *The American Statistician*, 42  
2101 (4):278–280, 1988.
- 2102 Michael M Zhang and Fernando Perez-Cruz. Accelerated inference for latent variable models. *arXiv*  
2103 *preprint arXiv:1705.07178*, 2017.
- 2104 Wei Zhang and Jana Křsecká. Nonparametric estimation of multiple structures with outliers. In  
2105 *Dynamical Vision*, pages 60–74. Springer, 2007.
- 2106 Mingyuan Zhou, Hongxia Yang, Guillermo Sapiro, David B Dunson, and Lawrence Carin. Dependent  
2107 hierarchical beta process for image interpolation and denoising. In *International conference on*  
2108 *artificial intelligence and statistics*, pages 883–891, 2011.





## PROBABILISTIC CONCEPTS

**TODO:** What follows down here is old and has to be adapted to the measure-theoretic realizations.

Let us introduce the notation of expectation for random variable  $X$ :

$$E_p[X] = \sum_i^k p(X = x_i)x_i \quad (\text{A.1})$$

For the continuous case, if  $f_X(x)$  is a properly defined probability density function, the expected value becomes:

$$E_f[X] = \int_{-\infty}^{\infty} x f_X(x) dx \quad (\text{A.2})$$

From the context it can be seen that the object  $X$  at the left is a different object than  $x$  at the right. The former is random variable (a one-dimensional function, an (in)finite vector), the latter is a value of a random variable (a scalar). The term  $dx$  is a measure, in this case it assigns volume to *subsets of random variables values*. A proper notation would incorporate this aspect, but not much will be gained by creating such complex notations.

Let us also introduce the conditional probability:

$$p(X|Y) = \frac{p(X, Y)}{p(Y)} \quad (\text{A.3})$$

Suppose  $X$  is a discrete random variable, then the object  $p(X)$  is a vector of finite size  $k$ ,  $p(Y)$  is a vector of finite size  $l$ , and  $p(X|Y)$  as well as  $p(X, Y)$  are matrices of size  $k \times l$ . A conditional probability hence ‘divides’ a matrix by a vector. It is a proper measure if  $p(Y) \neq 0$  (for all values of - or events in -  $Y$ ).

2125 In for example importance sampling (Sect. ??), the expectation is taken over a function of  
 2126 a random variable. A function, if measurable, can be taken the expectation over using the  
 2127 so-called law of the unconscious statistician:

$$E_f[g(X)] = \int_{-\infty}^{\infty} g(x)f_X(x)dx \quad (\text{A.4})$$

2128 Here  $g(X)$  is a general measurable function, and not restricted to a probability density func-  
 2129 tion.

2130 The notation above is an indefinite integral. We can approach this integral by Monte Carlo  
 2131 integration (Sect. ??).

$$E_f[g(X)] = \frac{1}{k} \sum_{i=1}^k g(x_i) \quad \text{with} \quad x_i \sim f_X(x) \quad (\text{A.5})$$

2132 Rather than summing over  $f_X(x_i)$ , we now sample  $x_i \sim f_X(x)$ .

## 2133 A.1 Common Inequalities

### 2134 A.1.1 Markov's Inequality

2135 Markov's inequality comes up with an upper bound for the probability that an non-negative  
 2136 random variable  $X$  exceeds some constant positive threshold  $a$ .

$$p(X \geq a) \leq \frac{E[X]}{a} \quad (\text{A.6})$$

2137 The proof in classical probability theory uses an indicator variable:

$$aI(X \geq a) \leq X \quad (\text{A.7})$$

2138 Here  $I(X \geq a) = 1$  if the event  $X \geq a$  occurs, setting the left-hand side to  $a$  (which is of  
 2139 course smaller than  $X$ ). And  $I(X \geq a) = 0$  on the event  $X < a$ , which is naturally smaller  
 2140 than the non-negative  $X$ .

2141 Expectations obey the inequality: if  $(X \leq Y)$ , then  $E[X] \leq E[Y]$ , hence:

$$E[aI(X \geq a)] \leq E[X] \quad (\text{A.8})$$

2142 And because expectations add up linearly:

$$E[aI(X \geq a)] = aE[I(X \geq a)] = a(1 \cdot p(X \geq a) + 0 \cdot p(X < a)) = a \cdot p(X \geq a) \quad (\text{A.9})$$

2143 So, we have Markov's inequality combining Eq. A.8 and A.9:

$$a \cdot p(X \geq a) \leq E[X] \quad (\text{A.10})$$

### 2144 A.1.2 Chebyshev's Inequality

2145 Now, Chebyshev's inequality defines in a similar way (Chebyshev was a teacher of Markov<sup>1</sup>)  
 2146 an upper bound on the deviation from the mean for a random variable. Recall the definition  
 2147 of the variance of  $X$  and assume it is finite:

$$\text{Var}(X) = E[(X - E[X])^2] = \sigma^2 \quad (\text{A.11})$$

2148 Consider now the random variable  $(X - E[X])^2$  and constant  $a = (\sigma k)^2$  and write down  
 2149 Markov's inequality:

$$p((X - E[X])^2 \geq (\sigma k)^2) \leq \frac{E[(X - E[X])^2]}{(\sigma k)^2} \quad (\text{A.12})$$

2150 Taking the square root of the inequality at the left, and using the definition of  $\sigma^2$  at the  
 2151 right, leads to:

$$p(|X - E[X]| \geq \sigma k) \leq \frac{1}{k^2} \quad (\text{A.13})$$

### 2152 A.1.3 Weak Law of Large Numbers

2153 Chebyshev's inequality can be used to prove the weak law of large numbers. Given that we  
 2154 have a series of random variables, all with the same finite expectation,  $E[X_i] = \mu$ , then this  
 2155 law states that the sample average  $\bar{X} = \frac{1}{n}(X_1 + \dots + X_n)$  converges in probability towards  
 2156 the expected value:

$$\bar{X} \xrightarrow{P} \mu \quad \text{for} \quad n \rightarrow \infty \quad (\text{A.14})$$

2157 We can use the independence assumption between variables  $X_i$  to write down the variance  
 2158 and expectation of  $\bar{X}$ :

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{n}(X_1 + \dots + X_n)\right) = \frac{1}{n^2} \text{Var}(X_1 + \dots + X_n) = \frac{\sigma^2}{n} E[\bar{X}] = \mu \quad (\text{A.15})$$

---

<sup>1</sup>There were many mathematically gifted Markov's. This is Andrey Andreyevich Markov Sr., known from the Markov chains and Markov processes. Jr. is known from Markov's principle, Markov's rule and the Markov algorithm.



2159 And now we can apply Chebyshev's inequality on  $\bar{X}$ :

$$p(|\bar{X} - \mu| \geq \epsilon) \leq \frac{\sigma^2}{n\epsilon^2} \quad (\text{A.16})$$

2160 Convergence in probability towards  $X$  is the case if for all  $\epsilon$ :

$$\lim_{n \rightarrow \infty} p(|\bar{X} - X| \geq \epsilon) = 0 \quad (\text{A.17})$$

2161 This is the case for  $n \rightarrow \infty$  indeed.

#### 2162 A.1.4 Strong Law of Large Numbers

2163 The strong law incorporates the weak law. Rather than convergence *in probability*, it states  
2164 convergence *almost surely* towards the expected value.

$$\bar{X} \xrightarrow{a.s.} \mu \quad \text{for} \quad n \rightarrow \infty \quad (\text{A.18})$$

2165 The strong law states that with probability 1, for any  $\epsilon > 0$ , the inequality  $|X - \mu| < \epsilon$   
2166 holds for large enough  $n$ . The weak law states only that the average  $\bar{X}$  is likely near  $\mu$ , but  
2167  $|X - \mu| \geq \epsilon$  can still happen, even for large  $n$ .

#### 2168 A.1.5 Common Distributions

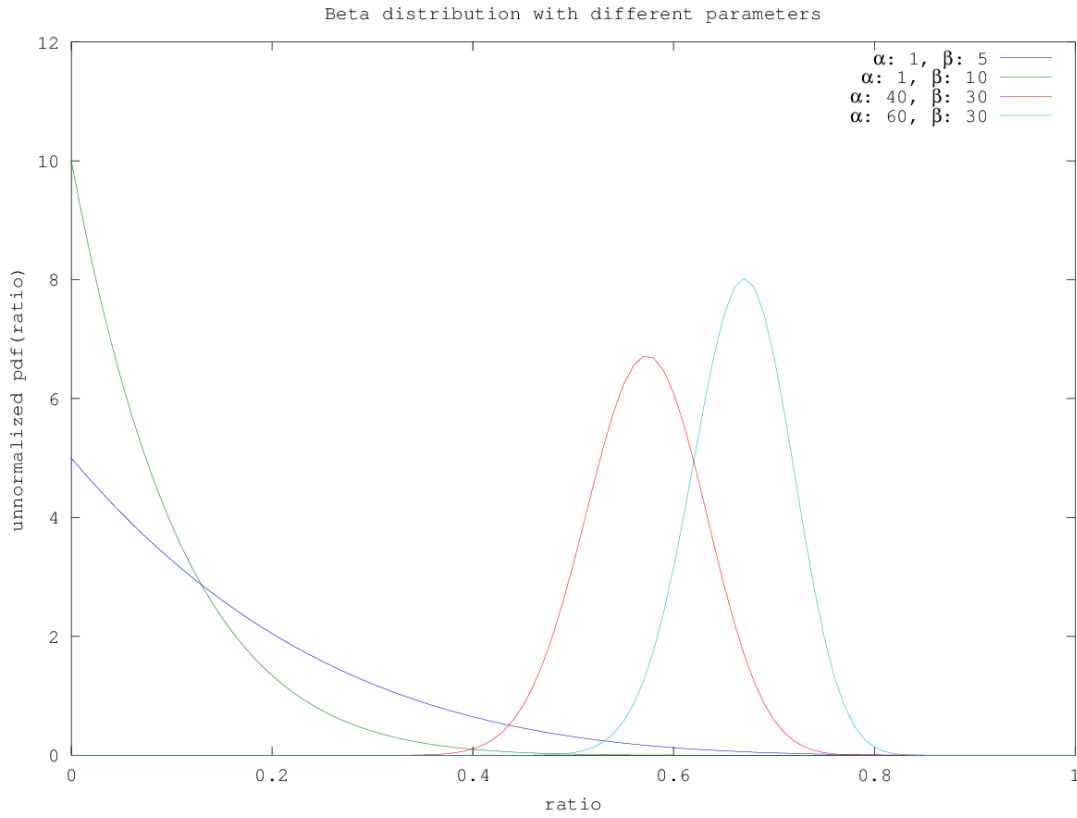
2169 One of the probability distributions that is interesting to us is the beta-distribution. A normal  
2170 distribution might be a reasonable prior for a continuous variable such as human heights in  
2171 a population. If this variable however is itself a probability, a reasonable prior is the beta-  
2172 distribution. The beta-distribution can be described as:

$$f(x, \alpha, \beta) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1} = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1} \quad (\text{A.19})$$

2173 Here  $B$  is the beta function and  $\Gamma$  is the gamma function, the continuous extension of the  
2174 factorial function:  $\Gamma(n) = (n-1)!$ . Naturally, there are many of such extensions. The  
2175 gamma function extends the factorial in a specific sense. It obeys the recurrence relation  
2176  $f(x+1) = xf(x)$  with  $f(1) = 1$ . Its description is defined with an improper integral:

$$\Gamma(t) = \int_0^\infty x^{t-1} e^{-x} dx \quad (\text{A.20})$$

2177 The expected value of a random variable  $X$  with a beta-distribution:



**Figure A.1:** The beta-distribution with different parameters. The x-axis is the quantity modelled, for example a ratio between wins and losses in a soccer season. The y-axis is the corresponding unnormalized density function. Setting  $\alpha = 1$  shows a monotonically decreasing density function. Having a variable  $\alpha$  allows probability mass to shift to the end.

$$E_f[X] = \int_0^1 x f(x, \alpha, \beta) dx = \int_0^1 \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^\alpha (1-x)^{\beta-1} dx = \frac{\alpha}{\alpha + \beta} \quad (\text{A.21})$$

2178 Let us illustrate the effects of the parameters of the beta-distribution.

2179 The stick-breaking presentation shows that a Dirichlet process consists of beta processes with  
 2180  $\alpha = 1$ . The Pitman-Yor process has  $\alpha$  left variable. Informally, the location where the stick  
 2181 will be broken (iteratively) for the Dirichlet process is ‘quite close to the beginning’ with high  
 2182 probability. In the case  $\alpha$  becomes larger, the breaking can also occur likely at the end of the  
 2183 stick. This means for breaking a stick, say 20 times, the distribution of stick lengths for the  
 2184 Pitman-Yor process has a much larger support. The difference between the large and small  
 2185 sticks is much larger.

2186

2187

## DIRICHLET-MULTINOMIAL

2188 There is confusion in the literature with respect to the Dirichlet-multinomial. There are three  
2189 ways to interpret this.

### 2190 B.1 Categorical Distribution

2191 Often, the Dirichlet-multinomial is actually not a compound Dirichlet and a multinomial,  
2192 but a compound Dirichlet and **categorical** distribution:

$$p(z|\theta) = \prod_i \theta_i^{z_i} \quad (\text{B.1})$$

2193 This means that this is about only one categorical variable, not a set. The notation of above  
2194 would for dice assign the vector  $[1, 0, 0, 0, 0, 0]$  to the face with one pip,  $[0, 1, 0, 0, 0, 0]$  to  
2195 the face with two pips, etc. Naturally, this means that  $\sum_i z_i = 1$ .

2196 This gets rid off the  $\frac{n!}{\prod_i z_i!}$  factor and leads to the much shorter:

$$p(z|\theta)p(\theta|\alpha) = \frac{1}{B(\alpha)} \prod_i \theta_i^{z_i + \alpha_i - 1} \quad (\text{B.2})$$

2197 To subsequently derive at the Dirichlet-multinomial, you'll have to integrate over:

$$\int p(z|\theta)p(\theta|\alpha)d\theta = \frac{1}{B(\alpha)} \int \prod_i \theta_i^{z_i + \alpha_i - 1} d\theta \quad (\text{B.3})$$

2198 Now, the Dirichlet didn't come from nowhere... The factor  $B(\alpha)$  is a normalization factor:

$$p(\theta|\alpha) = \frac{1}{B(\alpha)} \prod_i \theta_i^{\alpha_i-1} = \frac{\prod_i \theta_i^{\alpha_i-1}}{\int_{\Delta^n} \prod_i \theta_i^{\alpha_i-1} d\theta} \quad (\text{B.4})$$

2199 with  $\int_{\Delta^n}$  corresponding to the condition  $\sum_i \theta_i = 1$ .

2200 In other words, the multivariate Beta function is actually this integral directly from the def-  
2201 inition:

$$\int_{\Delta^n} \prod_i \theta_i^{\alpha_i-1} d\theta = B(\alpha) \quad (\text{B.5})$$

2202 And hence the integral:

$$\int \prod_i \theta_i^{z_i+\alpha_i-1} d\theta = B(\alpha + z) \quad (\text{B.6})$$

2203 Hence:

$$\int p(z|\theta)p(\theta|\alpha)d\theta = \frac{B(\alpha + z)}{B(\alpha)} \quad (\text{B.7})$$

2204 Or to end up with something commonly stated as the Dirichlet-multinomial:

$$\int p(z|\theta)p(\theta|\alpha)d\theta = \frac{\prod_i \Gamma(\alpha_i + z_i)}{\Gamma(\sum_i (\alpha_i + z_i))} \frac{\Gamma(\sum_i \alpha_i)}{\prod_i \Gamma(\alpha_i)} \quad (\text{B.8})$$

2205 Collecting terms:

$$\int p(z|\theta)p(\theta|\alpha)d\theta = \frac{\Gamma(\sum_i \alpha_i)}{\Gamma(\sum_i \alpha_i + \sum_i z_i)} \prod_i \frac{\Gamma(\alpha_i + z_i)}{\Gamma(\alpha_i)} \quad (\text{B.9})$$

2206 Note, however that we run  $i$  here over the entries in our categorical variable  $z$  represented  
2207 as a vector! This is very different from a multinomial distribution over a set of variables!

## 2208 B.2 Multinomial Distribution

2209 In case of an actual multinomial distribution, counts of  $z$ , let's write them  $n(z)$  are actually  
2210 the topic of consideration, not  $z$  itself.

$$p(z|\theta) = \frac{(\sum_k n(z_k))!}{\prod_k (n(z_k)!) } \prod_k \theta_k^{n(z_k)} \quad (\text{B.10})$$

2211 We now run over  $k$  unique variables, not over a vectorized categorical variable.

2212 Of course, we can know again multiply with a Dirichlet distribution and the derivation is  
 2213 along the lines as described before. The result:

$$\int p(z|\theta)p(\theta|\alpha)d\theta = \frac{(\sum_k n(z_k))!}{\prod_k (n(z_k)!) } \frac{\Gamma(\sum_k \alpha_k)}{\Gamma(\sum_k \alpha_k + \sum_k n(z_k))} \prod_k \frac{\Gamma(\alpha_k + n(z_k))}{\Gamma(\alpha_k)} \quad (\text{B.11})$$

2214 This might not be pretty, but it is the actual full Dirichlet-multinomial.

## 2215 B.3 N Categorical Distributions

2216 The third option, and this is meant most times is the distribution of a **sequence** of categorical  
 2217 variables. Recall that the multinomial assigns probabilities to the **number** of extracted balls  
 2218 (in an experiment getting n balls out of a bag with k ball types). A sequence of categorical  
 2219 variables assigns a probability to a **sequence** and has a form without the normalization  
 2220 factor:

$$p(z|\theta) = \prod_k \theta_k^{z_k} \quad (\text{B.12})$$

2221 Here  $k$  runs over the categories. We can now follow the derivation as with the single cate-  
 2222 gorical variable.



## GIBBS SAMPLING

2223

2224

2225 Notation:

$$\int dF(x) = F(x) \quad (\text{C.1})$$

2226 A mixture model:

$$L(x) = \int dF(x) \mu(x) \quad (\text{C.2})$$

2227 If we have a particular form of  $F(x)$ , namely it admits a decomposition of a sum of individual  
 2228 values  $x_i$ :

$$F(x) = \sum_i \delta_{x_i} = \delta(x = x_0) + \delta(x = x_1) + \dots \quad (\text{C.3})$$

2229 Then our mixture model can be written as:

$$L(x) = \int dF(x) \mu(x) = \sum_i \mu(x_i) \quad (\text{C.4})$$

2230 Let  $x_0 = 3$ ,  $x_1 = 4$ ,  $\mu(x) = x^2$ , then  $L(x) = 3^2 + 4^2 = 25$ .2231 Walker with  $P = F$ ,  $\mu(x) = N(y|\theta)$ ,  $i = j$  and giving each  $\theta_j$  a weight  $\omega_j$ :

$$f_P(y) = \int dP(\theta) N(y|\theta) \quad P = \sum_j \omega_j \delta_{\theta_j} \quad (\text{C.5})$$

2232 Then:

$$f_{\omega,j}(y) = \sum_j \omega_j N(y|\theta_j) \quad (\text{C.6})$$

2233 The likelihood:

$$L(w | \alpha, \lambda_0) = p(\phi | \alpha) \prod_{i=0}^{N-1} \int p(w_i | \theta_i, \phi) dG_0(\theta_i) \quad (\text{C.7})$$

2234 Here the index runs over all data points  $w_i$ . Each data point corresponds to a line with  
 2235 parameters  $\theta_i$ . Here the parameters  $\theta_i$  and  $\theta_j$  for data point  $w_i$  and  $w_j$  can be the same and  
 2236 thus reflect the same line. The index for  $\theta$  runs over the  $N$  data points, not over the  $K$  lines.

2237 The distribution  $G_0$  does have hyperparameters  $\lambda_0$ .

2238 We can also group all data points that belong to the same line  $k$  together by reordering the  
 2239 product terms:

$$L(w | \alpha, \lambda_0) = p(\phi | \alpha) \prod_k \prod_{i:z_i=k} \int p(w_i | \theta_i, \phi) dG_0(\theta_i) \quad (\text{C.8})$$

2240 Here the factors that belong to line  $k$  are multiplied. The index still runs over the data points.

2241 It is also possible not to limit the second product to only the data points  $i$  that are assigned  
 2242 to line  $k$ .

$$L(w | \alpha, \lambda_0) = p(\phi | \alpha) \prod_k \prod_i p(z_i | \phi) \int p(w_i | \theta_i, \phi) dG_0(\theta_i) \quad (\text{C.9})$$

2243 Now, we are gonna introduce the stick-breaking sum, which turns our integral into a discrete  
 2244 sum.

2245  $G = \sum_l p_l \delta_{Z_l}$  with  $Z_l$  iid from  $G_0$  and  $p_l$  defined as a product of beta distributions.

$$L(w | \alpha, \lambda_0) = p(\phi | \alpha) \sum_k \prod_i p(z_i | \phi) p(w_i | \theta_k) p(\theta_k | \lambda_0) \quad (\text{C.10})$$

2246 The first term at the right hand side,  $p(\phi | \alpha)$ , generates the partition  $\phi$  by the Dirichlet  
 2247 Process with concentration parameter  $\alpha$ . The second term  $p(z_i | \phi)$  defines indices  $z_0, \dots, z_N$   
 2248 to link observations  $w_0, \dots, w_N$  with the parameters  $\theta_0, \dots, \theta_K$ . The probability  $p(w_i | \theta_k)$   
 2249 corresponds to the likelihood equations 3.10 and 3.11 with  $w_i$  the tuple of  $x_i$  and  $y_i$  and  
 2250  $\theta_k$  the line parameters  $\sigma_k^2$  and  $\beta_k$ . The probability  $p(\theta_k | \lambda_0)$  corresponds to the prior from  
 2251 equation 3.14. The parameters  $\theta_k$  (that is,  $\sigma_k^2$  and  $\beta_k$ ) are generated from hyperparameters  
 2252  $\lambda_0$ . The hyperparameters  $\lambda_0 = \{\mu_0, \Lambda_0, a, b\}$  are the parameters from the Normal-Inverse-  
 2253 Gamma prior.

2254 The Dirichlet process can be used as a mixture model (Antoniak, 1974; Escobar and West,  
 2255 1995; MacEachern and Müller, 1998) in which it generates (non-unique) parameters that  
 2256 subsequently generate observations:

$$\begin{aligned} G &\sim DP(\alpha, G_0) \\ \theta_i &| G \sim G \\ w_i &| \theta_i \sim F(\theta_i) \end{aligned} \tag{C.11}$$

2257 Here  $F$  describes the mapping from parameters  $\theta_i$  to observations  $w_i$ . It is possible to inte-  
 2258 grate over  $G$  and sample the parameters directly from the base distribution  $G_0$ .

2259 It is possible to integrate over  $G$  and get a description in the form of conditionals over the  
 2260 parameters (Blackwell and MacQueen, 1973):

$$\theta_{n+1} | \theta_1 \dots \theta_{n-1} \sim \frac{1}{\alpha + n} (\alpha G_0 + \sum_{i=1}^n \delta_{\theta_i}) \tag{C.12}$$

## 2261 C.1 Gibbs Sampling of Parameters

2262 Algorithm, we will draw  $\theta_i | \theta_{-i}, y_i$  for all  $i$ .

2263 And that continuously. So, that's how we get theta.

2264 Gibbs sampling requires the conditional probabilities of all entities involved (Geman and  
 2265 Geman, 1984). Gibbs sampling just as other Markov chain Monte Carlo methods generates  
 2266 a sequence of correlated samples. Subsequently, if necessary, the Maximum A Posteriori  
 2267 estimation of a value can be found through picking the mode (most common occurring  
 2268 value) of a parameter.

2269 The derivation of the conditional probabilities of parameters with respect to the remain-  
 2270 ing parameters has been described in the literature (Neal, 2000). Such a derivation uses  
 2271 an important property of the Dirichlet process, namely that it is the conjugate prior of the  
 2272 multinomial distribution. Thanks to conjugacy the following equations have closed-form  
 2273 descriptions. The conditional probabilities are sampled from the base distribution  $G_0$  and  
 2274 the other parameters  $\theta_i$  in the following way:

$$\theta_{n+1} | \theta_1 \dots \theta_{n-1} \sim \frac{1}{\alpha + n} (\alpha G_0 + \sum_{i=1}^n \delta_{\theta_i}) \tag{C.13}$$

2275 If we include the observations themselves, we need to include the likelihood as well:

$$\theta_i | \theta_{-i}, w_i \sim C \left\{ \sum_{j, j \neq i} F(w_i, \theta_j) \delta_{\theta_j} + \alpha H_i \int F(w_i, \theta) dG_0(\theta) \right\} \tag{C.14}$$



2276 The constant  $C$  is a normalization factor to make the above a proper probability density  
2277 (summing to one). The entity  $H_i$  is the posterior density of  $\theta$  given  $G_0$  as prior and  $y_i$  as  
2278 observation. The notation  $\theta_{-i}$  describes the set of all parameters  $\Theta$  with  $\theta_i$  excluded. The  
2279 integral over  $dG_0(\theta)$  is a Lebesgue-Stieltjes integral that weighs the contribution of  $F(w_i, \theta)$   
2280 with the base distribution  $G_0(\theta)$ .

2281 Equation C.14 can be used to perform inference directly with all (non-unique) parameters  
2282  $\theta_i$  tied to observations  $w_i$ . Details on inference will be provided in Sect. ??.





## GLOSSARY

2283

2284

2285

2286 **burn-in** running a Markov chain Monte Carlo for a while before starting to sample from it,  
2287 so the results are not depending on its initial random starting position. 37

2288 **collapsed Gibbs sampling** Rao-Blackwellized Gibbs sampling. Certain sampling steps are  
2289 replaced by steps where one or more variables are integrated out. This is can be thanks  
2290 to analytic descriptions that arise from the use of conjugate priors. 37

## LIST OF FIGURES

2292	2.1	Probability measure . . . . .	12
2293	2.2	Generative vs Discriminative . . . . .	21
2294	2.3	Plate notation . . . . .	24
2295	2.4	Stick-breaking presentation . . . . .	26
2296	2.5	Beta Process . . . . .	27
2297	2.6	Gamma Process . . . . .	29
2298	2.7	Chinese Restaurant Process . . . . .	30
2299	2.8	Dirichlet Process . . . . .	30
2300	2.9	The difference visualized between a Dirichlet Process mixture and a Hierarchical Dirichlet Process. It illustrates also that the input of a Dirichlet process does not have to be a continuous function. If it is a continuous distribution it will become a discrete distributed almost surely. If it is a discrete distribution, it will have atoms at the locations where the discrete distribution had its probability mass concentrated. . . . .	33
2306	2.10	Rejection Sampling . . . . .	35
2307	3.1	The Infinite Line Model in the Chinese Restaurant Process representation (compare with Fig. 2.8). Top: $\alpha$ , the concentration parameter of the Dirichlet Process. Bottom, left to right: $w_i$ , the observation, an individual point in a 2D space; $\theta_i$ , the parameters (intercept, slope) of the line belonging to observation $w_i$ ; $H$ , the base distribution from which line parameter values are sampled. . . . .	43
2313	3.2	The Infinite Line Model in the stick-breaking representation (compare with Fig. 2.8). From left to right: $\alpha$ , the concentration parameter of the Dirichlet Process; $(\phi_1, \dots, \phi_k)$ , the partition of points over lines; $z_i$ , the assignment parameters that link observation $w_i$ with line $k$ ; $w_i$ , the observation, an individual point with $x$ and $y$ coordinates; $\theta_k$ , the parameters of line $k$ ; $\lambda_0$ , the base measure from which the line parameter values are sampled. . . . .	49
2319	3.3	The performance of Algorithm 1 with respect to clustering is measured using the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert metric. A figure of 1 means perfect clustering for all metrics, except Mirvin's where 0 denotes perfect clustering. . . . .	52
2323	3.4	The performance of Algorithm 2 with respect to clustering is measured using the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert metric. A figure of 1 means perfect clustering for all metrics, except Mirvin's where 0 denotes perfect clustering. . . . .	53

2327	3.5	One of the Gibbs steps in the inference of two particular lines. The points are more or less distributed according to the lines, but one line exists out of two large clusters. The line coordinates are visualized by a double circle. The x-coordinate is the y-intercept of the line, the y-coordinate is the slope. . . . .	54
2328			
2329			
2330			
2331	3.6	The assignment of a line to a single point. There are three clusters found, rather than only the obvious two. . . . .	55
2332			
2333	4.1	Sampling of Pareto pairs. The parameters are $\lambda_m = 2$ , $\lambda_n = -2$ , $k = 5$ , and we have sampled $N = 1000$ pairs. . . . .	58
2334			
2335	4.2	Consider the data uniformly distributed on a line segment and a symmetric Pareto prior for both the endpoints, then we can update the estimate for the endpoints given the data as visualized. Each subfigure shows an adjustment of the endpoints given more data points (1, 3, 10, and 100 data points). . .	59
2336			
2337			
2338			
2339	4.3	The Bayesian linear regression model for multiple line segments in plate notation is the same as for the Infinite Line Model. The Dirichlet process is defined at the left with concentration parameter $\alpha$ . It generates the partitions $(\pi_1, \dots, \pi_k)$ with assignment parameters $z_i$ that denote which observation $w_i$ belongs to which cluster $k$ . The cluster is summarized through the parameter set $\theta_k$ and has $\lambda_0$ as its hyperparameter. The parameter set $\theta_k$ includes parameters that signify the line itself such as slope and y-intercept, plus the parameters that denote the extend of the segment. . . . .	59
2340			
2341			
2342			
2343			
2344			
2345			
2346			
2347	4.4	Line segments generated through a Dirichlet Process. The Dirichlet Process itself is again the same. But now four parameters are generated. A normal-inverse-Wishart distribution is used to generate the center of the line segment, and an inverse-Wishart distribution to generate one of the endpoints of the line segments (the other end point is mirrored through its center). Points are generated normally over the line segments, with an additional Gaussian component to indicate the deviation from the line segment from the normal-inverse-Wishart. . . . .	60
2348			
2349			
2350			
2351			
2352			
2353			
2354			
2355	4.5	Line segments generated through a Dirichlet Process. Compared to Fig. 4.4 the points are generated uniformly over the line segments: points are not generated outside of the line segments. . . . .	61
2356			
2357			
2358	4.6	Bayesian point estimates of the sampling process with varying outcomes. . .	63
2359	4.7	Segment detection performs much worse than line detection across all three clustering performance indicators. Perfect clustering is indicated by 1.0 for Rand Index, Adjusted Rand Index, and Hubert. . . . .	64
2360			
2361			
2362	5.1	Dyadic vs triadic MCMC . . . . .	71
2363	5.2	Two examples of fitting a mixture of lines to data items scattered over a two-dimensional space. The lines drawn are inferred using one of the methods in this paper. The lines are not the ground truth, but are meant to demonstrate the typical errors made by fitting methods. Note for example that there are mistakes in both the assignment of points to lines as well as the line parameters (slope and intercept). . . . .	75
2364			
2365			
2366			
2367			
2368			
2369	5.3	The same results as in Table 5.1, but visualized in a violin plot. The distribution over metric values are displayed in a vertical fashion. From left to right the distribution shifts to one, signifying better clustering performance. . . . .	76
2370			
2371			
2372	6.1	Generative Adversarial Network . . . . .	81
2373	6.2	Varational Autoencoder . . . . .	81

2374	7.1	The likelihood function for the moments at which people decide to exercise during the day. On the horizontal axis time, on the vertical axis the frequency of exercising. . . . .	87
2375			
2376			
2377	7.2	The improved likelihood function for the moments at which people decide to exercise during the day using the Von Mises distribution. On the horizontal axis time, on the vertical axis the frequency of exercising. . . . .	88
2378			
2379			
2380	7.3	After running the algorithm, the above figures show the different categories of runners that have been found. . . . .	91
2381			
2382	A.1	Beta distribution . . . . .	108

LIST OF TABLES

2384 2.1 Structures and Processes . . . . . 23

2385 2.2 Levy measure . . . . . 25

2386 2.3 Exchangeable structures . . . . . 26

2387 5.1 The purity, rand index, and adjusted rand index establishing the quality of

2388 the clustering method. The closer the values to one, the better the method

2389 performed. The purity metric assigns high values to clusters that do not have

2390 data points from other clusters (but does not penalize the number of clus-

2391 ters). The rand index index computes similarity between clusters taking false

2392 negatives and false positives into account. The adjusted rand index accounts

2393 for chance. The adjusted rand index is most useful in our comparison. . . . . 76

2394 7.1 Example of the type of data about the timing of exercising. A person is repre-

2395 sented by row, her preferences by column. There is not a predefined number

2396 of users or groups of users. . . . . 86

2397 7.2 Top row: A sequence of cluster indices indicates the ground truth. Each clus-

2398 ter index represents a multi-modal Von-Mises-Uniform distribution with dif-

2399 ferent parameters  $\theta$ . Bottom row: A sequence of cluster indices that are the

2400 result of the described algorithm. Each cluster index represents again a multi-

2401 modal Von-Mises distribution. Errors would be represented by an inconsistent

2402 mapping from the top row to the bottom row. . . . . 90

## LIST OF ABBREVIATIONS

<b>AIC</b>	<b>A</b> kaike <b>I</b> nformation <b>C</b> riterion
<b>BIC</b>	<b>B</b> ayesian <b>I</b> nformation <b>C</b> riterion
<b>DP</b>	<b>D</b> irichlet <b>P</b> rocess





24 04

## SUMMARY

24 05 Summary...



2406

## SAMENVATTING

2407 Samenvatting...



2408

## ACKNOWLEDGMENTS

2409 I want to thank...

2410

## CURRICULUM VITAE

2411 Anne van Rossum ....

2412

## PUBLICATIONS

2413 The investigations performed during my Ph.D. research resulted in the following publica-  
2414 tions.

2415     ◦ A.C. van Rossum. ....



# SIKS DISSERTATION SERIES

2417	<b>1998</b>	2448	<b>2000</b>
2418	1 Johan van den Akker (CWI <sup>1</sup> ) <i>DEGAS - An Active,</i>	2449	1 Frank Niessink (VU) <i>Perspectives on Improving Soft-</i>
2419	<i>Temporal Database of Autonomous Objects</i>	2450	<i>ware Maintenance</i>
2420	2 Floris Wiesman (UM) <i>Information Retrieval by</i>	2451	2 Koen Holtman (TU/e) <i>Prototyping of CMS Storage</i>
2421	<i>Graphically Browsing Meta-Information</i>	2452	<i>Management</i>
2422	3 Ans Steuten (TUD) <i>A Contribution to the Linguistic</i>	2453	3 Carolien M.T. Metselaar (UvA) <i>Sociaal-</i>
2423	<i>Analysis of Business Conversations within the Lan-</i>	2454	<i>organisatorische Gevolgen van Kennistechnologie;</i>
2424	<i>guage/Action Perspective</i>	2455	<i>een Procesbenadering en Actorperspectief</i>
2425	4 Dennis Breuker (UM) <i>Memory versus Search in</i>	2456	4 Geert de Haan (VU) <i>ETAG, A Formal Model of Com-</i>
2426	<i>Games</i>	2457	<i>petence Knowledge for User Interface Design</i>
2427	5 Eduard W. Oskamp (RUL) <i>Computerondersteuning</i>	2458	5 Ruud van der Pol (UM) <i>Knowledge-Based Query</i>
2428	<i>bij Straftoemeting</i>	2459	<i>Formulation in Information Retrieval</i>
2429	<b>1999</b>	2460	6 Rogier van Eijk (UU) <i>Programming Languages for</i>
		2461	<i>Agent Communication</i>
		2462	7 Niels Peek (UU) <i>Decision-Theoretic Planning of Clin-</i>
2430	1 Mark Sloof (VU) <i>Physiology of Quality Change Mod-</i>	2463	<i>ical Patient Management</i>
2431	<i>elling; Automated Modelling of Quality Change of</i>	2464	8 Veerle Coupé (EUR) <i>Sensitivity Analysis of Decision-</i>
2432	<i>Agricultural Products</i>	2465	<i>Theoretic Networks</i>
2433	2 Rob Potharst (EUR) <i>Classification using Decision</i>	2466	9 Florian Waas (CWI) <i>Principles of Probabilistic</i>
2434	<i>Trees and Neural Nets</i>	2467	<i>Query Optimization</i>
2435	3 Don Beal (UM) <i>The Nature of Minimax Search</i>	2468	10 Niels Nes (CWI) <i>Image Database Management Sys-</i>
2436	4 Jacques Penders (UM) <i>The Practical Art of Moving</i>	2469	<i>tem Design Considerations, Algorithms and Architec-</i>
2437	<i>Physical Objects</i>	2470	<i>ture</i>
2438	5 Aldo de Moor (KUB) <i>Empowering Communities: A</i>	2471	11 Jonas Karlsson (CWI) <i>Scalable Distributed Data</i>
2439	<i>Method for the Legitimate User-Driven Specification</i>	2472	<i>Structures for Database Management</i>
2440	<i>of Network Information Systems</i>		
2441	6 Niek J.E. Wijngaards (VU) <i>Re-Design of Composi-</i>	2473	<b>2001</b>
2442	<i>tional Systems</i>		
2443	7 David Spelt (UT) <i>Verification Support for Object</i>	2474	1 Silja Renooij (UU) <i>Qualitative Approaches to Quan-</i>
2444	<i>Database Design</i>	2475	<i>tifying Probabilistic Networks</i>
2445	8 Jacques H.J. Lenting (UM) <i>Informed Gambling:</i>	2476	2 Koen Hindriks (UU) <i>Agent Programming Lan-</i>
2446	<i>Conception and Analysis of a Multi-Agent Mecha-</i>	2477	<i>guages: Programming with Mental Models</i>
2447	<i>nism for Discrete Reallocation</i>		

<sup>1</sup>Abbreviations: SIKS - Dutch Research School for Information and Knowledge Systems; CWI - Centrum voor Wiskunde en Informatica, Amsterdam; EUR - Erasmus Universiteit, Rotterdam; KUB - Katholieke Universiteit Brabant, Tilburg; KUN - Katholieke Universiteit Nijmegen; OU - Open Universiteit; RUL - Rijksuniversiteit Leiden; RUN - Radboud Universiteit Nijmegen; TUD - Technische Universiteit Delft; TU/e - Technische Universiteit Eindhoven; UL - Universiteit Leiden; UM - Universiteit Maastricht; UT - Universiteit Twente, Enschede; UU - Universiteit Utrecht; UvA - Universiteit van Amsterdam; UvT - Universiteit van Tilburg; VU - Vrije Universiteit, Amsterdam.

- 2478 3 Maarten van Someren (UvA) *Learning as Problem Solving* 2534  
2479 2535
- 2480 4 Evgueni Smirnov (UM) *Conjunctive and Disjunctive Version Spaces with Instance-Based Boundary Sets* 2536  
2481 2537
- 2482 5 Jacco van Ossenbruggen (VU) *Processing Structured Hypermedia: A Matter of Style* 2538  
2483 2539
- 2484 6 Martijn van Welie (VU) *Task-Based User Interface Design* 2540  
2485 2540
- 2486 7 Bastiaan Schonhage (VU) *Diva: Architectural Perspectives on Information Visualization* 2541  
2487 2541
- 2488 8 Pascal van Eck (VU) *A Compositional Semantic Structure for Multi-Agent Systems Dynamics* 2542  
2489 2543
- 2490 9 Pieter Jan 't Hoen (RUL) *Towards Distributed Development of Large Object-Oriented Models, Views of Packages as Classes* 2544  
2491 2545
- 2492 10 Maarten Sierhuis (UvA) *Modeling and Simulating Work Practice BRAHMS: a Multiagent Modeling and Simulation Language for Work Practice Analysis and Design* 2546  
2493 2547  
2494 2548  
2495 2549  
2496 2550
- 2497 11 Tom M. van Engers (VU) *Knowledge Management: The Role of Mental Models in Business Systems Design* 2551  
2498 2552  
2499 2553
- 2500 **2002** 2554  
2555  
2556
- 2501 1 Nico Lassing (VU) *Architecture-Level Modifiability Analysis* 2557  
2502 2558
- 2503 2 Roelof van Zwol (UT) *Modelling and Searching Web-based Document Collections* 2559  
2504 2560
- 2505 3 Henk Ernst Blok (UT) *Database Optimization Aspects for Information Retrieval* 2561  
2506 2562
- 2507 4 Juan Roberto Castelo Valdueza (UU) *The Discrete Acyclic Digraph Markov Model in Data Mining* 2563  
2508 2564
- 2509 5 Radu Serban (VU) *The Private Cyberspace Modeling Electronic Environments Inhabited by Privacy-Concerned Agents* 2565  
2510 2566  
2511 2567
- 2512 6 Laurens Mommers (UL) *Applied Legal Epistemology; Building a Knowledge-based Ontology of the Legal Domain* 2568  
2513 2569  
2514 2570
- 2515 7 Peter Boncz (CWI) *Monet: A Next-Generation DBMS Kernel For Query-Intensive Applications* 2571  
2516 2572
- 2517 8 Jaap Gordijn (VU) *Value Based Requirements Engineering: Exploring Innovative E-Commerce Ideas* 2573  
2518 2574
- 2519 9 Willem-Jan van den Heuvel (KUB) *Integrating Modern Business Applications with Objectified Legacy Systems* 2575  
2520 2576  
2521 2577
- 2522 10 Brian Sheppard (UM) *Towards Perfect Play of Scrabble* 2578  
2523 2579
- 2524 11 Wouter C.A. Wijngaards (VU) *Agent Based Modelling of Dynamics: Biological and Organisational Applications* 2580  
2525 2581  
2526 2582
- 2527 12 Albrecht Schmidt (UvA) *Processing XML in Database Systems* 2582  
2528 2583
- 2529 13 Hongjing Wu (TU/e) *A Reference Architecture for Adaptive Hypermedia Applications* 2584  
2530 2585
- 2531 14 Wieke de Vries (UU) *Agent Interaction: Abstract Approaches to Modelling, Programming and Verifying Multi-Agent Systems* 2586  
2532 2587  
2533 2588
- 15 Rik Eshuis (UT) *Semantics and Verification of UML Activity Diagrams for Workflow Modelling* 2589
- 16 Pieter van Langen (VU) *The Anatomy of Design: Foundations, Models and Applications* 2590
- 17 Stefan Manegold (UvA) *Understanding, Modeling, and Improving Main-Memory Database Performance* 2591
- 2003**
- 1 Heiner Stuckenschmidt (VU) *Ontology-Based Information Sharing in Weakly Structured Environments* 2592
- 2 Jan Broersen (VU) *Modal Action Logics for Reasoning About Reactive Systems* 2593
- 3 Martijn Schuemie (TUD) *Human-Computer Interaction and Presence in Virtual Reality Exposure Therapy* 2594
- 4 Milan Petkovic (UT) *Content-Based Video Retrieval Supported by Database Technology* 2595
- 5 Jos Lehmann (UvA) *Causation in Artificial Intelligence and Law – A Modelling Approach* 2596
- 6 Boris van Schooten (UT) *Development and Specification of Virtual Environments* 2597
- 7 Machiel Jansen (UvA) *Formal Explorations of Knowledge Intensive Tasks* 2598
- 8 Yong-Ping Ran (UM) *Repair-Based Scheduling* 2599
- 9 Rens Kortmann (UM) *The Resolution of Visually Guided Behaviour* 2600
- 10 Andreas Lincke (UT) *Electronic Business Negotiation: Some Experimental Studies on the Interaction between Medium, Innovation Context and Cult* 2601
- 11 Simon Keizer (UT) *Reasoning under Uncertainty in Natural Language Dialogue using Bayesian Networks* 2602
- 12 Roeland Ordelman (UT) *Dutch Speech Recognition in Multimedia Information Retrieval* 2603
- 13 Jeroen Donkers (UM) *Nosce Hostem – Searching with Opponent Models* 2604
- 14 Stijn Hoppenbrouwers (KUN) *Freezing Language: Conceptualisation Processes across ICT-Supported Organisations* 2605
- 15 Mathijs de Weerd (TUD) *Plan Merging in Multi-Agent Systems* 2606
- 16 Menzo Windhouwer (CWI) *Feature Grammar Systems - Incremental Maintenance of Indexes to Digital Media Warehouse* 2607
- 17 David Jansen (UT) *Extensions of Statecharts with Probability, Time, and Stochastic Timing* 2608
- 18 Levente Kocsis (UM) *Learning Search Decisions* 2609
- 2004**
- 1 Virginia Dignum (UU) *A Model for Organizational Interaction: Based on Agents, Founded in Logic* 2610
- 2 Lai Xu (UvT) *Monitoring Multi-party Contracts for E-business* 2611
- 3 Perry Groot (VU) *A Theoretical and Empirical Analysis of Approximation in Symbolic Problem Solving* 2612

- 2588 4 Chris van Aart (UvA) *Organizational Principles for* 2642  
2589 *Multi-Agent Architectures* 2643
- 2590 5 Viara Popova (EUR) *Knowledge Discovery and* 2644  
2591 *Monotonicity* 2645
- 2592 6 Bart-Jan Hommes (TUD) *The Evaluation of Busi-* 2646  
2593 *ness Process Modeling Techniques* 2647
- 2594 7 Elise Boltjes (UM) *Voorbeeld<sub>IG</sub> Onderwijs; Voor-* 2648  
2595 *beeldgestuurd Onderwijs, een Opstap naar Abstracte* 2649  
2596 *Denken, vooral voor Meisjes* 2650
- 2597 8 Joop Verbeek (UM) *Politie en de Nieuwe Inter-* 2651  
2598 *nationale Informatiemarkt, Grensregionale Politie* 2652  
2599 *Gegevensuitwisseling en Digitale Expertise* 2653
- 2600 9 Martin Caminada (VU) *For the Sake of the Argu-* 2654  
2601 *ment; Explorations into Argument-based Reasoning* 2655
- 2602 10 Suzanne Kabel (UvA) *Knowledge-rich Indexing of* 2656  
2603 *Learning-objects* 2657
- 2604 11 Michel Klein (VU) *Change Management for Dis-* 2658  
2605 *tributed Ontologies* 2659
- 2606 12 The Duy Bui (UT) *Creating Emotions and Facial Ex-* 2660  
2607 *pressions for Embodied Agents* 2661
- 2608 13 Wojciech Jamroga (UT) *Using Multiple Models of* 2662  
2609 *Reality: On Agents who Know how to Play* 2663
- 2610 14 Paul Harrenstein (UU) *Logic in Conflict. Logical Ex-* 2664  
2611 *plorations in Strategic Equilibrium* 2665
- 2612 15 Arno Knobbe (UU) *Multi-Relational Data Mining* 2666  
2613 16 Federico Divina (VU) *Hybrid Genetic Relational* 2667  
2614 *Search for Inductive Learning*
- 2615 17 Mark Winands (UM) *Informed Search in Complex* 2668  
2616 *Games*
- 2617 18 Vania Bessa Machado (UvA) *Supporting the Con-* 2669  
2618 *struction of Qualitative Knowledge Models* 2670
- 2619 19 Thijs Westerveld (UT) *Using generative probabilis-* 2671  
2620 *tic models for multimedia retrieval* 2672
- 2621 20 Madelon Evers (Nyenrode) *Learning from Design* 2673  
2622 *facilitating multidisciplinary design teams* 2674  
2675
- 2623 **2005** 2676
- 2624 1 Floor Verdenius (UvA) *Methodological Aspects of* 2679  
2625 *Designing Induction-Based Applications* 2680
- 2626 2 Erik van der Werf (UM) *AI techniques for the game* 2681  
2627 *of Go* 2682
- 2628 3 Franc Grootjen (RUN) *A Pragmatic Approach to the* 2683  
2629 *Conceptualisation of Language* 2684
- 2630 4 Nirvana Meratnia (UT) *Towards Database Support* 2685  
2631 *for Moving Object data* 2686
- 2632 5 Gabriel Infante-Lopez (UvA) *Two-Level Probabilis-* 2687  
2633 *tic Grammars for Natural Language Parsing* 2688
- 2634 6 Pieter Spronck (UM) *Adaptive Game AI* 2689
- 2635 7 Flavius Frasinca (TU/e) *Hypermedia Presentation* 2690  
2636 *Generation for Semantic Web Information Systems* 2691
- 2637 8 Richard Vdovjak (TU/e) *A Model-driven Approach* 2692  
2638 *for Building Distributed Ontology-based Web Appli-* 2693  
2639 *cations* 2694
- 2640 9 Jeen Broekstra (VU) *Storage, Querying and Infer-* 2695  
2641 *encing for Semantic Web Languages* 2696  
2697
- 10 Anders Bouwer (UvA) *Explaining Behaviour: Using*  
*Qualitative Simulation in Interactive Learning Envi-*  
*ronments*
- 11 Elth Ogston (VU) *Agent Based Matchmaking and*  
*Clustering - A Decentralized Approach to Search*
- 12 Csaba Boer (EUR) *Distributed Simulation in Indus-*  
*try*
- 13 Fred Hamburg (UL) *Een Computermodel voor het*  
*Ondersteunen van Euthanasiebeslissingen*
- 14 Borys Omelayenko (VU) *Web-Service configuration*  
*on the Semantic Web; Exploring how semantics*  
*meets pragmatics*
- 15 Tibor Bosse (VU) *Analysis of the Dynamics of Cog-*  
*nitive Processes*
- 16 Joris Graaumanns (UU) *Usability of XML Query Lan-*  
*guages*
- 17 Boris Shishkov (TUD) *Software Specification Based*  
*on Re-usable Business Components*
- 18 Danielle Sent (UU) *Test-selection strategies for prob-*  
*abilistic networks*
- 19 Michel van Dartel (UM) *Situated Representation*
- 20 Cristina Coteanu (UL) *Cyber Consumer Law, State*  
*of the Art and Perspectives*
- 21 Wijnand Derks (UT) *Improving Concurrency and*  
*Recovery in Database Systems by Exploiting Appli-*  
*cation Semantics*
- 2006**
- 1 Samuil Angelov (TU/e) *Foundations of B2B Elec-*  
*tronic Contracting*
- 2 Cristina Chisalita (VU) *Contextual issues in the de-*  
*sign and use of information technology in organiza-*  
*tions*
- 3 Noor Christoph (UvA) *The role of metacognitive*  
*skills in learning to solve problems*
- 4 Marta Sabou (VU) *Building Web Service Ontologies*
- 5 Cees Pierik (UU) *Validation Techniques for Object-*  
*Oriented Proof Outlines*
- 6 Ziv Baida (VU) *Software-aided Service Bundling -*  
*Intelligent Methods & Tools for Graphical Service*  
*Modeling*
- 7 Marko Smiljanic (UT) *XML schema matching – bal-*  
*ancing efficiency and effectiveness by means of clus-*  
*tering*
- 8 Eelco Herder (UT) *Forward, Back and Home Again*  
*- Analyzing User Behavior on the Web*
- 9 Mohamed Wahdan (UM) *Automatic Formulation of*  
*the Auditor's Opinion*
- 10 Ronny Siebes (VU) *Semantic Routing in Peer-to-Peer*  
*Systems*
- 11 Joeri van Ruth (UT) *Flattening Queries over Nested*  
*Data Types*
- 12 Bert Bongers (VU) *Interactivation - Towards an e-*  
*cology of people, our technological environment, and*  
*the arts*
- 13 Henk-Jan Lebbink (UU) *Dialogue and Decision*  
*Games for Information Exchanging Agents*

- 2698 14 Johan Hoorn (VU) *Software Requirements: Update, Upgrade, Redesign - towards a Theory of Requirements Change* 2752-2753-2754
- 2701 15 Rainer Malik (UU) *CONAN: Text Mining in the Biomedical Domain* 2755-2756
- 2702 16 Carsten Riggelsen (UU) *Approximation Methods for Efficient Learning of Bayesian Networks* 2757-2758
- 2703 17 Stacey Nagata (UU) *User Assistance for Multitasking with Interruptions on a Mobile Device* 2759-2760
- 2704 18 Valentin Zhizhkun (UvA) *Graph transformation for Natural Language Processing* 2761-2762
- 2705 19 Birna van Riemsdijk (UU) *Cognitive Agent Programming: A Semantic Approach* 2763-2764
- 2706 20 Marina Velikova (UvT) *Monotone models for prediction in data mining* 2765-2766
- 2707 21 Bas van Gils (RUN) *Aptness on the Web* 2767
- 2708 22 Paul de Vrieze (RUN) *Fundamentals of Adaptive Personalisation* 2768-2769
- 2709 23 Ion Juvina (UU) *Development of Cognitive Model for Navigating on the Web* 2770-2771
- 2710 24 Laura Hollink (VU) *Semantic Annotation for Retrieval of Visual Resources* 2772-2773
- 2711 25 Madalina Drugan (UU) *Conditional log-likelihood MDL and Evolutionary MCMC* 2774-2775
- 2712 26 Vojkan Mihajlovic (UT) *Score Region Algebra: A Flexible Framework for Structured Information Retrieval* 2776-2777-2778-2779
- 2713 27 Stefano Bocconi (CWI) *Vox Populi: generating video documentaries from semantically annotated media repositories* 2780-2781
- 2714 28 Borkur Sigurbjornsson (UvA) *Focused Information Access using XML Element Retrieval* 2782-2783-2784-2785
- 2715 2007
- 2716 1 Kees Leune (UvT) *Access Control and Service-Oriented Architectures* 2786
- 2717 2 Wouter Teepe (RUG) *Reconciling Information Exchange and Confidentiality: A Formal Approach* 2787-2788-2789
- 2718 3 Peter Mika (VU) *Social Networks and the Semantic Web* 2790-2791
- 2719 4 Jurriaan van Diggelen (UU) *Achieving Semantic Interoperability in Multi-agent Systems: a dialogue-based approach* 2792-2793-2794
- 2720 5 Bart Schermer (UL) *Software Agents, Surveillance and the Right to Privacy: a Legislative Framework for Agent-enabled Surveillance* 2795-2796-2797
- 2721 6 Gilad Mishne (UvA) *Applied Text Analytics for Blogs* 2798
- 2722 7 Natasa Jovanovic' (UT) *To Whom It May Concern Addressee Identification in Face-to-Face Meetings* 2799-2800
- 2723 8 Mark Hoogendoorn (VU) *Modeling of Change in Multi-Agent Organizations* 2801-2802
- 2724 9 David Mobach (VU) *Agent-Based Mediated Service Negotiation* 2803-2804
- 2725 10 Huib Aldewereld (UU) *Autonomy vs. Conformity: an Institutional Perspective on Norms and Protocols* 2805-2806-2807
- 2726 11 Natalia Stash (TU/e) *Incorporating Cognitive/Learning Styles in a General-Purpose Adaptive Hypermedia System*
- 2727 12 Marcel van Gerven (RUN) *Bayesian Networks for Clinical Decision Support: A Rational Approach to Dynamic Decision-Making under Uncertainty*
- 2728 13 Rutger Rienks (UT) *Meetings in Smart Environments; Implications of Progressing Technology*
- 2729 14 Niek Bergboer (UM) *Context-Based Image Analysis*
- 2730 15 Joyca Lacroix (UM) *NIM: a Situated Computational Memory Model*
- 2731 16 Davide Grossi (UU) *Designing Invisible Handcuffs. Formal investigations in Institutions and Organizations for Multi-agent Systems*
- 2732 17 Theodore Charitos (UU) *Reasoning with Dynamic Networks in Practice*
- 2733 18 Bart Orriens (UvT) *On the development and management of adaptive business collaborations*
- 2734 19 David Levy (UM) *Intimate relationships with artificial partners*
- 2735 20 Slinger Jansen (UU) *Customer Configuration Updating in a Software Supply Network*
- 2736 21 Karianne Vermaas (UU) *Fast diffusion and broadening use: A research on residential adoption and usage of broadband internet in the Netherlands between 2001 and 2005*
- 2737 22 Zlatko Zlatev (UT) *Goal-oriented design of value and process models from patterns*
- 2738 23 Peter Barna (TU/e) *Specification of Application Logic in Web Information Systems*
- 2739 24 Georgina Ramírez Camps (CWI) *Structural Features in XML Retrieval*
- 2740 25 Joost Schalken (VU) *Empirical Investigations in Software Process Improvement*
- 2741 2008
- 2742 1 Katalin Boer-Sorbán (EUR) *Agent-Based Simulation of Financial Markets: A modular, continuous-time approach*
- 2743 2 Alexei Sharpanskykh (VU) *On Computer-Aided Methods for Modeling and Analysis of Organizations*
- 2744 3 Vera Hollink (UvA) *Optimizing hierarchical menus: a usage-based approach*
- 2745 4 Ander de Keijzer (UT) *Management of Uncertain Data - towards unattended integration*
- 2746 5 Bela Mutschler (UT) *Modeling and simulating causal dependencies on process-aware information systems from a cost perspective*
- 2747 6 Arjen Hommersom (RUN) *On the Application of Formal Methods to Clinical Guidelines, an Artificial Intelligence Perspective*
- 2748 7 Peter van Rosmalen (OU) *Supporting the tutor in the design and support of adaptive e-learning*
- 2749 8 Janneke Bolt (UU) *Bayesian Networks: Aspects of Approximate Inference*
- 2750 9 Christof van Nimwegen (UU) *The paradox of the guided user: assistance can be counter-effective*

- 2808 10 Wauter Bosma (UT) *Discourse oriented Summariza-*2866  
2809 *tion* 2867
- 2810 11 Vera Kartseva (VU) *Designing Controls for Network*  
2811 *Organizations: a Value-Based Approach* 2868
- 2812 12 Jozsef Farkas (RUN) *A Semiotically oriented Cogni-*  
2813 *tive Model of Knowledge Representation* 2869
- 2814 13 Caterina Carraciolo (UvA) *Topic Driven Access to*2870  
2815 *Scientific Handbooks* 2871
- 2816 14 Arthur van Bunningen (UT) *Context-Aware Query-*2872  
2817 *ing; Better Answers with Less Effort* 2873
- 2818 15 Martijn van Otterlo (UT) *The Logic of Adaptive Be-*2874  
2819 *havior: Knowledge Representation and Algorithms*2875  
2820 *for the Markov Decision Process Framework in First-*2876  
2821 *Order Domains* 2877
- 2822 16 Henriette van Vugt (VU) *Embodied Agents from a*2878  
2823 *User's Perspective* 2879
- 2824 17 Martin Op't Land (TUD) *Applying Architecture and*2880  
2825 *Ontology to the Splitting and Allying of Enterprises*2881
- 2826 18 Guido de Croon (UM) *Adaptive Active Vision* 2882
- 2827 19 Henning Rode (UT) *From document to entity re-*2883  
2828 *trieval: improving precision and performance of fo-*2884  
2829 *cused text search* 2885
- 2830 20 Rex Arendsen (UvA) *Geen bericht, goed bericht. Een*2886  
2831 *onderzoek naar de effecten van de introductie van*2887  
2832 *elektronisch berichtenverkeer met een overheid op de*2888  
2833 *administratieve lasten van bedrijven* 2889
- 2834 21 Krisztian Balog (UvA) *People search in the enter-*2890  
2835 *prise* 2891
- 2836 22 Henk Koning (UU) *Communication of IT*2892  
2837 *architecture* 2893
- 2838 23 Stefan Visscher (UU) *Bayesian network models for*2894  
2839 *the management of ventilator-associated pneumonia*2895
- 2840 24 Zharko Aleksovski (VU) *Using background knowl-*2896  
2841 *edge in ontology matching* 2897
- 2842 25 Geert Jonker (UU) *Efficient and Equitable exchange*2898  
2843 *in air traffic management plan repair using spender-*2899  
2844 *signed currency* 2900
- 2845 26 Marijn Huijbregts (UT) *Segmentation, diarization*2901  
2846 *and speech transcription: surprise data unraveled* 2902
- 2847 27 Hubert Vogten (OU) *Design and implementation*2903  
2848 *strategies for IMS learning design* 2904
- 2849 28 Ildiko Flesh (RUN) *On the use of independence re-*2905  
2850 *lations in Bayesian networks* 2906
- 2851 29 Dennis Reidsma (UT) *Annotations and subjective*2907  
2852 *machines- Of annotators, embodied agents, users,*2908  
2853 *and other humans* 2909
- 2854 30 Wouter van Atteveldt (VU) *Semantic network anal-*2910  
2855 *ysis: techniques for extracting, representing and*2911  
2856 *querying media content* 2912
- 2857 31 Loes Braun (UM) *Pro-active medical information re-*2913  
2858 *trieval* 2914
- 2859 32 Trung B. Hui (UT) *Toward affective dialogue man-*2915  
2860 *agement using partially observable markov decision*2916  
2861 *processes* 2917
- 2862 33 Frank Terpstra (UvA) *Scientific workflow design,*2918  
2863 *theoretical and practical issues* 2919
- 2864 34 Jeroen de Knijf (UU) *Studies in Frequent Tree Min-*2920  
2865 *ing* 2921
- 35 Benjamin Torben-Nielsen (UvT) *Dendritic mor-*  
*phology: function shapes structure*
- 2009**
- 1 Rasa Jurgelenaite (RUN) *Symmetric Causal Inde-*  
*pendence Models*
- 2 Willem Robert van Hage (VU) *Evaluating Ontology-*  
*Alignment Techniques*
- 3 Hans Stol (UvT) *A Framework for Evidence-based*  
*Policy Making Using IT*
- 4 Josephine Nabukenya (RUN) *Improving the Qual-*  
*ity of Organisational Policy Making using Collabo-*  
*ration Engineering*
- 5 Sietse Overbeek (RUN) *Bridging Supply and De-*  
*mand for Knowledge Intensive Tasks - Based on*  
*Knowledge, Cognition, and Quality*
- 6 Muhammad Subianto (UU) *Understanding Classi-*  
*fication*
- 7 Ronald Poppe (UT) *Discriminative Vision-Based Re-*  
*covery and Recognition of Human Motion*
- 8 Volker Nannen (VU) *Evolutionary Agent-Based Pol-*  
*icy Analysis in Dynamic Environments*
- 9 Benjamin Kanagwa (RUN) *Design, Discovery and*  
*Construction of Service-oriented Systems*
- 10 Jan Wielemaker (UvA) *Logic programming for*  
*knowledge-intensive interactive applications*
- 11 Alexander Boer (UvA) *Legal Theory, Sources of Law*  
*& the Semantic Web*
- 12 Peter Massuthe (TU/e, Humboldt-Universität zu  
Berlin) *Operating Guidelines for Services*
- 13 Steven de Jong (UM) *Fairness in Multi-Agent Sys-*  
*tems*
- 14 Maksym Korotkiy (VU) *From ontology-enabled ser-*  
*vices to service-enabled ontologies (making onto-*  
*logies work in e-science with ONTO-SOA)*
- 15 Rinke Hoekstra (UvA) *Ontology Representation -*  
*Design Patterns and Ontologies that Make Sense*
- 16 Fritz Reul (UvT) *New Architectures in Computer*  
*Chess*
- 17 Laurens van der Maaten (UvT) *Feature Extraction*  
*from Visual Data*
- 18 Fabian Groffen (CWI) *Armada, An Evolving*  
*Database System*
- 19 Valentin Robu (CWI) *Modeling Preferences, Strate-*  
*gic Reasoning and Collaboration in Agent-Mediated*  
*Electronic Markets*
- 20 Bob van der Vecht (UU) *Adjustable Autonomy: Con-*  
*trolling Influences on Decision Making*
- 21 Stijn Vanderlooy (UM) *Ranking and Reliable Clas-*  
*sification*
- 22 Pavel Serdyukov (UT) *Search For Expertise: Going*  
*beyond direct evidence*
- 23 Peter Hofgesang (VU) *Modelling Web Usage in a*  
*Changing Environment*
- 24 Annerieke Heuvelink (VU) *Cognitive Models for*  
*Training Simulations*
- 25 Alex van Ballegooij (CWI) *"RAM: Array Database*  
*Management through Relational Mapping"*

- 2923 26 Fernando Koch (UU) *An Agent-Based Model for the* 2943  
 2924 *Development of Intelligent Mobile Services* 2944
- 2925 27 Christian Glahn (OU) *Contextual Support of social* 2945  
 2926 *Engagement and Reflection on the Web* 2946
- 2927 28 Sander Evers (UT) *Sensor Data Management with* 2947  
 2928 *Probabilistic Models* 2948
- 2929 29 Stanislav Pokraev (UT) *Model-Driven Semantic In-* 2949  
 2930 *tegration of Service-Oriented Applications* 2950
- 2931 30 Marcin Zukowski (CWI) *Balancing vectorized query* 2951  
 2932 *execution with bandwidth-optimized storage* 2952
- 2933 31 Sofiya Katrenko (UvA) *A Closer Look at Learning* 2953  
 2934 *Relations from Text* 2954
- 2935 32 Rik Farenhorst and Remco de Boer (VU) *Architec-* 2955  
 2936 *tural Knowledge Management: Supporting Archi-* 2956  
 2937 *tects and Auditors* 2957
- 2938 33 Khiet Truong (UT) *How Does Real Affect Affect Affect* 2958  
 2939 *Recognition In Speech?* 2959
- 2940 34 Inge van de Weerd (UU) *Advancing in Software* 2960  
 2941 *Product Management: An Incremental Method En-* 2961  
 2942 *gineering Approach*
- 35 Wouter Koelewijn (UL) *Privacy en Politiegegevens;*  
*Over geautomatiseerde normatieve informatie-*  
*uitwisseling*
- 36 Marco Kalz (OUN) *Placement Support for Learners*  
*in Learning Networks*
- 37 Hendrik Drachsler (OUN) *Navigation Support for*  
*Learners in Informal Learning Networks*
- 38 Riina Vuorikari (OU) *Tags and self-organisation: a*  
*metadata ecology for learning resources in a multi-*  
*lingual context*
- 39 Christian Stahl (TUE, Humboldt-Universität zu  
 Berlin) *Service Substitution – A Behavioral Ap-*  
*proach Based on Petri Nets*
- 40 Stephan Raaijmakers (UvT) *Multinomial Language*  
*Learning: Investigations into the Geometry of Lan-*  
*guage*
- 41 Igor Berezhnny (UvT) *Digital Analysis of Paintings*
- 42 Toine Bogers (UvT) *Recommender Systems for So-*  
*cial Bookmarking*