

1 **Nonparametric Bayesian Methods**
2 **in Robotics**

3 PROEFSCHRIFT

4 ter verkrijging van de graad van doctor
5 aan de Universiteit Leiden,
6 op gezag van de rector magnificus,
7 prof. mr. C.J.J.M. Stolker,
8 volgens besluit van het College voor Promoties
9 te verdedigen op ...2018
10 in de aula van de Universiteit
11 klokke ...

12 door

13 **Anne Cornelis van Rossum,**
14 **geboren op 30 december 1980**
15 **te Dirksland**

16 Samenstelling van de promotiecommissie

17

18

Promotoren: Prof. dr. H.J. van den Herik,
Prof. dr. ir. H. X. Lin

Copromotor: Dr. J.L.A. Dubbeldam

Promotiecommissie: Prof. dr. A. Plaat,
Prof. dr. J.N. Kok,
Prof. dr. ir. A.W. Heemink (TU Delft) ,
... (Buitenlands),
Prof. dr. ir. B.J.A. Kröse
Prof. dr. C.M. Jonker (SIKS)



19 SIKS Dissertation Series No. ...

20 The research reported in the thesis has been carried out under the auspices of SIKS, the
21 Dutch Research School for Information and Knowledge Systems.

22

23 ISBN

24 Copyright © 2018, A.C. van Rossum

25 *All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or*
26 *transmitted, in any form or by any means, electronically, mechanically, photocopying, recording*
27 *or otherwise, without prior permission of the author.*

“

28

The study of mental objects with reproducible properties is called mathematics.

29

”

30

The Mathematical Experience (Davis and Hersch, 1981)

“

31

The study of physical objects with reproducible properties is called science.

32

”

33

The dawning of the age of stochasticity, Mathematics: frontiers and perspectives

34

(Mumford, 2000)

CONTENTS

36	1 Introduction	1
37	1.1 Scope of the Thesis	1
38	1.2 Bayesian Nonparametrics	2
39	1.3 Problem Statement and Research Questions	2
40	1.4 Research Methodology	3
41	1.5 Main Contribution	4
42	1.6 Organization of the Thesis	5
43	2 Related Work	7
44	2.1 Probability Theory	8
45	2.1.1 Measure Theory	8
46	2.1.2 Bayesian Inference	16
47	2.1.3 Model Composition	21
48	2.1.4 General Random Elements	23
49	2.1.5 Plate Notation	24
50	2.1.6 Completely Random Measure and Lévy Measure	25
51	2.1.7 Exchangeability	26
52	2.1.8 Stick-breaking Representation	27
53	2.2 Five Random Processes	28
54	2.2.1 Dirichlet Process	28
55	2.2.2 Beta Process	31
56	2.2.3 Gamma Process	34
57	2.2.4 Pitman-Yor Process	36
58	2.2.5 Hierarchical Dirichlet Process	37
59	2.2.6 Comparison of Random Processes	38
60	2.3 Inference	39
61	2.3.1 Inverse Transform Sampling	39
62	2.3.2 Rejection Sampling	39
63	2.3.3 Approximate Bayesian Computation	41
64	2.3.4 Gibbs Sampling	42
65	2.3.5 Metropolis-Hastings Sampling	43
66	2.3.6 Split-Merge MCMC Sampling	44
67	2.3.7 Comparison of the Six Inference Methods	45
68	2.4 Chapter Conclusions	46

69	3 Nonparametric Bayesian Line Detection	47
70	3.1 Infinite Line Model	48
71	3.1.1 Posterior Predictive for a Line given Other Lines	49
72	3.1.2 Likelihood of Data given a Line	50
73	3.1.3 Conjugate Prior for a Line	51
74	3.1.4 Posterior Predictive for a Line given Data	52
75	3.2 Inference for the Infinite Line Model	53
76	3.3 Accelerating Inference for the Infinite Line Model	54
77	3.4 Performance of the Infinite Line Model	56
78	3.4.1 Clustering Performance	56
79	3.4.2 Two Examples	58
80	3.5 Chapter Conclusions	59
81	4 Nonparametric Bayesian Segment Estimation	61
82	4.1 Pareto Pairs	62
83	4.1.1 Pareto Prior	62
84	4.1.2 Posterior for a Pareto pair	64
85	4.2 Generative Process to Create a Line Segment	66
86	4.2.1 Dirichlet Process Prior	66
87	4.2.2 Likelihood of Data given Segment Parameters	66
88	4.3 Inference over a Line Segment	68
89	4.4 Results	68
90	4.5 Chapter Conclusions	70
91	5 Triadic Split-Merge Sampler	71
92	5.1 The Class of Split-Merge Samplers	71
93	5.1.1 Split and Merge Moves	72
94	5.1.2 Dirichlet process Prior	72
95	5.2 Conventional Split-Merge Sampler	73
96	5.2.1 Acceptance for the Split Step	74
97	5.2.2 Acceptance for the Merge Step	76
98	5.2.3 Sequentially-Allocated Merge-Split Sampler	76
99	5.3 Triadic split-merge sampler	77
100	5.3.1 Acceptance for the Split Step	79
101	5.3.2 Acceptance for the Merge Step	80
102	5.4 Results	80
103	5.4.1 Implementation	81
104	5.4.2 Comparison	82
105	5.5 Chapter Conclusions	83
106	6 Adversarially Trained MCMC Kernels	85
107	6.1 Data-Driven Inference	85
108	6.2 Learning the Transition Operator	86
109	6.2.1 Adversarial Training	86
110	6.2.2 Variational Autoencoders	87
111	6.2.3 Infusion Training	87
112	6.3 Volumetric Models	88

113	7 Recommender Engine	91
114	7.1 Application	91
115	7.2 Model of Individuals	92
116	7.2.1 Multi-modal Normal-Uniform Distribution Model	92
117	7.2.2 Multi-modal Von-Mises-Uniform Distribution Model	93
118	7.2.3 Hyperparameters	94
119	7.3 Model of Groups	95
120	7.4 Inference	95
121	7.5 Results	96
122	7.5.1 Artificial Dataset	96
123	7.5.2 Real-world Dataset	97
124	7.6 Discussion	98
125	8 Discussion and Conclusions	99
126	References	101
127	Appendices	
128	A Probabilistic Concepts	109
129	A.1 Common Inequalities	110
130	A.1.1 Markov's Inequality	110
131	A.1.2 Chebyshev's Inequality	111
132	A.1.3 Weak Law of Large Numbers	111
133	A.1.4 Strong Law of Large Numbers	112
134	A.1.5 Common Distributions	112
135	B Dirichlet-Multinomial Interpretations	115
136	B.1 Dirichlet-Categorical	116
137	B.2 Dirichlet-Multinomial Distribution	117
138	B.3 Dirichlet-N Categorical Distributions	118
139	C Gibbs Sampling	121
140	C.1 Gibbs Sampling of Parameters	123
141	D Glossary	127
142	List of Figures	129
143	List of Tables	133
144	Summary	136
145	Samenvatting	138
146	Acknowledgments	141
147	Curriculum Vitae	142

148	Publications	144
149	SIKS Dissertation Series	146

INTRODUCTION

Contents

The thesis addresses nonparametric Bayesian methods in robotic vision. Nonparametric Bayesian models can be simultaneously employed to perform inference over the number of entities observed and over the shape or nature of these entities. This chapter introduces nonparametric Bayesian models, the research methodology based on the Bayesian methodology, the main contribution towards robotic vision, and the general organization of the thesis.

Outline

The scope of thesis is to apply nonparametric Bayesian methods to robotic vision (Section 1.1). Bayesian nonparametric models define entities together with noise in such a way that inference can be performed in an optimal manner (Section 1.2). Particular problems in robotic vision that can benefit from Bayesian nonparametric methods are formulated and detailed (Section 1.3). The research methodology is described (Section 1.4). Our main contribution is to introduce nonparametric Bayesian models in robotic vision (Section 1.5). At the end of this chapter the organization of the thesis is given (Section 1.6).

1.1 Scope of the Thesis

In the thesis, modern Bayesian nonparametric methods are used to answer long-standing questions within computer vision and robotics. The following three challenging questions are typical examples. Is there a Bayesian form of line detection rather than applying the traditional Hough transform? Which of the nonparametric Bayesian priors can be used to detect multiple features simultaneously? What are efficient inference methods for these priors?

The scope of the thesis is the transfer of knowledge on Bayesian nonparametrics to well-described application domains. It will not establish a new body of work around a new family

of stochastic processes. The detailed application of complex models towards robotic vision is expected to help and encourage people in entirely different application domains, such as collaborative filtering, search engine optimization, and audio processing. All these different applications do not always need dedicated algorithms, but do deserve and can exploit the same optimal general inference techniques from Bayesian nonparametrics.

1.2 Bayesian Nonparametrics

In robotic vision (computer vision and depth perception) traditionally custom-made algorithms have been developed for a given task. There are specific methods to detect corners (e.g., Förstner and Gülch, 1987; Harris and Stephens, 1988; Shi and Tomasi, 1994), to detect edges (e.g., Sobel, 1970; Canny, 1986), to detect features (e.g., Hough, 1962), and to describe features (e.g., Lowe, 1999; Dalal and Triggs, 2005; Bay et al., 2006).

On the one hand, it is desirable that such sophisticated methods are generalizable to other application domains. On the other hand, it is important to take particular information about an application domain into account. The methods described in the previous paragraph are limited to their specific task. An example of limited generalizability can be found in the Hough transform. The Hough transform can be used to detect lines, but the way inference is performed in the algorithm does limit its application to basic forms of object detection. An example of limited specificity can be found in linear regression. Linear regression does not take into account real-world statistics.

Both generalization and specificity are formalized by a Bayesian model. A Bayesian model is general because it can be solved with general inference methods. One of such general inference methods is a Markov-Chain Monte Carlo method. It does not know anything about real-world statistics. A Bayesian model is also specific in that it can incorporate application-specific know-how by the definition of priors.

Typical problems in robotic vision will be about the recognition of several objects, multiple shapes, or objects that have multiple parts. Models that represent such objects do not have knowledge about the number of such objects, shapes, or parts. To incorporate application-specific know-how on the number of objects it is possible to define a prior that assigns a probability to this quantity. The number of objects can even be potentially infinite. The Bayesian models that define a prior on the number of objects, shapes, or parts are called nonparametric Bayesian models. This means that in contrast with conventional methods such as k -means clustering (Forgy, 1965; Lloyd, 1982) the number of objects does not need to be predefined.

1.3 Problem Statement and Research Questions

Many methods in robotics - and in particular in robotic vision - have been developed in times where computational resources were limited. Then, highly optimized algorithms have been developed, leveraging peculiarities of the application domain. Recent advances in Bayesian

214 methods, both with respect to concept development, as well as computational efficient so-
 215 lution strategies, now open up new ways to solve old problems. However, extending only
 216 the old methods themselves would lead to ad-hoc solution strategies that will miss benefits
 217 from potential optimal and more widely applicable algorithms.

218 This observation leads us to the formulation of our problem statement (PS).

219 **PS:** *How can robotic problems effectively be generalized and their structure
 exploited in a wider Bayesian framework?*

220 The problem statement is rather general. In our research, we focus on robotic vision, in the
 221 form of point cloud recognition and depth perception. In particular, we look at objects, lines,
 222 line segments, and more complex shapes.

223 The problem statement is divided into three research questions (RQs).

RQ 1 How can we estimate the number of objects simultaneously with the
 fitting of these objects?

224 **RQ 2** How can we estimate the number of lines simultaneously with line fitting
 in computer vision?

RQ 3 How can we recognize more general 3D objects?

225 1.4 Research Methodology

226 The research methodology advocated in the thesis follows the Bayesian methodology (cf.
 227 Savage, 1972; Jaynes, 2003). So, our research methodology exists out of two phases. In the
 228 first phase a Bayesian model is defined. This model exists of (1) a definition of parameters
 229 and relations between these parameters, (2) a definition of the noise, and (3) the data. In
 230 the second phase, the Bayesian method dictates all remaining unknowns, from the number
 231 of parameters to the values of the parameters. To perform Bayesian inference efficiently new
 232 methods are required if the model is complex (as is in the case of robotic vision).

233 The Bayesian methodology aims to establish the rationale for practical questions. The fol-
 234 lowing two questions are clear examples.

- 235 ◦ If we observe a single point in an image, can we expect it to be part of a line?
- 236 ◦ If we have two lines and we live in a world with squares, what are we able to infer?

237 The two questions tap into our capabilities to define models that makes our prior knowledge
 238 explicit. Moreover, if we are able to quickly assign (1) points to segments, (2) segments to
 239 lines, (3) objects to categories, we can enrich it with all corresponding group properties
 240 without the need to have them observed for this individual.

In robotic vision we take as an example the task of line detection. Both the Hough transform (Hough, 1962) and the RANSAC method (Bolles and Fischler, 1981) do detect lines, but they do not explicitly take noise into account. By applying Bayesian methodology to these tasks, the inference method becomes optimal in an information-theoretic sense. Also frequentist statisticians agree that nonparametric Bayesian models are consistent in the sense that they approach the underlying true distribution (Wasserman, 1998). There is no need to search for another method to infer lines in a line detection task. If someone would find a method that outperforms a Bayesian method it is either (1) because the signal or noise has not been correctly modeled after all, or (2) because the method overfits with respect to the available data. If approximations are used with respect to optimal Bayesian inference (either variational approximations or Markov-Chain Monte Carlo), there are theoretical guarantees on convergence.

A well known problem with nonparametric Bayesian models is the curse of dimensionality. Compared to maximum likelihood methods or other non-probabilistic methods that do not take noise into account at all, the nonparametric Bayesian models require significant computational resources. Our research methodology first establishes the correct models, even if solving them seems computationally infeasible. Our approach is to develop subsequently approximations using more sophisticated samplers, so that the theoretical guarantees on convergence are preserved.

Due to the fact that the models are optimal by construction, there are no experiments required to address the optimality in particular. However, experiments are still required to establish whether the models make sense. Yet, the methodology does also have limitations. For instance, we will not search over different noise models and limit priors to a particular hierarchical level.

1.5 Main Contribution

Our contribution to robotic vision can be subdivided into three parts that correspond with the three research questions.

The first part addresses the problem of inference about objects from a nonparametric Bayesian perspective. Contemporary methods in robotic vision do not allow for astute statements about their performance. In practice, this means that when using computer vision to detect cells under a microscope, someone cannot be confident about the number of detected cells. An autonomous cleaning robot in a supermarket cannot be confident about the aisle it is driving into. To be able to properly take into account models and uncertainty simultaneously, Bayesian models have found mainstream adoption. State-of-the-art Bayesian methods that reason about the number of objects alongside object models are a recent object of study (cf. Ferguson, 1973; Hjort, 1990; Lijoi and Prünster, 2010; Joho et al., 2011). The thesis applies such nonparametric Bayesian models towards the applications of robotic vision and depth perception. Models such as the infinite line model and the infinite line segment model are introduced.

280 The second part addresses the problem of high-dimensional data. To efficiently sample
281 more complex geometric structures, new MCMC (Markov-Chain Monte Carlo, Section 2.3.4)
282 methods are required. The thesis introduces such an MCMC sampler, namely a new Split-
283 Merge sampler, and applies it to complex geometric structures.

284 The third part addresses more complex robotic vision problems, in the form of object recog-
285 nition of point clouds in 3D. It combines nonparametric Bayesian inference with models
286 from deep learning.

287 1.6 Organization of the Thesis

288 **Chapter 1** (this chapter) introduces the problem of contemporary methods in computer
289 vision and depth perception. Due to the fact that these methods are not optimal
290 by construction, it is hard to articulate how they perform. The need for a Bayesian
291 methodology is sketched briefly. The problem statement and the research questions
292 are formulated. Moreover, the research methodology is described and the organization
293 of the thesis is outlined.

294 **Chapter 2** describes (1) probability theory using measure theory, (2) random measures
295 known as random processes of which five are described as nonparametric Bayesian
296 models, and (3) six inference methods that infer model parameters of such nonpara-
297 metric Bayesian models given the data. It is followed by a discussion that indicates
298 which parts will be most useful for chapters 3 and 4.

299 **Chapter 3** examines a first nonparametric Bayesian model, i.e., the infinite line model. The
300 infinite line model represents a countably infinite set of lines. Gibbs sampling is used
301 to perform simultaneous inference over (1) the number of lines and (2) line parameter
302 values such as slope and intercept.

303 **Chapter 4** examines a second nonparametric Bayesian model, i.e., the infinite line segment
304 model. The infinite line segment model represents a countably infinite set of line
305 segments. A split-merge MCMC sampling method is used to perform simultaneous
306 inference over (1) the number of line segments and (2) line segment parameter values
307 such as slope, intercept, and segment size. Chapters 2 to 4 answer the first research
308 question.

309 **Chapter 5** investigates a new MCMC method, the Triadic Split-Merge sampler. It is tailored
310 to clustering problems and accelerates inference of the models in Chapters 3 and 4.
311 This chapter answers the second research question.

312 **Chapter 6** examines a third nonparametric Bayesian model, particularly aimed at volumet-
313 ric inference. This chapter answers the third research question.

314 **Chapter 7** applies the hierarchical sampling method to the domain of recommender en-
315 gines. It estimates simultaneously the number of user types with a fitting procedure
316 for the individual user. I want to change this chapter to something relevant to point
317 clouds.

318 **Chapter 8** discusses the relevance of the developed models and inference methods. The
319 answers to the research questions are discussed. Then the problem statement is an-
320 swered and conclusions are formulated. Finally, recommendations are given and fu-
321 ture research is envisaged.

RELATED WORK

322

323

Contents

325

326

327

328

329

330

331

332

333

In robotics depth sensors generate point clouds. The tasks of robotic object recognition, positioning, and navigation require models that represent such point clouds. It is unclear whether the current methods that perform inference over point clouds are appropriate for these tasks. The current models do not model uncertainty explicitly. This chapter presents models that can be used for point cloud modeling and that represent uncertainty. This (partially) answers research question RQ1. The chapter concludes with recommendations for the development of point cloud inference models. They will be implemented in a new model for line inference in Chapter 3 and line segment inference in Chapter 4.

Outline

335

336

337

338

339

340

341

342

343

344

345

346

This chapter describes probability theory, and in particular, measure theory underlying random processes (Section 2.1). Five random processes are described, the Beta process, the Gamma process, the Dirichlet process, the Pitman-Yor process, and the hierarchical Dirichlet process. The random processes are presented as a Poisson process with a Lévy measure, a stick-breaking construction, and a sequential presentation (Section 2.2). These representations give rise to different inference methods. Six inference methods are described: Inverse transform sampling, rejection sampling, approximate Bayesian computation, Gibbs sampling, Metropolis-Hastings, and Split-Merge Markov chain Monte Carlo (Section 2.3). Inference about point clouds in the chapters to follow will use adaptations of the described models and inference methods for which some recommendations are given (Section 2.4).

2.1 Probability Theory

Modern probability is based on measure theory (Section 2.1.1). Measure theory will provide the means to formally describe random variables, random processes, and most generally, random measures. A model represented by random measures can be fitted to the data using Bayesian inference (Section 2.1.2). We give three typical examples of Bayesian model compositions, among which an infinite mixture model (Section 2.1.3). A number of processes are described that can be used with (for example as prior distribution) infinite mixture models (Section 2.1.4). We introduce plate notation which visualizes infinite models particularly well (Section 2.1.5). We investigate completely random measures and Lévy measures (Section 2.1.6), exchangeability (Section 2.1.7), and stick-breaking processes (Section 2.1.8), which will form the basis of the processes defined in Section 2.2.

2.1.1 Measure Theory

A random variable is a *function* that assigns values to a *set* of possible outcomes. The formal definition requires concepts such as “measurable function” and “probability space” from *measure theory* (Feller, 1950). Measure theory is used to generalize the notion of a random variable to that of a “random process”.

Informally, a measure generalizes the notion of size of Euclidean objects to sets and subsets. The definition of a measure is based on the definition of a σ -algebra. A σ -algebra ascribes a value to a sum of individual disjoint sets, even if they are infinite in number.

▼ Definition 2.1 — σ -algebra

A σ -algebra is a subset $\Sigma \in 2^X$, with X a set and 2^X its powerset, with three requirements:

- Σ is non-empty: at least one $A \in X$ is in Σ ;
 - Σ is closed under complementation: if A in Σ , so is its complement A^c ;
 - Σ is closed under countable unions: if A_1, A_2, \dots in Σ , so is $A = A_1 \cup A_2 \cup \dots$
-

The members of a σ -algebra are called *measurable sets*. Let $X = \{1, 2, 3, 4\}$ and let us define a σ -algebra $\Sigma = \{\emptyset, \{1\}, \{4\}, \{2, 3\}, \{1, 4\}, \{1, 2, 3\}, \{2, 3, 4\}, \{1, 2, 3, 4\}\}$. Here \emptyset denotes the empty set. The complement of A is defined with respect to X : $A \cup A^c = X$. An example of closure under complementation: let $A_1 = \{1\}$, then $A_1^c = \{2, 3, 4\}$ and A_1^c is indeed a member of Σ : $A_1^c \in \Sigma$. An example of closure under countable unions: let $A_1 = \{1\}$ and $A_2 = \{2, 3\}$, then $A_1 \cup A_2 = \{1, 2, 3\}$ and $A_1 \cup A_2 \in \Sigma$.

▼ Definition 2.2 — generated σ -algebra

A **generated σ -algebra**, with X a set and $B \in 2^X$, is the smallest σ -algebra $\sigma(B)$ that contains all sets of B .

Let $X = \{1, 2, 3, 4\}$ and $B = \{\{1\}, \{2, 3\}\}$, then the generated σ -algebra is the set $\sigma(B) = \{\emptyset, \{1\}, \{2, 3\}, \{1, 4\}, \{2, 3, 4\}, \{1, 2, 3, 4\}\}$. Here the sets in B are completed to the sets in $\sigma(B)$ by obeying the requirements of a σ -algebra of closure under complementation and countable unions by addition (e.g., $\{1, 4\}$ is added due to closure under completion with respect to $\{2, 3\}$).

The notion of a σ -algebra (Fremlin, 2000) can be applied to solve the so-called Banach-Tarski paradox (Banach and Tarski, 1924). This paradox describes how a unit-ball in \mathbf{R}^3 can be partitioned into a finite number of disjoint infinite sets (scattering of points) and then can be reassembled into two unit-balls again. This violates the intuitive notion of preservation of volume. If the measure μ of the union of two disjoint sets is equal to the sum of the measures of the two sets, this is called *finite additivity*: $\mu(\bigcup_{i=1}^N A_i) = \sum_{i=1}^N \mu(A_i)$. In probability theory σ -additivity extends this to infinite disjoint sets: $\mu(\bigcup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} \mu(A_i)$. Measure theory solves the Banach-Tarski paradox by only assigning a measure to subsets that are measurable sets (Tao, 2011).

A *measure* assigns values to measurable sets (as stated before, measurable sets are members or subsets of Σ).

▼ Definition 2.3 — *measure*

A **measure** μ is a function from Σ to $[-\infty, +\infty]$, with three requirements:

- μ is non-negative: $\mu(A) \geq 0$ for $\forall A \in \Sigma$;
- μ has a null empty set: $\mu(\emptyset) = 0$;
- μ is σ -additive: $\mu(\bigcup_{i \in I_{\Sigma}} A_i) = \sum_{i \in I_{\Sigma}} \mu(A_i)$ for A_i disjoint.

The first statement defines that a measure μ only assigns non-negative values to sets in Σ . The second statement equals the measure of the empty set \emptyset to 0. The third statement defines that σ -additivity is required. For any two sets in Σ the measure of the union of the sets equals the sum of the measures of the individual sets. Here I_{Σ} defines an index over sets in Σ .

Informally, a measure relates the concepts of *sets* and *subsets* to notions of size. A measure can be seen as a *monotonically* increasing function. Let the set A in X be the interval $[0, 1)$, an uncountable (infinite) set of real numbers. Define the σ -algebra $\{\emptyset, A\}$. The empty set has measure 0, the set A has measure 1. Let us define the σ -algebra $\{\emptyset, A_{0,0.5}, A_{0.5,1}, A\}$. The set $A_{0,0.5}$ corresponds to the interval $[0, 0.5)$ and $A_{0.5,1}$ to $[0.5, 1)$. Both sets are assigned measure 0.5 and their union has measure 1. This examples shows that with σ -additive unions, measures can be assigned to sets that are uncountable.

A *measurable space* (X, Σ) is defined as a pair.

▼ Definition 2.4 — measurable space

A **measurable space** (X, Σ) is a pair with:

- X a set;
 - Σ a σ -algebra over X .
-

A **measure space** (X, Σ, μ) is defined as a triple.

▼ Definition 2.5 — measure space

A **measure space** (X, Σ, μ) is a triple with:

- X a set;
 - Σ a σ -algebra over X ;
 - μ a measure from Σ to $[-\infty, \infty]$.
-

A finite measure μ assigns a finite real number to all A .

▼ Definition 2.6 — finite measure

A **finite measure** μ is a measure from Σ to $[0, \infty)$:

- μ is non-negative: $\mu(A) \geq 0$ for $\forall A \in \Sigma$;
 - μ has a null empty set: $\mu(\emptyset) = 0$;
 - μ is σ -additive: $\mu(\bigcup_{i \in I_\Sigma} A_i) = \sum_{i \in I_\Sigma} \mu(A_i)$ for A_i disjoint;
 - μ for the whole sample space, X , is finite: $\mu(X) = N$.
-

A σ -finite measure allows A to be a countable union of sets with finite measure.

▼ Definition 2.7 — σ -finite measure

A **σ -finite measure** μ is a finite measure with:

- X is a countable union of sets with finite measures.
-

We will now define five measures: (A) the *probability measure* (Definition 2.8), (B) the *counting measure* (Definition 2.10), (C) the *Borel measure* (Definition 2.12), (D) the *Lebesgue measure* (Definition 2.17), and (E) the *random measure* (Definition 2.18). These measures are important because they are fundamental to different branches of mathematics. In probability theory a σ -algebra is interpreted as a collection of events to which probabilities are assigned. Counting measures play a fundamental role in discrete probability distributions. In integration theory a σ -algebra corresponding to the Borel and Lebesgue measures are relevant for integration in the Euclidean space \mathcal{R}^n . In statistics a σ -algebra formally defines

the concept of sufficient statistics and generalizes random variables to random functions and measures.

A: Probability measure

A *probability measure*, \mathbb{P} , is a finite measure that assigns non-negative values \mathbb{P} , called probabilities, to sets A , called events (see Definition 2.8).

▼ Definition 2.8 — *probability measure*

A **probability measure** \mathbb{P} is a measure μ with:

- \mathbb{P} is non-negative: $\mathbb{P}(A) \geq 0$ for $\forall A \in \Sigma$;
- \mathbb{P} has a null empty set: $\mathbb{P}(\emptyset) = 0$;
- \mathbb{P} is σ -additive: $\mathbb{P}(\bigcup_{i \in I_\Sigma} A_i) = \sum_{i \in I_\Sigma} \mu(A_i)$ for A_i disjoint;
- \mathbb{P} for the whole sample space, X , is unity: $\mathbb{P}(X) = 1$.

The four requirements are called the Kolmogorov axioms (Kolmogorov, 1933). The probability measure is an actual *measure*. It therefore obeys the three requirements: (1) non-negativity for any set, (2) the existence of a null empty set, and (3) σ -additivity. Here we note that a *probability* measure compared to a general measure obeys a fourth requirement, namely the restriction of the measure for the whole space X to 1. This can be seen as some kind of normalization. It influences how two probability measures have to be summed to become again a probability measure.

In Figure 2.1 the probability measure is visualized as a mapping from the probability space to the unit interval $[0, 1]$.

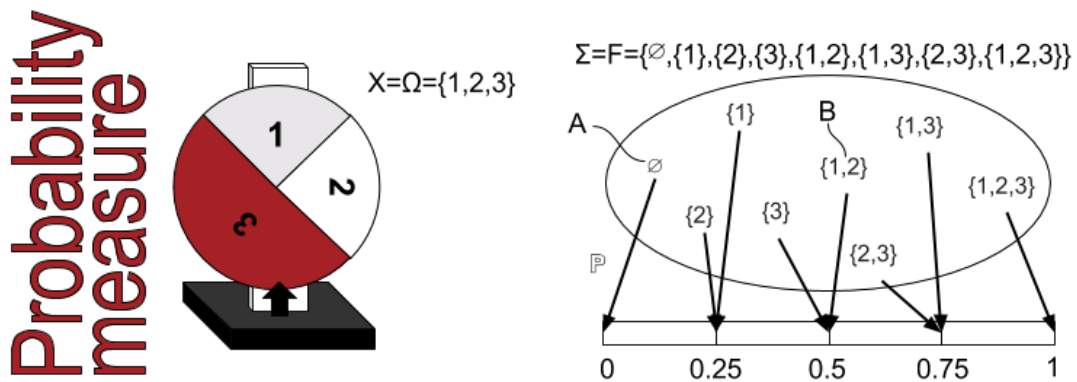


Figure 2.1: A probability measure \mathbb{P} mapping the probability space for 3 events to the unit interval. Left: a turning wheel representing three possible outcomes of which the third is twice as likely as the other two outcomes. Right: a probability measure \mathbb{P} assigned to each outcome. The empty set, $A = \emptyset$, has probability measure 0. The set of encountering either 1 or 2, $B = \{1, 2\}$, has probability measure 0.5. Taken from Wikipedia.

A probability space (X, Σ, \mathbb{P}) is a measure space (X, Σ, μ) with the probability measure \mathbb{P} as its measure μ .

▼ Definition 2.9 — *probability space*

A probability space (X, Σ, \mathbb{P}) is a triple with:

- X a set;
 - Σ a σ -algebra over X ;
 - \mathbb{P} a probability measure from Σ to $[0, 1]$.
-

The triple for the probability space (X, Σ, \mathbb{P}) is also written as $(\Omega, \mathbb{F}, \mathbb{P})$. The space X is the event space Ω , the set of *elementary outcomes*. The σ -algebra over subsets of Ω is denoted by \mathbb{F} . The probability measure \mathbb{P} assigns a value on the unit interval $[0, 1]$ to every event in \mathbb{F} .

B: Counting measure

The *counting measure* forms the basis for the definition of discrete probabilities (Schilling, 2005).

▼ Definition 2.10 — *counting measure*

A counting measure ν on a space X is a measure μ with:

- ν is non-negative and integer-valued for $\forall A \in \Sigma$;
 - $\nu < \infty$ for $\forall A \in \Sigma$ if A bounded (of finite size);
 - $\nu = \infty$ if $\exists A \in \Sigma$ with A unbounded (infinite).
-

A counting measure is a measure that is integer-valued. Every set A has a measure that is a positive integer or zero. The set A is unbounded if and only if its counting measure is infinite.

C: Borel measure

The *Borel σ -algebra* defines a σ -algebra for the real line \mathbb{R} .

▼ Definition 2.11 — *Borel σ -algebra*

A Borel σ -algebra \mathbb{B}_σ on \mathbb{R} is the smallest σ -algebra that contains all open subsets of \mathbb{R} :

- $\mathbb{B} = \Sigma(U)$ with $U = \bigcup U \subseteq \mathbb{R}$: U is open.
-

500 The Borel σ -algebra contains all open subsets of \mathbb{R} . The property of closure under comple-
 501 mentation of a σ -algebra means that it also contains the closed subsets of \mathbb{R} . If $A = (0, 1)$,
 502 then $A^c = \{[-\infty, 0], [1, \infty]\}$.

503 A Borel measure assigns values to subsets of \mathbb{B}_σ .

504 ▼ Definition 2.12 — Borel measure

505 A Borel measure μ is a function from $\Sigma = \mathbb{B}_\sigma$ to $[-\infty, +\infty]$, with the three measure
 506 requirements:
 507

- 508 ◦ μ is non-negative: $\mu(A) \geq 0$ for $\forall A \in \Sigma$;
- 509 ◦ μ has a null empty set: $\mu(\emptyset) = 0$;
- 510 ◦ μ is σ -additive: $\mu(\bigcup_{i \in I_\Sigma} A_i) = \sum_{i \in I_\Sigma} \mu(A_i)$ for A_i disjoint.

512 The Borel space is a measurable space with a Borel σ -algebra rather than a general σ -
 513 algebra.

514 ▼ Definition 2.13 — Borel space

515 A Borel space (X, \mathbb{B}_σ) is a pair with:
 516

- 517 ◦ X a set;
- 518 ◦ \mathbb{B}_σ a Borel σ -algebra over X .

520 A complete measure space is a measure space in which every subset of every null set is mea-
 521 surable.

522 ▼ Definition 2.14 — complete measure space

523 A complete measure space (X, Σ, μ) :
 524

- 525 ◦ $S \subseteq N \in \Sigma$ and $\mu(N) = 0 \Rightarrow S \in \Sigma$.

527 The Borel space is not a complete measure space. There are sets in the Borel σ -algebra that
 528 are of measure zero and that contain subsets that are undefined.

529 D: Lebesgue measure

530 The Lebesgue measure defines a size to subsets of \mathbb{R}^n that completes the Borel measure
 531 (Lebesgue, 1902). It makes use of the notion of an outer measure.

532 ▼ Definition 2.15 — outer measure

533 An outer measure ϕ on a space \mathbb{R} is a measure μ with:
 534

- ϕ is non-negative and real-valued for $\forall A \in \Sigma$;
- ϕ has a null empty set: $\phi(\emptyset) = 0$;
- ϕ is σ -subadditive: $\phi(\bigcup_{i \in I_\Sigma} A_i) \leq \sum_{i \in I_\Sigma} \phi(A_i)$ for $\forall A_i$;
- ϕ is monotone: $A \subseteq B$ implies $\phi(A) \leq \phi(B)$;
- ϕ is translation-invariant: $\phi(A + x) = \phi(A)$ for $\forall A \in \Sigma$ and $\forall x \in \mathbb{R}$.

An outer measure relaxes σ -additivity of disjoint sets of X to σ -subadditivity for any sequence of sets. Intuitively, the outer measure of a set is an upper bound on the size of a set.

▼ **Definition 2.16 — Lebesgue outer measure**

A **Lebesgue outer measure** λ on a space \mathbb{R}^n is an outer measure ϕ with:

- $\lambda(A) = \inf \left\{ \sum_{k=1}^{\infty} l(I_k) : (I_k)_{k \in \mathbb{N}} \text{ is a sequence of open intervals with } A \subseteq \bigcup_{k=1}^{\infty} I_k \right\}$.

Here $A \subseteq \mathbb{R}$ is a subset of the real line. The Lebesgue outer measure λ is the infimum (greatest lower bound) of the sum of the lengths $l(I) = b - a$ of the intervals $I = [a, b]$.

The *Lebesgue measure* is defined through the Lebesgue outer measure.

▼ **Definition 2.17 — Lebesgue measure**

A **Lebesgue measure** m on a space \mathbb{R}^n is a Lebesgue outer measure λ with:

- $m(B) = \lambda(B \cup A) + \lambda(B \cup A^c)$.

E: Random measure and random process

A measurable function is defined between two measurable spaces.

▼ **Definition 2.18 — measurable function**

A **measurable function** $f : X \rightarrow Y$ fulfills:

- $f^{-1}(E) \in \Sigma$ for $\forall E \in T$,

with both (X, Σ) and (Y, T) measurable spaces.

A measurable function *preserves the structure* of the corresponding measurable spaces (captured through the σ -algebras).

A random variable is a measurable function between two measurable spaces, with as domain a measurable space that is a probability space.

▼ **Definition 2.19 — general random variable**

A (X, Σ) -valued random variable X is a measurable function from probability space $(\Omega, \mathbb{F}, \mathbb{P})$ to measurable space (X, Σ) .

A random variable is a (X, Σ) -valued random variable with a choice for the codomain and σ -algebra (see Definition 2.20).

▼ **Definition 2.20 — random variable**

A random variable X is a measurable function from probability space $(\Omega, \mathbb{F}, \mathbb{P})$ to the real line with the Borel σ -algebra $(\mathbb{R}, \mathbb{B}_{\mathbb{R}})$.

The codomain is the real line \mathbb{R} and the Borel σ -algebra.

Random variables can be generalized to complex random variables or random elements of any type. A *complex random variable* is a measurable function from Ω to \mathbb{C} . A *random elephant* is a measurable function from Ω to a suitable space of elephants (Kingman, 1993).

▼ **Definition 2.21 — random measure**

A random measure is a function $\xi : \Omega \times X \rightarrow [0, +\infty]$ from probability space $(\Omega, \mathbb{F}, \mathbb{P})$ to measurable space (X, Σ) such that $\xi(\cdot, X)$ is a random variable on $(\Omega, \mathbb{F}, \mathbb{P})$ and $\xi(\omega, \cdot)$ is a measure on Σ .

We have encountered a random variable, and a probability measure \mathbb{P} on the original probability space. Now, one might wonder whether probabilities are logically assigned to elements on the measurable space that is the codomain of this random variable. Why does it map to a measurable space and not a measure space actually? This is because (through the σ -algebras of both spaces, or more precisely the random variable itself) the probability measure is *induced* on the target space. This is known as a *probability distribution*:

▼ **Definition 2.22 — probability distribution**

Given a random variable X from $(\Omega, \mathbb{F}, \mathbb{P})$ to $(\mathbb{R}, \mathbb{B}_{\sigma})$, the **probability distribution** μ is the induced probability measure: $\mu(B) = \mathbb{P}(X^{-1}(B))$ for all Borel sets $B \in \mathbb{B}_{\sigma}$.

The measurable function X is inverted: $X^{-1}(\cdot)$. The measure μ exists on $(\mathbb{R}, \mathbb{B}_{\sigma})$ just as \mathbb{P} exists on (Ω, \mathbb{F}) . The notation for the measure μ does not include the original probability space or σ -algebra. The complete notation for the probability distribution μ can be written as a function f of X :

$$f_X(x) = f_{X,(\Omega,\mathbb{F},\mathbb{P}),(\mathbb{R},\mathbb{B}_{\sigma})}(x). \quad (2.1)$$

At the left X denotes the random variable, $x \in \Omega$ are the (elementary) outcomes on the sample space Ω . At the right the complete notation adds \mathbb{F}, \mathbb{P} and $\mathbb{R}, \mathbb{B}_\sigma$. The shorthand notation at the left will be used to indicate the real line with a Borel σ -algebra as codomain.

A random variable X is *distributed as* $f_X(x)$, notation:

$$X \sim f_X(x). \quad (2.2)$$

A *random process* is an *ordered* set of random variables. The set can be a sequence of random variables in a time series. It can be a series of steps in the spatial domain, called a random field.

▼ **Definition 2.23 — random process**

A **random process** X is a collection $\{X_t : t \in T\}$ with X_t an (S, Σ) -valued random variable on Ω and $(\Omega, \mathbb{F}, \mathbb{P})$ a probability space, (S, Σ) a measurable space, and T a totally ordered set.

A random process is a probability distribution with a domain that is a set of probability distributions. A random process is a distribution over distributions, a hierarchy over distribution.

2.1.2 Bayesian Inference

Let x be a (S, Σ_S, μ_S) -valued random variable¹, y a (T, Σ_T, μ_T) -valued random variable, then we can construct z , a (C, Σ_C, μ_C) -valued random variable with the latter being a subset of the product set of x and y : $C \in S \otimes T$.

▼ **Definition 2.24 — product space**

A **product space** $(S \otimes T, \Sigma_{S \otimes T})$ has σ -algebra $\Sigma_{S \otimes T} = \sigma(F \otimes G : F \in \Sigma_S, G \in \Sigma_T)$ with (S, Σ_S, μ_S) and (T, Σ_T, μ_T) two σ -finite measure spaces.

▼ **Definition 2.25 — product measure**

A **product measure** $\mu_{S \otimes T}$ is a measure $\mu_{S \otimes T}(F \otimes G) = \mu_S(F) \otimes \mu_T(G)$ with (S, Σ_S, μ_S) and (T, Σ_T, μ_T) two σ -finite measure spaces.

The **joint probability distribution** P_C is a probability measure on the product σ -algebra Σ_C with $C \in S \otimes T$. As function of the random variables x and y the joint probability distribution is written as $x_{,Y}(x, y)$, $f(x, y)$, or $p(x, y)$.

A σ -algebra is *independent* in the following sense.

¹The lowercase x is used instead of X in the context of probability distributions as in Eq. 2.1.

▼ **Definition 2.26 — independent σ -algebra**

Let (Ω, \mathbb{F}, P) be a probability space and \mathbb{A} and \mathbb{B} be sub- σ -algebras of \mathbb{F} . \mathbb{A} and \mathbb{B} are **independent σ -algebras** if:

$$\circ P(A \cap B) = P(A)P(B) \quad \forall A \in \mathbb{A} \text{ and } B \in \mathbb{B}.$$

Two random variables x and y are independent if and only if the σ -algebras that they generate are independent.

▼ **Definition 2.27 — conditional probability distribution**

Let (Ω, \mathbb{F}, P) be a probability space, $\mathbb{G} \subseteq \mathbb{F}$ a sub- σ -algebra of \mathbb{F} , and $X : \Omega \rightarrow \mathbb{R}$ a real-valued random variable (\mathbb{F} -measurable with respect to the Borel σ -algebra \mathbb{B}_{σ} on \mathbb{R}). There exists a function $\mu : \mathbb{B}_{\sigma} \times \Omega \rightarrow \mathbb{R}$ such that $\mu(\cdot, \omega)$ is a probability measure on \mathbb{B}_{σ} for each $\omega \in \Omega$ and $\mu(H, \cdot) = P(X \in H | \mathbb{G})$ (almost surely) for every $H \in \mathbb{B}_{\sigma}$. For any $\omega \in \Omega$, the function $\mu(\cdot, \omega) : \mathbb{B}_{\sigma} \rightarrow \mathbb{R}$ is called a **conditional probability distribution** of X given \mathbb{G} .

Informally, a conditional probability is described with a sub- σ -algebra which only presents part of the structure of the full σ -algebra. As function of the random variables x and y the conditional probability distribution of y given x is written as $f_{Y|X}(y|x)$, $f(y|x)$, or $p(y|x)$.

A typical conditional probability distribution is that of the data given parameters, or parameters given the data. Another often used conditional probability distribution is that of the data given a statistic (summary) of that data. This statistic can be a so-called sufficient statistic.

▼ **Definition 2.28 — sufficient statistic**

A conditional probability distribution of the data X given a *sufficient statistic* $t = T(X)$ does not depend on parameter θ :

$$\circ P(x|t, \theta) = P(x|t)$$

The random variables x and θ define a Bayesian model with observations x and parameters θ .

▼ **Definition 2.29 — Bayesian model**

A **Bayesian model** $f(x, \theta)$ defines a function between observations x and parameters θ with both x and θ random variables.

In a **supervised learning** task both x and θ are known. In an **unsupervised learning** task x is known, but θ is unknown. The random variable θ is called a hidden or latent variable.

The random variable θ can be any random element: a random vector, a random matrix, a random process.

Let the observations x be a sequence x_0, x_1, \dots , then the observations x_i can be distributed *independent and identically*.

▼ **Definition 2.30 — *independent and identically distributed***

A collection of random variables $x = \{x_0, x_1, \dots\}$ is **independent and identically distributed (i.i.d.)** if:

- the probability distribution $p(x_i)$ is the same for $\forall x_i \in x$;
 - each x_i is independent with respect to x_j with $i \neq j$.
-

The observations x_i can be distributed in an *exchangeable* sequence in which any order is equally likely.

▼ **Definition 2.31 — *exchangeable***

A sequence of random variables $x = \{x_0, x_1, \dots\}$ is **exchangeable** if for any finite permutation ρ of the indices $0, 1, \dots$:

- the joint probability distribution of the permuted sequence $p(x_{\rho(0)}, x_{\rho(1)}, \dots)$ equals that of the original sequence $p(x_0, x_1, \dots)$.
-

The joint probability distribution of i.i.d. observations given parameters can be written as a product:

$$p(x_0, \dots, x_{k-1} | \theta) = \prod_{i=0}^{k-1} p(x_i | \theta). \quad (2.3)$$

▼ **Definition 2.32 — *likelihood function***

The **likelihood function** is defined as:

$$\mathcal{L}(\theta | x) = p(x | \theta). \quad (2.4)$$

The likelihood of the parameters θ given observations x is the probability of these observations given the parameter values.

The definition of the likelihood function allows us to find an optimal set of parameter values given the observations. The probability $p(x | \theta)$ can be maximized (Aldrich and Others, 1997).

▼ Definition 2.33 — *maximum likelihood*

Maximum likelihood is defined as:

$$\theta^* \in \operatorname{argmax}_{\theta} \prod_{i=0}^{k-1} p(x_i | \theta). \quad (2.5)$$

The maximum likelihood method finds the maximum of $\prod_{i=0}^{k-1} p(x_i | \theta)$ for all possible parameter values θ . The maximum in maximum likelihood does not need to be unique (Steel, 1994). The notation makes this explicit by writing θ^* as a member (denoted by the \in symbol) of the outcomes of the argmax operation (and does not use the equal sign).

A function $f(\cdot)$ and the logarithm of a function $\log f(\cdot)$ have the same maxima. This is due to the fact that the logarithm is a monotonic function (a monotonically increasing function). The log of a product of logarithms is equal to the sum of the individual logarithms.

▼ Definition 2.34 — *maximum log-likelihood*

Maximum log-likelihood is defined as:

$$\theta^* \in \operatorname{argmax}_{\theta} \sum_{i=0}^{k-1} \log p(x_i | \theta). \quad (2.6)$$

In the case we have information about the parameters θ we can model this with a probability distribution.

▼ Definition 2.35 — *prior probability distribution*

A **prior probability distribution** defines a probability distribution $p(\theta)$ to parameters θ without a dependency on the observations x .

Given the definition of a prior probability distribution, we can define *maximum a posteriori* estimation.

▼ Definition 2.36 — *maximum a posteriori*

maximum a posteriori estimation is defined as:

$$\theta^* \in \operatorname{argmax}_{\theta} \sum_{i=0}^{k-1} \log p(x_i | \theta) + \log p(\theta). \quad (2.7)$$

If we are not only interested in the parameter θ^* that maximizes $p(x|\theta)$ and $p(\theta)$, but in the complete distribution for $p(\theta)$ we need Bayes' theorem described by Laplace (1820).

▼ **Definition 2.37 — Bayesian inference**

Bayesian inference using Bayes' theorem is defined as:

$$f(\theta|x) = p(\theta|x) = \frac{\overbrace{p(x|\theta)}^{\text{likelihood}} \overbrace{p(\theta)}^{\text{prior}}}{\underbrace{p(x)}_{\text{normalization constant}}} = \frac{p(x|\theta)p(\theta)}{\int p(x|\theta)p(\theta)d\theta}. \quad (2.8)$$

Bayes' theorem describes the posterior probability $p(\theta|x)$ as the likelihood times the prior probability distribution divided by a normalization constant, also called the evidence. The normalization constant is not a function of the parameters θ . If a function is known except for the normalization constant, it is indicated by the “proportional to” symbol \propto .

$$f(\theta|x) \propto p(x|\theta)p(\theta) \quad (2.9)$$

In Bayesian inference $p(\theta|x)$ is calculated. In contrast, in maximum likelihood and maximum a posteriori only parts of Eq. 2.8 are calculated, respectively $p(x|\theta)$ and $p(x|\theta)p(\theta)$. In Section 2.3 inference methods will be described that approximate Bayesian inference. Approximation is required in the case closed-form expressions are not available. If the inference task only requires maximum a posteriori, approximation methods are also available (Daume, 2007), but this is outside of the scope of the current thesis.

There are two supervised learning models, a generative model and a discriminative model. Below we provide their definitions and in Figure 2.2 we give three examples for each model.

▼ **Definition 2.38 — generative model**

A **generative** model defines the joint probability distribution $p(x, \theta)$ and uses Bayes rule to define $p(x|\theta)$.

▼ **Definition 2.39 — discriminative model**

A **discriminative** model defines the conditional probability distribution $p(x|\theta)$ directly.

Figure 2.2 shows three generative and three discriminative models. They are chosen for their structure; from left to right, in both cases, the structure is between the random variables is enriched, first in the form of a sequence structure, then in the form of a graph structure. Figure 2.2 visualizes three generative models: (1) the Naive Bayes Model (Russell et al., 1995), (2) the Hidden Markov Model (Baum and Petrie, 1966), and (3) the Directional

Model (Koller and Friedman, 2009). It shows also three discriminative models: (1) Logistic Regression, (2) Linear-chain Conditional Random Fields, and (3) general Conditional Random Fields.

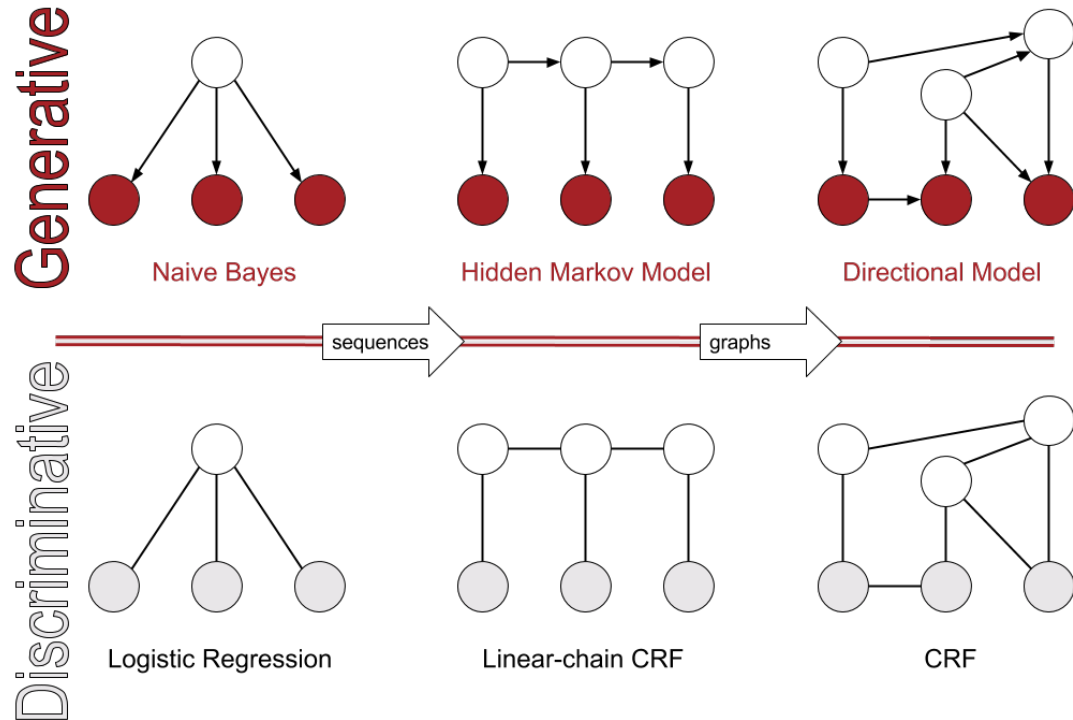


Figure 2.2: Generative models: Naive Bayes Model, Hidden Markov Model, and Directional Model. Discriminative models: Logistic Regression, Linear-chain Conditional Random Fields, and general Conditional Random Fields. Figure adapted from Sutton and McCallum (2011).

There is no definitive reason to use a generative model rather than a discriminative model or vice-versa. Here we confine ourselves to two remarks. First, a discriminative model has lower asymptotic error, but a generative model approaches its asymptotic error faster in the case of a Naive Bayes classifier versus Logistic Regression (Jordan, 2002). However, Xue and Titterton (2008) doubt the existence of such precisely defined regimes. According to them the asymptotic error denotes the error with an increasing number of samples. Second, the prior $p(\theta)$ in the generative model provides a principled way to handle missing information, while the direct modeling of decision boundaries in a discriminative model often leads to better performance in a classification task (Jaakkola et al., 1999). Apart from generative models and discriminative models, there are also hybrid models (Bouchard and Triggs, 2004; Raina et al., 2003; Bosch et al., 2008). In the thesis we will limit ourselves to generative models.

2.1.3 Model Composition

A model can be composed out of a set of probability distributions. We list three of such possible compositions. The Naive Bayes model is a *product* of probability distributions with a

prior distribution (Definition 2.40). The finite mixture model is a *sum* over a finite number of probability distributions where each one is weighted (Definition 2.41). The infinite mixture model is a *sum* over an infinite number of probability distributions where each one is weighted (Definition 2.42).

▼ **Definition 2.40 — naive Bayes model**

The **naive Bayes model** is a product over a finite number $k \neq \infty$ of probability distributions $p(x_i|\theta)$ multiplied by the prior distribution $p(\theta)$:

$$p(\theta|x) \propto p(\theta) \prod_{i=0}^{k-1} p(x_i|\theta). \quad (2.10)$$

A finite mixture model is a sum over a finite number of probability distributions.

▼ **Definition 2.41 — finite mixture model**

A **finite mixture model** is a sum over a finite number $k \neq \infty$ of probability distributions $p(x_i)$, with each distribution weighted by a factor w_i with $\sum_i w_i = 1$.

$$p(x) = \sum_{i=0}^{k-1} w_i p(x_i). \quad (2.11)$$

The mixture model is finite in the sense that there are only $k \neq \infty$ distributions summed up. The weights of the individual distributions $p(x_i)$ are normalized (sum up to one) such that the weighted sum over the probability distributions is itself a probability distribution.

An infinite mixture model is a sum over an infinite number of probability distributions.

▼ **Definition 2.42 — infinite mixture model**

A **infinite mixture model** is a sum over an infinite number of probability distributions $p(x_i)$, with each distribution weighted by a factor w_i with $\sum_i w_i = 1$.

$$p(x) = \sum_{i=0}^{\infty} w_i p(x_i). \quad (2.12)$$

The infinite mixture model is a sum over an infinite number of probability distributions with weights that sum up to one. In this way it assigns a finite value to a countably infinite set of functions. In the thesis we will encounter infinite mixture models in Chapters 3 and 4.

A mixture model relates one latent cause (cluster) to each data point. A factorial model or feature allocation model assigns multiple factors to a data point.

▼ **Definition 2.43 — *feature model***

A **feature model** is a sum over a number k of sets of probability distributions $p(x_{C_i})$.

$$p(x) = \sum_{i=0}^{k-1} \sum_{j=0}^{C_i} p(x_j). \quad (2.13)$$

A counting model allows each feature to occur multiple times.

▼ **Definition 2.44 — *counting model***

A **counting model** is a sum over a number k of sets of probability distributions $p(x_{C_i})$ with each feature occurring a number c_j of times.

$$p(x) = \sum_{i=0}^{k-1} \sum_{j=0}^{C_i} c_j p(x_j). \quad (2.14)$$

2.1.4 General Random Elements

In section 2.1.1 random elements were described in general. Random elements can vary from random vectors, random distributions, random clusters (partitions), to random trees. Table 2.1 describes the random elements and the corresponding examples of random processes in the literature. Below we mention them with the appropriate references.

The Gaussian Process (Rasmussen and Williams, 2006) describes a distribution on functions. The Beta Process (Hjort, 1990), the Gamma Process (Ferguson, 1974), the Dirichlet Process and the Polya Tree (Ferguson, 1973) describe a distribution on distributions. The Chinese Restaurant Process (Aldous, 1985) and Pitman-Yor Process (Pitman and Yor, 1997) describe a distribution on partitions (in the form of cluster assignments). The Stick-breaking Process describes a distribution on partition sizes (with no information on assignments themselves). The Dirichlet Diffusion Tree (Neal, 2001) and Kingman's coalescence (Kingman, 1965) describe a distribution on hierarchical partitions. The Indian Buffet Process (Ghahramani and Griffiths, 2005) describes a distribution over sparse binary matrices. The Gamma-Poisson Process (Titsias, 2008) describes a distribution over integer-valued matrices. The Mondrian Process (Roy and Teh, 2009) describes a distribution over kd-trees.

Table 2.1: A list of seven mathematical structures and for each of these structures one or more random processes that can generate the structure. For example, a distribution on distributions can be generated by a Beta Process, Gamma Process, Dirichlet Process, or a Polya Tree.

Structure	Example
Distribution on functions	Gaussian Process
Distribution on distributions	Beta Process
	Gamma Process
	Dirichlet Process
	Polya Tree
Distribution on partition assignments	Chinese Restaurant Process Pitman-Yor Process
Distribution on partition sizes	Stick-breaking Process
Distribution on hierarchical partitions	Dirichlet Diffusion Tree Kingman's coalescence
Distribution on sparse binary matrices	Indian Buffet Process
Distribution on integer-valued matrices	Gamma-Poisson Process
Distribution on kd-trees	Mondrian Process

2.1.5 Plate Notation

Random processes and mixture models are visually represented by a method called *plate notation* (cf. Buntine, 1994; Koller and Friedman, 2009). Sets of variables are represented in a plate, a rectangular region (see Figure 2.3).

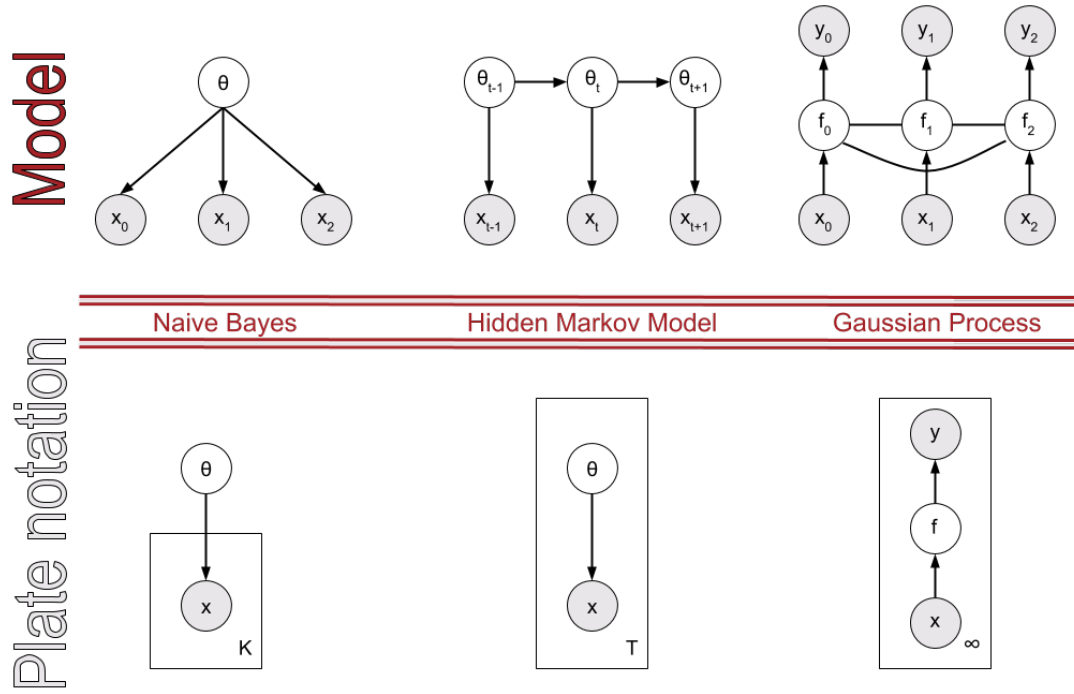


Figure 2.3: Top: graphical model of a Naive Bayes, hidden Markov model, and Gaussian process. Bottom: corresponding plate notation of the Naive Bayes, hidden Markov model, and Gaussian process. Observed variables are denoted by a circle that is shaded.

854 Plate notation is a representation that does not preserve all dependencies between variables.
 855 For example, the dependencies between the states in the Hidden Markov Model (e.g., be-
 856 tween θ_0 and θ_1) are not represented.

857 2.1.6 Completely Random Measure and Lévy Measure

858 Some random process are mathematically represented by a completely random measure
 859 (Kingman, 1967), which is defined as follows.

860 ▼ Definition 2.45 — *completely random measure*

862 A **completely random measure** is a random measure $\mu : \Omega \times X \rightarrow [0, +\infty]$ from
 863 probability space $(\Omega, \mathbb{F}, \mathbb{P})$ to measurable space (X, Σ) with

- 864 ◦ for any collection of disjoint sets $A_1, \dots, A_k \in \Sigma$ and $A_i \cap A_j = \emptyset$ for $i \neq j$ a mutual
 865 independency between $\mu(A_1), \dots, \mu(A_k)$.

867 Kingman (1967) shows that a completely random measure can be decomposed into three
 868 components:

- 869 1. a deterministic function;

2. a countable set of non-negative random masses at deterministic locations;
3. a countable set of non-negative random masses at random locations.

The first component is a deterministic function. The second component has non-negative random masses, also called atoms, on deterministic locations. The third component is the one of interest. It has a set of random masses (atoms) that can be represented as a Poisson random measure on $\mathbb{R}^+ \otimes X$ with mean measure ν which is known as the Lévy intensity measure (Favaro et al., 2013).

Table 2.2: Lévy measure of the Beta Process (Wang and Carin, 2012), Gamma Process (Knowles et al., 2014), the Dirichlet Process (Lijoi and Prünster, 2010) (indirectly through $F = 1 - e^{-\nu}$).

Random Process	Lévy measure
Beta Process	$\nu(da, dw) = H(da)\alpha w^{-1}(1-w)^{\alpha-1}dw$
Gamma Process	$\nu(da, dw) = H(da)w^{-1}e^{-\alpha w}dw$
Dirichlet Process	$\nu(da, dw) = H(da)e^{-w\alpha(x, \infty)}(1-e^{-w})^{-1}dw$

For Lévy measure decompositions of other processes such as the Indian buffet process, we refer to Wang and Carin (2012).

2.1.7 Exchangeability

Here we recall Definition 2.31 for exchangeable sequences. De Finetti's theorem states that there is parameter θ such that the data x_i is conditionally independent given this parameter for exchangeable sequences (cf. De Finetti, 1937).

▼ Definition 2.46 — De Finetti's theorem

A sequence $\{x_0, x_1, \dots\}$ of (X, Σ_X) -valued random variables is an infinitely exchangeable sequence if and only if there exist a measure $\mu(d\theta)$ on θ such that

$$p(x_0, \dots, x_{k-1}) = \int_{\Sigma_X(X)} \prod_{i=0}^{k-1} p(x_i|\theta) \mu(d\theta) \quad \forall k \geq 1. \quad (2.15)$$

In words, de Finetti's theorem states that if we have *exchangeable* data, we have a parameter θ , a likelihood $p(x|\theta)$, and some measure μ on θ , such that the data (x_0, \dots, x_{k-1}) is *conditionally independent*. Hence, although the data is not i.i.d., there are underlying, unobservable, quantities that are i.i.d. and exchangeable sequences are mixtures of these quantities. The theorem proofs that if the observations are exchangeable, they must be a random sample from some model and there must exist a prior probability distribution over the parameters of that model, hence requiring a Bayesian approach.

896 The theorem is not limited to exchangeable *sequences*. In contrast, there are similar theo-
 897 rems for other exchangeable objects (Orbanz and Roy, 2015). Five examples (see Table 2.3)
 898 of exchangeable structures that have a theorem that describes an underlying measure that
 899 can be sampled i.i.d. are: (1) exchangeable sequences (de Finetti, 1930), (2) increments
 900 (Bühlmann, 1960), (3) partitions (Kingman, 1978), (4) arrays (Aldous, 1981), and (5)
 901 Markov chains (Diaconis and Freedman, 1980).

Table 2.3: Five exchangeable structures and their theorems.

Mathematical Object	Theorem
Exchangeable Sequence	de Finetti
Exchangeable Increment	Bühlmann
Exchangeable Partition	Kingman
Exchangeable Array	Aldous-Hoover
Exchangeable Markov Chain	Diaconis-Freedman

902 2.1.8 Stick-breaking Representation

903 Now we introduce the *stick-breaking representation* by Freedman and Diaconis (1983), also
 904 known as the residual allocation model (Sawyer and Hartl, 1985; Hoppe, 1986).

▼ Definition 2.47 — *stick-breaking*

An infinite sequence of random variables $\phi = \{\phi_0, \phi_1, \dots\}$ has a **stick-breaking representation** with parameters α and β denoted by $\phi \sim GEM(\alpha, \beta)$.

$$w_k \stackrel{i.i.d.}{\sim} \text{Beta}(1 - \beta, \alpha + k\beta) \quad k = 1, \dots, K \quad (2.16)$$

$$\phi_k = w_k \prod_{i=1}^{k-1} (1 - w_i) \quad (2.17)$$

905

906

907 The stick-breaking process samples repeatedly from a $\text{Beta}(1 - \beta, \alpha + k\beta)$ distribution. The
 908 result of the process is a vector of k weights ϕ_k . The abbreviation *GEM* stands for Griffiths,
 909 Engen, and McCloskey (Ewens, 1990; Ethier, 1990). There is also a variant of GEM with a
 910 single parameter α which can be obtained by setting $\beta = 0$. In that case w_k are drawn from
 911 a $\text{Beta}(1, \alpha)$ distribution.

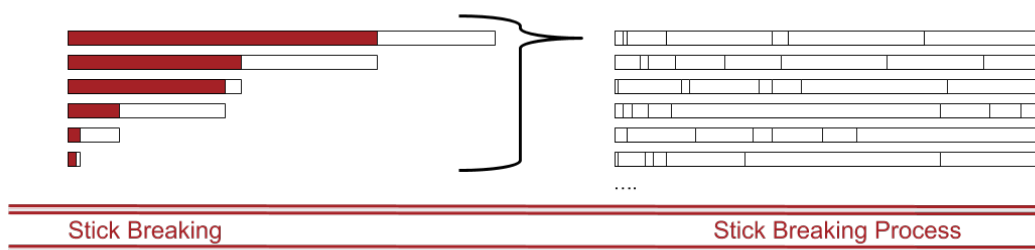


Figure 2.4: The stick-breaking representation. Left: at the first row, the stick is broken at x_0 , at the next rows the remaining part of the stick is broken x_i with $i > 0$. Only six iterations are shown. Right: samples of a stick-breaking process. The first row shows the stick ratios from the stick-breaking representation at the left. The next rows show other samples from the same process.

Figure 2.4 visualizes the stick-breaking process. A stick of fixed length 1 gets broken at a position w_0 sampled from a Beta distribution. The remainder of the stick is broken again at position $w_1(1 - w_0)$. This process continues for an infinite number of times. In this manner generates a stick-breaking process a sequence of non-negative values that sum up to one. The stick-breaking representation can on itself give rise to more sophisticated stochastic processes (Dunson et al., 2012). Computationally it can also fulfill a useful rule. Namely, it is possible to approximate a distribution over partitions by truncating a stick-breaking process. The stick-breaking is then only performed a limited number of times (Kurihara et al., 2007).

In Section 2.2.1 the relevance of the stick-breaking process for the Dirichlet process will be shown. In that case the values generated by the stick-breaking process represent the weights of the partitions induced by the Dirichlet Process.

2.2 Five Random Processes

In this section we investigate five random processes. The Dirichlet process (Section 2.2.1), the Beta process (Section 2.2.2), the Gamma process (Section 2.2.3), the Pitman-Yor process (Section 2.2.4), and the hierarchical Dirichlet process (Section 2.2.5). We compare the random processes in Section 2.2.6.

2.2.1 Dirichlet Process

The Dirichlet process is presented as a measure (Section A), is shown to have a sequential representation in the form of the Chinese restaurant process (Section B), and a stick-breaking representation (Section C).

933 A: Dirichlet Process as a Measure

934 The Dirichlet process (DP) is a distribution over distributions (Ferguson, 1973).

▼ Definition 2.48 — Dirichlet process

A Dirichlet process DP over a set S can be used to draw sample paths X :

$$X \sim DP(\alpha, H)$$

with α the dispersion parameter and H a measure on S and for which any measurable partition $\{B_0, \dots, B_{n-1}\} \in S$ is drawn from a Dirichlet distribution:

$$(X(B_0), \dots, X(B_{n-1})) \sim \text{Dirichlet}(\alpha H(B_0), \dots, \alpha H(B_{n-1}))$$

935

936

937 The Lévy intensity of the Dirichlet process is complicated, because it is a so-called normalized
938 process, see Regazzini et al. (2003).

939 The Dirichlet process can be used as a prior for a mixture model (Definition 2.42). This is
940 visualized in (Figure 2.5).

	θ_0	θ_1	θ_2	θ_3	θ_4	θ_5	θ_6	θ_7	θ_8	θ_9	sum
data 0	0	0	1	0	0	0	0	0	0	...	1
data 1	1	0	0	0	0	0	0	0	0	...	1
data 2	1	0	0	0	0	0	0	0	0	...	1
data 3	0	1	0	0	0	0	0	0	0	...	1
data 4	1	0	0	0	0	0	0	0	0	...	1
data 5	0	1	0	0	0	0	0	0	0	...	1
data 6	0	0	0	0	1	0	0	0	0	...	1

Figure 2.5: Matrix representation of a mixture model. At the horizontal axis the latent variables (potentially infinite clusters). At the vertical axis the data items. The rows sum up to one. A data item is only assigned to one cluster.

941 B: Chinese Restaurant Process

942 De Finetti's theorem (Definition 2.46) can be used to establish the existence of an infinitely
943 exchangeable sequence. In the particular case of the Dirichlet process the sequence is an
944 exchangeable *distribution over partitions* and is called the CRP.

▼ **Definition 2.49 — Chinese restaurant process**

A **Chinese restaurant process** is a sequential process that is an exchangeable distribution over partitions:

$$p(z_i = k | z_0, \dots, z_{i-1}) = \begin{cases} \frac{n_k}{\alpha + i} & \text{if } k \leq K_+ \\ \frac{\alpha}{\alpha + i} & \text{if } k > K_+ \end{cases} \quad (2.18)$$

The conditional probability of a cluster assignment z_i for data item y_i given the cluster assignments z_0, \dots, z_{i-1} is proportional to the number of data items n_k assigned to an existing cluster k , or proportional to α for a new cluster.

The Chinese Restaurant Process is visualized in Figure 2.6.




Figure 2.6: The Chinese Restaurant Process with i customers already sitting down. A new customer $y_{i=7}$ arrives and gets assigned, z_i . This is an existing table $\{1, 2, 3\}$ with a probability proportional to the number of customers n_i sitting at that table: $n_i/(\alpha + i)$, or a new, empty table 4 with probability $1/(\alpha + i)$. In the visualized Chinese restaurant process the dispersion factor $\alpha = 1$.

C: Stick-breaking Representation of the Dirichlet process

▼ **Definition 2.50 — stick-breaking representation of the Dirichlet process**

The **stick-breaking representation** of the Dirichlet process states that if

$$\phi_k \sim GEM(\alpha, 0) \quad (2.19)$$

$$\theta_k \sim H \quad (2.20)$$

$$G = \sum_{k=1}^{\infty} \phi_k \delta(\theta, \theta_k) \quad (2.21)$$

then $G \sim DP(\alpha, H)$.

The weights ϕ_k are sampled from the stick-breaking process $GEM(\alpha, 0)$ (see Definition 2.47). The parameter values θ_k are sampled from the base measure H . To sample from the Dirichlet Process we have to sample these parameters with the given weights.

If the stick-breaking process is used as a prior for a mixture, then the cluster assignments z_i are sampled according to the mixing proportions ϕ :

$$\phi \sim GEM(\alpha, 0) \quad (2.22)$$

$$\theta_k \sim H \quad (2.23)$$

$$G = \sum_{k=1}^{\infty} \phi_k \delta(\theta, \theta_k) \quad (2.24)$$

$$z_i \sim Mult(\phi) \quad (2.25)$$

$$x_i \sim F(\theta_{z_i}) \quad (2.26)$$

963 Here $\theta_k = \theta_{z_i}$ for observation with index i and cluster assignment k : $z_i = k$.

964 The Dirichlet process, the Dirichlet process as prior for a mixture model, the Chinese restaurant
965 process and the stick-breaking representation can be compared in Figure 2.7.

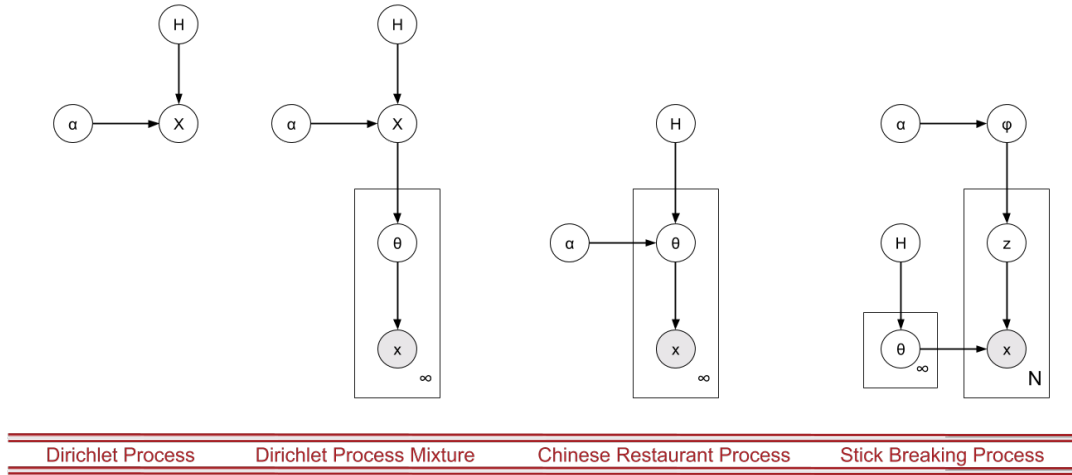


Figure 2.7: From left to right: (1) The Dirichlet process $X \sim DP(\alpha, H)$; (2) The Dirichlet mixture model with X the prior for a sum $\sum_i w_i p(x_i | \theta_i)$; (3) the Chinese restaurant process with X marginalized out; and (4) the Stick-breaking process with a distribution over partition sizes π and indicator variables z_i .

966 2.2.2 Beta Process

967 The Beta process is presented as a measure (Section A), is shown to have a sequential rep-
968 resentation in the form of the Indian buffet process (Section B), and a stick-breaking repre-
969 sentation (Section C).

970 A: Beta Process as a Measure

971 A Beta process (Hjort, 1990) is a random process with a countably infinite collection of
972 weighted atoms in a space (X, \mathbb{B}) with weights that are in between $[0, 1]$.

▼ Definition 2.51

Let (X, \mathbb{B}) be a Borel space, ν a finite measure, and $\alpha > 0$ a scale parameter, then a **Beta process** is a Lévy process on (X, \mathbb{B}) with its Lévy measure ν corresponding to the density:

$$\nu(dw) = \alpha w^{-1}(1-w)^{\alpha-1}dw \quad (2.27)$$

with $w > 0$.

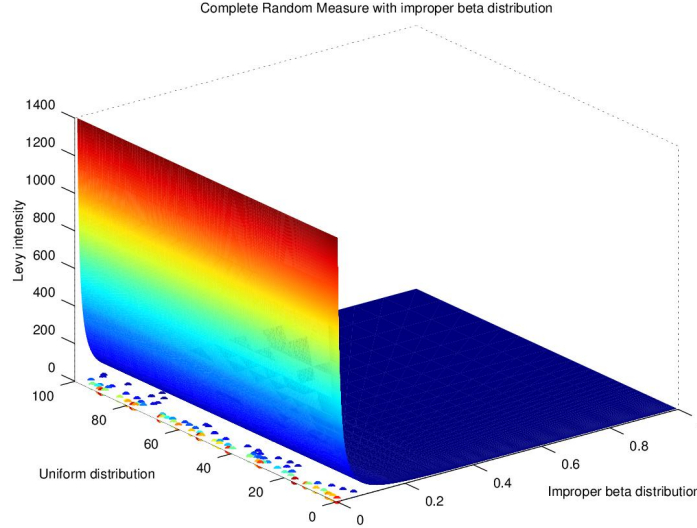


Figure 2.8: A Completely Random Measure with a Lévy intensity defined on the product space $\Omega \otimes (0, 1)$. Here Ω is a bounded interval on which the base measure $B_0 = U(0, 100)$ is defined. On $(0, 1)$ we define an improper beta distribution $\alpha w^{-1}(1-w)^{\alpha-1}$. In this example $\alpha = 10$. This is how a Beta Process can be generated from a nonhomogeneous spatial Poisson point process. This has been visualized before (Jordan, 2010). The image is produced by rejection sampling using a homogeneous Poisson point process at $\max(\nu)$ over $w = [0.01, 0.9]$. For $w \rightarrow 0$ this maximum would go to ∞ and all points would be rejected. Hence, the points should be denser for w around 0 and should be seen as an approximation of the actual process.

In Figure 2.8 the Beta Process is generated from a Completely Random Measure (see Section 2.1.6) with a Lévy intensity defined on $\Omega \otimes (0, 1)$ (Thibaux and Jordan, 2007). In this case Ω is the so-called base measure B_0 and is assumed uniform over a bounded region. The $(0, 1)$ space is equipped with an improper Beta distribution. It is called improper or degenerate because the scale parameter of the standard Beta distribution is set to zero. This has the consequence that the integral is infinite: $\nu(\Omega \otimes (0, 1)) = \infty$. It is due to the fact that the density $w^{-1}(1-w)^{\alpha-1}$ goes to infinity for $w \rightarrow 0$. That means that a countable infinite number of points can be obtained from the Poisson process.

The Beta process can be used as a prior for a feature or factorial model (Definition 2.43). This is visualized in (Figure 2.9).

	θ_0	θ_1	θ_2	θ_3	θ_4	θ_5	θ_6	θ_7	θ_8	θ_9	sum
data 0	1	0	1	0	0	0	0	0	0	...	2
data 1	1	1	0	0	1	0	0	0	0	...	3
data 2	1	0	0	0	0	0	0	0	0	...	1
data 3	0	0	0	0	0	0	1	1	1	...	3
data 4	0	1	0	1	1	0	0	0	1	...	4
data 5	0	0	0	0	1	0	0	0	1	...	2
data 6	0	0	1	0	0	0	0	0	0	...	1

Figure 2.9: Matrix representation of a factorial model. At the horizontal axis the latent variables (potentially infinite number of features). At the vertical axis the data items. The rows sum up to nonnegative integers. A data item has multiple features.

990 B: Indian Buffet Process

991 The Beta process has a sequential representation in the form of the Indian Buffet Process:

992 ▼ Definition 2.52 — *Indian buffet process*

993
994 An **Indian buffet process** is a sequential process that is an exchangeable distribution
995 over sparse binary matrices:

$$p(z_{i,j} = k | z_{0,0}, \dots, z_{i-1,K_+}) = \begin{cases} \frac{n_{-i,k}}{i} & \text{if } k \leq K_+ \\ \frac{\lambda^{k_{new}} e^{-\lambda}}{k_{new}!} & \text{if } k > K_+ \end{cases} \quad (2.28)$$

996

997

998 Here $\lambda = \alpha/i$, $k_{new} = K_+ - k$. The i 'th data item samples an existing column with a probability
999 of the number of times it has been sampled before divided by its index, $n_{-i,k}/i$. It samples
1000 a new column with a probability according to a Poisson distribution, $\lambda^{k_{new}} e^{-\lambda}/k_{new}!$. The
1001 conditional form of the sequential presentation describes a closed-form solution for Gibbs
1002 sampling, section 2.3.4 (Ghahramani and Griffiths, 2005).

1003 C: Stick-breaking Representation of the Beta process

▼ Definition 2.53 — *stick-breaking representation of the Beta process*

The **stick-breaking representation** of the Beta process states that if

$$\phi_{k,j} \sim GEM(\alpha, 0) \quad (2.29)$$

$$C_k \sim \text{Poisson}(\gamma) \quad (2.30)$$

$$\theta_{k,j} \sim \frac{1}{\gamma} H \quad (2.31)$$

$$G = \sum_{k=1}^{\infty} \sum_{j=1}^{C_k} \phi_{k,j} \delta(\theta, \theta_{k,j}) \quad (2.32)$$

then $G \sim BP(\alpha, H)$.

The Beta process is used in linguistics (He et al., 2013; Vanhainen and Salvi, 2012), computer vision (Zhou et al., 2011; Gao and Sun, 2013), risk assessment (Li et al., 2014), and medicine (Ross et al., 2014).

2.2.3 Gamma Process

The Gamma process is presented as a measure (Section A), is shown to have a sequential representation in the form of a multi-scoop Indian buffet process (Section B), and a stick-breaking representation (Section C).

A: Gamma Process as a Measure

A Gamma process is a random process with independent gamma distributed increments (Ferguson, 1974). Below we provide a formal definition.

▼ Definition 2.54 — Gamma process

Let (X, \mathbb{B}) be a Borel space, ν a finite measure, and $\alpha > 0$ a scale parameter, then a **Gamma process** is a Lévy process on (X, \mathbb{B}) with its Lévy measure ν corresponding to the density:

$$\nu(dw) = w^{-1} e^{-\alpha w} dw \quad (2.33)$$

with $w > 0$.




Figure 2.10: A Completely Random Measure with a Lévy intensity defined on the product space $\Omega \otimes \mathbb{R}$. Here Ω is a bounded interval on which the base measure $B_0 = U(0, 100)$ is defined. On \mathbb{R} we define an improper gamma distribution $w^{-1}e^{-\alpha w}$. In this example $\alpha = 1$. This is how a Gamma Process can be generated from a nonhomogeneous spatial Poisson point process. The image is produced by rejection sampling in the same way as Figure 2.8.

1023 The Gamma process can be used as a prior for a counting model (Definition 2.44). This is
 1024 visualized in (Figure 2.11).

	θ_0	θ_1	θ_2	θ_3	θ_4	θ_5	θ_6	θ_7	θ_8	θ_9	sum
data 0	4	0	3	0	0	0	0	0	0	...	7
data 1	1	1	0	0	1	0	0	0	0	...	3
data 2	2	0	0	0	0	0	0	0	0	...	2
data 3	0	0	0	0	0	4	1	0	0	...	5
data 4	0	0	0	2	0	0	0	0	0	...	2
data 5	0	0	0	0	0	4	0	0	6	...	10
data 6	0	0	0	0	0	4	0	0	0	...	4

Figure 2.11: Matrix representation of a counting model. At the horizontal axis the latent variables (potentially infinite number of features). At the vertical axis the data items. Each data item has one or more features and these features can occur multiple times.

1025 B: Multi-Scoop Indian Buffet Process

1026 Find sequential representation. It is an Indian Buffet Process where the customers can pick
 1027 up multiple scoops from the same dish.

1028 (Zhou et al., 2012)

1029 C: Stick-breaking Representation of the Gamma process

1030 The stick-breaking representation of the Gamma process introduces an additional Gamma
1031 distribution that defines the number of times a feature is represented in a data object com-
1032 pared to the Beta process, see Definition 2.55, (Roychowdhury and Kulis, 2015).

▼ Definition 2.55 — *stick-breaking representation of the Gamma process*

The **stick-breaking representation** of the Gamma process states that if

$$\phi_{k,j} \sim GEM(\alpha, 0) \quad (2.34)$$

$$G_{k,j} \sim Gamma(\alpha + 1, c) \quad (2.35)$$

$$C_k \sim Poisson(\gamma) \quad (2.36)$$

$$\theta_{k,j} \sim \frac{1}{\gamma} H \quad (2.37)$$

$$G = \sum_{k=1}^{\infty} \sum_{j=1}^{C_k} G_{k,j} \phi_{k,j} \delta(\theta, \theta_{k,j}) \quad (2.38)$$

1033 then $G \sim GamP(\alpha, H)$.

1034

1035 The Gamma process is used in risk theory (Dufresne et al., 1991), spatial statistics (Wolpert
1036 and Ickstadt, 1998; Rao and Teh, 2009), erosion (Singpurwalla, 1997; Abdel-Hameed, 2012),
1037 and finance (Madan and Seneta, 1990; Küchler and Tappe, 2008).

1038 2.2.4 Pitman-Yor Process

1039 The Pitman-Yor process (PYP) introduces another parameter d with respect to the Dirichlet
1040 process. It has been developed by Pitman and Yor as the two-parameter Poisson-Dirichlet
1041 distribution (Pitman and Yor, 1997). The Pitman-Yor process has the following definition.

▼ Definition 2.56 — *Pitman-Yor process*

A Pitman-Yor process PY over a set S can be used to draw sample paths X :

$$X \sim PY(d, \alpha, H)$$

1042 with $\alpha > -d$ a strength parameter, $0 \leq d < 1$ a discount parameter, and H a measure
1043 on S .

1044

1045 The Pitman-Yor process generalizes the Dirichlet process. The Pitman-Yor process has a stick-
1046 breaking representation in which sticks are drawn from $GEM(\alpha, \beta)$. The Dirichlet process
1047 has a stick-breaking representation in which sticks are drawn from $GEM(\alpha, 0)$, see Def. 2.53.

▼ Definition 2.57

The **stick-breaking representation** of the PYP states that if

$$\phi_k \sim GEM(\alpha, \beta) \quad (2.39)$$

$$\theta_k \sim H \quad (2.40)$$

$$G = \sum_{k=1}^{\infty} \phi_k \delta(\theta, \theta_k) \quad (2.41)$$

1048 then $G \sim PYP(\alpha, H)$.

1049

1050 The Pitman-Yor process is used in quite a few applications, such as language models (Teh
1051 et al., 2006), scene segmentation (Sudderth and Jordan, 2009), speech induction (Blunsom
1052 and Cohn, 2011), and time series (Bassetti et al., 2014).

1053 2.2.5 Hierarchical Dirichlet Process

1054 The HDP extends the Dirichlet mixture model with a hierarchical structure (Teh et al., 2006).

▼ Definition 2.58

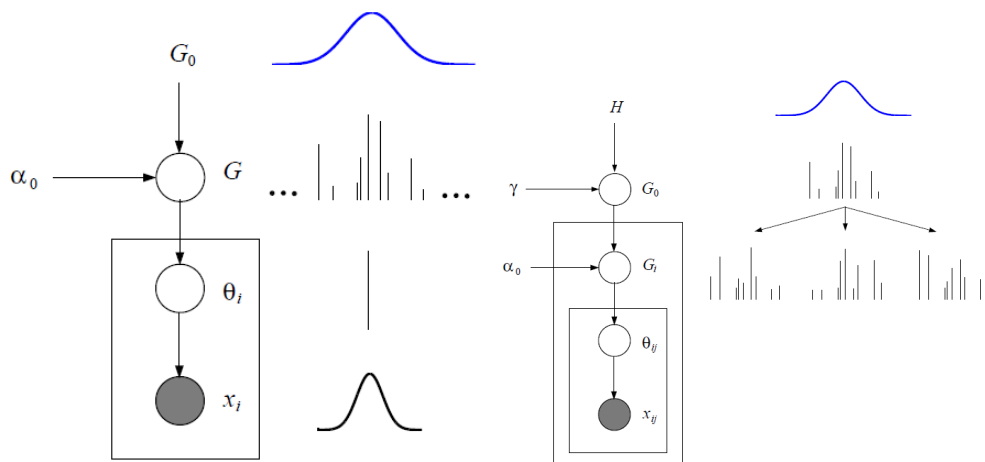
A Hierarchical Dirichlet process HDP over a set S can be used to draw sample paths X :

$$G_0 \sim DP(\gamma, H)$$

$$X_i \sim DP(\alpha_0, G_0) \text{ for each group } i$$

1055 with a Dirichlet Process with a general γ dispersion parameter and base distribution H
1056 as a measure on S of which the generated distributions G_0 are used as base distribution
1057 for each group distribution X_i .

1058



(a) Dirichlet Process Mixture Model. Each draw from the process corresponds to a parameter. Each parameter is associated with a distribution (in this case a Gaussian). (b) Hierarchical Dirichlet Process. Observe that the location of the atoms are fixed through the highest layer G_0 . The only freedom left to express by G_i is in the weights of those atoms. This reflects a decomposition in a structural and non-structural component.

Figure 2.12: The difference visualized between a Dirichlet Process mixture and a hierarchical Dirichlet process. It illustrates also that the input of a Dirichlet process does not have to be a continuous function. If it is a continuous distribution it will become a discrete distributed almost surely. If it is a discrete distribution, it will have atoms at the locations where the discrete distribution had its probability mass concentrated.

The hierarchical Dirichlet process uses the outcome of a Dirichlet Process as a starting point to define multiple distributions with atoms at the same locations, while they come equipped with different weights. So, the Dirichlet process on the lower level uses not a continuous distribution as input, but a discrete one, generated by the DP at the top layer. Note, that the Dirichlet Process will *generate* an almost surely (a.s.) discrete distribution, but it can also have a discrete distribution as *prior* H .

2.2.6 Comparison of Random Processes

The Dirichlet process (Section 2.2.1), the Beta process (Section 2.2.2), the Gamma process (Section 2.2.3), the Pitman-Yor process (Section 2.2.4), and the hierarchical Dirichlet process (Section 2.2.5) are not be compared on performance or computationally efficiency. The processes each represent different assumptions on the model structure. The Dirichlet process is suited for clustering problems where observations have to be assigned to a single class. The Beta process is a good model for applications with combinatorial structure: features that are shared among multiple objects. The Gamma process is a model for applications with counting: features are shared among multiple objects and there is a number of features per object. The Pitman-Yor process fits clustering problems where the distribution over clusters sizes obeys a power law. The hierarchical Dirichlet process is a typical example of a process that can model additional structure within a cluster.

1077 2.3 Inference

1078 There will be six inference methods described, all sampling methods. Inverse transform sam-
 1079 pling is described in Section 2.3.1. Rejection sampling in Section 2.3.2. Approximate Bayes-
 1080 ian computation in Section 2.3.3. Gibbs sampling in Section 2.3.4. Metropolis-Hastings in
 1081 Section 2.3.5. Split-Merge MCMC in Section 2.3.6. We rapport for every inference method
 1082 the corresponding algorithm in pseudo code. We compare the inference methods in Sec-
 1083 tion 2.3.7.

1084 2.3.1 Inverse Transform Sampling

1085 Let $p(x)$ be a discrete probability distribution with two possible values $x = f$ and $x = g$.
 1086 The probability distribution sums up to one: $\sum_v p(x = v) = 1$. Sample from a uniform
 1087 distribution $u \sim U(0, 1)$. If $u < p(x = f)$ generate f , else generate g . This procedure
 1088 samples f with probability $p(x = f)$ and g with probability $p(x = g)$. This can be readily
 1089 generalized to more than two values by making use of the cumulative distribution function.
 1090 In Algorithm 1 we sample from $f(x)$ by making use of the inverse cumulative distribution.

Algorithm 1 Inverse transform sampling for $f(x)$

```

1: procedure INVERSE TRANSFORM SAMPLING( $f(x)$ )           ▷ Distribution to sample from.
2:    $F(x) = P(X \leq x) \quad \forall x \in X$                      ▷ Create cumulative distribution function  $F(x)$ .
3:   for  $t = 1 \rightarrow T$  do
4:      $u \sim U(0, 1)$                                        ▷ Sample from uniform distribution.
5:      $x \sim F^{-1}(u)$                                        ▷ Sample  $x$  from (the inverse)  $F^{-1}(x)$ .
6:      $X = X \cup x$ 
7:   end for
8:   return  $X$                                              ▷  $X$  will have the distribution of  $f(x)$ .
9: end procedure

```

1091 The term “inverse” stems from the fact that we return x (or $f(x)$) given u . Inverse trans-
 1092 form sampling is a common component in sampling methods. When one of the steps in
 1093 an algorithm samples from a uniform distribution, it is often an inverse transform sampling
 1094 step.

1095 2.3.2 Rejection Sampling

1096 Let $f(x)$ be a complicated function from which it is hard to take samples. Let $g(x)$ be a
 1097 simple function that is easy to sample from. Then we can sample from $f(x)$ by making
 1098 sure $Mg(x) \geq f(x)$. The function $Mg(x)$ is an *envelope* function. This sampling method S
 1099 generates the sample set X using $f(x)$ and $g(x)$.

$$X = S(f(x), g(x)) \quad (2.42)$$

1100 The rejection sampling method (Halperin and Burrows, 1960) for $f(x)$ is described in Algo-
 1101 rithm 2.

Algorithm 2 Rejection sampling for $f(x)$

```

1: procedure REJECTION SAMPLING( $f(x), g(x)$ )      ▷ Target and proposal distribution.
2:   for  $t = 1 \rightarrow T$  do
3:      $x^t \sim g(x)$                                 ▷ Generate  $x^t$  from  $g(x)$ 
4:      $u \sim U(0, 1)$                                 ▷ Inverse transform sampling
5:      $p_0 = f(x)/(Mg(x))$ 
6:     if  $u < p_0$  then
7:        $X = X \cup x^t$                                 ▷ Accept
8:     end if
9:   end for
10:  return  $X$                                 ▷  $X$  will have the distribution of  $f(x)$ 
11: end procedure

```

1102 We can use rejection sampling to *sample* from the *posterior* $f(\theta|x)$ given that we know
 1103 the *exact* likelihood function and that we can *sample* from the prior. We know that we can
 1104 sample from the posterior by sampling from $p(\theta)p(x|\theta)$. Moreover, we know that the prior
 1105 $p(\theta)$ necessarily has to be larger than $p(\theta)p(x|\theta)$ for any observation, because $p(x|\theta)$ is a
 1106 probability density function, hence for each x and θ it is smaller than one. Hence we can
 1107 use rejection sampling with $Mg(x) \geq f(x)$ with $M = 1$, $p(\theta) = g(x)$ and $p(x|\theta) = f(x)$.

1108 We introduce the following notation. We make explicit that we need $p(x|\theta)$ for each com-
 1109 bination of observations and parameters, but that we only need to *sample* from the prior,
 1110 which we indicate by a tilde, $\sim p(\theta)$.

$$\Theta = S(\sim p(\theta), p(x|\theta), x) \quad (2.43)$$

Algorithm 3 Rejection sampling for $f(\theta|x)$

```

1: procedure REJECTION SAMPLING( $p(\theta), p(x|\theta), x$ )  ▷ Requires prior, likelihood and
   observations.
2:   for  $t = 1 \rightarrow T$  do
3:      $\theta^t \sim p(\theta)$                                 ▷ Generate  $\theta^t$  from prior
4:      $u \sim U(0, 1)$                                 ▷ Inverse transform sampling
5:      $p_0 = p(x|\theta)$ 
6:     if  $u < p_0$  then
7:        $\Theta = \Theta \cup \theta^t$                                 ▷ Accept
8:     end if
9:   end for
10:  return  $\Theta$                                 ▷  $\Theta$  will have the distribution of  $f(\theta|x)$ 
11: end procedure

```

1111 In Algorithm 3 the envelope distribution $p(\theta)$ and the target distribution $p(\theta)p(x|\theta)$, cancel
 1112 in such way that only $p(x|\theta)$ remains.

1113 Most examples illustrate rejection sampling by estimating the area of a circle, but let us
 1114 visualize the method in the context of sampling (Figure 2.13).

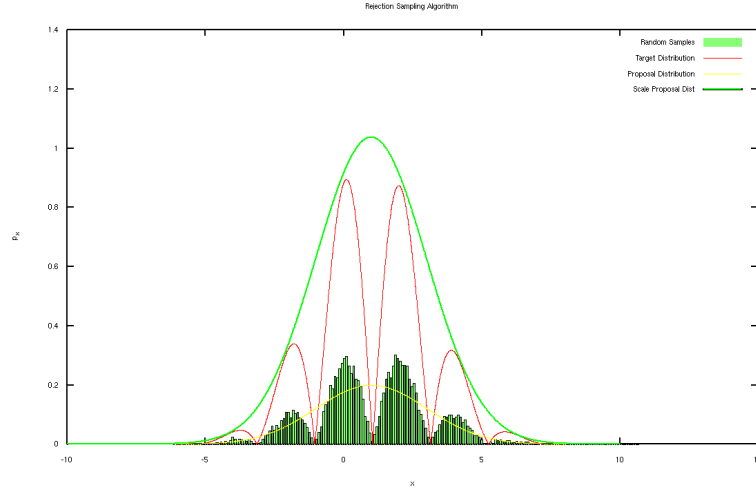


Figure 2.13: A Gaussian is placed over the complex target probability density function. Subsequently the samples that fall in between these two ‘envelopes’ are rejected. This results in a sampling scheme that follows exactly the more complicated probability density function. Note that if the function is scaled by a factor, the sampling scheme stays the same. Such a scaling factor is only important if we want, for example, to know the area under the graph.

1115 2.3.3 Approximate Bayesian Computation

1116 In approximate Bayesian computation (ABC) (Rubin and Others, 1984) the likelihood func-
 1117 tion does not need to be calculated² (Sisson and Fan, 2011). In contrast, it is assumed
 1118 that there is a model available that allows to generate observations given the (searched for)
 1119 parameters. In ABC for each configuration of parameters a set of observations is generated.

$$\Theta = S(\sim p(\theta), X, \sim M(\theta), d(X^t, X), \epsilon) \quad (2.44)$$

1120 Approximate Bayesian computation uses many tuning parameters. Its most salient charac-
 1121 teristic though, is that it generates pseudo-observations through $M(\theta)$ (see Algorithm 4).

²ABC is also called likelihood-free computation

Algorithm 4 Approximate Bayesian computation

```

1: procedure APPROXIMATE BAYESIAN COMPUTATION( $p(\theta), X, M, d, \epsilon$ )    ▷ Requires prior,
   observations, model, distance function, and threshold.
2:   for  $t = 1 \rightarrow T$  do
3:      $\theta^t \sim p(\theta)$                                               ▷ Generate  $\theta$  from prior
4:      $X^t \sim M(\theta)$                                               ▷ Simulate observations  $X^t$  from model  $M$ 
5:      $\rho = d(X^t, X)$     ▷ Calculate distance between simulated and actual observations
6:     if  $\rho \leq \epsilon$  then
7:        $\Theta = \Theta \cup \theta^t$     ▷ Accept  $\theta^t$  if distance falls under threshold  $\epsilon$ .
8:     end if
9:   end for
10:  return  $\Theta$     ▷  $\Theta$  will have the distribution of  $f(\theta|X)$ 
11: end procedure

```

The term Bayesian reflects the fact that a prior is involved. The weight of this prior can be manipulated by the threshold ϵ . If this threshold is set very low, the prior plays no role and only observations are taken into account. If ϵ is set extremely high, all θ coming from the prior will be accepted, and the actual observations are not used in the process. There are several disadvantages to approximate Bayesian computation.

- A set of simulated observations has to be compared with the actual observations. This becomes unwieldy if there are many observations.
- It is possible to use summary statistics rather than the observations themselves. If these are sufficient statistics there will be no information loss. If not, there will be information loss in practice.
- The distance function suffers from the curse of dimensionality. In the case that the dimensionality of the individual observations becomes high, or the number of parameters becomes large, it gets increasingly difficult to come up with a distance function which is efficient and accurate at the same time.

2.3.4 Gibbs Sampling

Gibbs sampling (Geman and Geman, 1984) is similar to the *coordinate descent* optimization algorithm (Wright, 2015). In coordinate descent a local minimum of a function is found by iteratively performing a line search along one coordinate direction at a time. Gibbs sampling optimizes over one variate in the multivariate probability distribution at a time. The update value is set and fixed. Then, the next variate is chosen in a round-robin like manner.

$$\Theta = S(X, \sim p(\theta_i | \theta_{-i}, X), \sim p(\theta), B) \quad (2.45)$$

The Gibbs algorithm is given in Algorithm 5. Some explanation on the notation is as follows. The multiple parameters in the multivariate probability distribution are denoted by θ . The parameters are denoted individually with θ_i . The set of all parameters except for i is denoted

1145 by θ_{-i} . If we sample a parameter we write θ^t with t the iteration or sampling round. The
 1146 set of parameter samples has capital letter Θ .

Algorithm 5 Gibbs sampling

```

1: procedure GIBBS SAMPLING( $p(\theta_i|\theta_{-i}, X), p(\theta), X, B$ )           ▷ Requires parameters,
   observations and burn-in.
2:    $\theta^0 \sim p(\theta)$                                            ▷ Set parameters to some initial value
3:   for  $t = 1 \rightarrow T$  do
4:     for  $i = 1 \rightarrow k$  do
5:        $\theta_i^t \sim p(\theta_i^{t-1}|\theta_{-i}^t, X)$            ▷ Generate  $\theta_i^t$  from the full conditional probability
6:     end for
7:      $\Theta = \Theta \cup \theta^t$ 
8:   end for
9:    $\Theta_{B:T} \in \Theta$                                            ▷ Get  $\Theta_T$  set, from burn-in  $B$  to end of run  $T$ 
10:   $\Theta \sim \Theta_{B:T}$                                            ▷ Sample  $\Theta$  from correlated  $\Theta_{B:T}$ 
11:  return  $\Theta$ 
12: end procedure

```

1147 Gibbs samples are Markovian. This means that the conditional probability only takes into
 1148 account values at the previous time step $t - 1$. When running the Gibbs sampling algorithm
 1149 long enough, it will visit all possible states eventually. The Markovian property has an un-
 1150 desired side effect. It makes subsequent steps correlated. Hence when finally extracting the
 1151 parameter probabilities, it is important to skip multiple steps to remove the temporal cor-
 1152 relations. It is also important to run the algorithm for a while after its start. In that case it
 1153 does not suffer from a bad choice of initial parameter values. Disregarding the first samples
 1154 is called burn-in. In words, Gibbs sampling works by having the algorithm spend time in
 1155 parts of the space proportionally to the probability of getting into that part of the space.

1156 In the physics literature Gibbs sampling is known as Glauber dynamics or the heat bath al-
 1157 gorithm. First, observe that Gibbs sampling does not necessary require an actual calculation
 1158 of the conditional probability in all cases. The obvious exception is for the observations,
 1159 which are already known. Second, observe that a neat optimization procedure arises when
 1160 conjugate priors are used. A conjugate prior leads to a posterior distribution that can be de-
 1161 scribed analytically. In such a case it is computationally unnecessary to perform sampling.
 1162 It is much faster to use the actual available analytic description. This is commonly called
 1163 collapsed Gibbs sampling.

1164 2.3.5 Metropolis-Hastings Sampling

1165 Metropolis-Hastings (Metropolis et al., 1953) is one of the most well-known Markov chain
 1166 Monte Carlo (MCMC) algorithms. An MCMC algorithm uses a Markov chain (see Gibbs
 1167 sampling, Section 2.3.4) and combines this with a stochastic (Monte Carlo) component.
 1168 This sampling method can be used for high-dimensional distributions. Metropolis-Hastings
 1169 calculates an acceptance factor α which takes into account if a step should be taken according

1170 to a predefined proposal distribution. In case this step is not accepted, the current sample is
 1171 resampled (see Algorithm 6).

$$\Theta = S(X, \theta^0, Q(\theta^{t+1}|\theta^t), f(\theta, X)) \quad (2.46)$$

1172 Here we need $Q(\theta^{t+1}|\theta^t)$ explicitly as well as samples from it.

Algorithm 6 Metropolis-Hastings sampling

```

1: procedure METROPOLIS-HASTINGS SAMPLING( $\theta^0, X, Q, f$ )  $\triangleright$  Requires initial parameters,
   observations, proposal distribution, and function proportional to desired distribution
2:   for  $t = 1 \rightarrow T$  do
3:      $\theta^{t+1} \sim Q(\theta^{t+1}|\theta^t)$   $\triangleright$  Sample from proposal distribution  $Q$ 
4:      $\alpha = \frac{f(\theta^{t+1}, X^{t+1})Q(\theta^t|\theta^{t+1})}{f(\theta^t, X^t)Q(\theta^{t+1}|\theta^t)}$   $\triangleright$  Calculate acceptance
5:      $u \sim U(0, 1)$   $\triangleright$  Inverse transform sampling
6:     if  $\alpha > u$  then
7:        $\Theta = \Theta \cup \theta^{t+1}$   $\triangleright$  Accept  $\theta^{t+1}$ 
8:     else
9:        $\Theta = \Theta \cup \theta^t$   $\triangleright$  Reuse previous sample (note, different from rejection)
10:    end if
11:  end for
12:  return  $\Theta$   $\triangleright$   $\Theta$  will be samples from the distribution  $f(\theta|x)$ 
13: end procedure

```

1173 A particular choice of a Metropolis-Hastings step is that of a proposal distribution that does
 1174 not depend on the state of the chain. This is already suggested by Hastings and is called the
 1175 independence sampler.

1176 2.3.6 Split-Merge MCMC Sampling

1177 The discussed sampling methods do not assume much structure in the model. This means
 1178 that in hierarchical models sampling either occurs through updating the to-be-estimated
 1179 quantities by iterating over every single observation or over every single cluster. This has
 1180 a disadvantage, the procedure in which a cluster is split into two clusters is very slow by
 1181 moving data points one by one from an old to a new cluster. Much more efficient sampling
 1182 methods can be designed if we would be able to handle large chunks of cluster assignments
 1183 at once.

1184 Split-merge samplers are such methods that can update cluster assignments for multiple ob-
 1185 servations at once. These samples adjust the acceptance method in the Metropolis-Hastings
 1186 algorithm. Split-Merge sampling is described in Algorithm 7.

Algorithm 7 Split-Merge MCMC sampling

```

1: procedure SPLIT-MERGE MCMC SAMPLING( $\theta^0, X, Q, f$ )    ▷ Requires initial parameters,
   observations, proposal distribution, and function proportional to desired distribution
2:   for  $t = 1 \rightarrow T$  do
3:      $i \sim D(0, N - 1)$                                 ▷ Sample observation  $i$  discretely
4:      $j \sim D(0, N - 1)$                                 ▷ Sample observation  $j$  discretely
5:     if  $c_i == c_j$  then
6:        $c_{old} = c_i$ 
7:        $\theta_{c_{new}}^{t+1} \sim Q(\theta^{t+1} | \theta^t)$           ▷ Sample from proposal distribution  $Q$ 
8:       for  $k \in c_{old}$  do
9:          $c_k \sim C(c_{old}, c_{new})$                     ▷ Assign to new cluster categorically
10:      end for
11:    else
12:       $c_{merge} = c_i$ 
13:      for  $k \in c_j$  do
14:         $c_k = c_{merge}$                                 ▷ Assign all to first cluster
15:      end for
16:    end if
17:     $\alpha = \frac{f(\theta^{t+1}, X^{t+1})Q(\theta^{t+1} | \theta^t)}{f(\theta^t, X^t)Q(\theta^t | \theta^{t+1})}$           ▷ Calculate acceptance
18:     $u \sim U(0, 1)$                                     ▷ Inverse transform sampling
19:    if  $\alpha > u$  then
20:       $\Theta = \Theta \cup \theta^{t+1}$                           ▷ Accept  $\theta^{t+1}$ 
21:    else
22:       $\Theta = \Theta \cup \theta^t$                             ▷ Reuse previous sample (note, different from rejection)
23:    end if
24:  end for
25:  return  $\Theta$                                           ▷  $\Theta$  will be samples from the distribution  $f(\theta|x)$ 
26: end procedure

```

1187 The exact acceptance probability depends on the model. For the mixture model with a
 1188 Dirichlet Process as prior, its performance is further improved by adjusting the assignment
 1189 process from random to observation-supported by introducing intermediate restricted Gibbs
 1190 sampling steps (Jain and Neal, 2004, 2007). Similarly, there are other variants that incor-
 1191 porate data fit to the splitting step. Labels can for example be calculated sequentially (Dahl,
 1192 2003) or methods can be used that postulate subcluster structure within clusters to optimize
 1193 inference over split and merge sets (Chang and Fisher III, 2013).

1194 2.3.7 Comparison of the Six Inference Methods

1195 In robotic vision the type of data we are obtaining from depth sensors are point clouds. To
 1196 perform inference over objects made out of point clouds, clustering algorithms benefit from
 1197 two sampling strategies. If conjugate probability densities are used, Gibbs sampling, or col-
 1198 lapsed Gibbs sampling can be used (Section 2.3.4). If the model becomes more complicated
 1199 and nonconjugate split-merge sampling will likely accelerate the inference (Section 2.3.6).

1200 2.4 Chapter Conclusions

1201 Chapter 3 describes Gibbs sampling to perform inference over an infinite set of lines. Gibbs
1202 sampling requires conditional probabilities. These are given in closed-form because there
1203 is a conjugate description of the line parameters given the points that form the lines in this
1204 application.

1205 Chapter 4 describes Split-Merge MCMC sampling to perform inference over an infinite set
1206 of line segments. The parameters for line segments do not have a conjugate description.
1207 Metropolis-Hastings can be used to perform inference over the line segments, but the search
1208 space is quite large. The Split-Merge MCMC method performs faster inference than Metropolis-
1209 Hastings because it is able to split and merge line segments (with multiple points ascribed
1210 to them) at once.

NONPARAMETRIC BAYESIAN LINE DETECTION

Contents

In this chapter the nonparametric Bayesian models from the literature (Chapter 2) are applied to perform inference over point clouds. The point cloud under study will be a point cloud distributed over lines in a two-dimensional space. Traditionally, RANSAC and the Hough transform have been used to perform inference over such lines. We use a nonparametric Bayesian model to perform inference over a countably infinite number of lines. Given a prior with respect to the noise and the distribution of points over the lines, Bayesian inference describes the optimal procedure to perform line fitting.

Published in

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. Nonparametric Bayesian Line Detection. *International Conference on Pattern Recognition and Methods*, ICPRAM 2016, Rome, Italy, February 24-26, 2016. Best paper award in theory and methods track.

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. *Fundamentals of Nonparametric Bayesian Line Detection*. Springer, 2017.

Outline

The infinite line model describes a collection of lines with a Dirichlet process as prior (Section 3.1). Inference in the infinite line model is performed through Gibbs sampling (Section 3.2). As is known, Gibbs sampling over *parameters* converges slowly, however it can be accelerated through sampling over *clusters* (Section 3.3). The results by the inference method are assessed using clustering performance measures (Section 3.4). The chapter summarizes the findings (Section 3.5) and introduces extensions which will be handled in the next chapters.

In computer vision and particularly in robotics, traditionally the task of line detection has been performed through sophisticated, but ad-hoc methods. Here we mention two examples

of such methods, RANSAC and the Hough transform. RANSAC (Bolles and Fischler, 1981) is a method that iteratively tests a hypothesis. A line is fitted through a subset of points. Then other points that are in consensus with this line (according to a certain loss function) are added to the subset. This procedure is repeated till a certain performance level is obtained. The Hough transform (Hough, 1962) is a deterministic approach which maps points in the image space to curves in the so-called Hough space of slopes and intercepts. A line is extracted by getting the maximum in the Hough space.

There are four main problems with these methods. First, the extension of RANSAC or Hough to the detection of multiple lines is nontrivial (Zhang and Křsecká, 2007; Gallo et al., 2011; Chen et al., 2001). Second, the noise level is hardcoded into model parameters and it is not possible to incorporate knowledge about the nature of the noise. Third, it is hard to extend the model to hierarchical forms, for example, to lines that form more complicated structures such as squares or volumetric forms. Fourth, there are no results known with respect to any form of optimality of the mentioned algorithms.

In this chapter we postulate a method to perform inference over the number of lines and over the fitting of points on that line using the nonparametric Bayesian methods from chapter 2. The method aims at overcoming the four main problems mentioned above.

3.1 Infinite Line Model

The Dirichlet process has been previously described as prior for a mixture distribution (in Figure 2.6, see Section 2.2.1). It will be used in our model as a prior for the *distribution of points over a countably infinite set of lines*. From now on we will refer to this model as the infinite line model (ILM).

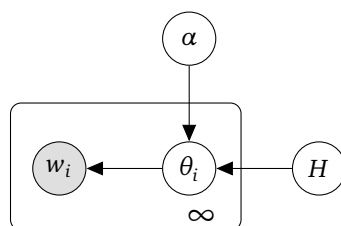


Figure 3.1: The infinite line model using the Chinese restaurant process representation (compare with Figure 2.6). Top: α , the concentration parameter of the Dirichlet process. Bottom, left to right: w_i , the observation, an individual point in a 2D space; θ_i , the parameters (intercept, slope) of the line belonging to observation w_i ; H , the base distribution from which line parameter values are sampled.

The infinite line model is visualized in Figure 3.1 using plate notation (see Section 2.1.5). In Section 3.1.1 it is described how θ_i is sampled from H and α . In Section 3.1.2 it is described how w_i is sampled from θ_i . In Section 3.1.3 the prior H for θ_i is described. In Section 3.1.4 it is described how line parameters θ_i can be updated given the data w_i .

1265 3.1.1 Posterior Predictive for a Line given Other Lines

1266 Let us start reiterating the definition of the Dirichlet process. Let H be a distribution over Θ
 1267 and let α be scalar. The Dirichlet process generates a distribution $G \sim DP(\alpha, H)$:

$$G(\theta_1, \dots, \theta_\infty) \sim DP(\alpha, H(\theta_1, \dots, \theta_\infty)). \quad (3.1)$$

1268 A Dirichlet process assigns a Dirichlet distribution to every parameter partition $\Theta_1, \dots, \Theta_r$:

$$(G(\Theta_1), \dots, G(\Theta_r)) \sim Dir(\alpha H(\Theta_1), \dots, \alpha H(\Theta_r)). \quad (3.2)$$

1269 The Dirichlet is conjugate to the categorical:

$$(G(\Theta_1), \dots, G(\Theta_r)) \mid \theta_1, \dots, \theta_n \sim Dir(\alpha H(\Theta_1) + n_1, \dots, \alpha H(\Theta_r) + n_r),$$

$$n_k = \sum_{j=1}^n \delta_{\theta_j}(\Theta_k). \quad (3.3)$$

1270 To study the details about the conjugacy of the Dirichlet distribution with respect to multi-
 1271 nomial and categorical distributions we refer to the derivations in Appendix B.

1272 In the above notation, $\delta_{\theta_j}(\Theta_k)$ is a Dirac measure (a generalization of the Dirac delta func-
 1273 tion), also known as an indicator function. Given a set Θ_k with a σ -algebra over subsets of
 1274 Θ :

$$\delta_{\theta_j}(\Theta_k) = 1_{\Theta_k}(\theta_j) = \begin{cases} 1 & \text{if } \theta_j \in \Theta_k \\ 0 & \text{if } \theta_j \notin \Theta_k \end{cases}. \quad (3.4)$$

1275 The posterior for the Dirichlet process base distribution and dispersion parameter is a Dirich-
 1276 let process with adjusted parameters:

$$G(\cdot) \mid \theta_1, \dots, \theta_n \sim DP\left(\alpha + n, \frac{\alpha}{\alpha + n} H(\cdot) + \frac{n}{\alpha + n} \frac{\sum_{j=1}^n \delta_{\theta_j}(\cdot)}{n}\right). \quad (3.5)$$

1277 The posterior base distribution G is a weighted average between the prior base distribution
 1278 H and the empirical distribution $n^{-1} \sum_{j=1}^n \delta_{\theta_j}$ with the weights respectively α and n (nor-
 1279 malized). The dispersion parameter α is updated to $\alpha + n$. Note that $\delta_{\theta_j}(\cdot)$ is a distribution,
 1280 the Dirac measure Eq. 3.4.

1281 The posterior predictive for a new parameter θ_{n+1} has the form:

$$P(\theta_{n+1} \in \Theta_k \mid \theta_1, \dots, \theta_n) = \frac{1}{\alpha + n} \left(\alpha H(\Theta_k) + \sum_{j=1}^n \delta_{\theta_j}(\Theta_k) \right). \quad (3.6)$$

In other words, the posterior predictive of θ_{n+1} given the parameters $\theta_1, \dots, \theta_n$ in Eq. 3.6 has exactly the same form as the posterior base distribution G given the parameters $\theta_1, \dots, \theta_n$ (Blackwell and MacQueen, 1973) in Eq. 3.5, namely:

$$\theta_{n+1} \mid \theta_1, \dots, \theta_n \sim \frac{1}{\alpha + n} \left(\alpha H(\theta_{n+1}) + \sum_{j=1}^n \delta(\theta_j - \theta_{n+1}) \right). \quad (3.7)$$

A normal Dirac delta function $\delta(\theta_j - \theta_{n+1})$ can be used here, which is only non-zero when θ_j is equal to θ_{n+1} .

Equivalently, if we describe θ_n conditioned on $\theta_1, \dots, \theta_{n-1}$ we have to run over $n-1$ rather than n parameters:

$$\theta_n \mid \theta_1, \dots, \theta_{n-1} \sim \frac{1}{\alpha + n - 1} \left(\alpha H(\theta_n) + \sum_{j=1}^{n-1} \delta(\theta_j - \theta_n) \right). \quad (3.8)$$

Due to the exchangeability property we can also consider any other parameter update (Neal, 2000):

$$\theta_i \mid \theta_{-i} \sim \frac{1}{\alpha + n - 1} \left(\alpha H(\theta_i) + \sum_{j \neq i} \delta(\theta_j - \theta_i) \right). \quad (3.9)$$

The notation θ_{-i} means every parameter θ except for the one equal to θ_i .

3.1.2 Likelihood of Data given a Line

The likelihood of data given line parameters is defined to be according to the **Bayesian linear regression** model. The Bayesian linear regression model for a single line (Box and Tiao, 2011) assumes a linear relationship between the independent x_i and dependent variables y_i with Gaussian noise added in the y -direction. The individual points i are drawn from a Normal distribution:

$$y_i \sim \mathcal{N}(x_i \beta, \sigma^2). \quad (3.10)$$

The (column) vector β maps the (row) vector with independent variables x_i to the dependent variable y_i . The noise is normally distributed with standard deviation σ along the dimension of the dependent variable.

In a 2D point cloud the point p is represented by (x_p, y_p) . The points are mapped into an intercept-slope representation through defining $X_i = [1 \ x_p]$ and $y_i = y_p$. The vector β will then contain the y -intercept as the first value, the slope as the second value.

1304 All observations that belong to the same single line lead to a likelihood function that corre-
 1305 sponds to a normally distributed random variable with y and X as parameters:

$$p(y | X, \beta, \sigma^2) \propto \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2}(y - X\beta)^T(y - X\beta)\right). \quad (3.11)$$

1306 The dependent variable is now a column vector of values y and each observation has a row
 1307 of independent variables in X . The vector β and the standard deviation σ are shared across
 1308 all observations. The term $y - X\beta$ is written out like this:

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} - \begin{pmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix}. \quad (3.12)$$

1309 Note that Eq. 3.11 has exactly the same form for a single point or for multiple points that be-
 1310 long to the same line. Hence, we have the probability of a point w_i given the line parameters
 1311 $\theta_k = (\beta_k, \sigma_k^2)$:

$$F(w_i, \theta_k) = p(w_i | \theta_k) = p(w_i | \beta_k, \sigma_k^2) = p(y_i | X_i, \beta_k, \sigma_k^2). \quad (3.13)$$

1312 To get the full distribution $p(w_i, \beta, \sigma^2)$ we will need also $p(\beta, \sigma^2)$.

1313 3.1.3 Conjugate Prior for a Line

1314 The conjugate prior for the likelihood in Eq. 3.11 is a product of a prior for the standard
 1315 deviation $p(\sigma)$ and the conditional probability of the line coefficients given the standard
 1316 deviation $p(\beta | \sigma^2)$.

$$p(\sigma^2, \beta) = p(\sigma^2)p(\beta | \sigma^2). \quad (3.14)$$

1317 The standard deviation σ is sampled from an Inverse-Gamma (IG) distribution:

$$p(\sigma) \propto (\sigma^2)^{-(\nu_0/2+1)} \exp\left(-\frac{1}{2\sigma^2} \nu_0 s_0^2\right). \quad (3.15)$$

1318 This is an $IG(a_0, b_0)$ with $a_0 = \nu_0/2$ and $b_0 = 1/2 \nu_0 s_0^2$. The conditional with respect to the
 1319 line coefficients has a normal distribution as prior:

$$p(\beta | \sigma^2) \propto \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2}(\beta - \mu_0)^T \Lambda_0(\beta - \mu_0)\right). \quad (3.16)$$

1320 Let us collect $\Lambda_0, \mu_0, a_0, b_0$ into λ_0 , we have now a description of our base distribution H :

$$H(\theta_k) = NIG(\theta_k; \lambda_0). \quad (3.17)$$

1321 The Normal-Inverse-Gamma (NIG) is a distribution that combines a Normal and an Inverse
 1322 Gamma distribution. The line coefficients are sampled from a Normal distribution and the
 1323 standard deviation is sampled from the Gamma distribution with a_0 and b_0 as hyperparam-
 1324 eters.

$$\begin{aligned} \sigma_k &= \tau_k^{-1/2} & \tau_k &\sim \mathcal{G}(a_0, b_0), \\ \mu_k &\sim \mathcal{N}(\mu_0, \sigma^2 \Lambda_0^{-1}). \end{aligned} \quad (3.18)$$

1325 3.1.4 Posterior Predictive for a Line given Data

1326 Due to the fact that the NIG is a conjugate prior with respect to the normal distribution (with
 1327 unknown mean and variance), we have a simplified description for updating the hyperpa-
 1328 rameters, given a set of observations. The hyperparameters are updated¹ according to (c.f.
 1329 Denison, 2002):

$$\begin{aligned} \Lambda_n &= \Lambda_0 + X^T X, \\ \mu_n &= \Lambda_n^{-1}(\Lambda_0 \mu_0 + X^T y), \\ a_n &= a_0 + n/2, \\ b_n &= b_0 + 1/2(y^T y + \mu_0^T \Lambda_0 \mu_0 - \mu_n^T \Lambda_n \mu_n). \end{aligned} \quad (3.19)$$

1330 Let us collect $\Lambda_0, \mu_0, a_0, b_0$ into λ and $\Lambda_n, \mu_n, a_n, b_n$ into λ^* . Let us collect a set of our obser-
 1331 vations and $(X, y)_k$ into w_k . The update for the hyperparameters can then be summarized
 1332 as:

$$\lambda^* = U_{up}(\lambda, w_k). \quad (3.20)$$

1333 Removing observations does lead to similar updates (“downdates”) for the hyperparameters:

$$\begin{aligned} \Lambda_n &= \Lambda_0 - X^T X, \\ \mu_n &= \Lambda_n^{-1}(\Lambda_0 \mu_0 - X^T y), \\ a_n &= a_0 - n/2, \\ b_n &= b_0 - 1/2(y^T y + \mu_n^T \Lambda_n \mu_n - \mu_0^T \Lambda_0 \mu_0). \end{aligned} \quad (3.21)$$

1334 The downdate for the hyperparameters can then be summarized as:

$$\lambda^* = U_{down}(\lambda, w_k). \quad (3.22)$$

¹In comparison with the notation of Denison (2002), we update Λ rather than $V = \Lambda^{-1}$ and subsequently use Λ_n at the right-hand side to simplify the notation for μ_n and b_n .

1335 If we combine this update with sampling θ_k from λ_n according to Eq. 3.17, then we obtain:

$$p(\theta_k | \lambda_0, w_k) \propto F(w_k, \theta_k)H(\theta_k; \lambda_0) = p(\theta_k | \lambda_n) = NIG(\theta_k; \lambda_n). \quad (3.23)$$

1336 Sampling of $NIG(\theta_k; \lambda_n)$ is as in Eq. 3.18, but with λ_n rather than λ_0 .

$$\begin{aligned} \sigma_k &= \tau_k^{-1/2} & \tau_k &\sim \mathcal{G}(a_n, b_n), \\ \mu_k &\sim \mathcal{N}(\mu_n, \sigma^2 \Lambda_n^{-1}). \end{aligned} \quad (3.24)$$

1337 Let us integrate over θ (through the function H):

$$Q(w_k, \lambda_0) = \int_{\Theta} F(w_k, \theta) dH(\theta; \lambda_0). \quad (3.25)$$

1338 3.2 Inference for the Infinite Line Model

1339 The posterior predictive for parameters (see Eq. 3.9) combined with observations w_i is de-
1340 scribed by:

$$p(\theta_i | \theta_{-i}, w_i) \propto r_i H_i(\theta_i) + \sum_{j \neq i} L_{i,j} \delta(\theta_j - \theta_i). \quad (3.26)$$

1341 Eq. 3.26 defines the posterior as proportional (indicated by \propto) to three terms. First, the
1342 α -weighted posterior r_i for a new *cluster*. Second, the posterior with respect to the prior
1343 distribution over the parameter values $H_i(\theta_i)$ given observation w_i . Third, a term that sums
1344 over the likelihood of w_i given existing line θ_j , indicated by $L_{i,j}$.

1345 The α -weighted posterior r_i defines the probability that a new cluster will be sampled:

$$r_i = \alpha Q(w_i, \lambda_0) = \alpha \int_{\Theta} F(w_i, \theta) dH(\theta). \quad (3.27)$$

1346 The posterior $H_i(\theta)$ is the normalized product of the prior distribution $H(\theta_i)$ with the like-
1347 lihood $F(w_i, \theta_i)$ for a single observation w_i .

$$H_i(\theta_i) \propto H(\theta_i) F(w_i, \theta_i). \quad (3.28)$$

1348 The likelihood $F(w_i, \theta_i)$ we will denote $L_{i,i}$ or, in general, for observation w_i and line θ_j :

$$L_{i,j} = F(w_i, \theta_j). \quad (3.29)$$

1349 Sampling a new cluster parameter from $H_i(\theta_i)$ is done with probability:

$$p(\theta_{new}) = \frac{r_i}{r_i + \sum_{j \neq i} L_{i,j}}. \quad (3.30)$$

1350 We can use these equations to derive the parameters θ_i with Gibbs sampling.

Algorithm 8 Gibbs sampling over parameters θ_i

```

1: procedure GIBBS ALGORITHM 1( $w, \lambda_0, \alpha$ )           ▷ Accepts points  $w$ , hyperparameters  $\lambda_0, \alpha$  and
   returns  $k$  line coordinates
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:        $r_i = \alpha Q(w_i, \lambda_0)$            ▷ Posterior predictive of  $w_i$  given hyper parameters (Eq. 3.27)
5:       for all  $j = 1 : N, j \neq i$  do
6:          $L_{i,j} = F(w_i, \theta_j)$            ▷ Likelihood for a line given observation (Eq. 3.29)
7:       end for
8:        $p(\theta_{new}) = \frac{r_i}{r_i + \sum_{j \neq i} L_{i,j}}$    ▷ Probability of sampling a new parameter (Eq. 3.30)
9:        $u \sim U(0, 1)$ 
10:      if  $p(\theta_{new}) > u$  then                               ▷ Sample with probability  $p(\theta_{new})$ 
11:         $\lambda_n = U_{up}(w_i, \lambda_0)$            ▷ Update hyperparameters with  $w_i$  (Eq. 3.20)
12:         $\theta_i \sim NIG(\theta_i; \lambda_n)$            ▷ Sample  $\theta_i$  from NIG (Eq. 3.24)
13:      else
14:         $\theta_i$  sampled from existing clusters           ▷ Sample old cluster
15:      end if
16:    end for
17:  end for
18:  return summary on  $\theta_k$  for  $k$  lines
19: end procedure

```

1351 This Gibbs algorithm is earlier described in its general form (see algorithm 1 in Neal, 2000).
1352 As shown in Algorithm 8 we perform a loop in which for T iterations each θ_i belonging to
1353 observation w_i is updated in succession. The loop consists of four steps. First, the posterior
1354 predictive for w_i given the hyperparameters $p(w_i | \lambda_0)$ is calculated. Second, the likelihood
1355 $L_{i,j}$ for all θ_j given w_i (with $j \neq i$) is calculated. Third, the fraction with r_i defines the
1356 probability for θ_i to be sampled from a new or existing cluster. Fourth, depending on the
1357 probability u , (1) a new cluster is sampled, the hyperparameters are updated with informa-
1358 tion on w_i and thereafter θ is sampled from a Normal-Inverse-Gamma distribution with the
1359 updated hyperparameters, or (2) an existing cluster is sampled.

1360 3.3 Accelerating Inference for the Infinite Line Model

1361 Gibbs sampling of this model might be accelerated. In Figure 3.2 we use plate notation to
1362 show the stick-breaking representation of the infinite line model.

1363 In the previous section we sampled over individual parameters. It is possible to iterate only
1364 over the clusters. The derivation takes a few steps (Neal, 2000) but leads to a simple update
1365 for the component indices that only depends on the number of data items per cluster, the
1366 parameter α , and the available data.

1367 The probability to sample from an existing cluster depends on the number of items in that
1368 cluster (the current data item excluded). This is expressed in equation 3.31.

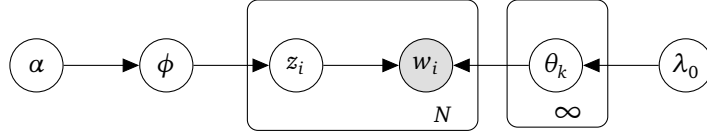


Figure 3.2: The infinite line model in the stick-breaking representation (compare with Figure 3.1). From left to right: α , the concentration parameter of the Dirichlet process; (ϕ_1, \dots, ϕ_k) , the partition of points over lines; z_i , the assignment parameters that link observation w_i with line k ; w_i , the observation, an individual point with x and y coordinates; θ_k , the parameters of line k ; λ_0 , the base measure from which the line parameter values are sampled.

Algorithm 9 Gibbs sampling over clusters c_k

```

1: procedure GIBBS ALGORITHM 2( $w, \lambda_0, \alpha$ )    ▷ Accepts points  $w$  and hyperparameters  $\lambda_0$  and  $\alpha$ ,
   returns  $k$  line coordinates
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:        $c = \text{cluster}(w_i)$                     ▷ Get cluster  $c$  currently assigned to observation  $w_i$ 
5:        $\lambda_c = U_{\text{down}}(w_i, \lambda_c)$           ▷ Adjust cluster hyperparameters on removing  $w_i$  (Eq. 3.22)
6:        $m_c = m_c - 1$                         ▷ Adjust cluster size  $m_c$ 
7:       for all  $k = 1 : K$  do
8:          $L_k = m_k F(w_i, \theta_k)$             ▷ Likelihood for cluster  $k$  given  $w_i$  (Eq. 3.33)
9:       end for
10:       $r_i = Q(w_i, \lambda_0)$                     ▷ Posterior predictive of  $w_i$  given hyper parameters
11:       $p(\text{new}) = \frac{r_i}{r_i + \sum_k L_k}$           ▷ Sample new or old?
12:      if  $p(\text{new})$  then
13:         $\lambda_k = U_{\text{up}}(w_i, \lambda_0)$           ▷ Update hyperparameters with observation  $w_i$ 
14:         $\theta_i \sim \text{NIG}(\lambda)$                 ▷ Sample  $\theta_i$  from NIG
15:      else
16:         $k$  sampled from existing clusters
17:         $\lambda_k = U_{\text{up}}(w_i, \lambda_k)$           ▷ Restore hyperparameters with observation  $w_i$ 
18:      end if
19:       $m_k = m_k + 1$                         ▷ Increment cluster size  $m_k$ 
20:    end for
21:    for all  $k = 1 : K$  do
22:       $\theta_k \sim \text{NIG}(\lambda_k)$                 ▷ Sample  $\theta_k$  from NIG
23:    end for
24:  end for
25:  return summary on  $\theta_k$  for  $k$  lines
26: end procedure

```

$$p(c_i = c \text{ and } c_i = c_j \text{ and } i \neq j \mid c_{-i}, w_i, \alpha, \theta) \propto \frac{n_{c,-i}}{\alpha + n - 1} F(w_i \mid \theta_i). \quad (3.31)$$

1369 The probability to sample a new cluster only depends on α and the total number of data
 1370 items. This is formally described in equation 3.32.

$$p(c_i \in \Omega(c) \text{ and } c_i \neq c_j \text{ and } i \neq j \mid c_{-i}, \alpha) \propto \frac{\alpha}{\alpha + n - 1} \int F(w_i \mid \theta_i) dH(\theta). \quad (3.32)$$

1371 Here $\Omega(c)$ denotes all admitted values for c_i . The importance of conjugacy is obvious from
 1372 Eq. 3.32, it will lead to an analytic form of the integral. The inference method using Eqs. 3.31

and 3.32 is described in Section 3.1.

One benefit of iterating over clusters rather than non-unique parameters is that we can calculate the likelihood by multiplying it with the number of observations at that cluster (rather than per parameter). If we write the number of observations as $n_{c,-i} = m_k$, we can update the likelihood on a cluster level like this:

$$L_k = m_k F(w_i, \theta_k). \quad (3.33)$$

Directly sampling over the clusters is described in its general form (see algorithm 2 in Neal, 2000). Rather than updating each θ_i per observation w_i , an entire cluster θ_k is updated. In Algorithm 8 the update of a cluster would require a first observation to generate a new cluster at θ_j and then moving all observations of the old cluster θ_i to θ_j . In contrast, in Algorithm 9 when a data item either is added or deleted from a cluster, the cluster parameters are updated for all data items in that cluster at once. For this algorithm this means that when w_i is excluded from calculating the likelihood we have to “downdate” the corresponding hyperparameters (as previously mentioned). In Algorithm 9 after all observations have been iterated over and assigned the corresponding cluster k , an outer loop iterates over all clusters to obtain new parameters θ from the NIG prior.

3.4 Performance of the Infinite Line Model

The infinite line model (see Section 3.1) is able to fit an infinite number of lines through a point cloud in two dimensions. These lines are no line segments, but infinite lines. However, to test the model a variable number of lines are generated of a length that is considerably larger compared to the spread caused by the standard deviation of points from that line.

As described earlier, Gibbs sampling leads to correlated samples. Our choice is to get the Maximum A Posterior estimates for the clusters by picking the median values for all the parameters involved. In Section 3.4.1 we discuss the clustering performance and in Section 3.4.2 we provide two clustering examples.

3.4.1 Clustering Performance

The results of the clustering algorithms are measured using conventional metrics. For instance, we may use the Rand Index. It describes the accuracy of cluster assignments (Rand, 1971) by:

$$R = \frac{a + b}{a + b + c + d}. \quad (3.34)$$

Here a counts the pair of points that belong to the same cluster, both at ground truth as well as after the inference procedure. Likewise b numbers the pair of points that belong to

different clusters in both sets. The values c and d describe discrepancies between the ground truth and the results after inference. A Rand Index of one means that there have been no mistakes.

The clustering performance is quite different from the line estimation performance. If the points are not properly assigned, the line will not be estimated correctly. Due to the fact that line estimation has this secondary effect, line estimation performance is not taken into account. Moreover, from lines that generated only a single, or very few points, we can extract point assignments, but line coefficients are impossible to derive. In fact, any derivation would lead to introducing a threshold for the number of points per cluster. Then the performance would need to be measured by weighting the fitting versus the assignment.

The performance of Algorithm 8 can be seen in Figure 3.3 and is rather disappointing. On average the inference procedure agrees upon the ground truth for 75% of the cases considering the Rand Index. Even worse, if we adjust for chance as with the Adjusted Rand Index, the performance would then drop to only having 25% correct cases!

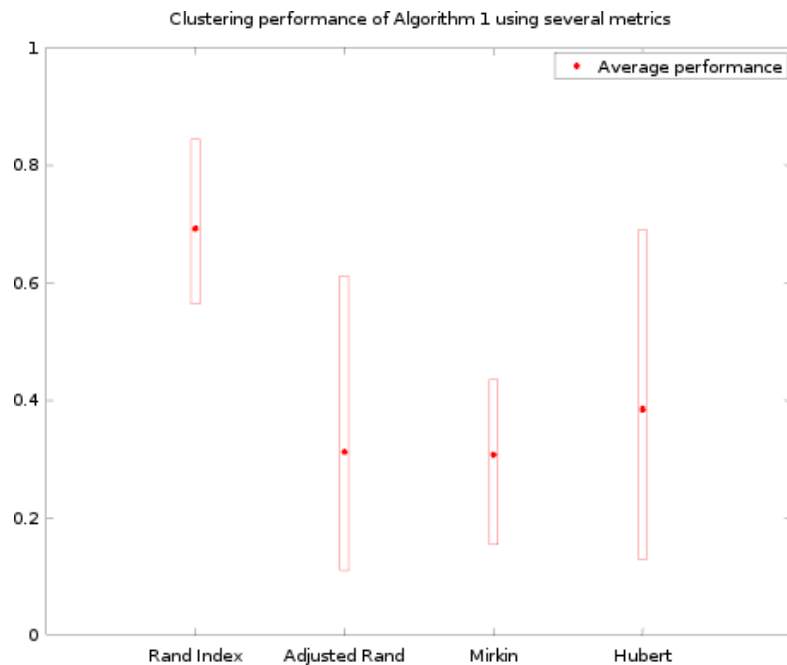


Figure 3.3: The performance of Algorithm 8 with respect to clustering is measured using the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert metric. A score of 1 means perfect clustering for all metrics, except Mirvin's where 0 denotes perfect clustering.

Algorithm 9 leads to stellar performance measures (Figure 3.4). Apparently, updating entire clusters at once with respect to their parameter values leads at times to perfect clustering, bringing the performance metrics close to their optimal values (see also van Rossum et al., 2016b).

The lack of performance of Algorithm 8 is not only caused by slow mixing. Even when allowing it ten times the number of iterations of Algorithm 8, it does not reach the same performance levels. A line seems to form local regions of high probability, making it difficult for points to postulate slightly changed line coordinates.

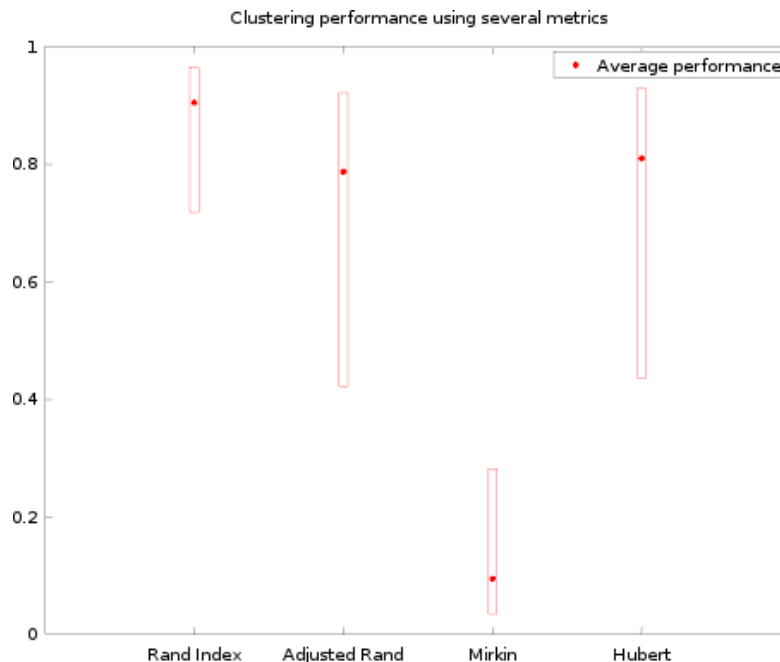


Figure 3.4: The performance of Algorithm 9 with respect to clustering is measured using the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert metric. A score of 1 means perfect clustering for all metrics, except Mirvin's where 0 denotes perfect clustering.

3.4.2 Two Examples

In the following we show two examples to understand the inference process better. The first example is seen in Figure 3.5. It shows the assignment after a single Gibbs step in Algorithm 8. There is a single line that is represented by two clusters. Algorithm 8 does not have merge or split steps to perform inference about sets of data points, it thus has to move each data point one by one. In passing we mention that there are split-merge algorithms that take these more sophisticated Gibbs steps into account (Jain and Neal, 2004) and we will see these in the following two chapters.

The second example in Figure 3.6 shows that a single point as an outlier is not a problem for our method. A single point might throw off Bayesian linear regression, but because there are multiple lines to be estimated in our Infinite Line Mixture Model, this single point is assigned its own line.

The extension to more points as outliers would, of course, require us to postulate a distribution for these outlier points as well. For instance, a uniform distribution might be used in tandem with the proposed model. However, this would lead to a non-conjugate model and hence it would require different inference methods.

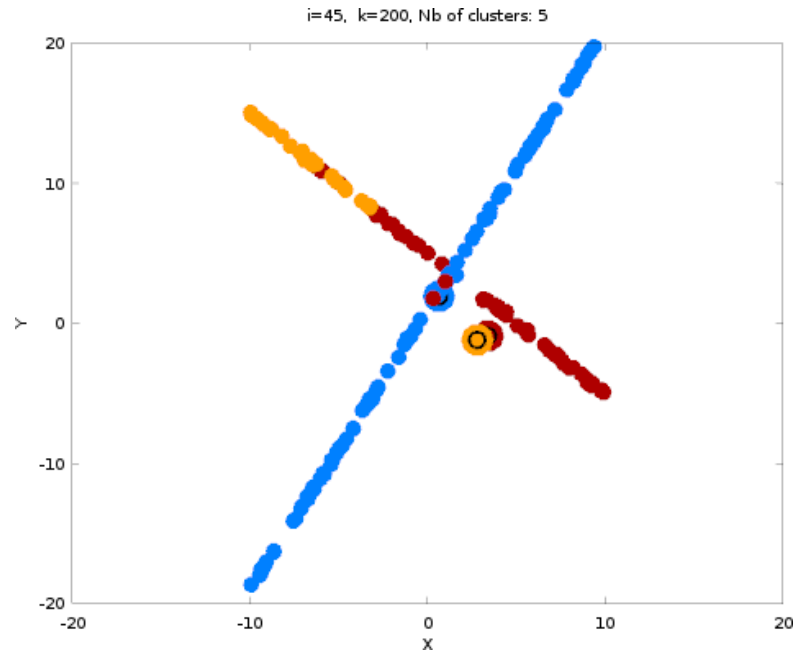


Figure 3.5: One of the Gibbs steps in the inference of two particular lines. The points are roughly distributed according to the lines, but one line exists out of two large clusters. The line coordinates are visualized by a double circle. The x-coordinate is the y-intercept of the line, the y-coordinate is the slope.




Figure 3.6: The assignment of a line to a single point. In the figure, there are three clusters found, rather than only the two obvious clusters.

1441 3.5 Chapter Conclusions

1442 The infinite line model that is proposed extends the familiar Bayesian linear regression model
 1443 to an infinite number of lines using a Dirichlet Process as prior. The model is a full Bayesian
 1444 method to detect multiple lines. A full Bayesian method, in contrast to ad-hoc methods

such as RANSAC or the Hough transform, means optimal inference (Zellner, 1988) given the model and noise definition.

Results in section 3.4 show high values for different performance metrics for clustering, such as the Rand Index, the Adjusted Rand Index, and other metrics (van Rossum et al., 2016b,a). The Bayesian model is solved through two types of algorithms. Algorithm 8 iterates over all observations and suffers from slow mixing. The individual updates makes it hard to reassign large number of points at the same time. Algorithm 9 iterates over entire clusters. This allows updates for groups of points leading to much faster mixing. Note, that even optimal inference may occasionally result in misclassifications. The dataset is generated by a random process. Hence, occasionally two lines are generated with almost the same slope and intercept. Points on these lines are impossible to assign to the proper line.

The essential contribution of this chapter is the introduction of a fully Bayesian method to infer lines. For such a model, it holds that there are two ways in which it can be extended for full-fledged inference in computer vision as required in robotics. First, the extension of lines in 2D to planes in 3D. This is quite a trivial extension that does not change anything of the model except for the dimension of the data points. Second, somehow a prior needs to be incorporated to limit the lines of infinite length, to line segments. To restrict points on the lines to a uniform distribution of points over a line segment, a symmetric Pareto distribution can be used as prior (see next Chapter). This would subsequently allow for a hierarchical model in which these end points are on their turn part of more complicated objects. Hence, the Infinite Line Mixture Model is an essential step towards the use of Bayesian methods (and thus properly formulated priors) for robotic computer vision.

1467

1468

1469

NONPARAMETRIC BAYESIAN SEGMENT ESTIMATION

1470

Contents

1471

1472

1473

1474

1475

1476

1477

The nonparametric Bayesian model for line estimation (Chapter 3) does not take into account lines that are of finite length. In this chapter, we introduce a Bayesian method to perform inference over such line segments. In this model our prior for the length of the line segment is a symmetric Pareto distribution. Due to the fact that the prior and likelihood do not form a conjugate pair, a more general inference method is used (than the inference methods for the conjugate model in Chapter 3), namely Gibbs sampling with auxiliary variables.

1478

Published in

1479

1480

1481

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. Non-parametric Segment Detection. *Proceedings of the Eighth European Starting AI Researcher Symposium, STAIRS 2016*, the Hague, the Netherlands, August 26-27, 2016.

1482

Outline

1483

1484

1485

1486

1487

1488

1489

1490

The model is using both a Normal-Inverse-Gamma distribution and a Pareto distribution as priors for an individual line segment (Section 4.1.1). The parameters for the line segments are generated through a Dirichlet process (Section 4.2). The generative Dirichlet process is used to perform inference using Gibbs sampling over auxiliary variables (Section 4.3). The results for inference over line segments are compared with those for lines (Section 4.4). Finally, weak aspects of the current MCMC method are established (Section 4.5). They will form the basis for new inference methods in the next chapters.

4.1 Pareto Pairs

Lines in a two-dimensional space are mathematical objects that can be described by *two* parameters. To restrict a line to a line segment, a total of *four* parameters are required. Two parametrizations will then come to mind. First, a center-point parametrization, in which parameters describe the center of a line segment, the slope of the line through the center, and the size of the line segment. Second, an endpoint parametrization, in which parameters describe the locations of the two endpoints. The two parametrizations are *equivalent*, but generalizations can either be intuitive or cumbersome. The reason is that the generalization to a line segment from a two-dimensional space to a three-dimensional space, requires the endpoints to be positions in a 3D space. Here we see that the center-point parametrization would require a nonintuitive description of the angles in particular directions. In contrast, the generalization to squares and rectangles or shapes with many endpoints, might benefit from the center-point parametrization.

As far as we know there is no statistical description of data points distributed over a line segment that has a conjugate prior form. A line segment itself, however, has a conjugate form! Assume that we have a prior for the location of endpoints on the x-axis. Given the data, we then update the location of the endpoints. By disregarding the distribution of the data points over the segment, we can update the location of the endpoints by a conjugate Bayesian construction.

4.1.1 Pareto Prior

Assume that the data is distributed according to a symmetric uniform distribution. Hence, the likelihood is given by:

$$p(x | a) \sim \mathcal{U}(-a, a) = \begin{cases} \frac{1}{2a} & \text{for } x \leq |a| \\ 0 & \text{otherwise} \end{cases} . \quad (4.1)$$

Here the uniform distribution is centered around 0 and extends with size a in both directions. It is possible to shift the entire distribution over a distance b . For now, let us continue with one endpoint at a and one endpoint at $-a$.

A prior for the (endpoints of a) symmetric uniform distribution is a symmetric Pareto distribution:

$$p(a) \sim \mathcal{P}_s(\lambda, k) = \begin{cases} \frac{1}{2} k \lambda^k |a|^{-k-1} & |a| \geq \lambda \\ 0 & \text{otherwise} \end{cases} . \quad (4.2)$$

1518 The factor $\frac{1}{2}$ stems from the fact that the symmetric Pareto distribution is now mirrored
 1519 across the y-axis. Hence, the probability density is half of that of the normal Pareto distri-
 1520 bution for the positive x-axis.

1521 If we would just sample from a symmetric Pareto distribution, we can sample multiple times
 1522 from the positive x-axis. To actually sample endpoints of segments we have to sample pairs
 1523 of points.

$$p(a, b) \sim \mathcal{P}_p(\lambda_m, \lambda_n, k) \quad (4.3)$$

1524 This process can be described in two steps. First, we sample a and b from a categorical
 1525 distribution to decide which one will be the left endpoint and which one the right endpoint.
 1526 Second, we sample the right endpoint from a normal Pareto distribution and the left endpoint
 1527 from a mirrored Pareto distribution.

1528 The sampling of Pareto pairs is visualized in Fig. 4.1.

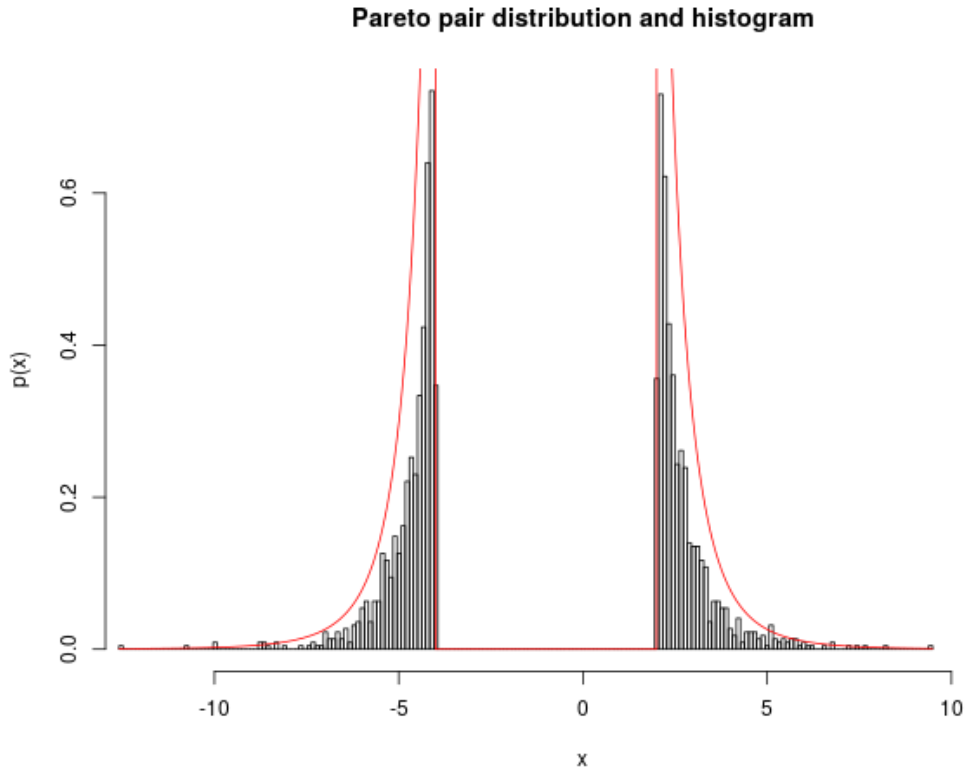


Figure 4.1: Sampling of Pareto pairs. The parameters are $\lambda_m = -4$, $\lambda_n = 2$, $k = 5$, and we have sampled $N = 1000$ pairs. The position parameters λ_m and λ_n define the positions of the endpoints. The shape parameter k defines the variance in the exact positions on both sides. The Pareto pairs always sample endpoints both at the "left" and the "right" (not both at the left or the right side).

1529 4.1.2 Posterior for a Pareto pair

1530 The Pareto distribution is a conjugate prior for the uniform distribution, with updated hy-
 1531 perparameters:

$$p(a \mid D) = \mathcal{P}(c, N + k). \quad (4.4)$$

1532 The data is denoted by $D = \{x_0, \dots, x_{N-1}\}$, the parameter k is adjusted with the number of
 1533 data points N , and the parameter c is the maximum of $\{m, \lambda\}$ with m the maximum value
 1534 in D .

1535 The posterior for a Pareto pair can be found by sampling in parallel for the endpoint at
 1536 the "right" and the one at the "left". The endpoint at the right is sampled from a Pareto
 1537 distribution $\mathcal{P}(c_n, N + k_n)$ with (1) c_n the maximum of the data points D and λ_n , (2) N
 1538 the number of Pareto pairs, and (3) k_n the scale hyperparameter. The endpoint at the left
 1539 is sampled from a Pareto distribution $\mathcal{P}(c_m, N + k_m)$ with (1) c_m the minimum of the data
 1540 points D and λ_m , (2) N the number of Pareto pairs, and k_m the scale hyperparameter.

1541 If $k_n \neq -k_m$ the distribution is shifted such that $k'_n = -k'_m$. This makes the form of the
 1542 probability distribution symmetric with respect to the y axis. In the end, the results are
 1543 shifted back. This transformation makes sense for pairs of points. We do not want the two
 1544 scale parameters of the Pareto distribution to influence the symmetry (the shape) of the
 1545 overall distribution.

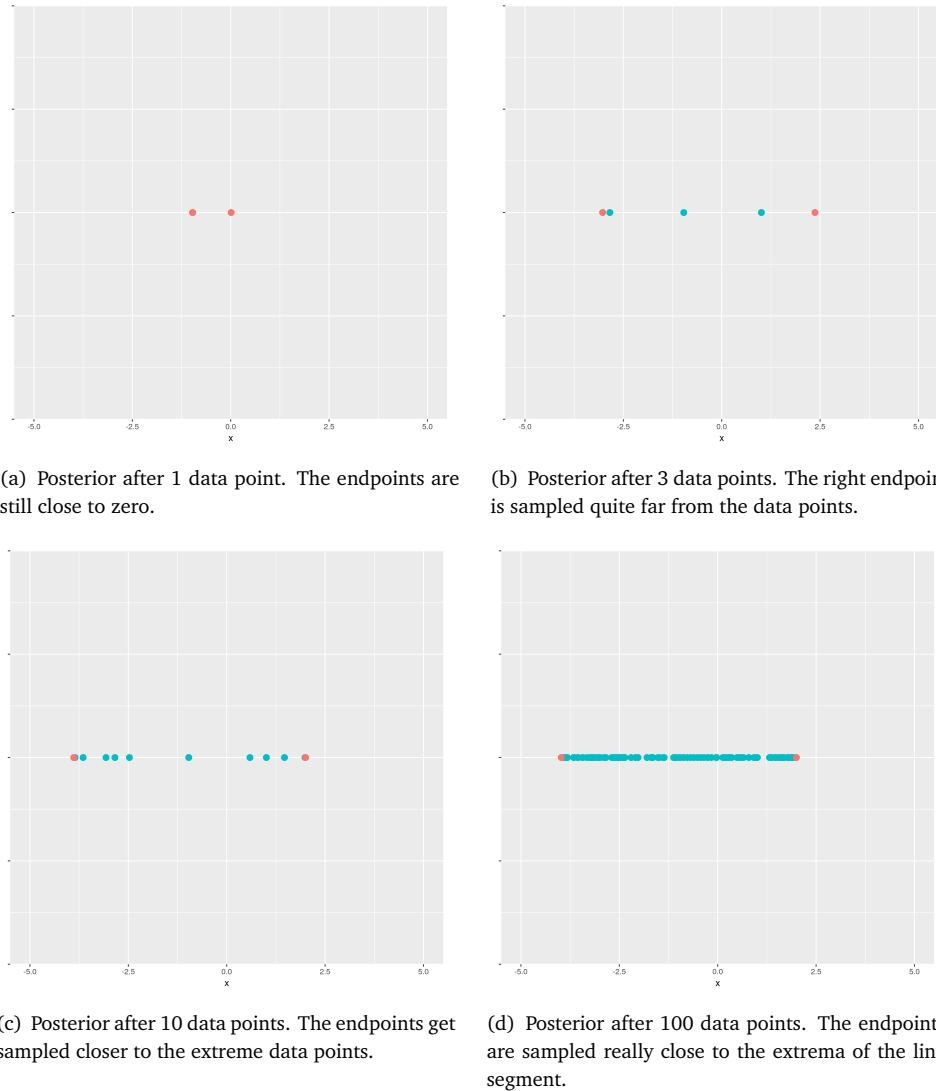


Figure 4.2: Consider (1) the data uniformly distributed on a line segment and (2) a symmetric Pareto prior for both endpoints, then we can update the estimate for the endpoints given the data as visualized. Each subfigure shows an adjustment of the endpoints given more data points (1, 3, 10, and 100 data points). The y-axis does not have a significance in these plots.

1546 Sampling from the Pareto distribution is through inverse transform sampling. By sampling
 1547 from $U(0, 1)$ with 1 included, we transform according to $k/U^{1/a}$.

1548 Figure 4.2 shows how the endpoints are updated given the data. An uninformative prior is
 1549 used. In this case the hyperparameters k_n and k_m are set close to 0, thus the data will wash
 1550 out the prior immediately. Naturally, it is possible to define large k_n and k_m . In that case it
 1551 should be noted that the data will never be able to “correct for” this prior: if the true length
 1552 of the segment is smaller than $|k_n - k_m|$ this will not be inferred properly. Here we note also
 1553 that the maximum and minimum operators are quite sensitive to outliers as well.

4.2 Generative Process to Create a Line Segment




Figure 4.3: The Bayesian linear regression model for multiple line segments in plate notation is the same as for the Infinite Line Model. The Dirichlet process is defined at the left with concentration parameter α . It generates the partitions (π_1, \dots, π_k) with assignment parameters z_i that denote which observation w_i belongs to which cluster k . The cluster is summarized through the parameter set θ_k and has λ_0 as its hyperparameter. The parameter set θ_k includes parameters that signify the line itself such as slope and y-intercept, plus the parameters that denote the extent of the segment.

To be able to perform inference over a line segment in a two-dimensional space, we will have to map somehow these points to a one-dimensional space (see Figure 4.2).

In the case of a *line* we can sample θ_i from a Normal-Inverse-Gamma distribution with hyperparameter λ_{temp} . The update for the hyperparameter we obtain in closed form given observations (described in the previous chapter, Sections 3.1.3 and 3.1.4). However, in the case of a *line segment* there is no known conjugate prior available. Here we will introduce a non-conjugate model to describe line segments. It uses a Dirichlet process prior for the assignment of points unto multiple segments (Section 4.2.1). It uses a uniform distribution (Section 4.2.2) to describe the distribution of points across an individual line segment.

4.2.1 Dirichlet Process Prior

Let us reiterate the Dirichlet process for our nonparametric line segment model:

$$\begin{aligned} G &\sim DP(\alpha, H), \\ \theta_i &| G \sim G, \\ w_i &| \theta_i \sim F(w_i, \theta_i). \end{aligned} \tag{4.5}$$

The likelihood function F describes the mapping from parameters θ_i to observations w_i . In the previous chapter this has been a likelihood function that describes observations on lines.

4.2.2 Likelihood of Data given Segment Parameters

The likelihood $F(w_i, \theta_i)$ describes the mapping from parameters θ_i to observations w_i . We will again use an intercept-slope representation through defining $w_i = (X_i, y_i)$ with $X_i = [1, x_i]$. The column vector $\beta = [\beta_0, \beta_1]$ contains two parameters: the y-intercept β_0 and the slope parameter β_1 (compare Section 3.1.2).

We assume a normally distributed random variable across $y - X\beta$, the same as in the line model (Eq. 3.11):

$$p(y | X, \beta, \sigma^2) \propto \sigma^{-n} \exp\left(-\frac{1}{2\sigma^2}(y - X\beta)^T(y - X\beta)\right). \quad (4.6)$$

1575 However, in contrast to the line model we restrict the variable $X_i = [1 \ x_i]$ to a subset of
 1576 all possible values. The data points on the segment are to be generated from a uniform
 1577 distribution:

$$\begin{aligned} F(y_i, X_i | \beta_i, \sigma_i, a_i, b_i) &= \mathcal{N}(y_i - X_i \beta_i, \sigma_i), \\ X_i &= [1 \ x_i], \\ x_i &\sim \mathcal{U}(a_i, b_i). \end{aligned} \quad (4.7)$$

1578 Fig. 4.4 displays the adjustment with points generated uniformly over the line segment.

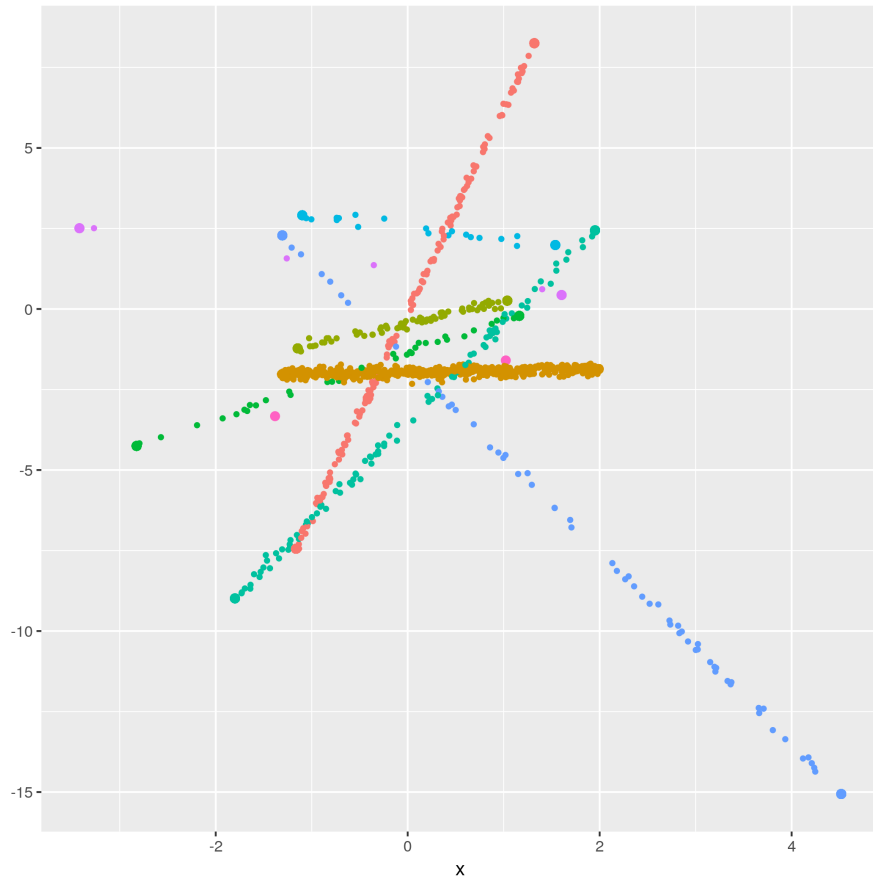


Figure 4.4: Line segments generated through a Dirichlet Process. The line parameters are generated from a Normal-Inverse-Gamma distribution. The endpoints of the line segment are generated from a symmetric Pareto distribution. The points are generated uniformly over the line segments.

1579 Finally, we remark that the description in Eq. 4.7 is not conjugate. This means that infer-
 1580 ence over line segments requires more complicated sampling strategies than the inference
 1581 methods in the previous chapter that require conjugacy.

4.3 Inference over a Line Segment

As a sequel to the last remark of section Section 4.2 we are aware that the line segment model is not conjugate. This therefore requires a sampling algorithm that does not make use of conjugacy. One algorithm that does not assume conjugacy is earlier described in its general form by Neal (2000) and is called *Gibbs sampling over auxiliary variables*. The sampling process proposes m new values for the parameters from the hyperparameters. The m values are called auxiliary parameters. Now, to establish to which cluster a certain observation w_i needs to be assigned, the likelihood of each existing and new clusters alike are compared. The weight of an old cluster is defined through the number of data points assigned to it. The weight of a new cluster is defined through α/m .

After every data item is assigned a cluster, the cluster parameters themselves are updated given the assigned data items. In a conjugate model the sufficient statistics (Definition 2.28) can be updated at once, given such observations. In a nonconjugate model we will need to update θ_j by sampling from $p(\theta_j | y)$.

Algorithm 10 Gibbs sampling over auxiliary variables (a θ_i)

```

1: procedure GIBBS ALGORITHM WITH AUXILIARY VARIABLES( $w, \lambda_0, \alpha$ )  $\triangleright$  Accepts points  $w$ ,
   hyperparameters  $\lambda_0, \alpha$ , number of auxiliary variables  $m$ , and returns  $k$  line coordinates.
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:       for all  $j = 1 : m$  do
5:          $\theta_j \sim \text{NIG}(\lambda_0)$   $\triangleright$  Sample  $\theta_j$  from NIG.
6:       end for
7:       for all  $j = 1 : K + m, j \neq i$  do
8:          $L_j = \text{likelihood}(w_i, \theta_j)$   $\triangleright$  Update likelihood for all  $\theta$  except  $\theta_i$  given  $w_i$ .
9:       end for
10:       $P_{-i=1:K} = b \sum_{-i} L_{-i}$   $\triangleright$  Calculate probability of existing cluster.
11:       $P_{-i=K:K+m} = b\alpha/m L_m L_{-i}$   $\triangleright$  Calculate probability of new cluster.
12:       $\theta_i = \theta_j$  according to above  $P_{-i}$   $\triangleright$  Sample  $\theta_i$  accord. to above prob.
13:      Remove unused clusters.
14:    end for
15:    for all  $j = 1 : K$  do
16:       $\theta_j \sim p(\theta_j | y)$   $\triangleright$  Update  $\theta_j$ .
17:    end for
18:  end for
19:  return summary on  $\theta_k$  for  $k$  line segments.
20: end procedure

```

4.4 Results

In Figure 4.5 we show four Bayesian point estimates of the sampling process. These are examples that demonstrate the type of errors that are made in the inference process. In example (a) the segments are correctly sampled. In (b) the type of error is that of recognizing multiple segments where there is only one segment to the human observer. In (c) the error

1601 is due to the fact that some segments contain very few points. In (d) the error stems from
 1602 line segments being chosen orthogonal to the actual segment.

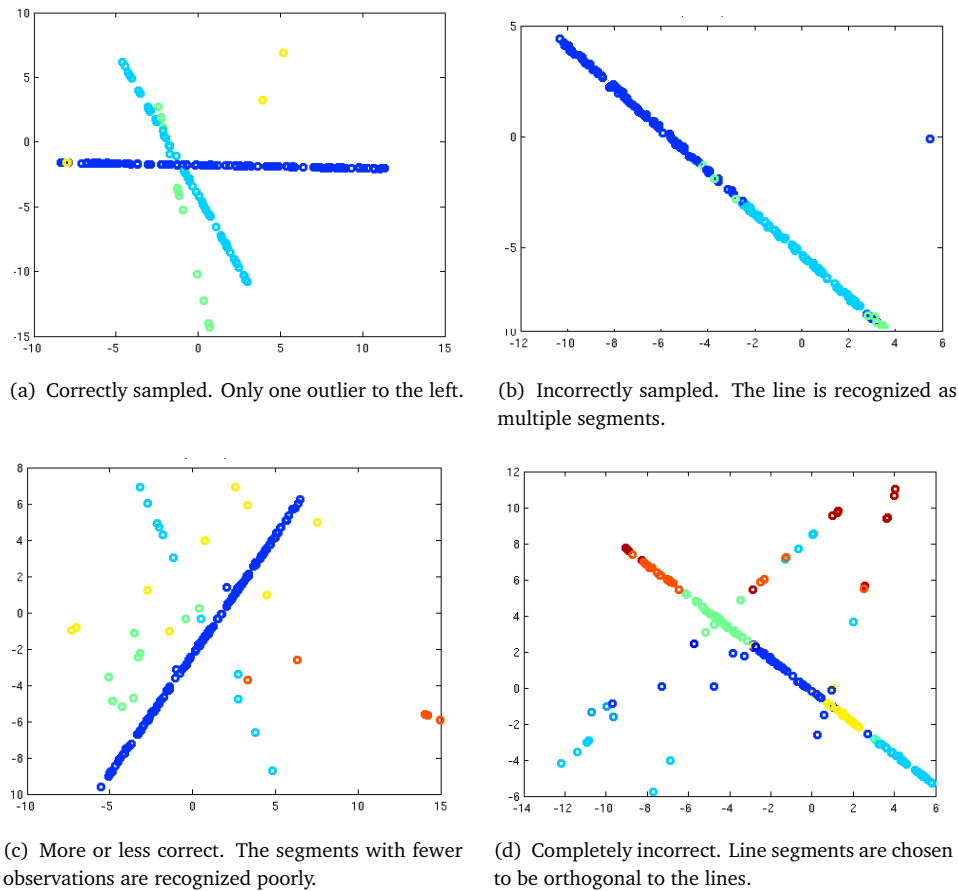


Figure 4.5: Bayesian point estimates of the sampling process with varying outcomes.

1603 The results over a larger dataset can be measured with clustering metrics as visualized in
 1604 Figure 4.6. The clustering performance of the segment detection algorithm, measured by
 1605 the clustering index, such as the Rand Index, the Adjusted Rand Index, and the Hubert
 1606 metric, show all reduced performance (see Figure 4.6) compared to line detection (without
 1607 constraints on segment size).

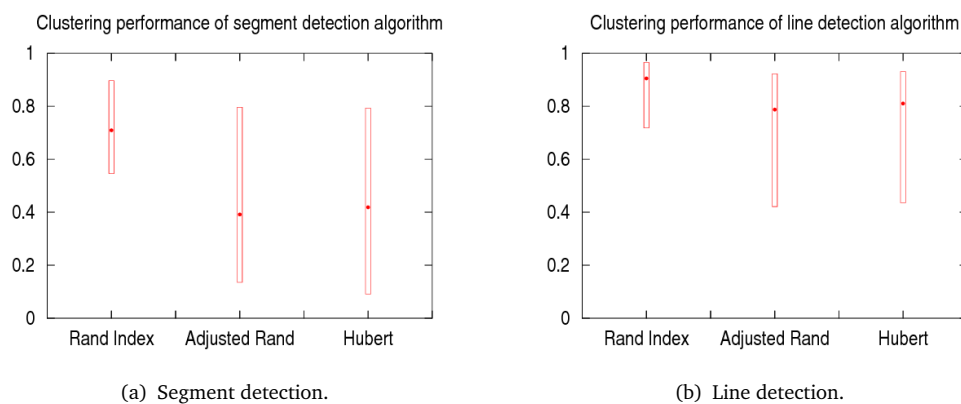


Figure 4.6: Segment detection performs much worse than line detection across all three clustering performance indicators. Perfect clustering is indicated by 1.0 for Rand Index, Adjusted Rand Index, and Hubert.

4.5 Chapter Conclusions

From Chapter 3 we know that segment estimation is a much harder problem than line estimation. In this chapter we used an advanced method, namely Gibbs sampling with auxiliary variables to perform inference over an infinite set of line segments. The auxiliary variable Gibbs sampling method converges faster than the ordinary Metropolis-Hastings algorithm by postulating multiple segments rather than only one.

However, the segment estimation problem remains a challenge for the current inference methods. The target probability density has modes that each needs to be found and tend to be separated by very low probability regions. In Chapter 5 we will introduce new sampling methods that will cope with this challenge.

1618

1619

TRIADIC SPLIT-MERGE SAMPLER

Contents

1621

1622

1623

1624

1625

1626

1627

1628

1629

1630

This chapter introduces a new sampling method called the triadic split-merge sampler. The reason is that naive implementations of MCMC methods suffer from slow convergence in machine vision due to the complexity of the parameter space. Towards this blocked Gibbs and split-merge samplers have been developed that assign multiple data points to clusters at once. The triadic split-merge sampler improves on these samplers by defining split and merge steps between two and three clusters. This has two advantages. First, it reduces the asymmetry between the split and merge steps. Second, it is able to propose a new cluster that is composed out of data points from two different clusters. Both advantages speed up the convergence of the sampler on a line estimation problem.

Published in

1632

1633

A.C. van Rossum, H.X. Lin, J. Dubbeldam, and H.J. van den Herik. Triadic Split-Merge Sampler. *The 10th International Conference on Machine Vision, ICMV 2017, Vienna, Austria, November 13-November 15, 2017.*

Outline

1635

1636

1637

1638

1639

1640

1641

We introduce the class of split-merge samplers as part of the MCMC samplers (Section 5.1). A conventional split-merge sampler, labeled the dyadic split-merge sampler, is detailed (Section 5.2). A new split-merge sampler, the triadic split-merge sampler is introduced (Section 5.3). The results for inference over lines are compared between the conventional dyadic sampler and the new triadic sampler (Section 5.4). Finally, we give the chapter conclusions and we describe how we can further improve the inference procedure (Section 5.5). They will be the basis of the next chapter.

5.1 The Class of Split-Merge Samplers

1642

In clustering models there is a hierarchical structure. At the lowest level there are individual data points. At a higher level the points are grouped into clusters. Split-merge samplers

are samplers that take into account such structure as already described in Section 2.3.6. Rather than moving data points one by one from an old to a cluster, split-merge samplers can perform moves that operate on partitions of the dataset. For example, a cluster can be split into two clusters in one single move or two clusters can be merged into one cluster in another single move.

5.1.1 Split and Merge Moves

One of the first split-merge samplers has been defined for so-called point sources in nuclear imaging (Stawinski et al., 1998). This split-merge sampler proposes split and merge moves that are defined locally. For instance, two clusters that are close to each other are a candidate for a merge step into one cluster. By defining pairs of split and merge steps with the right probabilities, a properly balanced Metropolis-Hastings step can be performed (Section 2.3.6).

In the hierarchical models of lines and segments we have used a Dirichlet process as prior. Split-merge samplers have been defined with a Dirichlet prior (Dahl, 2003; Jain and Neal, 2004). These split-merge samplers operate on one or two clusters and are defined in the thesis as dyadic split-merge samplers (Section 5.2). The split step is different per sampler: (1) the simple random split procedure does not take into account the data distribution, (2) the sequentially allocated merge-split sequentially assigns data towards one of the two clusters that is the better fit.

Other examples of split-merge samplers are a sampler that uses sub-cluster splits (Chang and Fisher III, 2013) a sampler that uses data-driven jumps besides split-merge steps (Hughes et al., 2012), a sampler that uses data-driven jumps, split-merge steps, and operates on a hierarchical Dirichlet process (Bryant and Sudderth, 2012), and a sampler that generalizes split-merge steps to birth-death steps (with multiple clusters generated simultaneously) (Hughes and Sudderth, 2013). We will introduce a sampler that generalizes the dyadic sampler to moves over two and three clusters (Section 5.3).

5.1.2 Dirichlet process Prior

Let us first reiterate the Dirichlet process. We will consider a Dirichlet process as a prior on the distribution over parameters G . The form of this model is:

$$\begin{aligned} y_i | \theta_i &\sim F(\theta_i) \\ \theta_i | G &\sim G \\ G &\sim DP(H, \alpha) \end{aligned} \tag{5.1}$$

The split and merge steps dictate the simultaneous assignment of observations y_i unto parameters θ_i .

1675 5.2 Conventional Split-Merge Sampler

1676 The conventional split-merge sampler (see Jain and Neal, 2004) splits a single cluster into
 1677 two clusters, and merges two clusters into a single cluster. Hence, this split-merge sampler
 1678 operates on two clusters at each time step. Therefore we will call their algorithm a dyadic
 1679 split-merge sampler in contrast with our approach (van Rossum et al., 2017). Below we
 1680 describe this dyadic split-merge sampler in pseudo-code (see Algorithm 11).

Algorithm 11 Dyadic split-merge sampler

```

1: procedure DYADIC SPLIT-MERGE SAMPLER( $c$ )                                 $\triangleright$  Accepts cluster assignments
    $c$  of length  $N$  (besides Metropolis-Hastings acceptance factors  $a(c', c)$  and a split procedure e.g.
   SIMPLERANDOMSPLIT) and returns a (potentially) updated cluster assignment vector  $c'$ .
2:    $i \sim U(1, N)$                                                           $\triangleright$  Sample  $i$  random uniformly over cluster assignments.
3:    $j \sim U(1, N) \cap i$                                                   $\triangleright$  Sample  $j$  also random uniformly, but with  $j \neq i$ .
4:    $S_R = \{c_i, c_j\}$                                                         $\triangleright$  Sampled clusters  $c_i, c_j$ .
5:    $S_I = \{c_x\}$  with  $c_x \in S_R$  for  $x \in \{1, \dots, N\}$                  $\triangleright$  All data in clusters  $c_i, c_j$ .
6:    $S_E = S \cap S_R$                                                         $\triangleright$  All data in clusters  $c_i, c_j$  excluding  $S_R$ .
7:    $N_S = \text{unique}(S_R)$ 
8:   if  $N_S = 1$  then                                                          $\triangleright$  Case:  $i, j$  belong to the same cluster.
9:      $c_i^{(2)} = c_k$  with  $c_k \notin \{c_1, \dots, c_N\}$                      $\triangleright$  Sample new cluster for  $c_i^{(2)}$ .
10:     $c_j^{(2)} = c_j^{(1)}$                                                         $\triangleright$  Keep  $c_j$  the same.
11:     $c_e^{(2)} = \text{SPLITPROCEDURE}(S_E, c_i^{(2)}, c_j^{(2)})$                  $\triangleright$  After  $c_i^{(2)}, c_j^{(2)}$  assign  $S_E$ .
12:    for all  $m \notin S_I$  do
13:       $c_m^{(2)} = c_m^{(1)}$                                                         $\triangleright$  Data points in clusters other than  $c_i, c_j$  are not adjusted.
14:    end for
15:     $c' = \{c_i^{(2)}, c_j^{(2)}, c_e^{(2)}, c_m^{(2)}\}$ 
16:     $a = a_{\text{split}}(c', c)$  according to Eq. ??                             $\triangleright$  MH acceptance for a split.
17:  else                                                          $\triangleright$  Case:  $i, j$  belong to different clusters  $c_i \neq c_j$  ( $N_S = 2$ ).
18:    for all  $q \in S_I$  do
19:       $c_q^{(1)} = c_j^{(1)}$                                                         $\triangleright$  Assign all data points in  $c_i$  and  $c_j$  to  $c_j$ .
20:    end for
21:    for all  $m \notin S_I$  do
22:       $c_m^{(1)} = c_m^{(2)}$                                                         $\triangleright$  Data points in clusters other than  $c_i, c_j$  are not adjusted.
23:    end for
24:     $c' = \{c_q^{(1)}, c_m^{(1)}\}$ 
25:     $a = a_{\text{merge}}(c', c)$  according to Eq. ??                             $\triangleright$  MH acceptance for a merge.
26:  end if
27:   $u \sim U(0, 1)$                                                           $\triangleright$  Sample  $u$  between 0 or 1 uniformly.
28:  if  $a < u$  then
29:     $c' = c$                                                           $\triangleright$  Reject  $c'$  by setting it to  $c$ 
30:  end if
31:  return  $c'$ , the (updated) cluster assignment vector:  $c \rightarrow c'$ .
32: end procedure

```

1681 In Algorithm 11 the notation $c_i^{(2)}$ is used to signify that the cluster assignment c_i has 2 clusters
 1682 under consideration. In the dyadic algorithm we could have used c_i^{merge} and c_i^{split} , however
 1683 in the triadic algorithm (see Algorithm 14) with multiple split and merge operations the
 1684 latter notation would become confusing.

1685 The dyadic split-merge sampler in Algorithm 11 samples two distinct data items. If the data
 1686 items belong to the same cluster a split step is attempted. If the data items belong to different

Algorithm 12 Simple random split

```

1: procedure SIMPLERANDOMSPLIT( $S, c_0, c_1$ )  $\triangleright$  Accepts unassigned set  $S$  and cluster indices  $c_0, c_1$ ,
   returns cluster assignment  $c'_m$ .
2:   for all  $m \in S$  do
3:      $c'_m \sim \text{Cat}(c_0, c_1)$  with equiprobable  $p(c_0) = p(c_1) = \frac{1}{2}$ .
4:   end for
5:   return  $c'_m$ , the cluster assignment for  $S$ .
6: end procedure

```

clusters a merge step is attempted. The split procedure itself is the so-called simple random split (Algorithm 12) that assigns data items with the same probability to one of the parts of the split cluster without any consideration for a proper data fit.

5.2.1 Acceptance for the Split Step

The acceptance ratio contains the Metropolis ratio to step from c to c' :

$$\frac{P(c')L(c'|y)}{P(c)L(c|y)}. \quad (5.2)$$

Additionally, the Hastings correction is applied because of the asymmetry of the proposal distribution in the form of $q(c|c')/q(c'|c)$:

$$a_{split}(c^{(2)}, c^{(1)}) = \min \left[1, \frac{q(c^{(1)}|c^{(2)})}{q(c^{(2)}|c^{(1)})} \frac{P(c^{(2)})}{P(c^{(1)})} \frac{L(c^{(2)}|y)}{L(c^{(1)}|y)} \right]. \quad (5.3)$$

The notation $c^{(2)}$ is used to indicate that the cluster index vector is referencing 2 unique clusters (in this case after the split step).

The prior distribution is represented by a Chinese Restaurant Process with concentration parameter α and no discount factor. Data not yet assigned is assigned (1) with probability $\alpha/(n + \alpha)$ to a new cluster and (2) with probability $n_c/(n + \alpha)$ to an existing cluster c . Here n is the total number of assigned data points, n_c is the number of data points assigned to cluster c . There are D clusters. Hence, the prior over clusters will be:

$$P(c) = \frac{\Gamma(\alpha)}{\Gamma(\alpha + n)} \alpha^D \prod_{c_l} \Gamma(n_{c_l}) = \alpha^D \frac{\prod_{c_l} (n_{c_l} - 1)!}{\prod_{k=1}^n (\alpha + k - 1)}. \quad (5.4)$$

In the prior distribution ratio before and after the split step many of the factors drop out. There is one factor α remaining and the number of data points in the split cluster is part of the equation. There is no dependency on other clusters or the total number of data points. We can simplify the formula using the Beta function $B(a, b) = \Gamma(a)\Gamma(b)/\Gamma(a + b)$ with $\Gamma(x) = (x - 1)!$ the Euler-Gamma function:

$$\frac{P(c^{(2)})}{P(c^{(1)})} = \alpha \frac{(n_{c_i^{(2)}} - 1)!(n_{c_j^{(2)}} - 1)!}{(n_{c_i^{(1)}} - 1)!} = \alpha B(n_{c_i^{(2)}}, n_{c_j^{(2)}}). \quad (5.5)$$

1706 The likelihood can be written as a product over all observations y_i or as a product over
 1707 clusters with each cluster a product over its observations y_k :

$$L(c|y) = \prod_{c=1}^D \prod_{k:c_k=c} p(y_k|\phi). \quad (5.6)$$

1708 Here we write $p(y_k|\theta_k)$ rather than assuming conjugacy between the likelihood $F(\theta_k)$ and
 1709 the prior distribution $H(\theta_k)$ (see Dahl, 2005). (In the case of conjugacy we can analyti-
 1710 cally calculate $\int F(\theta_k)dH(\theta_k)$ which speeds up inference, but which restricts our choice of
 1711 likelihoods and priors).

1712 With the above formula for the likelihood, we can calculate the likelihood ratio of two clus-
 1713 ters versus a single cluster:

$$\frac{L(c^{(2)}|y)}{L(c^{(1)}|y)} = \frac{\prod_{k:c_k^{(2)}=c_i^{(2)}} p(y_k|\phi) \prod_{k:c_k^{(2)}=c_j^{(2)}} p(y_k|\phi)}{\prod_{k:c_k^{(1)}=c_i^{(1)}} p(y_k|\phi)}. \quad (5.7)$$

1714 The split step determines the probability of a particular split. Algorithm 11 commences with
 1715 picking two random points. These two points are henceforth are already assigned to distinct
 1716 clusters. Only the remaining points have to be assigned (see Figure 5.1).

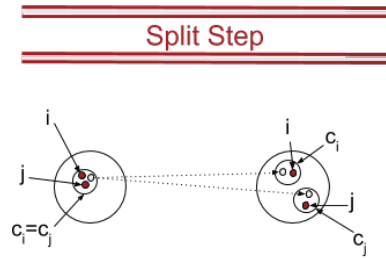


Figure 5.1: A split step. Points i and j are already assigned to separate clusters c_i and c_j . Each next point is assigned to one of the clusters with probability $\frac{1}{2}$, indicated by dotted arrows.

1717 The remaining points are assigned with equal probability $\frac{1}{2}$ to $c_i^{(2)}$ and $c_j^{(2)}$:

$$q(c^{(2)}|c^{(1)}) = \left(\frac{1}{2}\right)^{-2+n_{c_i^{(2)}}+n_{c_j^{(2)}}} = \left(\frac{1}{2}\right)^{-2+n_{c_i^{(1)}}}. \quad (5.8)$$

1718 The probability of the reverse of the split operation is exactly 1. There is only one way in
 1719 which a single cluster can be the starting state for a split operation. It must have had all
 1720 points assigned to it. This means that the ratio with respect to the assignment of points over
 1721 clusters in the split step becomes:

$$\frac{q(c^{(1)}|c^{(2)})}{q(c^{(2)}|c^{(1)})} = \frac{1}{\left(\frac{1}{2}\right)^{n_{c_i^{(2)}}+n_{c_j^{(2)}}-2}} = 2^{-2+n_{c_i^{(1)}}}. \quad (5.9)$$

Only basic identities are used and the fact that the number of data items does not change after a split, $n_{c_i^{(2)}} + n_{c_j^{(2)}} = n_{c_i^{(1)}}$. Note that the reverse split transition looks like a ‘merging’ operation. The merge step, however, is defined independently and its description can be found in the next section.

5.2.2 Acceptance for the Merge Step

Acceptance of a merge step consists of the same components as that of the split step.

$$a_{merge}(c^{(1)}, c^{(2)}) = \min \left[1, \frac{q(c^{(2)}|c^{(1)})}{q(c^{(1)}|c^{(2)})} \frac{P(c^{(1)})}{P(c^{(2)})} \frac{L(c^{(1)}|y)}{L(c^{(2)}|y)} \right]. \quad (5.10)$$

$$\frac{P(c^{(1)})}{P(c^{(2)})} = \alpha^{-1} \frac{(n_{c_i^{(1)}} - 1)!}{(n_{c_i^{(2)}} - 1)!(n_{c_j^{(2)}} - 1)!} = \frac{1}{\alpha B(n_{c_i^{(2)}}, n_{c_j^{(2)}})}. \quad (5.11)$$

$$\frac{L(c^{(1)}|y)}{L(c^{(2)}|y)} = \frac{\prod_{k:c_k^{(1)}=c_i^{(1)}} p(y_k|\phi)}{\prod_{k:c_k^{(2)}=c_i^{(2)}} p(y_k|\phi) \prod_{k:c_k^{(2)}=c_j^{(2)}} p(y_k|\phi)}. \quad (5.12)$$

$$\frac{q(c^{(2)}|c^{(1)})}{q(c^{(1)}|c^{(2)})} = \left(\frac{1}{2} \right)^{-2+n_{c_i^{(1)}}} = 2^{2-n_{c_i^{(1)}}}. \quad (5.13)$$

The ratios of the merge step are the inverse of the ratios of the split step. That is, Eq. 5.11 is the inverse of Eq. 5.5, Eq. 5.12 is the inverse of Eq. 5.7, and Eq. 5.13 is the inverse of Eq. 5.9.

5.2.3 Sequentially-Allocated Merge-Split Sampler

A variant on the conventional split-merge sampler is the sequentially allocated merge-split (SAMS) sampler¹ (Dahl, 2003). The simple random split procedure of Algorithm 12 is replaced by a procedure that sequentially assigns observations to clusters rather than splitting the data random uniformly over the split clusters.

Algorithm 13 Sequentially Allocated Merge-Split

```

1: procedure SAMS( $S, c_0, c_1$ )    ▷ Accepts unassigned set  $S$ , cluster indices  $c_i$ , and  $p(y_k|\theta_{c_i})$  with
    $i = 0, 1$ , returns cluster assignment  $c'_m$ .
2:    $T = \text{random\_shuffle}(S)$ 
3:   for all  $m \in T$  do
4:      $p(c_m = c_0 | c_0, c_1, \theta_{c_0}, \theta_{c_1}) = \frac{N_0 p(y_m | \theta_{c_0})}{N_0 p(y_m | \theta_{c_0}) + N_1 p(y_m | \theta_{c_1})}$ 
5:      $p(c_m = c_1 | c_0, c_1, \theta_{c_0}, \theta_{c_1}) = 1 - p(c_m = c_0 | c_0, c_1, \theta_{c_0}, \theta_{c_1})$ 
6:      $c'_m \sim p(c_m | c_0, c_1, \theta_{c_0}, \theta_{c_1})$ 
7:   end for
8:   return  $c'_m$ , the cluster assignment for  $S$ .
9: end procedure

```

In contrast to the simple random split, observations y_k are used in the SAMS to obtain cluster assignments that correspond with the data rather than cluster assignments independent of the data.

¹In the naming of split-merge or merge-split samplers, the order of merge split does not bear any significance.

5.3 Triadic split-merge sampler

The triadic split-merge sampler uses up to three clusters for a split or merge step (Fig. 5.2).




Figure 5.2: Right: dyadic MCMC picks two data items i, j random uniformly. If both are in the same cluster a split towards two clusters is attempted. If both are in distinct clusters a merge towards one cluster is attempted. Left: triadic MCMC picks three data items i, j, k random uniformly. If all three are in the same cluster a split towards two clusters is attempted. If the three items are in two clusters either a split into three (with probability $1 - \beta$) or a merge into a single cluster (with probability β) is attempted. If the three data items are in three distinct clusters a merge is attempted. There are no direct transitions from a single cluster to three clusters or the other way around.

The intuition behind the triadic split-merge sampler is twofold:

- In the dyadic sampler there is a large asymmetry between split and merge steps. There is only one way in which two clusters can be merged into one single cluster, while there are many ways in which one single cluster can be split into two clusters. This asymmetry is reduced by transitioning between two and three clusters. This is a straightforward improvement in balancing split and merge steps (for alternatives, see Wang and Russell (2015)).
- In practical optimization problems it might be useful to form a third cluster out of subsets of two other clusters. The dyadic MCMC sampler requires immediate steps in which (1) one of these clusters is split into two, (2) the other is split into two, and (3) the two new clusters are merged. This means that (a) mixing and hence convergence will be slow and (b) the intermediate steps might have very low probability and function as an unnecessary barrier between high probable states.

The algorithm is detailed in Algorithm 14. Compare with Algorithm 11. It starts with sampling three distinct points i, j , and k . These points can originate from one, two, or three distinct clusters with non-unique indices c_i, c_j , and c_k . Depending on the number of distinct clusters N_S , either a dyadic or triadic step is performed. If all points belong to the same cluster, $N_S = 1$, a split step into two clusters is attempted. If the points belong to two clusters $N_S = 2$, with probability β a merge step into one cluster is tried, with probability $1 - \beta$ a triadic split into three clusters will be considered. If the points belong to three distinct clusters, $N_S = 3$, a triadic merge step into two clusters will be attempted. All steps will be accepted or rejected according to the Metropolis-Hastings probability ratios in the following section. Note that there are no steps with which a single cluster is immediately split into three, nor is there is a step that merges three clusters into one.

Algorithm 14 Triadic split-merge sampler

```

1: procedure TRIADIC SPLIT-MERGE SAMPLER( $c$ )                                 $\triangleright$  Accepts
   cluster assignments  $c$  of length  $N$  (besides Metropolis-Hastings acceptance factors  $a(c', c)$  and a
   split procedure) and returns a (potentially) updated cluster assignment vector  $c'$ .
2:    $i \sim U(1, N)$                                  $\triangleright$  Sample  $i$  random uniformly over cluster assignments.
3:    $j \sim U(1, N) \cap i$                              $\triangleright$  Sample  $j$  also random uniformly, but with  $j \neq i$ .
4:    $k \sim U(1, N) \cap \{i, j\}$                      $\triangleright$  Sample  $k$  random uniformly, but with  $k \neq j, k \neq i$ .
5:    $S_R = \{c_i, c_j, c_k\}$                              $\triangleright$  Sampled clusters  $c_i, c_j, c_k$ .
6:    $S_I = \{c_x\}$  with  $c_x \in S_R$  for  $x \in \{1, \dots, N\}$   $\triangleright$  All data in clusters  $c_i, c_j, c_k$ .
7:    $S_E = S_I \cap S_R$                                  $\triangleright$  All data in clusters  $c_i, c_j, c_k$  excluding  $S_R$ .
8:    $N_S = \text{unique}(S_R)$ 
9:    $u \sim U(0, 1)$                                  $\triangleright$  Sample  $u$  between 0 or 1 uniformly.
10:  if  $N_S = 1$  then                                 $\triangleright$  Case:  $i, j, k$  belong to the same cluster.
11:    return  $c' = \text{DYADIC SPLIT-MERGE SAMPLER}(c)$ 
12:  else if  $N_S = 2$  and  $u < \beta$  then  $\triangleright$  Case: a cluster with one item and one with two items and
    $u < \beta$ .
13:    return  $c' = \text{DYADIC SPLIT-MERGE SAMPLER}(c)$ 
14:  else if  $N_S = 2$  and  $u \geq \beta$  then  $\triangleright$  Case: a cluster with one item and one with two items and
    $u \geq \beta$ .
15:     $c_i^{(3)} = c_k$  with  $c_k \notin \{c_1, \dots, c_N\}$   $\triangleright$  Sample new cluster for  $c_i^{(3)}$ .
16:     $c_j^{(3)} = c_j^{(2)}$   $\triangleright$  Keep  $c_j$  the same.
17:     $c_e^{(3)} = \text{SPLITPROCEDURE}(S_E, c_i^{(3)}, c_j^{(3)})$   $\triangleright$  After  $c_i^{(3)}, c_j^{(3)}$  assign  $S_E$ .
18:    for all  $m \notin S_I$  do
19:       $c_m^{(3)} = c_m^{(2)}$   $\triangleright$  Data points in clusters other than  $c_i, c_j$  are not adjusted.
20:    end for
21:     $c' = \{c_i^{(3)}, c_j^{(3)}, c_e^{(3)}, c_m^{(3)}\}$ 
22:     $a = a_{\text{split}}(c', c)$  according to Eq. ??  $\triangleright$  MH acceptance for a split.
23:  else  $\triangleright$  Case:  $i, j, k$  belong to three different clusters  $c_i \neq c_j \neq c_k$  ( $N_S = 3$ ).
24:     $S_L = S_I \cap \{c_i^{(3)}, c_j^{(3)}\}$   $\triangleright$  Data in clusters  $c_i, c_j, c_k$  except for  $i$  and  $j$  itself.
25:     $\{c_i^{(2)}, c_j^{(2)}\} = \text{SAMS}(S_L, c_i^{(3)}, c_j^{(3)})$   $\triangleright$  Assign data points in  $c_i, c_j, c_k$  to  $c_i, c_j$ .
26:    for all  $m \notin S_L$  do
27:       $c_m^{(2)} = c_m^{(3)}$   $\triangleright$  Data points in clusters other than  $S_L$  are not adjusted.
28:    end for
29:     $c' = \{c_i^{(2)}, c_j^{(2)}, c_m^{(2)}\}$ 
30:     $a = a_{\text{merge}}(c', c)$  according to Eq. 5.21  $\triangleright$  MH acceptance for a merge.
31:  end if
32:   $u \sim U(0, 1)$   $\triangleright$  Sample  $u$  between 0 or 1 uniformly.
33:  if  $a < u$  then
34:     $c' = c$   $\triangleright$  Reject  $c'$  by setting it to  $c$ 
35:  end if
36:  return  $c'$ , the (updated) cluster assignment vector:  $c \rightarrow c'$ .
37: end procedure

```

1767 Sampling random uniformly for three unique items is implemented through a random shuf-
 1768 fle algorithm, in particular the modern version of the Fisher-Yates shuffle introduced by
 1769 Durstenfeld (1964) and picking the first three items.

1770 The Metropolis-Hastings probabilities for the triadic split and merge steps are calculated in
 1771 the following Sections 5.3.1 and 5.3.2.

1772 5.3.1 Acceptance for the Split Step

1773 In the triadic split-merge sampler there are two splitting steps. It is possible to split according
 1774 to the dyadic split-merge sampler. However, given two clusters there are (split) jumps to
 1775 three states as well as (merge) jumps to single states again. To account for this asymmetry
 1776 another Hastings correction is applied to establish detailed balance.

$$a_{split}(c^{(2)}, c^{(1)}) = \min \left[1, \frac{r(c^{(1)}|c^{(2)}) q(c^{(1)}|c^{(2)}) P(c^{(2)}) L(c^{(2)}|y)}{r(c^{(2)}|c^{(1)}) q(c^{(2)}|c^{(1)}) P(c^{(1)}) L(c^{(1)}|y)} \right]. \quad (5.14)$$

1777 Here we have one additional term compared to the split step from one cluster to two clusters:

$$\frac{r(c^{(1)}|c^{(2)})}{r(c^{(2)}|c^{(1)})} = \frac{\beta}{1}. \quad (5.15)$$

1778 The parameter β is free to control, as long as $0 < \beta < 1$ (to maintain ergodicity). The
 1779 transition from two states to three states is another split step:

$$a_{split}(c^{(3)}, c^{(2)}) = \min \left[1, \frac{r(c^{(2)}|c^{(3)}) q(c^{(2)}|c^{(3)}) P(c^{(3)}) L(c^{(3)}|y)}{r(c^{(3)}|c^{(2)}) q(c^{(3)}|c^{(2)}) P(c^{(2)}) L(c^{(2)}|y)} \right]. \quad (5.16)$$

1780 The fraction with r reads as follows:

$$\frac{r(c^{(2)}|c^{(3)})}{r(c^{(3)}|c^{(2)})} = \frac{1}{1 - \beta}. \quad (5.17)$$

1781 The fraction with q uses the total number of data points n_c in the clusters:

$$\frac{q(c^{(2)}|c^{(3)})}{q(c^{(3)}|c^{(2)})} = \frac{\left(\frac{1}{2}\right)^{n_c-2}}{\left(\frac{1}{3}\right)^{n_c-3}} = (3^{n_c-3}) (2^{2-n_c}) = \left(\frac{3}{2}\right)^{n_c} \frac{2^2}{3^3}. \quad (5.18)$$

1782 To move from 2 clusters to 3 clusters the probability is a $1/3$ for each cluster index in vector
 1783 c (except for the three data items already selected randomly, hence $n_c - 3$). To move back,
 1784 the probability is a $1/2$ and there are only two data items randomly assigned beforehand.
 1785 The fraction with P uses the number of data points in each of the clusters before and after
 1786 the step:

$$\frac{P(c^{(3)})}{P(c^{(2)})} = \alpha \frac{(n_{c_i^{(3)}} - 1)!(n_{c_j^{(3)}} - 1)!(n_{c_k^{(3)}} - 1)!}{(n_{c_i^{(2)}} - 1)!(n_{c_j^{(2)}} - 1)!} = \alpha \frac{B(n_{c_i^{(3)}}, n_{c_j^{(3)}}, n_{c_k^{(3)}})}{B(n_{c_i^{(2)}}, n_{c_j^{(2)}})}. \quad (5.19)$$

1787 Here we introduced a generalized Beta function $B(a, b, c) = \Gamma(a)\Gamma(b)\Gamma(c)/\Gamma(a + b + c)$ with
 1788 $\Gamma(x) = (x - 1)!$ the Gamma function. The likelihood ratio becomes:

$$\frac{L(c^{(3)}|y)}{L(c^{(2)}|y)} = \frac{\prod_{m:c_m^{(3)}=c_i^{(3)}} P(y_m|\phi) \prod_{m:c_m^{(3)}=c_j^{(3)}} P(y_m|\phi) \prod_{m:c_m^{(3)}=c_k^{(3)}} P(y_m|\phi)}{\prod_{m:c_m^{(2)}=c_i^{(2)}} P(y_m|\phi) \prod_{m:c_m^{(2)}=c_j^{(2)}} P(y_m|\phi)}. \quad (5.20)$$

5.3.2 Acceptance for the Merge Step

The merge step from two to one cluster is analogous to the split step:

$$a_{merge}(c^{(1)}, c^{(2)}) = \min \left[1, \frac{r(c^{(2)}|c^{(1)}) q(c^{(2)}|c^{(1)}) P(c^{(1)}) L(c^{(1)}|y)}{r(c^{(1)}|c^{(2)}) q(c^{(1)}|c^{(2)}) P(c^{(2)}) L(c^{(2)}|y)} \right]. \quad (5.21)$$

The merge step from three clusters to two clusters is:

$$a_{merge}(c^{(2)}, c^{(3)}) = \min \left[1, \frac{r(c^{(3)}|c^{(2)}) q(c^{(3)}|c^{(2)}) P(c^{(2)}) L(c^{(2)}|y)}{r(c^{(2)}|c^{(3)}) q(c^{(2)}|c^{(3)}) P(c^{(3)}) L(c^{(3)}|y)} \right]. \quad (5.22)$$

Note that all the fractions in Eq. 5.22 are the inverse of the fractions in Eq. 5.16. Inverting Eq. 5.17–5.20 will be left to the reader.

One additional issue we have to consider. When merging three clusters into two we can (1) distribute the data over all three clusters or (2) alternatively, keep the data in two clusters assigned to these clusters and only distribute the data in the third cluster over the other two clusters. The second and alternative option however would introduce unnecessary asymmetry with the merge step. In other words, Eq. 5.23 is not the inverse of Eq. 5.18. In contrast, the equation is similar to splitting one cluster across two as in Eq. 5.9:

$$\frac{q_{alt}(c^{(3)}|c^{(2)})}{q_{alt}(c^{(2)}|c^{(3)})} = 2^{-2+n_c}. \quad (5.23)$$

Hence the first option is entertained and the q -fraction is exactly the inverse of Eq. 5.18.

A second issue has to be considered, namely the inclusion or exclusion of direct operations between a single cluster and three clusters. This is because factors such as

$$\frac{P(c^{(3)})}{P(c^{(1)})} = \alpha^2 \frac{(n_{c_i^{(3)}} - 1)!(n_{c_j^{(3)}} - 1)!(n_{c_k^{(3)}} - 1)!}{(n_{c_i^{(1)}} - 1)!}, \quad (5.24)$$

become very small and although compensated by a large q fraction, remain further away from an acceptance factor of 1. Note that by the ability to split a single cluster into two and then into three, there is no ergodic argument to introduce also the immediate step.

5.4 Results

The problem we use to test our sampler is a well-known problem in computer vision, namely that of the inference of line parameters (slope and intercept) given data points. Rather than ordinary linear regression, in computer vision there is a mixture of lines that have to be estimated. Moreover, the number of lines is not known in advance. To solve this problem we use the Dirichlet process mixture (Eq. 5.1) with a normal distribution $N(0, \sigma_0)$ to generate the line parameters and a likelihood function that defines points to be uniformly distributed across a line of length 20 and deviating from the line according to a normal distribution $N(0, \sigma_1)$.

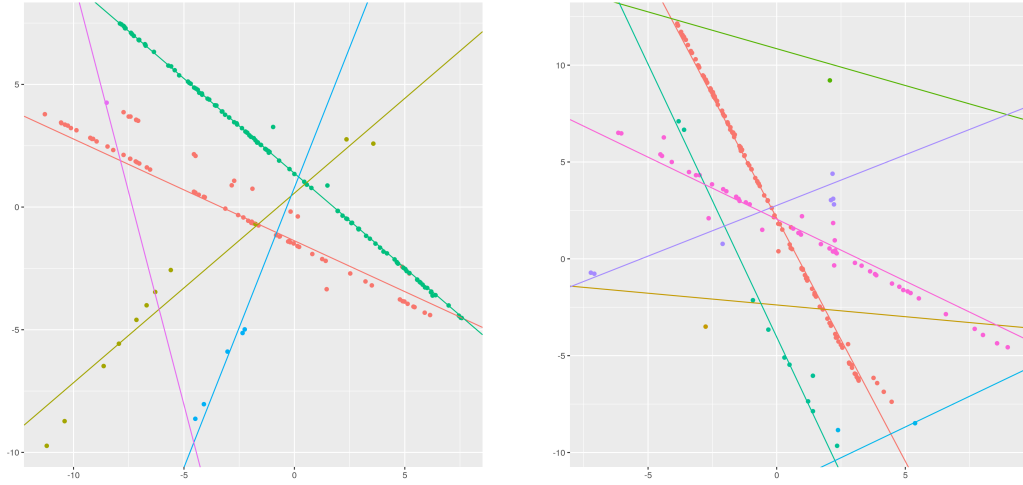


Figure 5.3: Two examples of fitting a mixture of lines to data items scattered over a two-dimensional space. The lines drawn are inferred using the triadic sampler. The lines are not the ground truth, but are meant to demonstrate the typical errors made by fitting methods. Note, for example, that there are mistakes in both the assignment of points to lines as well as the line parameters (slope and intercept). Left: In example 1 two lines with similar slope are seen as the same line. Right: In example 2 points on one vertical line are assigned to multiple lines. In example 1 and 2 slopes are not always through the points.

1815 5.4.1 Implementation

1816 The sampler is open-source² implemented in C++ which means that (a) it is computationally
 1817 fast, (b) it can be run on embedded devices if a cross-compiler is available and the Eigen3
 1818 library is ported. Note, that due to the fact that the simulator uses a lot of random numbers
 1819 the system should use a modern compiler (g++-6 or newer) and should have enough entropy
 1820 available³. Rather than a random scan, the implementation uses a fixed scan as advocated
 1821 in the literature (MacEachern, 2007).

1822 To speed up the sampler most calculations are done in log-space. Consider $v = u + 1$. The
 1823 ratio with probabilities (Eq. 5.5 and 5.19) becomes:

$$\log \frac{P(c^{(v)})}{P(c^{(u)})} = \log(\alpha) + \sum_i \log \Gamma(n_{c_i^{(v)}}) - \sum_i \log \Gamma(n_{c_i^{(u)}}). \quad (5.25)$$

1824 The fraction with $q(\cdot)$ (Eq. 5.9 and 5.18) becomes:

$$\log \frac{q(c^{(v-1)}|c^{(v)})}{q(c^{(v)}|c^{(v-1)})} = (v - n_c - 1) \log(v - 1) - (v - n_c) \log(v). \quad (5.26)$$

1825 The fraction with r becomes, for example, (Eq. 5.17):

$$\log \frac{r(c^{(2)}|c^{(3)})}{r(c^{(3)}|c^{(2)})} = -\log(1 - \beta). \quad (5.27)$$

²Code can be found at <https://code.annevanrossum.nl/noparama>.

³On Linux this can be checked in `/proc/sys/kernel/random/entropy_avail`.

The log-probability to calculate the likelihood given by a multivariate Normal distribution is well-known.

5.4.2 Comparison

The Triadic sampler using SAMS is compared with the Jain-Neal Dyadic sampler using SAMS and an auxiliary variable sampler with $m = 3$ (see algorithm 8 in Neal (2000)).

Method	Purity	Rand Index	Adjusted Rand Index
Dyadic sampler	0.80960	0.80580	0.56382
Auxiliary variables	0.87235	0.85879	0.68224
Triadic sampler	0.86405	0.87188	0.71067

Table 5.1: The purity, rand index, and adjusted rand index establishing the quality of the clustering method. The closer the values to one, the better the method performed. The purity metric assigns high values to clusters that do not have data points from other clusters (but does not penalize the number of clusters). The rand index computes similarity between clusters taking false negatives and false positives into account. The adjusted rand index accounts for chance. The adjusted rand index is most useful in our comparison.

In Table 5.1 the line estimation problem is compared for the dyadic sampler, an auxiliary variables sampler, and the proposed triadic sampler. The simulation is run with $\beta = 0.1$ so that a significant number of steps are tried between two and three clusters (rather than only between one and two clusters).

In Figures 5.4 to 5.6 the different metrics are visualized in the form of violin plots.

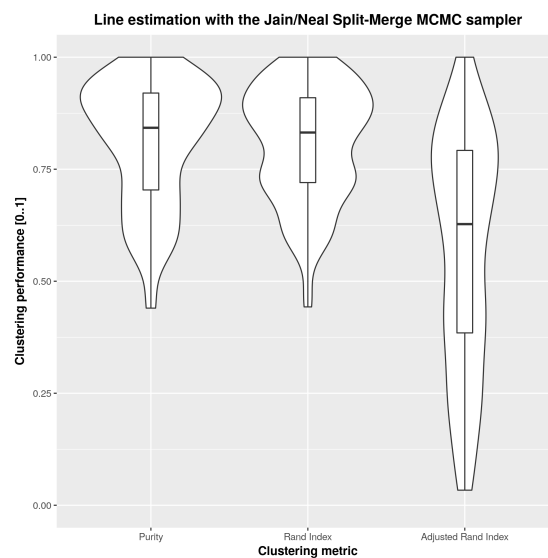


Figure 5.4: Line estimation with the diadic split-merge sampler. The same results as in Table 5.1, but visualized in a violin plot. The distribution over metric values are displayed in a vertical fashion.




Figure 5.5: Line estimation with the auxiliary variable sampler.

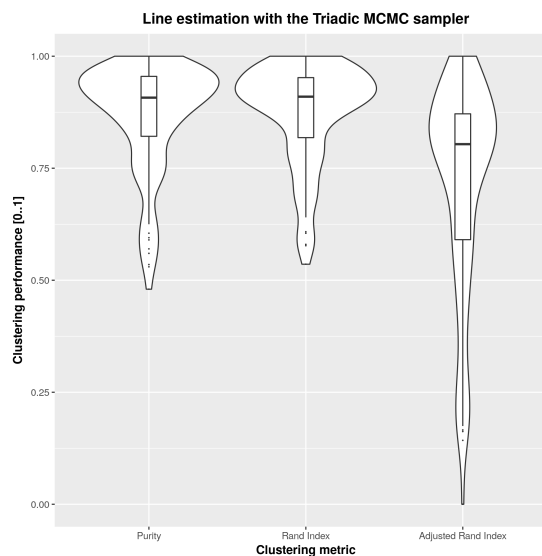


Figure 5.6: Line estimation with the triadic split-merge sampler (our inference method). The values are all shifted up towards one.

1836 The improvement in clustering is especially visible with the adjusted rand index.

1837 5.5 Chapter Conclusions

1838 A new split-merge sampler has been introduced, implemented, and applied to the computer
 1839 vision problem of line estimation. The sampler outperforms existing samplers, such as the
 1840 ordinary (dyadic) split-merge sampler (Jain and Neal, 2004) and auxiliary variable sampler
 1841 (Neal, 2000).

1842 This chapter answers our second research question: "How can we estimate the number of
 1843 lines simultaneously with line fitting in computer vision." The triadic split-merge sampler has

been used with likelihood functions that correspond to line fitting. It therefore estimates the number of lines and simultaneously performs line fitting. Moreover, the sampler is optimized to reassign points from three lines to two lines and the other way around. This means a hypothesized third line can be composed at once from two existing lines. These triadic steps accelerate the inference process as shown in Section 5.4.

Although the proposed split-merge sampler is able to mix considerably faster through a mixture model, it does not use global jumps directly based on the data. It is reasonable to suggest that MCMC methods benefit from combining the local jumps with global jumps, for example by a mixture of the local Metropolis-Hastings sampler with a Metropolized independence sampler (Jampani et al., 2015). We will introduce such a sampler in chapter 6.

ADVERSARIALLY TRAINED MCMC KERNELS

Contents

To use MCMC for volumetric inference it is necessary to be able to accelerate the algorithms even further. Volumetric objects exhibit more structure, which is reflected by symmetry.

Outline

We describe MCMC methods that cope with symmetric objects.

6.1 Data-Driven Inference

There are three aspects we would like to address in our inference engine.

The first aspect aims to have structure within our inference engine. The proposal distribution in a Markov chain, although moderately complex in the previous chapter, does not have much knowledge about the model. An artificial border is maintained that does not allow the inference engine to have knowledge about the model. The purpose of this is never articulated in particular. However, it is logical from a separation of concern. Such an inference engine (1) does not need to receive any information about the model and (2) is guaranteed to be general in the sense that it is not tailored to a particular model. This is nicely articulated by Tran et al. (2017) from which we quote.

“

Many existing probabilistic programming languages treat the inference engine as a black box, abstracted away from the model. These cannot capture probabilistic inferences that reuse the model's representation - a key idea in recent advances in variational inference, generative adversarial networks, and also in more classic inference.

”

The second aspect concerns the data. In MCMC the position for the chain is driven by (1) the prior, (2) the prior and the likelihood, (3) a sequence of priors and likelihood, (4) a

sequence of priors, likelihood and proposal distributions, basically anything, except for the data itself. Data-driven approaches would namely destroy the convergence of the Markov chain. To start an MCMC sampler in a data-driven manner and continue in a data-oblivious manner is a possible solution (Zhang and Perez-Cruz, 2017). Even better, it is possible to use a Metropolized independence sampler Jampani et al. (2015). Such a sampler samples independently from the previous state and uses global information. However, to work well its proposal distribution needs to match the target distribution quite well. Although, when combined with a local sampler, it might be sufficient to just be able to match the modes of the target distribution well.

The third aspect concerns the way we build our MCMC engine. The split-merge sampler of the previous chapter has been meticulously designed. If we admit a data-driven approach, we might as well adjust our MCMC engine using training samples. Note, that this training will be across a set of line, box, or scenes mixtures. The MCMC engine will not be able to learn just the parameters of a particular visual object. It will learn how to jump around (optionally, adaptively) from one visual object to the next or from one cluster configuration to the next. In other words, it will be able to teach itself to become a Triadic Split-Merge sampler if that happens to be a good engine. Is it possible to constrain the search through MCMC kernels such that its result is always converging in an MCMC sense? If we aim to learn the transition operator of our Markov chains, there is new literature that makes use of deep nets.

6.2 Learning the Transition Operator

There are multiple methods that can be used in a generative setting. We will discuss the three most prominent ones: (1) generative adversarial networks, (2) variational autoencoders, and (3) infusion training. This is far from an extensive categorization, worth studying are variational walkback (Goyal et al., 2017), stacked generative adversarial networks (Huang et al., 2016), generative latent optimization (Bojanowski et al., 2017), deep learning through the use of non-equilibrium thermodynamics (Sohl-Dickstein et al., 2015), denoising autoencoders, or generative stochastic networks, to name just a few.

6.2.1 Adversarial Training

Adversarial training has been extensively studied since the article on generative adversarial networks by Goodfellow et al. (2014). A particular adversarial setup for training an MCMC has been suggested as well (Song et al., 2017). The generator samples from a Markov chain. A discriminator subsequently needs to judge if its incoming data comes from the generator or if it is sampled from the actual data set. To start the process, the generator can run the chain from the model as well as from the data.

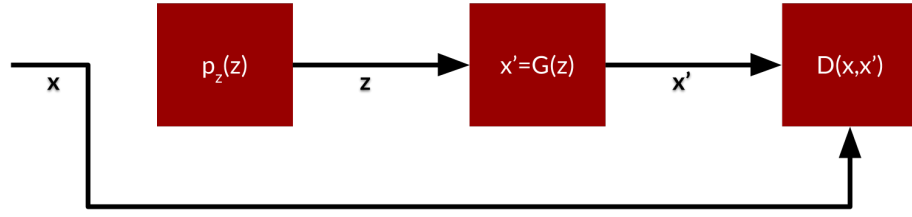


Figure 6.1: Left: $p_z(z)$ is a (prior) random distribution that generates random variables z . Middle: the generator G maps the random variables z to simulated data points x' . Right: the discriminator $D(x, x')$ compares the simulated data x' with the real data x . The generator tries to generate samples in such way that the discriminator has difficulties distinguishing the simulated from the real data.

1914 6.2.2 Variational Autoencoders

1915 Variational autoencoders (Kingma and Welling, 2013; Rezende et al., 2014) are ordinary
 1916 autoencoders with additional constraints on the latent variables. The latent variables in
 1917 autoencoder parlance are called the code. In a variational autoencoder the latent variables
 1918 are forced to approximately describe a unit Gaussian distribution. The autoencoder is trained
 1919 using a loss function that is composed out of (1) a generative loss, a mean squared error
 1920 that measures how accurately the network reconstructs its input, and (2) a latent loss, a KL-
 1921 divergence that measures how closely the latent variables match a unit Gaussian. To optimize
 1922 the KL divergence a reparameterization trick is applied. The encoder does not generate a
 1923 vector with real values, but generates a vector with means and standard deviations instead.

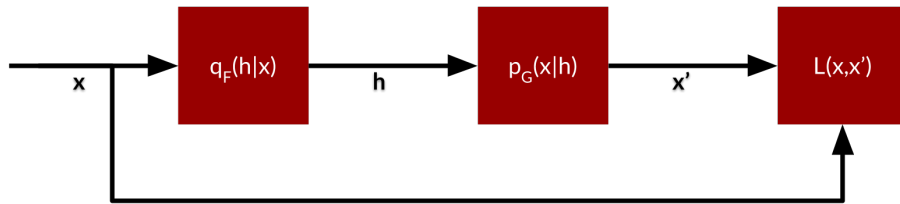


Figure 6.2: Left: $q_F(h|x)$ maps the data x to (hidden) random variables h . Middle: $p_G(x|h)$ maps the hidden random variables to reconstructed data x' . Right: $L(x, x')$ measures the similarity between x and x' .

1924 6.2.3 Infusion Training

1925 The transition operator can also be learned directly through infusion training. In infusion
 1926 training we gradually adjust totally unstructured noise to a target distribution as well. In this
 1927 method a particular data point is ‘infused’ into the Markov chain to bias the model sampling
 1928 to move towards this particular data point (and not another). In contrast to a generative
 1929 loss that is a mean squared error, this promises to have less blurry reconstructions.

6.3 Volumetric Models

In image processing autoencoders have been used for 2D shape recognition. To apply the same type of models to 3D point clouds, these point clouds are represented through voxels. The application of data-driven deep learning techniques, be it autoencoders, generative adversarial networks, or adversarial autoencoders promises similar good results in these 3D settings than in the current computer vision tasks.

The 3D ShapeNet model (Wu et al., 2015) exists of 3D voxel input that is piped through several stages with an increasing number of filters. The used voxel representation is a binary tensor. It assigns a value of 1 to each voxel that is inside the 3D object mesh and a value 0 to each voxel outside the mesh (empty space). The voxel sizes are fixed as well as the grid size (in this particular model the grid exists of 30x30x30 voxels). The inference model is a Deep Belief Network (DBN). Convolution operators, in the form of filters over small neighbourhoods, are used to reduce the number of model parameters (30x30x30 fully connected would be really many weights). The DBN is used in a supervised setting where shapes are trained with object labels. The model subsequently learns to generate shapes given an object label.

An unsupervised method in the form of a convolutional (volumetric) autoencoder (Sharma et al., 2016) has been applied to the same type of data. This (denoising) autoencoder, coined VConv-DAE maps from an entire voxel grid to another voxel grid. This work uses a combination of standard techniques, a dropout layer, a deconvolution layer, ReLu as well as sigmoid activation, but it is not in particular tailored to 3D point clouds.

Other representations than voxels are used. For example collections of 2D views and transformation parameters (Dosovitskiy et al., 2017). The most interesting are methods that work with raw data, the point cloud themselves. This alleviate the need to process the data and does not inadvertently increase the data dimensions, for example by artificially introduce voxels where there is no object present.

PointNet (?) directly operates on point clouds. To handle the input as a set of points (unordered), it uses a symmetric function over n input vectors and outputs a vector that is invariant to the input order. Typically sum and multiplication operators are such symmetric functions. After input and feature transforms by multi-layer perceptrons, a max pooling operator is used to map the input to a global feature. In the ModelNet40 shape classification benchmark there are more than 12000 CAD models from 40 object categories. PointNet achieves state of the art results compared to volumetric methods for a fraction of the computational costs.

Point clouds are also directly used in so-called deep kd-networks (Klokov and Lempitsky, 2017). A kd-tree is constructed by recursively picking the coordinate axis with the largest range of point coordinates and splitting the set of points into two subsets of equal size. These subsets are recursed into successively. The recursion stops at a particular level, depth D . The kd-networks are purported to outperform for example PointNet amongst other model architectures.

1970 A deep permutation equivariant (for semisupervised learning) and permutation invariant
1971 (for supervised learning) network has also been directly applied to point clouds (Ravanbakhsh
1972 et al., 2016). It does not reach the ModelNet40 accuracy levels from PointNet or the kd-
1973 networks though.

1974 PointNet++ Qi et al. (2017) introduces hierarchical structure to PointNet. This fits better
1975 non-uniform point distributions and seems to surpass kd-nets again on the ModelNet40 task.

1976

1977

RECOMMENDER ENGINE

Contents

1979

1980

1981

The described nonparametric Bayesian models (Chapter 3, 4, 5, and 6) are not limited to computer vision tasks. This chapter describes a recommender engine in which groups of runners are extracted from data collected from social media.

Outline

1983

1984

1985

1986

1987

We (1) introduce the form of the data at hand, (2) describe a multi-modal Von Mises-Uniform distribution to model the individual runners, (3) use a Dirichlet Process prior to group people, (4) use the previously described MCMC methods to perform inference, (5) show the results on an artificial and real-world data set, and (6) discuss ways with which the model can be expanded.

7.1 Application

1988

The data of people exercising can be considered binary (someone is either exercising or not in a particular timeslot). We do have however more information available. We know how often people have been exercising in a timeslot. This data has the form as visualized in Table 7.1.

Table 7.1: Example of the type of data about the timing of exercising. A person is represented by row, her preferences by column. There is not a predefined number of users or groups of users.

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
0	0	0	0	4	23	38	9	12	6	2	7	2	3	2	7	5	3	2	0	3	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	4	11	4	4	2	1	2	0	0	0
0	0	0	0	0	1	0	0	1	2	0	0	0	1	0	2	3	1	1	3	2	0	0	0
0	0	0	0	0	0	5	4	1	3	11	3	7	3	3	4	3	10	23	3	0	0	0	0
0	0	0	0	0	0	8	35	23	12	4	41	14	11	8	7	14	38	36	6	5	1	2	0
0	0	0	0	0	0	0	1	9	2	0	3	0	0	0	1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	29	9	8	7	5	3	0	2	0	1	8	6	1	4	1	2	0	0
0	0	0	0	0	2	14	12	7	1	0	2	0	2	0	3	4	6	4	3	9	2	0	1
0	0	0	0	0	0	2	0	0	0	0	0	0	0	0	0	3	3	2	0	1	0	0	0
0	0	0	0	0	0	0	1	1	0	1	1	0	0	0	2	1	2	3	0	1	0	0	0
0	0	0	0	0	0	1	0	0	0	0	0	0	2	0	4	5	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	3	2	1	0	1	1	0	0	0	0	3	11	12	0	0	0
0	0	0	0	0	0	0	0	0	2	1	0	0	0	1	0	3	5	4	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	18	4	0	0	0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	2	0	0	1	9	14	14	2	1	4	15	4	0	0
0	0	0	0	0	0	0	0	2	4	2	1	3	5	2	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	0	1	1	0	0	0	1	1	3	3	1	3	2	3	0	0
0	0	0	0	0	0	0	1	2	1	3	2	0	1	2	0	0	0	0	1	2	0	0	0
0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	8	4	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	4	6	2	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	0	1	1	6	2	1	0	0	0	0	1	1	0	0	0	0
0	0	0	0	0	0	2	2	3	12	18	12	12	2	4	3	5	6	43	37	16	1	2	0

1993 The first row defines the time of day, starting at midnight till the last timeslot from 23.00 till
 1994 00.00. In this particular case it shows that nobody is running from midnight to four o'clock
 1995 in the morning, understandably so.

1996 7.2 Model of Individuals

1997 7.2.1 Multi-modal Normal-Uniform Distribution Model

1998 We first postulate the likelihood function for the moments at which people exercise through
 1999 the day as defined in Eq. 7.1.

$$f(x|\theta) = w_0 \mathcal{U}(a, b) + \sum_{i=1}^2 w_i \mathcal{N}(\mu_i, \sigma_i) \quad (7.1)$$

2000 The likelihood (Eq. 7.1) is built up out of three probability density functions: one Uniform dis-
 2001 tribution $\mathcal{U}(a, b)$ with a and b as parameters and two Normal distributions $\mathcal{N}(\mu_i, \sigma_i)$ with
 2002 mean μ_i and σ_i . The distributions are weighted by the factors w_0, w_1, w_2 . The collection of
 2003 parameters for the likelihood function is referred to by $\theta = \{a, b, \mu_1, \sigma_1, \mu_2, \sigma_2, w_0, w_1, w_2\}$.

2004 This probability density function $f(x|\theta)$ will have the form as in Fig. 7.1. The uniform
 2005 distribution generates values here between 00:00 and 24:00. There are on top of that the
 2006 two Normal distributions that form peaks at certain moments during the day.



Figure 7.1: The likelihood function for the moments at which people decide to exercise during the day. On the horizontal axis time, on the vertical axis the frequency of exercising.

There is something noticable in Fig. 7.1, namely that the 24 hours of a day cause the Normal distribution to be cut off. Especially in Fig. 7.1 (a) there should be some considerable likelihood of running in the wee hours of the morning between 00:00 and 01:00.

7.2.2 Multi-modal Von-Mises-Uniform Distribution Model

There are several options to define a distribution over a limited range T . A so-called wrapped distribution is a distribution defined over the unity circle. By just multiplying it with $T/(2\pi)$ it can be used to define a probability density function over a day ($T = 24$).

$$f(x|\theta) = w_0 \mathcal{U}(a, b) + \sum_{i=1}^2 w_i \mathcal{VM}(\mu_i, \kappa_i) \quad (7.2)$$

The likelihood (Eq. 7.2) is again built up out of three probability density functions: one Uniform distribution $\mathcal{U}(a, b)$ with a and b as parameters and two Von Mises distributions $\mathcal{VM}(\mu_i, \kappa_i)$ with mean μ_i and κ_i . The parameters μ_i will be scaled and shifted with $[a, b]$ so all variables within this range fall on the unity circle. The parameter κ_i plays the same role as σ_i for the Normal distribution. The distributions are weighted by the factors w_0, w_1, w_2 . The collection of parameters for the likelihood function is referred to by $\theta = \{a, b, \mu_1, \kappa_1, \mu_2, \kappa_2, w_0, w_1, w_2\}$.

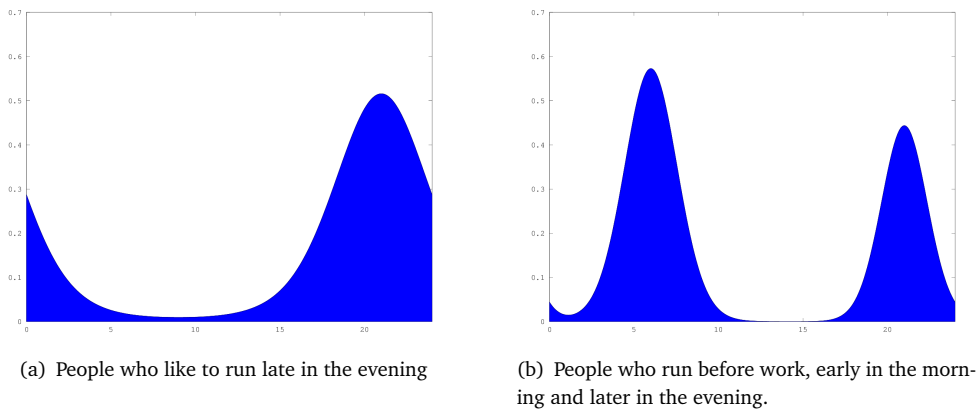


Figure 7.2: The improved likelihood function for the moments at which people decide to exercise during the day using the Von Mises distribution. On the horizontal axis time, on the vertical axis the frequency of exercising.

The likelihood with Von Mises distributions rather than Normal distributions is visualized in Fig. 7.2. The Von Mises distributions capture behavior in deviations from a standard exercise time. It does not take into account structured deviations, for example running late from work every Thursday, or weekends, or a person either running in the morning or in the evening, but never both at the same day.

7.2.3 Hyperparameters

The parameters θ for each user j are either fixed or generated from prior distributions. The hyperparameters a and b for the Uniform distribution are set to 0 and 24. People can run potentially any time of the day. The hyperparameters μ_i, κ_i for the Von Mises distributions are generated from a Uniform-Exponential distribution (Eq. 7.3). The Uniform distribution reflects the fact that if people run on regular times, this time can be any time of the day. The Exponential distribution defines a prior on how much people deviate from such a regular time to exercise.

$$\begin{aligned} f(\mu_i|a, b) &= \mathcal{U}(a, b) \\ f(\kappa_i|\lambda) &= \mathcal{E}(\lambda) \end{aligned} \tag{7.3}$$

The Uniform distribution generates μ_i between a and b . The parameter κ_i is generated from an Exponential distribution with hyperparameter λ . If λ is set to be small (< 0.5) we have a high likelihood that κ can be large and we have pronounced peaks. In contrary, if λ is set to be large, the Von Mises distribution likely approaches the Uniform distribution due to a higher chance of sampling a small value for κ .

The weights we sample from a normalized product of a zero-deflated Bernoulli distribution and a Dirichlet distribution (Eq. 7.4).

$$f(w_i|p, \alpha_i) = \mathcal{B}(p)\mathcal{D}(\alpha_i)/Z \quad (7.4)$$

The Bernoulli distribution samples zeros and ones with probability of $p = 0.5$, leading to 3-vectors like 001, 101, etc. The distribution is corrected in such way that the chance to sample 000 is zero. The Bernoulli distribution only gives weights $w_i = 0$ or $w_i = 1$, hence it is multiplied with a Dirichlet distribution. The Dirichlet samples a 3-vector with weights between zero and one where the weights $\sum_i w_i = 1$. A symmetric Dirichlet distribution with $\alpha = 1$ is similar to the Uniform distribution over the simplex. It is set slightly more towards favoring particular distributions with $\alpha_i = 1/3$ (we assume that people sample one or two distributions, and rarely from all three). The product with the zero-deflated Bernoulli distribution is made up to sum up to one again by normalizing the result.

We can combine Eq. 7.3 and Eq. 7.4 in Eq. 7.5:

$$\theta \sim \mathcal{U}(a, b)\mathcal{E}(\lambda)\mathcal{B}(p)\mathcal{D}(\alpha_i)/Z \quad (7.5)$$

To sample the parameters θ this is the base distribution we will encounter in the next section 7.3.

7.3 Model of Groups

Each person's exercise schedule is represented by a Von-Mises-Uniform distribution. People that are similar do have exercise schedules that can be represented by the same Von-Mises-Uniform distribution. To group similar schedules we define a nonparametric discrete distribution over a potentially infinite number of groups with each person assigned to a group.

A Dirichlet Process (Eq. 7.6) is a distribution over distributions that can be used as a prior for such a nonparametric discrete distribution.

$$DP(\alpha, H) \quad (7.6)$$

The Dirichlet Process has (1) a hyperparameter α , which defines the likelihood that there are many clusters versus few clusters (although it doesn't say anything about its actual count), and (2) a base distribution H , the distribution that generates θ (Eq.7.5).

7.4 Inference

The implementation of the model makes use of Gibbs sampling with auxiliary variables Jain and Neal (2007).

Details on this algorithm can be found in Jain and Neal (2007) and previous work of the authors van Rossum et al. (2016b,c).

Algorithm 15 Gibbs sampling over auxiliary variables

```

1: procedure GIBBS ALGORITHM WITH AUXILIARY VARIABLES( $w, \lambda_0, \alpha$ )      ▷ Accepts schedule  $w$ ,
   hyperparameters  $\lambda_0, \alpha$ , number of auxiliary variables  $m$ , and returns  $k$  groups
2:   for all  $t = 1 : T$  do
3:     for all  $i = 1 : N$  do
4:       for all  $j = 1 : m$  do
5:          $\theta_j \sim H(\lambda_0)$       ▷ Sample  $\theta_j$  from base distribution  $H$  in Eq. 7.5
6:       end for
7:       for all  $j = 1 : K + m, j \neq i$  do
8:          $L_j = \text{likelihood}(w_i, \theta_j)$       ▷ Update likelihood for all theta (except  $\theta_i$ ) given  $w_i$ 
9:       end for
10:       $P_{-i=1:K} = b \sum_{-i} L_{-i}$       ▷ Calculate probability of existing cluster
11:       $P_{-i=K:K+m} = b\alpha / mL_m L_{-i}$       ▷ Calculate probability of new cluster
12:       $\theta_i = \theta_j$  according to above  $P_{-i}$       ▷ Sample  $\theta_i$  accord. to above prob
13:      Remove unused clusters
14:    end for
15:    for all  $j = 1 : K$  do
16:       $\theta_j \sim p(\theta_j | y)$       ▷ Update  $\theta_j$ 
17:    end for
18:  end for
19:  return summary on  $\theta_k$  for  $k$  groups of runners
20: end procedure

```

2067 7.5 Results

2068 The algorithm is run first on an artificial dataset of which we know the ground truth (Sect. 7.5.1)
 2069 and next on the real-world dataset from Twitter (Sect. 7.5.2).

2070 7.5.1 Artificial Dataset

2071 The algorithm has been used on a self-generated dataset (generated from the probability
 2072 density function as in Fig. 7.2). In this case we have the ground truth that establishes which
 2073 exercise schedule comes from which probability density function. Using this ground truth
 2074 we can calculate how often our algorithm makes a mistake, grouping a person with people
 2075 that belong to another group. The results with this dataset are perfect (Rand Index equal to
 2076 one: perfect clustering). The results are of the form of Table 7.2 (except by a permutation
 2077 of indices), hence are not shown (the indices have no intrinsic meaning).

Table 7.2: Top row: A sequence of cluster indices indicates the ground truth. Each cluster index represents a multi-modal Von-Mises-Uniform distribution with different parameters θ . Bottom row: A sequence of cluster indices that are the result of the described algorithm. Each cluster index represents again a multi-modal Von-Mises distribution. Errors would be represented by an inconsistent mapping from the top row to the bottom row.

1	4	4	5	2	2	2	1	3	2	1	5	1	1	4	4	3	1	3	1	2	...
4	5	5	1	2	2	2	4	3	2	4	1	4	4	5	5	3	4	3	4	2	...

2078 7.5.2 Real-world Dataset

2079 The results when we actually use the collected Twitter dataset can be best visualized (Fig. 7.3).
 2080 This dataset consists of around 4000 moments at which people decide to run. The dataset
 2081 is subsequently filtered on regulars, people that at least have run a few times.

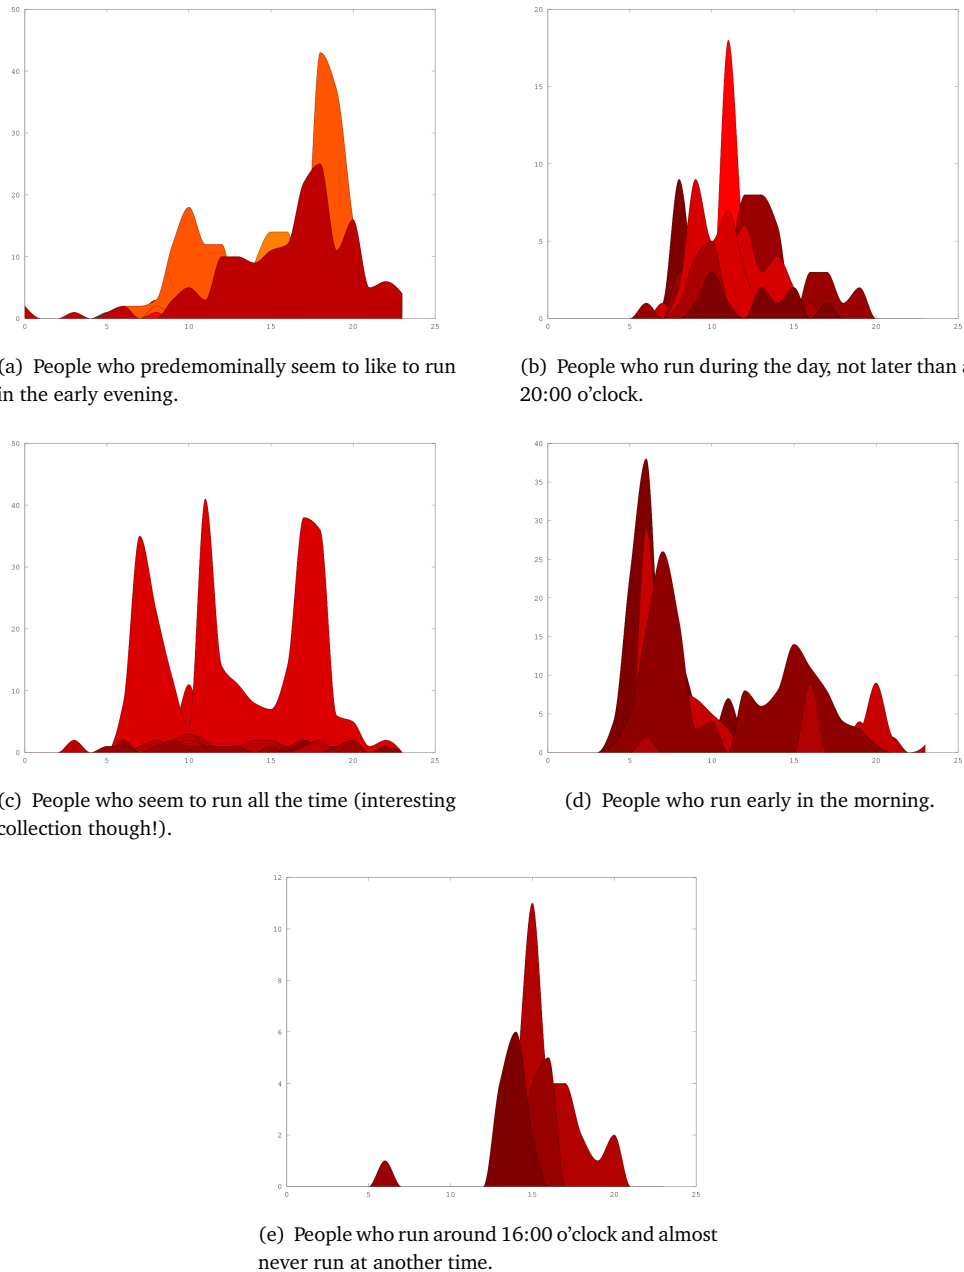


Figure 7.3: After running the algorithm, the above figures show the different categories of runners that have been found.

2082 The results in Fig. 7.3 show that there is a categorization of people indeed. After the al-
 2083 gorithm has done its work it is possible to assign a certain label in a post-hoc manner, e.g.
 2084 “People who run around 16:00 o'clock and almost never run at another time.”.

7.6 Discussion

The postulated model allows to reason about groups of people performing exercising, without predefining what these groups constitute apart from very general characteristics such as that there might be preferred times of day to exercise. A Dirichlet process is used as a nonparametric Bayesian model in which both the number of groups and the assignment of people exercising are learnt from the data.

There are several directions in which this research can be extended. First, more data would be very helpful. Only a limited number of people are posting consistently their training data online. Tapping into the data of current fitness promoting companies would allow the model to wash out the prior a bit more and adjust to the data. Second, we also collected weather data over this time period and also expect the day of the week to have significant influence. When there will be more data available these are logically dimensions to include in the dataset. Third, it would be interesting to study if people relate to the group of people they have been categorized with. Can this help or support their exercise regime?

2099

CHAPTER
8

2100

DISCUSSION AND CONCLUSIONS

REFERENCES

- 2103 M Abdel-Hameed. Optimal replacement policies for devices subject to a gamma wear process. *The*
 2104 *theory and applications of reliability*, pages 397–412, 2012.
- 2105 David J Aldous. Representations for partially exchangeable arrays of random variables. *Journal of*
 2106 *Multivariate Analysis*, 11(4):581–598, 1981.
- 2107 David J Aldous. *Exchangeability and related topics*. Springer, 1985.
- 2108 John Aldrich and Others. R.A. Fisher and the making of Maximum Likelihood 1912-1922. *Statistical*
 2109 *Science*, 12(3):162–176, 1997.
- 2110 Charles E Antoniak. Mixtures of Dirichlet processes with applications to Bayesian nonparametric
 2111 problems. *The annals of statistics*, pages 1152–1174, 1974.
- 2112 Stefan Banach and Alfred Tarski. Sur la décomposition des ensembles de points en parties respec-
 2113 tivement congruentes. *Fund. math*, 6(1):924, 1924.
- 2114 Federico Bassetti, Roberto Casarin, and Fabrizio Leisen. Beta-product dependent Pitman–Yor pro-
 2115 cesses for Bayesian inference. *Journal of Econometrics*, 180(1):49–72, 2014.
- 2116 Leonard E Baum and Ted Petrie. Statistical inference for probabilistic functions of finite state Markov
 2117 chains. *The annals of mathematical statistics*, 37(6):1554–1563, 1966.
- 2118 Herbert Bay, Tinne Tuytelaars, and Luc Van Gool. Surf: Speeded up robust features. In *European*
 2119 *conference on computer vision*, pages 404–417. Springer, 2006.
- 2120 Christopher M Bishop. *Pattern Recognition and Machine Learning*. Springer-Verlag New York, 2016.
- 2121 David Blackwell and James B MacQueen. Ferguson distributions via Pólya urn schemes. *The annals*
 2122 *of statistics*, pages 353–355, 1973.
- 2123 Phil Blunsom and Trevor Cohn. A hierarchical Pitman-Yor process HMM for unsupervised part of
 2124 speech induction. In *Proceedings of the 49th Annual Meeting of the Association for Computational*
 2125 *Linguistics: Human Language Technologies-Volume 1*, pages 865–874. Association for Computa-
 2126 tional Linguistics, 2011.
- 2127 Piotr Bojanowski, Armand Joulin, David Lopez-Paz, and Arthur Szlam. Optimizing the latent space
 2128 of generative networks. *arXiv preprint arXiv:1707.05776*, 2017.
- 2129 Robert C Bolles and Martin A Fischler. A RANSAC-Based Approach to Model Fitting and Its Application
 2130 to Finding Cylinders in Range Data. In *IJCAI*, volume 1981, pages 637–643, 1981.
- 2131 Anna Bosch, Andrew Zisserman, and Xavier Muñoz. Scene classification using a hybrid generative/dis-
 2132 criminative approach. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 30(4):712–
 2133 727, 2008.

- Guillaume Bouchard and Bill Triggs. The tradeoff between generative and discriminative classifiers. In *16th IASC International Symposium on Computational Statistics (COMPSTAT'04)*, pages 721–728, 2004.
- George E P Box and George C Tiao. *Bayesian inference in statistical analysis*, volume 40. John Wiley & Sons, 2011.
- Michael Bryant and Erik B Sudderth. Truly nonparametric online variational inference for hierarchical dirichlet processes. In *Advances in Neural Information Processing Systems*, pages 2699–2707, 2012.
- H Bühlmann. *Austauschbare stochastische Variablen und ihre Grenzwertsätze*. PhD thesis, ETH Zürich, 1960.
- Wray L Buntine. Operations for learning with graphical models. *JAIR*, 2:159–225, 1994.
- John Canny. A computational approach to edge detection. *IEEE Transactions on pattern analysis and machine intelligence*, (6):679–698, 1986.
- Jason Chang and John W Fisher III. Parallel sampling of dp mixture models using sub-cluster splits. In *Advances in Neural Information Processing Systems*, pages 620–628, 2013.
- Haifeng Chen, Peter Meer, and David E Tyler. Robust regression for data with multiple structures. In *Computer Vision and Pattern Recognition, 2001. CVPR 2001. Proceedings of the 2001 IEEE Computer Society Conference on*, volume 1, pages I—1069. IEEE, 2001.
- David B Dahl. An Improved Merge-Split Sampler for Conjugate Dirichlet Process Mixture Models. Technical report, University of Wisconsin–Madison, November 2003.
- David B Dahl. Sequentially-Allocated Merge-Split Sampler for Conjugate and Nonconjugate Dirichlet Process Mixture Models. Technical report, Texas A&M University, November 2005.
- Navneet Dalal and Bill Triggs. Histograms of oriented gradients for human detection. In *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*, volume 1, pages 886–893. IEEE, 2005.
- Hal Daume. Fast search for Dirichlet process mixture models. In *International Conference on Artificial Intelligence and Statistics*, pages 83–90, 2007.
- B de Finetti. Funzione caratteristica di un fenomeno aleatorio. *Atti Reale Accademia Nazionale dei Lincei*, VI:86–133, 1930.
- Bruno De Finetti. La prévision: ses lois logiques, ses sources subjectives. In *Annales de l'institut Henri Poincaré*, volume 7, pages 1–68, 1937.
- David GT Denison. *Bayesian methods for nonlinear classification and regression*, volume 386. John Wiley & Sons, 2002.
- Persi Diaconis and David Freedman. de Finetti’s theorem for Markov chains. *The Annals of Probability*, pages 115–130, 1980.
- Alexey Dosovitskiy, Jost Tobias Springenberg, Maxim Tatarchenko, and Thomas Brox. Learning to generate chairs, tables and cars with convolutional networks. *IEEE transactions on pattern analysis and machine intelligence*, 39(4):692–705, 2017.
- François Dufresne, Hans U Gerber, and Elias S W Shiu. Risk theory with the gamma process. *Astin Bulletin*, 21(02):177–192, 1991.
- David B Dunson, Ya Xue, and Lawrence Carin. The matrix stick-breaking process. *Journal of the American Statistical Association*, 2012.
- Richard Durstenfeld. Algorithm 235: Random Permutation. *Communications of the ACM*, 7(7):420, July 1964. ISSN 0001-0782. doi: 10.1145/364520.364540. URL <http://doi.acm.org/10.1145/364520.364540>.
- Michael D Escobar and Mike West. Bayesian density estimation and inference using mixtures. *Journal of the american statistical association*, 90(430):577–588, 1995.
- Stewart N Ethier. The distribution of the frequencies of age-ordered alleles in a diffusion model. *Advances in Applied Probability*, pages 519–532, 1990.

- Warren John Ewens. Population genetics theory-the past and the future. In *Mathematical and statistical developments of evolutionary theory*, pages 177–227. Springer, 1990.
- Stefano Favaro, Yee Whye Teh, and Others. MCMC for normalized random measure mixture models. *Statistical Science*, 28(3):335–359, 2013.
- William Feller. An introduction to probability theory and its applications. Vol. I. 1950.
- Thomas S Ferguson. A Bayesian analysis of some nonparametric problems. *The annals of statistics*, pages 209–230, 1973.
- Thomas S Ferguson. Prior distributions on spaces of probability measures. *The annals of statistics*, pages 615–629, 1974.
- E.W. Forgy. Cluster analysis of multivariate data: Efficiency vs. interpretability of classifications. *Biometrics*, 21:768–769, 1965.
- Wolfgang Förstner and Eberhard Gülch. A fast operator for detection and precise location of distinct points, corners and centres of circular features. In *Proc. ISPRS intercommission conference on fast processing of photogrammetric data*, pages 281–305, 1987.
- David Freedman and Persi Diaconis. On inconsistent Bayes estimates in the discrete case. *The Annals of Statistics*, pages 1109–1118, 1983.
- David Heaven Fremlin. *Measure theory*, volume 4. Torres Fremlin, 2000.
- Orazio Gallo, Roberto Manduchi, and Abbas Rafii. CC-RANSAC: Fitting planes in the presence of multiple surfaces in range data. *Pattern Recognition Letters*, 32(3):403–410, 2011.
- Qing-Bin Gao and Shi-Liang Sun. Human activity recognition with beta process hidden Markov models. In *Machine Learning and Cybernetics (ICMLC), 2013 International Conference on*, volume 2, pages 549–554. IEEE, 2013.
- Stuart Geman and Donald Geman. Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, (6):721–741, 1984.
- Zoubin Ghahramani and Thomas L Griffiths. Infinite latent feature models and the Indian buffet process. In *Advances in neural information processing systems*, pages 475–482, 2005.
- Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in neural information processing systems*, pages 2672–2680, 2014.
- Anirudh Goyal, Nan Rosemary Ke, Surya Ganguli, and Yoshua Bengio. Variational walkback: Learning a transition operator as a stochastic recurrent net. *arXiv preprint arXiv:1711.02282*, 2017.
- Max Halperin and G L Burrows. The Effect of Sequential Batching for Acceptance–Rejection Sampling Upon Sample Assurance of Total Product Quality. *Technometrics*, 2(1):19–26, 1960.
- Chris Harris and Mike Stephens. A combined corner and edge detector. In *Alvey vision conference*, volume 15, page 50. Citeseer, 1988.
- Li He, Hairong Qi, and Russell Zaretzki. Beta process joint dictionary learning for coupled feature spaces with application to single image super-resolution. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 345–352, 2013.
- Nils Lid Hjort. Nonparametric Bayes estimators based on beta processes in models for life history data. *The Annals of Statistics*, pages 1259–1294, 1990.
- Fred M Hoppe. Size-biased filtering of Poisson-Dirichlet samples with an application to partition structures in genetics. *Journal of Applied Probability*, pages 1008–1012, 1986.
- Paul V.C. Hough. Method and Means for Recognizing Complex Patterns, Dec 1962. URL <https://www.google.com/patents/US3069654>. Patent US 3069654 A.
- Xun Huang, Yixuan Li, Omid Poursaeed, John Hopcroft, and Serge Belongie. Stacked generative adversarial networks. *arXiv preprint arXiv:1612.04357*, 2016.
- Michael C Hughes and Erik Sudderth. Memoized online variational inference for dirichlet process mixture models. In *Advances in Neural Information Processing Systems*, pages 1133–1141, 2013.

- 2231 Michael C Hughes, Emily Fox, and Erik B Sudderth. Effective split-merge monte carlo methods for
2232 nonparametric models of sequential data. In *Advances in neural information processing systems*,
2233 pages 1295–1303, 2012.
- 2234 Tommi S Jaakkola, David Haussler, and Others. Exploiting generative models in discriminative clas-
2235 sifiers. *Advances in neural information processing systems*, pages 487–493, 1999.
- 2236 Sonia Jain and Radford M. Neal. A Split-Merge Markov Chain Monte Carlo Procedure for the Dirichlet
2237 Process Mixture Model. *Journal of Computational and Graphical Statistics*, 13(1):158–182, 2004.
2238 ISSN 10618600. URL <http://www.jstor.org/stable/1391150>.
- 2239 Sonia Jain and Radford M Neal. Splitting and Merging Components of a Nonconjugate Dirichlet
2240 Process Mixture Model. *Bayesian Analysis*, 2(3):445–472, 2007.
- 2241 Varun Jampani, Sebastian Nowozin, Matthew Loper, and Peter V Gehler. The informed sampler: A
2242 discriminative approach to bayesian inference in generative computer vision models. *Computer*
2243 *Vision and Image Understanding*, 136:32–44, 2015.
- 2244 Edwin T Jaynes. *Probability theory: The logic of science*. Cambridge university press, 2003.
- 2245 Dominik Joho, Gian Diego Tipaldi, Nikolas Engelhard, Cyrill Stachniss, Wolfram Burgard, Mar-
2246 tin Senk, Felix Faber, Maren Bennewitz, Clemens Eppner, Attila Görög, and Others. Unsuper-
2247 vised Scene Analysis and Reconstruction Using Nonparametric Bayesian Models. *Robotics and*
2248 *Autonomous Systems (RAS)*, 59(5):319–328, 2011.
- 2249 A Jordan. On discriminative vs. generative classifiers: A comparison of logistic regression and naive
2250 Bayes. *Advances in neural information processing systems*, 14:841, 2002.
- 2251 Michael I Jordan. Hierarchical models, nested models and completely random measures. *Frontiers of*
2252 *Statistical Decision Making and Bayesian Analysis: in Honor of James O. Berger*. New York: Springer,
2253 2010.
- 2254 Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint*
2255 *arXiv:1312.6114*, 2013.
- 2256 J F C Kingman. *Poisson processes*, volume 3 of *Oxford Studies in Probability*. The Clarendon Press
2257 Oxford University Press, New York, 1993. ISBN 0-19-853693-3.
- 2258 J F C Kingman. Some further analytical results in the theory of regenerative events. *Journal of*
2259 *Mathematical Analysis and Applications*, 11:422–433, 1965.
- 2260 J F C Kingman. Random partitions in population genetics. In *Proceedings of the Royal Society of*
2261 *London A: Mathematical, Physical and Engineering Sciences*, volume 361, pages 1–20. The Royal
2262 Society, 1978.
- 2263 John Kingman. Completely random measures. *Pacific Journal of Mathematics*, 21(1):59–78, 1967.
- 2264 Roman Klovov and Victor Lempitsky. Escape from cells: Deep kd-networks for the recognition of
2265 3d point cloud models. In *2017 IEEE International Conference on Computer Vision (ICCV)*, pages
2266 863–872. IEEE, 2017.
- 2267 David Knowles, Zoubin Ghahramani, and Konstantina Palla. A reversible infinite HMM using nor-
2268 malised random measures. In *Proceedings of the 31st International Conference on Machine Learning*
2269 *(ICML-14)*, pages 1998–2006, 2014.
- 2270 Daphne Koller and Nir Friedman. *Probabilistic graphical models: principles and techniques*. MIT press,
2271 2009.
- 2272 A Kolmogorov. *Grundbegriffe der wahrscheinlichkeitsrechnung*, volume 2 of *Ergebnisse der Mathematik*
2273 *und ihrer Grenzgebiete*. Springer-Verlag, 1933.
- 2274 Uwe Küchler and Stefan Tappe. Bilateral Gamma distributions and processes in financial mathemat-
2275 ics. *Stochastic Processes and their Applications*, 118(2):261–283, 2008.
- 2276 Kenichi Kurihara, Max Welling, and Yee Whye Teh. Collapsed Variational Dirichlet Process Mixture
2277 Models. In *IJCAI*, volume 7, pages 2796–2801, 2007.
- 2278 Pierre-Simon Laplace. *Théorie analytique des probabilités*. V. Courcier, 1820.

- Henri Lebesgue. Intégrale, longueur, aire. *Annali di Matematica Pura ed Applicata (1898-1922)*, 7 (1):231–359, 1902.
- Zhidong Li, Bang Zhang, Yang Wang, Fang Chen, Ronnie Taib, Vicky Whiffin, and Yi Wang. Water pipe condition assessment: a hierarchical beta process approach for sparse incident data. *Machine learning*, 95(1):11–26, 2014.
- Antonio Lijoi and Igor Prünster. Models beyond the Dirichlet process. *Bayesian nonparametrics*, 28: 80, 2010.
- S. Lloyd. Least squares quantization in pcm. *IEEE Trans. Inform. Theory*, 28:129–137, 1982. Originally as an unpublished Bell laboratories Technical Note (1957).
- David G Lowe. Object recognition from local scale-invariant features. In *Computer vision, 1999. The proceedings of the seventh IEEE international conference on*, volume 2, pages 1150–1157. Ieee, 1999.
- Steven N MacEachern. Comment on "Splitting and Merging Components of a Nonconjugate Dirichlet Process Mixture Model" by Jain and Neal. *Bayesian Analysis*, 2(3):483–494, 2007.
- Steven N MacEachern and Peter Müller. Estimating mixture of Dirichlet process models. *Journal of Computational and Graphical Statistics*, 7(2):223–238, 1998.
- Dilip B Madan and Eugene Seneta. The variance gamma (VG) model for share market returns. *Journal of business*, pages 511–524, 1990.
- Nicholas Metropolis, Arianna W. Rosenbluth, Marshall N. Rosenbluth, Augusta H. Teller, and Edward Teller. Equation of State Calculations by Fast Computing Machines. *The Journal of Chemical Physics*, 21(6):1087–1092, 1953. doi: 10.1063/1.1699114.
- Radford M Neal. Defining Priors for Distributions Using Dirichlet Diffusion Trees. Technical report, University of Toronto, Toronto, Ontario, Canada, 2001.
- Radford M Neal. Markov Chain Sampling Methods for Dirichlet Process Mixture Models. *Journal of Computational and Graphical Statistics*, 9(2):249–265, 2000.
- Peter Orbanz and Daniel M Roy. Bayesian models of graphs, arrays and other exchangeable random structures. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 37(2):437–461, 2015.
- Jim Pitman and Marc Yor. The two-parameter Poisson-Dirichlet distribution derived from a stable subordinator. *The Annals of Probability*, pages 855–900, 1997.
- Charles Ruizhongtai Qi, Li Yi, Hao Su, and Leonidas J Guibas. Pointnet++: Deep hierarchical feature learning on point sets in a metric space. In *Advances in Neural Information Processing Systems*, pages 5105–5114, 2017.
- Rajat Raina, Yirong Shen, Andrew McCallum, and Andrew Y Ng. Classification with hybrid generative/discriminative models. In *Advances in neural information processing systems*, page None, 2003.
- William M Rand. Objective criteria for the evaluation of clustering methods. *Journal of the American Statistical association*, 66(336):846–850, 1971.
- Vinayak Rao and Yee W Teh. Spatial normalized gamma processes. In *Advances in neural information processing systems*, pages 1554–1562, 2009.
- Carl Edward Rasmussen and Christopher K I Williams. *Gaussian processes for machine learning*, volume 2. 2006.
- Siamak Ravanbakhsh, Jeff Schneider, and Barnabas Poczos. Deep learning with sets and point clouds. *arXiv preprint arXiv:1611.04500*, 2016.
- Eugenio Regazzini, Antonio Lijoi, and Igor Prünster. Distributional results for means of normalized random measures with independent increments. *Annals of Statistics*, pages 560–585, 2003.
- Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and approximate inference in deep generative models. *arXiv preprint arXiv:1401.4082*, 2014.
- James C Ross, Peter J Castaldi, Michael H Cho, and Jennifer G Dy. Dual beta process priors for latent cluster discovery in chronic obstructive pulmonary disease. In *Proceedings of the 20th ACM SIGKDD*

- international conference on Knowledge discovery and data mining, pages 155–162. ACM, 2014.
- Daniel M Roy and Yee W Teh. The mondrian process. In *Advances in neural information processing systems*, pages 1377–1384, 2009.
- Anirban Roychowdhury and Brian Kulis. Gamma processes, stick-breaking, and variational inference. In *Artificial Intelligence and Statistics*, pages 800–808, 2015.
- Donald B Rubin and Others. Bayesianly justifiable and relevant frequency calculations for the applied statistician. *The Annals of Statistics*, 12(4):1151–1172, 1984.
- Stuart Russell, Peter Norvig, and Artificial Intelligence. Artificial Intelligence: A modern approach. *Artificial Intelligence*. Prentice-Hall, Egnlewood Cliffs, 25:27, 1995.
- Leonard J Savage. *The foundations of statistics*. Courier Corporation, 1972.
- Stanley Sawyer and Daniel Hartl. A sampling theory for local selection. *Journal of Genetics*, 64(1): 21–29, 1985.
- René L Schilling. *Measures, integrals and martingales*, volume 13. Cambridge University Press, 2005.
- Abhishek Sharma, Oliver Grau, and Mario Fritz. Vconv-dae: Deep volumetric shape learning without object labels. In *European Conference on Computer Vision*, pages 236–250. Springer, 2016.
- Jianbo Shi and Carlo Tomasi. Good features to track. In *Computer Vision and Pattern Recognition, 1994. Proceedings CVPR’94., 1994 IEEE Computer Society Conference on*, pages 593–600. IEEE, 1994.
- Nozer Singpurwalla. Gamma processes and their generalizations: an overview. In *Engineering probabilistic design and maintenance for flood protection*, pages 67–75. Springer, 1997.
- Scott A Sisson and Yanan Fan. Likelihood-free MCMC. *Handbook of Monte Carlo*, ed. Brooks, A., Gelman, A., Jones, GL, Meng XL, pages 313–335, 2011.
- I Sobel. *Camera Models and Perception*. PhD thesis, Stanford University, Stanford, CA, 1970.
- Jascha Sohl-Dickstein, Eric A Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. *arXiv preprint arXiv:1503.03585*, 2015.
- Jiaming Song, Shengjia Zhao, and Stefano Ermon. A-nice-mc: Adversarial training for mcmc. *arXiv preprint arXiv:1706.07561*, 2017.
- Guillaume Stawinski, Arnaud Doucet, and Patrick Duvaut. Reversible jump markov chain monte carlo for bayesian deconvolution of point sources. In *Bayesian Inference for Inverse Problems*, volume 3459, pages 179–191. International Society for Optics and Photonics, 1998.
- Mike Steel. The maximum likelihood point for a phylogenetic tree is not unique. *Systematic Biology*, 43(4):560–564, 1994.
- Erik B Sudderth and Michael I Jordan. Shared segmentation of natural scenes using dependent Pitman-Yor processes. In *Advances in Neural Information Processing Systems*, pages 1585–1592, 2009.
- Charles Sutton and Andrew McCallum. An introduction to conditional random fields. *Machine Learning*, 4(4):267–373, 2011.
- Terence Tao. *An introduction to measure theory*, volume 126. American Mathematical Soc., 2011.
- Yee Whye Teh, Michael I Jordan, Matthew J Beal, and David M Blei. Hierarchical Dirichlet processes. *Journal of the American statistical association*, 101(476), 2006.
- Romain Thibaux and Michael I Jordan. Hierarchical beta processes and the Indian buffet process. In *International Conference on Artificial Intelligence and Statistics*, pages 564–571, 2007.
- Michalis K Titsias. The infinite gamma-Poisson feature model. In *Advances in Neural Information Processing Systems*, pages 1513–1520, 2008.
- Dustin Tran, Matthew D Hoffman, Rif A Saurous, Eugene Brevdo, Kevin Murphy, and David M Blei. Deep probabilistic programming. *arXiv preprint arXiv:1701.03757*, 2017.
- Anne C van Rossum, Hai Xiang Lin, Johan Dubbeldam, and H Jaap van den Herik. Fundamentals of

- nonparametric bayesian line detection. In *International Conference on Pattern Recognition Applications and Methods*, pages 175–193. Springer, 2016a.
- Anne C. van Rossum, Hai Xiang Lin, Johan Dubbeldam, and H. Jaap van den Herik. Nonparametric Bayesian Line Detection - Towards Proper Priors for Robotic Computer Vision. In *Proceedings of the 5th International Conference on Pattern Recognition Applications and Methods*, pages 119–127, Feb 2016b. ISBN 978-989-758-173-1. doi: 10.5220/0005673301190127.
- Anne C. van Rossum, Hai Xiang Lin, Johan Dubbeldam, and H. Jaap van den Herik. Nonparametric Segment Detection. In *Proceedings of the 8th European Starting AI Researcher Symposium (STAIRS)*, pages 203–208, Aug 2016c. ISBN 978-989-758-173-1. doi: 10.5220/0005673301190127.
- Anne C. van Rossum, Hai Xiang Lin, Johan Dubbeldam, and H. Jaap van den Herik. Triadic split-merge sampler. Nov 2017.
- Niklas Vanhainen and Giampiero Salvi. Word Discovery with Beta Process Factor Analysis. In *INTER-SPEECH*, pages 799–802, 2012.
- Wei Wang and Stuart Russell. A Smart-dumb/Dumb-smart Algorithm for Efficient Split-merge MCMC. In *Proceedings of the Thirty-First Conference on Uncertainty in Artificial Intelligence, UAI15*, pages 902–911, Arlington, Virginia, United States, 2015. AUAI Press. ISBN 978-0-9966431-0-8. URL <http://dl.acm.org/citation.cfm?id=3020847.3020940>.
- Yingjian Wang and Lawrence Carin. Levy measure decompositions for the beta and gamma processes. *arXiv preprint arXiv:1206.4615*, 2012.
- Larry Wasserman. Asymptotic properties of nonparametric Bayesian procedures. In *Practical non-parametric and semiparametric Bayesian statistics*, pages 293–304. Springer, 1998.
- Robert L Wolpert and Katja Ickstadt. Poisson/gamma random field models for spatial statistics. *Biometrika*, 85(2):251–267, 1998.
- Stephen J Wright. Coordinate descent algorithms. *Mathematical Programming*, 151(1):3–34, 2015.
- Zhirong Wu, Shuran Song, Aditya Khosla, Fisher Yu, Linguang Zhang, Xiaoou Tang, and Jianxiong Xiao. 3d shapenets: A deep representation for volumetric shapes. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1912–1920, 2015.
- Jing-Hao Xue and D Michael Titterton. Comment on “On discriminative vs. generative classifiers: A comparison of logistic regression and naive Bayes”. *Neural processing letters*, 28(3): 169–187, 2008.
- Arnold Zellner. Optimal information processing and Bayes’s theorem. *The American Statistician*, 42(4):278–280, 1988.
- Michael M Zhang and Fernando Perez-Cruz. Accelerated inference for latent variable models. *arXiv preprint arXiv:1705.07178*, 2017.
- Wei Zhang and Jana Kösecká. Nonparametric estimation of multiple structures with outliers. In *Dynamical Vision*, pages 60–74. Springer, 2007.
- Mingyuan Zhou, Hongxia Yang, Guillermo Sapiro, David B Dunson, and Lawrence Carin. Dependent hierarchical beta process for image interpolation and denoising. In *International conference on artificial intelligence and statistics*, pages 883–891, 2011.
- Mingyuan Zhou, Lauren A Hannah, David B Dunson, and Lawrence Carin. Beta-negative binomial process and poisson factor analysis. *Journal of Machine Learning Research*, 2012.



PROBABILISTIC CONCEPTS

TODO: What follows down here is old and has to be adapted to the measure-theoretic realizations.

Let us introduce the notation of expectation for random variable X :

$$E_p[X] = \sum_i^k p(X = x_i)x_i \quad (\text{A.1})$$

For the continuous case, if $f_X(x)$ is a properly defined probability density function, the expected value becomes:

$$E_f[X] = \int_{-\infty}^{\infty} x f_X(x) dx \quad (\text{A.2})$$

From the context it can be seen that the object X at the left is a different object than x at the right. The former is random variable (a one-dimensional function, an (in)finite vector), the latter is a value of a random variable (a scalar). The term dx is a measure, in this case it assigns volume to *subsets of random variables values*. A proper notation would incorporate this aspect, but not much will be gained by creating such complex notations.

Let us also introduce the conditional probability:

$$p(X|Y) = \frac{p(X, Y)}{p(Y)} \quad (\text{A.3})$$

Suppose X is a discrete random variable, then the object $p(X)$ is a vector of finite size k , $p(Y)$ is a vector of finite size l , and $p(X|Y)$ as well as $p(X, Y)$ are matrices of size $k \times l$. A conditional probability hence ‘divides’ a matrix by a vector. It is a proper measure if $p(Y) \neq 0$ (for all values of - or events in - Y).

In for example importance sampling (Sect. ??), the expectation is taken over a function of a random variable. A function, if measurable, can be taken the expectation over using the so-called law of the unconscious statistician:

$$E_f[g(X)] = \int_{-\infty}^{\infty} g(x)f_X(x)dx \quad (\text{A.4})$$

Here $g(X)$ is a general measurable function, and not restricted to a probability density function.

The notation above is an indefinite integral. We can approach this integral by Monte Carlo integration (Sect. ??).

$$E_f[g(X)] = \frac{1}{k} \sum_{i=1}^k g(x_i) \quad \text{with} \quad x_i \sim f_X(x) \quad (\text{A.5})$$

Rather than summing over $f_X(x_i)$, we now sample $x_i \sim f_X(x)$.

A.1 Common Inequalities

A.1.1 Markov's Inequality

Markov's inequality comes up with an upper bound for the probability that an non-negative random variable X exceeds some constant positive threshold a .

$$p(X \geq a) \leq \frac{E[X]}{a} \quad (\text{A.6})$$

The proof in classical probability theory uses an indicator variable:

$$aI(X \geq a) \leq X \quad (\text{A.7})$$

Here $I(X \geq a) = 1$ if the event $X \geq a$ occurs, setting the left-hand side to a (which is of course smaller than X). And $I(X \geq a) = 0$ on the event $X < a$, which is naturally smaller than the non-negative X .

Expectations obey the inequality: if $(X \leq Y)$, then $E[X] \leq E[Y]$, hence:

$$E[aI(X \geq a)] \leq E[X] \quad (\text{A.8})$$

And because expectations add up linearly:

$$E[aI(X \geq a)] = aE[I(X \geq a)] = a(1 \cdot p(X \geq a) + 0 \cdot p(X < a)) = a \cdot p(X \geq a) \quad (\text{A.9})$$

2450 So, we have Markov's inequality combining Eq. A.8 and A.9:

$$a \cdot p(X \geq a) \leq E[X] \quad (\text{A.10})$$

2451 A.1.2 Chebyshev's Inequality

2452 Now, Chebyshev's inequality defines in a similar way (Chebyshev was a teacher of Markov¹)
 2453 an upper bound on the deviation from the mean for a random variable. Recall the definition
 2454 of the variance of X and assume it is finite:

$$\text{Var}(X) = E[(X - E[X])^2] = \sigma^2 \quad (\text{A.11})$$

2455 Consider now the random variable $(X - E[X])^2$ and constant $a = (\sigma k)^2$ and write down
 2456 Markov's inequality:

$$p((X - E[X])^2 \geq (\sigma k)^2) \leq \frac{E[(X - E[X])^2]}{(\sigma k)^2} \quad (\text{A.12})$$

2457 Taking the square root of the inequality at the left, and using the definition of σ^2 at the
 2458 right, leads to:

$$p(|X - E[X]| \geq \sigma k) \leq \frac{1}{k^2} \quad (\text{A.13})$$

2459 A.1.3 Weak Law of Large Numbers

2460 Chebyshev's inequality can be used to prove the weak law of large numbers. Given that we
 2461 have a series of random variables, all with the same finite expectation, $E[X_i] = \mu$, then this
 2462 law states that the sample average $\bar{X} = \frac{1}{n}(X_1 + \dots + X_n)$ converges in probability towards
 2463 the expected value:

$$\bar{X} \xrightarrow{P} \mu \quad \text{for} \quad n \rightarrow \infty \quad (\text{A.14})$$

2464 We can use the independence assumption between variables X_i to write down the variance
 2465 and expectation of \bar{X} :

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{n}(X_1 + \dots + X_n)\right) = \frac{1}{n^2} \text{Var}(X_1 + \dots + X_n) = \frac{\sigma^2}{n} E[\bar{X}] = \mu \quad (\text{A.15})$$

¹There were many mathematically gifted Markov's. This is Andrey Andreyevich Markov Sr., known from the Markov chains and Markov processes. Jr. is known from Markov's principle, Markov's rule and the Markov algorithm.

2466 And now we can apply Chebyshev's inequality on \bar{X} :

$$p(|\bar{X} - \mu| \geq \epsilon) \leq \frac{\sigma^2}{n\epsilon^2} \quad (\text{A.16})$$

2467 Convergence in probability towards X is the case if for all ϵ :

$$\lim_{n \rightarrow \infty} p(|\bar{X} - X| \geq \epsilon) = 0 \quad (\text{A.17})$$

2468 This is the case for $n \rightarrow \infty$ indeed.

2469 **A.1.4 Strong Law of Large Numbers**

2470 The strong law incorporates the weak law. Rather than convergence *in probability*, it states
2471 convergence *almost surely* towards the expected value.

$$\bar{X} \xrightarrow{a.s.} \mu \quad \text{for} \quad n \rightarrow \infty \quad (\text{A.18})$$

2472 The strong law states that with probability 1, for any $\epsilon > 0$, the inequality $|X - \mu| < \epsilon$
2473 holds for large enough n . The weak law states only that the average \bar{X} is likely near μ , but
2474 $|X - \mu| \geq \epsilon$ can still happen, even for large n .

2475 **A.1.5 Common Distributions**

2476 One of the probability distributions that is interesting to us is the beta-distribution. A normal
2477 distribution might be a reasonable prior for a continuous variable such as human heights in
2478 a population. If this variable however is itself a probability, a reasonable prior is the beta-
2479 distribution. The beta-distribution can be described as:

$$f(x, \alpha, \beta) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1} = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1} \quad (\text{A.19})$$

2480 Here B is the beta function and Γ is the gamma function, the continuous extension of the
2481 factorial function: $\Gamma(n) = (n-1)!$. Naturally, there are many of such extensions. The
2482 gamma function extends the factorial in a specific sense. It obeys the recurrence relation
2483 $f(x+1) = xf(x)$ with $f(1) = 1$. Its description is defined with an improper integral:

$$\Gamma(t) = \int_0^\infty x^{t-1} e^{-x} dx \quad (\text{A.20})$$

2484 The expected value of a random variable X with a beta-distribution:




Figure A.1: The beta-distribution with different parameters. The x-axis is the quantity modelled, for example a ratio between wins and losses in a soccer season. The y-axis is the corresponding unnormalized density function. Setting $\alpha = 1$ shows a monotonically decreasing density function. Having a variable α allows probability mass to shift to the end.

$$E_f[X] = \int_0^1 x f(x, \alpha, \beta) dx = \int_0^1 \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^\alpha (1-x)^{\beta-1} dx = \frac{\alpha}{\alpha + \beta} \quad (\text{A.21})$$

2485 Let us illustrate the effects of the parameters of the beta-distribution.

2486 The stick-breaking presentation shows that a Dirichlet process consists of beta processes with
 2487 $\alpha = 1$. The Pitman-Yor process has α left variable. Informally, the location where the stick
 2488 will be broken (iteratively) for the Dirichlet process is ‘quite close to the beginning’ with high
 2489 probability. In the case α becomes larger, the breaking can also occur likely at the end of the
 2490 stick. This means for breaking a stick, say 20 times, the distribution of stick lengths for the
 2491 Pitman-Yor process has a much larger support. The difference between the large and small
 2492 sticks is much larger.

DIRICHLET-MULTINOMIAL INTERPRETATIONS

There are multiple interpretations of the compound compound "Dirichlet-multinomial" distribution. We will derive different related compound distributions and show how each of these compound distributions can be used as a likelihood function that is conjugate to the Dirichlet distribution. Let us recall the Dirichlet distribution (Eq. B.1).

$$p(\theta|\alpha) = \frac{1}{B(\alpha)} \prod_i \theta_i^{\alpha_i-1}. \quad (\text{B.1})$$

The normalizing constant $1/B(\alpha)$ contains the Beta function:

$$B(\alpha) = \frac{\prod_i \Gamma(\alpha_i)}{\Gamma(\sum_i \alpha_i)}. \quad (\text{B.2})$$

Regarding notation, the Γ function can be written as a factorial $\Gamma(x+1) = x!$. Steps like $\Gamma(1+x) = x\Gamma(x)$ and $\frac{\Gamma(x+n)}{\Gamma(x)} = x^n$ will not be made explicit in the following text.

We would like to find a conjugate distribution for the Dirichlet that will result in a Dirichlet distribution after performing Bayesian inference. The shape of this conjugate distribution is of the form:

$$p(z|\theta) = C \prod_i \theta_i^{z_i}. \quad (\text{B.3})$$

Here C is a normalization constant that depends on the interpretation of the random variable z .

There are three ways to interpret the Dirichlet-multinomial, namely as (1) a compound Dirichlet-categorical distribution (Appendix B.1), (2) a compound Dirichlet-multinomial distribution (Appendix B.2), and as (3) compound Dirichlet with N categorical distributions (Appendix B.3). Each of these interpretations are different. However, if they are used in a Bayesian context, they will all result in a posterior that is a Dirichlet distribution.

B.1 Dirichlet-Categorical

Often, the Dirichlet-multinomial is actually not a compound Dirichlet and *multinomial* distribution, but a compound Dirichlet and *categorical* distribution. Informally, the categorical distribution is a distribution over categories. More formally, it is a generalization of a Bernoulli distribution that assigns probabilities to a random variable with two possible outcomes to a random variable with more than two possible outcomes. The categorical distribution has the form as in Eq. B.4.

$$p(z|\theta) = \prod_i \theta_i^{z_i}. \quad (\text{B.4})$$

This notation defines only *a single* categorical variable, not a set of variables. It is using a 1-of- K encoded representation (cf. Bishop, 2016). Assume a regular 6-faced dice. The face with one pip is represented by $z_0 = [1, 0, 0, 0, 0, 0]$. The face with two pips $z_1 = [0, 1, 0, 0, 0, 0]$. There is only one category non-zero, $\sum_j z_{i,j} = 1$. The probabilities of each category are represented by θ_i and are one sixth for a regular dice.

Let us take the product of the Dirichlet distribution in Eq. B.1 and the categorical distribution in Eq. B.4:

$$p(z|\theta)p(\theta|\alpha) = \frac{1}{B(\alpha)} \prod_i \theta_i^{z_i + \alpha_i - 1}. \quad (\text{B.5})$$

The compound distribution is created by integrating over θ :

$$p(z|\alpha) = \int p(z|\theta)p(\theta|\alpha)d\theta = \frac{1}{B(\alpha)} \int \prod_i \theta_i^{z_i + \alpha_i - 1} d\theta. \quad (\text{B.6})$$

The multivariate Beta function (Eq. B.2) stems from normalizing the Dirichlet distribution and can also be written as an integral over the unnormalized Dirichlet distribution:

$$B(\alpha) = \int \prod_i \theta_i^{\alpha_i - 1} d\theta. \quad (\text{B.7})$$

Now we can use the 1-of- K representation to write the integral as follows:

$$\int \prod_i \theta_i^{z_i + \alpha_i - 1} d\theta = B(\alpha + z). \quad (\text{B.8})$$

2531 The compound distribution of the Dirichlet-categorical henceforth can be written as:

$$p(z|\alpha) = \int p(z|\theta)p(\theta|\alpha)d\theta = \frac{B(\alpha + z)}{B(\alpha)}. \quad (\text{B.9})$$

2532 We can now perform Bayes' rule that shows the benefit of the closed-form description of the
2533 compound distribution in the conjugate context:

$$p(\theta|z, \alpha) = \frac{p(z|\theta, \alpha)p(\theta|\alpha)}{p(z|\alpha)} \quad (\text{B.10})$$

2534 In this case the likelihood $p(z|\theta, \alpha) = p(z|\theta)$ has observations z that only depend on θ (and
2535 θ depends on α).

$$p(\theta|z, \alpha) = \frac{p(z|\theta)p(\theta|\alpha)}{p(z|\alpha)} \quad (\text{B.11})$$

2536 Hence the posterior of the Dirichlet prior with the categorical likelihood is a Dirichlet poste-
2537 rior (use Eq. B.5 and Eq. B.9 in Eq. B.11).

$$p(\theta|z, \alpha) = \frac{1}{B(\alpha + z)} \prod_i \theta_i^{z_i + \alpha_i - 1}. \quad (\text{B.12})$$

2538 Note again that we run i over the entries in our categorical variable z represented as a vector.
2539 This is different from a multinomial distribution over a set of variables!

2540 B.2 Dirichlet-Multinomial Distribution

2541 In case of an actual multinomial distribution, the topic of consideration are *counts* of z . Let's
2542 write the counts as $n(z)$.

$$p(z|\theta) = \frac{(\sum_k n(z_k))!}{\prod_k (n(z_k))!} \prod_k \theta_k^{n(z_k)}. \quad (\text{B.13})$$

2543 We now run over k unique variables, not over a vectorized categorical variable. The fol-
2544 lowing example clarifies the multinomial. Suppose we have a bag with k colored balls. The
2545 multinomial is the random variable that assigns a probability to the vector $[n(z_0), \dots, n(z_{k-1})]$
2546 with $n(z_i)$ the number of balls of color i . The distribution also represents the task of throwing
2547 a k -faced dice n times and counting the number of times $n(z)$ category z is occurring.

2548 The product of the Dirichlet distribution with the multinomial can be performed along the
 2549 same lines as in Appendix B.1.

$$p(z|\theta)p(\theta|\alpha) = \frac{\sum_k n(z_k) \Gamma(\sum_k n(z_k))}{\prod_k n(z_k) \prod_k \Gamma(n(z_k))} \frac{1}{B(\alpha)} \prod_k \theta_k^{n(z_k)+\alpha_k-1}. \quad (\text{B.14})$$

2550 This we can simplify to:

$$p(z|\theta)p(\theta|\alpha) = \frac{\sum_k n(z_k)}{\prod_k n(z_k)} \frac{1}{B(n(z_k))} \frac{1}{B(\alpha)} \prod_k \theta_k^{n(z_k)+\alpha_k-1}. \quad (\text{B.15})$$

2551 We will get the following result for the compound distribution (using Eq. B.7 with the Beta
 2552 function as known result for the integral):

$$p(z|\alpha) = \int p(z|\theta)p(\theta|\alpha)d\theta = \frac{\sum_k n(z_k)}{\prod_k n(z_k)} \frac{1}{B(n(z_k))} \frac{1}{B(\alpha)} B(\alpha + n(z_k)). \quad (\text{B.16})$$

2553 The Dirichlet prior with the multinomial likelihood results in a Dirichlet posterior as well:

$$p(\theta|z, \alpha) = \frac{p(z|\theta)p(\theta|\alpha)}{p(z|\alpha)} = \frac{1}{B(\alpha + n(z_k))} \prod_k \theta_k^{n(z_k)+\alpha_k-1}. \quad (\text{B.17})$$

2554 **B.3 Dirichlet-N Categorical Distributions**

2555 The third interpretation of a Dirichlet-multinomial distribution is the one in which the term
 2556 "multinomial" refers to a distribution of a *sequence* of categorical variables. Recall that the
 2557 multinomial assigns probabilities to the *number* of extracted balls (in an experiment getting
 2558 n balls out of a bag with k ball types). A sequence of categorical variables has a form for its
 2559 probability distribution that has no normalization factor:

$$p(z|\theta) = \prod_k \theta_k^{z_k}. \quad (\text{B.18})$$

2560 Here k runs over the categories. Note that this form is exactly the same as that of Eq. B.4.
 2561 The compound distribution will henceforth have the same form as the Dirichlet-categorical
 2562 distribution:

$$p(z|\alpha) = \int p(z|\theta)p(\theta|\alpha)d\theta = \frac{B(\alpha + z)}{B(\alpha)}. \quad (\text{B.19})$$

2563 The Dirichlet prior with the n-categorical distribution as likelihood results in a Dirichlet
 2564 posterior:

$$p(\theta|z, \alpha) = \frac{p(z|\theta)p(\theta|\alpha)}{p(z|\alpha)} = \frac{1}{B(\alpha + z)} \prod_k \theta_k^{z_k + \alpha_k - 1}. \quad (\text{B.20})$$



GIBBS SAMPLING

2565

2566

2567 Notation:

$$\int dF(x) = F(x) \quad (\text{C.1})$$

2568 A mixture model:

$$L(x) = \int dF(x) \mu(x) \quad (\text{C.2})$$

2569 If we have a particular form of $F(x)$, namely it admits a decomposition of a sum of individual
 2570 values x_i :

$$F(x) = \sum_i \delta_{x_i} = \delta(x = x_0) + \delta(x = x_1) + \dots \quad (\text{C.3})$$

2571 Then our mixture model can be written as:

$$L(x) = \int dF(x) \mu(x) = \sum_i \mu(x_i) \quad (\text{C.4})$$

2572 Let $x_0 = 3$, $x_1 = 4$, $\mu(x) = x^2$, then $L(x) = 3^2 + 4^2 = 25$.2573 Walker with $P = F$, $\mu(x) = N(y|\theta)$, $i = j$ and giving each θ_j a weight ω_j :

$$f_P(y) = \int dP(\theta) N(y|\theta) \quad P = \sum_j \omega_j \delta_{\theta_j} \quad (\text{C.5})$$

2574 Then:

$$f_{\omega,j}(y) = \sum_j \omega_j N(y|\theta_j) \quad (\text{C.6})$$

2575 The likelihood:

$$L(w|\alpha, \lambda_0) = p(\phi|\alpha) \prod_{i=0}^{N-1} \int p(w_i|\theta_i, \phi) dG_0(\theta_i) \quad (\text{C.7})$$

2576 Here the index runs over all data points w_i . Each data point corresponds to a line with
 2577 parameters θ_i . Here the parameters θ_i and θ_j for data point w_i and w_j can be the same and
 2578 thus reflect the same line. The index for θ runs over the N data points, not over the K lines.

2579 The distribution G_0 does have hyperparameters λ_0 .

2580 We can also group all data points that belong to the same line k together by reordering the
 2581 product terms:

$$L(w|\alpha, \lambda_0) = p(\phi|\alpha) \prod_k \prod_{i:z_i=k} \int p(w_i|\theta_i, \phi) dG_0(\theta_i) \quad (\text{C.8})$$

2582 Here the factors that belong to line k are multiplied. The index still runs over the data points.

2583 It is also possible not to limit the second product to only the data points i that are assigned
 2584 to line k .

$$L(w|\alpha, \lambda_0) = p(\phi|\alpha) \prod_k \prod_i p(z_i|\phi) \int p(w_i|\theta_i, \phi) dG_0(\theta_i) \quad (\text{C.9})$$

2585 Now, we are gonna introduce the stick-breaking sum, which turns our integral into a discrete
 2586 sum.

2587 $G = \sum_l p_l \delta_{Z_l}$ with Z_l iid from G_0 and p_l defined as a product of beta distributions.

$$L(w|\alpha, \lambda_0) = p(\phi|\alpha) \sum_k \prod_i p(z_i|\phi) p(w_i|\theta_k) p(\theta_k|\lambda_0) \quad (\text{C.10})$$

2588 The first term at the right hand side, $p(\phi|\alpha)$, generates the partition ϕ by the Dirichlet
 2589 process with concentration parameter α . The second term $p(z_i|\phi)$ defines indices z_0, \dots, z_N
 2590 to link observations w_0, \dots, w_N with the parameters $\theta_0, \dots, \theta_K$. The probability $p(w_i|\theta_k)$
 2591 corresponds to the likelihood equations 3.10 and 3.11 with w_i the tuple of x_i and y_i and
 2592 θ_k the line parameters σ_k^2 and β_k . The probability $p(\theta_k|\lambda_0)$ corresponds to the prior from
 2593 equation 3.14. The parameters θ_k (that is, σ_k^2 and β_k) are generated from hyperparameters
 2594 λ_0 . The hyperparameters $\lambda_0 = \{\mu_0, \Lambda_0, a, b\}$ are the parameters from the Normal-Inverse-
 2595 Gamma prior.

2596 The Dirichlet process can be used as a mixture model (Antoniak, 1974; Escobar and West,
2597 1995; MacEachern and Müller, 1998) in which it generates (non-unique) parameters that
2598 subsequently generate observations:

$$\begin{aligned} G &\sim DP(\alpha, G_0) \\ \theta_i &| G \sim G \\ w_i &| \theta_i \sim F(\theta_i) \end{aligned} \tag{C.11}$$

2599 Here F describes the mapping from parameters θ_i to observations w_i . It is possible to inte-
2600 grate over G and sample the parameters directly from the base distribution G_0 .

2601 It is possible to integrate over G and get a description in the form of conditionals over the
2602 parameters (Blackwell and MacQueen, 1973):

$$\theta_{n+1} | \theta_1 \dots \theta_{n-1} \sim \frac{1}{\alpha + n} (\alpha G_0 + \sum_{i=1}^n \delta_{\theta_i}) \tag{C.12}$$

2603 C.1 Gibbs Sampling of Parameters

2604 Algorithm, we will draw $\theta_i | \theta_{-i}, y_i$ for all i .

2605 And that continuously. So, that's how we get theta.

2606 Gibbs sampling requires the conditional probabilities of all entities involved (Geman and
2607 Geman, 1984). Gibbs sampling just as other Markov chain Monte Carlo methods generates
2608 a sequence of correlated samples. Subsequently, if necessary, the Maximum A Posteriori
2609 estimation of a value can be found through picking the mode (most common occurring
2610 value) of a parameter.

2611 The derivation of the conditional probabilities of parameters with respect to the remain-
2612 ing parameters has been described in the literature (Neal, 2000). Such a derivation uses
2613 an important property of the Dirichlet process, namely that it is the conjugate prior of the
2614 multinomial distribution. Thanks to conjugacy the following equations have closed-form
2615 descriptions. The conditional probabilities are sampled from the base distribution G_0 and
2616 the other parameters θ_i in the following way:

$$\theta_{n+1} | \theta_1 \dots \theta_{n-1} \sim \frac{1}{\alpha + n} (\alpha G_0 + \sum_{i=1}^n \delta_{\theta_i}) \tag{C.13}$$

2617 If we include the observations themselves, we need to include the likelihood as well:

$$\theta_i | \theta_{-i}, w_i \sim C \left\{ \sum_{j, j \neq i} F(w_i, \theta_j) \delta_{\theta_j} + \alpha H_i \int F(w_i, \theta) dG_0(\theta) \right\} \tag{C.14}$$

2618 The constant C is a normalization factor to make the above a proper probability density
2619 (summing to one). The entity H_i is the posterior density of θ given G_0 as prior and y_i as
2620 observation. The notation θ_{-i} describes the set of all parameters Θ with θ_i excluded. The
2621 integral over $dG_0(\theta)$ is a Lebesgue-Stieltjes integral that weighs the contribution of $F(w_i, \theta)$
2622 with the base distribution $G_0(\theta)$.

2623 Equation C.14 can be used to perform inference directly with all (non-unique) parameters
2624 θ_i tied to observations w_i . Details on inference will be provided in Sect. ??.



GLOSSARY

2625

2626

2627

2628 **burn-in** running a Markov chain Monte Carlo for a while before starting to sample from it,
2629 so the results are not depending on its initial random starting position. 43

2630 **collapsed Gibbs sampling** Rao-Blackwellized Gibbs sampling. Certain sampling steps are
2631 replaced by steps where one or more variables are integrated out. This is can be thanks
2632 to analytic descriptions that arise from the use of conjugate priors. 43

2633 **slow mixing** a high degree of correlation between subsequent samples in a Markov chain
2634 leading to a long time before the chain converges. 57

LIST OF FIGURES

2636	2.1	Probability measure	11
2637	2.2	Generative vs Discriminative	21
2638	2.3	Plate notation	25
2639	2.4	Stick-breaking representation	28
2640	2.5	Matrix Representation	29
2641	2.6	Chinese Restaurant Process	30
2642	2.7	Dirichlet Process	31
2643	2.8	Beta Process	32
2644	2.9	Matrix Representation	33
2645	2.10	Gamma Process	35
2646	2.11	Matrix Representation	35
2647	2.12	The difference visualized between a Dirichlet Process mixture and a hierar-	
2648		chical Dirichlet process. It illustrates also that the input of a Dirichlet process	
2649		does not have to be a continuous function. If it is a continuous distribution	
2650		it will become a discrete distributed almost surely. If it is a discrete distribu-	
2651		tion, it will have atoms at the locations where the discrete distribution had its	
2652		probability mass concentrated.	38
2653	2.13	Rejection Sampling	41
2654	3.1	The infinite line model using the Chinese restaurant process representation	
2655		(compare with Figure 2.6). Top: α , the concentration parameter of the Dirich-	
2656		let process. Bottom, left to right: w_i , the observation, an individual point in	
2657		a 2D space; θ_i , the parameters (intercept, slope) of the line belonging to ob-	
2658		servation w_i ; H , the base distribution from which line parameter values are	
2659		sampled.	48
2660	3.2	The infinite line model in the stick-breaking representation (compare with	
2661		Figure 3.1). From left to right: α , the concentration parameter of the Dirich-	
2662		let process; (ϕ_1, \dots, ϕ_k) , the partition of points over lines; z_i , the assignment	
2663		parameters that link observation w_i with line k ; w_i , the observation, an indi-	
2664		vidual point with x and y coordinates; θ_k , the parameters of line k ; λ_0 , the	
2665		base measure from which the line parameter values are sampled.	55
2666	3.3	The performance of Algorithm 8 with respect to clustering is measured using	
2667		the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert	
2668		metric. A score of 1 means perfect clustering for all metrics, except Mirvin's	
2669		where 0 denotes perfect clustering.	57

2670	3.4	The performance of Algorithm 9 with respect to clustering is measured using	
2671		the Rand Index, the Adjusted Rand Index, the Mirvin metric, and the Hubert	
2672		metric. A score of 1 means perfect clustering for all metrics, except Mirvin's	
2673		where 0 denotes perfect clustering.	58
2674	3.5	One of the Gibbs steps in the inference of two particular lines. The points	
2675		are roughly distributed according to the lines, but one line exists out of two	
2676		large clusters. The line coordinates are visualized by a double circle. The	
2677		x-coordinate is the y-intercept of the line, the y-coordinate is the slope. . . .	59
2678	3.6	The assignment of a line to a single point. In the figure, there are three clusters	
2679		found, rather than only the two obvious clusters.	59
2680	4.1	Sampling of Pareto pairs. The parameters are $\lambda_m = -4$, $\lambda_n = 2$, $k = 5$, and	
2681		we have sampled $N = 1000$ pairs. The position parameters λ_m and λ_n define	
2682		the positions of the endpoints. The shape parameter k defines the variance in	
2683		the exact positions on both sides. The Pareto pairs always sample endpoints	
2684		both at the "left" and the "right" (not both at the left or the right side). . . .	63
2685	4.2	Consider (1) the data uniformly distributed on a line segment and (2) a sym-	
2686		metric Pareto prior for both endpoints, then we can update the estimate for	
2687		the endpoints given the data as visualized. Each subfigure shows an adjust-	
2688		ment of the endpoints given more data points (1, 3, 10, and 100 data points).	
2689		The y-axis does not have a significance in these plots.	65
2690	4.3	The Bayesian linear regression model for multiple line segments in plate no-	
2691		tation is the same as for the Infinite Line Model. The Dirichlet process is de-	
2692		defined at the left with concentration parameter α . It generates the partitions	
2693		(π_1, \dots, π_k) with assignment parameters z_i that denote which observation w_i	
2694		belongs to which cluster k . The cluster is summarized through the param-	
2695		eter set θ_k and has λ_0 as its hyperparameter. The parameter set θ_k includes	
2696		parameters that signify the line itself such as slope and y-intercept, plus the	
2697		parameters that denote the extent of the segment.	66
2698	4.4	Line segments generated through a Dirichlet Process. The line parameters are	
2699		generated from a Normal-Inverse-Gamma distribution. The endpoints of the	
2700		line segment are generated from a symmetric Pareto distribution. The points	
2701		are generated uniformly over the line segments.	67
2702	4.5	Bayesian point estimates of the sampling process with varying outcomes. . .	69
2703	4.6	Segment detection performs much worse than line detection across all three	
2704		clustering performance indicators. Perfect clustering is indicated by 1.0 for	
2705		Rand Index, Adjusted Rand Index, and Hubert.	70
2706	5.1	Dyadic Split Step	75
2707	5.2	Dyadic vs triadic MCMC	77
2708	5.3	Two examples of fitting a mixture of lines to data items scattered over a two-	
2709		dimensional space. The lines drawn are inferred using the triadic sampler.	
2710		The lines are not the ground truth, but are meant to demonstrate the typical	
2711		errors made by fitting methods. Note, for example, that there are mistakes	
2712		in both the assignment of points to lines as well as the line parameters (slope	
2713		and intercept). Left: In example 1 two lines with similar slope are seen as	
2714		the same line. Right: In example 2 points on one vertical line are assigned to	
2715		multiple lines. In example 1 and 2 slopes are not always through the points.	81

2716	5.4	Line estimation with the diadic split-merge sampler. The same results as in	
2717		Table 5.1, but visualized in a violin plot. The distribution over metric values	
2718		are displayed in a vertical fashion.	82
2719	5.5	Line estimation with the auxiliary variable sampler.	83
2720	5.6	Line estimation with the triadic split-merge sampler (our inference method).	
2721		The values are all shifted up towards one.	83
2722	6.1	Generative Adversarial Network	87
2723	6.2	Varational Autoencoder	87
2724	7.1	The likelihood function for the moments at which people decide to exercise	
2725		during the day. On the horizontal axis time, on the vertical axis the frequency	
2726		of exercising.	93
2727	7.2	The improved likelihood function for the moments at which people decide to	
2728		exercise during the day using the Von Mises distribution. On the horizontal	
2729		axis time, on the vertical axis the frequency of exercising.	94
2730	7.3	After running the algorithm, the above figures show the different categories	
2731		of runners that have been found.	97
2732	A.1	Beta distribution	113

LIST OF TABLES

2734 2.1 Structures and Processes 24

2735 2.2 Levy measure 26

2736 2.3 Exchangeable structures 27

2737 5.1 The purity, rand index, and adjusted rand index establishing the quality of

2738 the clustering method. The closer the values to one, the better the method

2739 performed. The purity metric assigns high values to clusters that do not have

2740 data points from other clusters (but does not penalize the number of clus-

2741 ters). The rand index index computes similarity between clusters taking false

2742 negatives and false positives into account. The adjusted rand index accounts

2743 for chance. The adjusted rand index is most useful in our comparison. 82

2744 7.1 Example of the type of data about the timing of exercising. A person is repre-

2745 sented by row, her preferences by column. There is not a predefined number

2746 of users or groups of users. 92

2747 7.2 Top row: A sequence of cluster indices indicates the ground truth. Each clus-

2748 ter index represents a multi-modal Von-Mises-Uniform distribution with dif-

2749 ferent parameters θ . Bottom row: A sequence of cluster indices that are the

2750 result of the described algorithm. Each cluster index represents again a multi-

2751 modal Von-Mises distribution. Errors would be represented by an inconsistent

2752 mapping from the top row to the bottom row. 96

2753 Acronyms

2754 **ABC** approximate Bayesian computation

2755 **a.s.** almost surely

2756 **BP** Beta process

2757 **CRP** Chinese restaurant process

2758 **DMM** Dirichlet mixture model

2759 **DP** Dirichlet process

2760 **GP** Gamma process

2761 **GEM** Griffiths, Engen, and McCloskey

2762 **HDP** hierarchical Dirichlet process

2763 **IBP** Indian buffet process

2764 **ILM** infinite line model

2765 **MAP** maximum a posteriori

2766 **MCMC** Markov chain Monte Carlo

2767 **ML** maximum likelihood

2768 **MLL** maximum log-likelihood

2769 **NB** naive Bayes

2770 **NIG** Normal-Inverse-Gamma

2771 **PYP** Pitman-Yor process

2772 **SAMS** sequentially allocated merge-split

2773

SUMMARY

2774 Summary....

2775

SAMENVATTING

2776 Samenvatting...

2777

ACKNOWLEDGMENTS

2778 I want to thank...

2779

CURRICULUM VITAE

2780 Anne van Rossum

2781

PUBLICATIONS

2782 The investigations performed during my Ph.D. research resulted in the following publica-
2783 tions.

2784 ◦ A.C. van Rossum.

SIKS DISSERTATION SERIES

1998

- 1 Johan van den Akker (CWI¹) *DEGAS - An Active Temporal Database of Autonomous Objects*
- 2 Floris Wiesman (UM) *Information Retrieval by Graphically Browsing Meta-Information*
- 3 Ans Steuten (TUD) *A Contribution to the Linguistic Analysis of Business Conversations within the Language/Action Perspective*
- 4 Dennis Breuker (UM) *Memory versus Search in Games*
- 5 Eduard W. Oskamp (RUL) *Computerondersteuning bij Straftoemeting*

1999

- 1 Mark Sloof (VU) *Physiology of Quality Change Modelling; Automated Modelling of Quality Change of Agricultural Products*
- 2 Rob Potharst (EUR) *Classification using Decision Trees and Neural Nets*
- 3 Don Beal (UM) *The Nature of Minimax Search*
- 4 Jacques Penders (UM) *The Practical Art of Moving Physical Objects*
- 5 Aldo de Moor (KUB) *Empowering Communities: A Method for the Legitimate User-Driven Specification of Network Information Systems*
- 6 Niek J.E. Wijngaards (VU) *Re-Design of Compositional Systems*
- 7 David Spelt (UT) *Verification Support for Object Database Design*
- 8 Jacques H.J. Lenting (UM) *Informed Gambling: Conception and Analysis of a Multi-Agent Mechanism for Discrete Reallocation*

2000

- 1 Frank Niessink (VU) *Perspectives on Improving Software Maintenance*
- 2 Koen Holtman (TU/e) *Prototyping of CMS Storage Management*
- 3 Carolien M.T. Metselaar (UvA) *Sociaal-organisatorische Gevolgen van Kennistechnologie; een Procesbenadering en Actorperspectief*
- 4 Geert de Haan (VU) *ETAG, A Formal Model of Competence Knowledge for User Interface Design*
- 5 Ruud van der Pol (UM) *Knowledge-Based Query Formulation in Information Retrieval*
- 6 Rogier van Eijk (UU) *Programming Languages for Agent Communication*
- 7 Niels Peek (UU) *Decision-Theoretic Planning of Clinical Patient Management*
- 8 Veerle Coupé (EUR) *Sensitivity Analysis of Decision-Theoretic Networks*
- 9 Florian Waas (CWI) *Principles of Probabilistic Query Optimization*
- 10 Niels Nes (CWI) *Image Database Management System Design Considerations, Algorithms and Architecture*
- 11 Jonas Karlsson (CWI) *Scalable Distributed Data Structures for Database Management*

2001

- 1 Silja Renooij (UU) *Qualitative Approaches to Quantifying Probabilistic Networks*
- 2 Koen Hindriks (UU) *Agent Programming Languages: Programming with Mental Models*

¹Abbreviations: SIKS - Dutch Research School for Information and Knowledge Systems; CWI - Centrum voor Wiskunde en Informatica, Amsterdam; EUR - Erasmus Universiteit, Rotterdam; KUB - Katholieke Universiteit Brabant, Tilburg; KUN - Katholieke Universiteit Nijmegen; OU - Open Universiteit; RUL - Rijksuniversiteit Leiden; RUN - Radboud Universiteit Nijmegen; TUD - Technische Universiteit Delft; TU/e - Technische Universiteit Eindhoven; UL - Universiteit Leiden; UM - Universiteit Maastricht; UT - Universiteit Twente, Enschede; UU - Universiteit Utrecht; UvA - Universiteit van Amsterdam; UvT - Universiteit van Tilburg; VU - Vrije Universiteit, Amsterdam.

- 2847 3 Maarten van Someren (UvA) *Learning as Problem Solving* 2903
2848 2904
- 2849 4 Evgueni Smirnov (UM) *Conjunctive and Disjunctive Version Spaces with Instance-Based Boundary Sets* 2905
2850 2906
- 2851 5 Jacco van Ossenbruggen (VU) *Processing Structured Hypermedia: A Matter of Style* 2907
2852 2908
- 2853 6 Martijn van Welie (VU) *Task-Based User Interface Design* 2909
2854 2910
- 2855 7 Bastiaan Schonhage (VU) *Diva: Architectural Perspectives on Information Visualization* 2910
2856 2911
- 2857 8 Pascal van Eck (VU) *A Compositional Semantic Structure for Multi-Agent Systems Dynamics* 2912
2858 2913
- 2859 9 Pieter Jan 't Hoen (RUL) *Towards Distributed Development of Large Object-Oriented Models, Views of Packages as Classes* 2913
2860 2914
- 2861 10 Maarten Sierhuis (UvA) *Modeling and Simulating Work Practice BRAHMS: a Multiagent Modeling and Simulation Language for Work Practice Analysis and Design* 2915
2862 2916
- 2863 11 Tom M. van Engers (VU) *Knowledge Management: The Role of Mental Models in Business Systems Design* 2917
2864 2918
- 2865 2919
- 2866 11 Tom M. van Engers (VU) *Knowledge Management: The Role of Mental Models in Business Systems Design* 2920
2867 2921
- 2868 2922
- 2869 **2002** 2923
2924
2925
- 2870 1 Nico Lassing (VU) *Architecture-Level Modifiability Analysis* 2926
2871 2927
- 2872 2 Roelof van Zwol (UT) *Modelling and Searching Web-based Document Collections* 2928
2873 2929
- 2874 3 Henk Ernst Blok (UT) *Database Optimization Aspects for Information Retrieval* 2930
2875 2931
- 2876 4 Juan Roberto Castelo Valdueza (UU) *The Discrete Acyclic Digraph Markov Model in Data Mining* 2932
2877 2933
- 2878 5 Radu Serban (VU) *The Private Cyberspace Modeling Electronic Environments Inhabited by Privacy-Concerned Agents* 2934
2879 2935
- 2880 2936
- 2881 6 Laurens Mommers (UL) *Applied Legal Epistemology; Building a Knowledge-based Ontology of the Legal Domain* 2937
2882 2938
- 2883 2939
- 2884 7 Peter Boncz (CWI) *Monet: A Next-Generation DBMS Kernel For Query-Intensive Applications* 2940
2885 2941
- 2886 8 Jaap Gordijn (VU) *Value Based Requirements Engineering: Exploring Innovative E-Commerce Ideas* 2942
2887 2943
- 2888 9 Willem-Jan van den Heuvel (KUB) *Integrating Modern Business Applications with Objectified Legacy Systems* 2944
2889 2945
- 2890 2946
- 2891 10 Brian Sheppard (UM) *Towards Perfect Play of Scrabble* 2947
2892 2948
- 2893 11 Wouter C.A. Wijngaards (VU) *Agent Based Modelling of Dynamics: Biological and Organisational Applications* 2949
2894 2950
- 2895 2951
- 2896 12 Albrecht Schmidt (UvA) *Processing XML in Database Systems* 2951
2897 2952
- 2898 13 Hongjing Wu (TU/e) *A Reference Architecture for Adaptive Hypermedia Applications* 2952
2899 2953
- 2900 14 Wieke de Vries (UU) *Agent Interaction: Abstract Approaches to Modelling, Programming and Verifying Multi-Agent Systems* 2954
2901 2955
- 2902 2956
- 15 Rik Eshuis (UT) *Semantics and Verification of UML Activity Diagrams for Workflow Modelling*
- 16 Pieter van Langen (VU) *The Anatomy of Design: Foundations, Models and Applications*
- 17 Stefan Manegold (UvA) *Understanding, Modeling, and Improving Main-Memory Database Performance*
- 2003**
- 1 Heiner Stuckenschmidt (VU) *Ontology-Based Information Sharing in Weakly Structured Environments*
- 2 Jan Broersen (VU) *Modal Action Logics for Reasoning About Reactive Systems*
- 3 Martijn Schuemie (TUD) *Human-Computer Interaction and Presence in Virtual Reality Exposure Therapy*
- 4 Milan Petkovic (UT) *Content-Based Video Retrieval Supported by Database Technology*
- 5 Jos Lehmann (UvA) *Causation in Artificial Intelligence and Law – A Modelling Approach*
- 6 Boris van Schooten (UT) *Development and Specification of Virtual Environments*
- 7 Machiel Jansen (UvA) *Formal Explorations of Knowledge Intensive Tasks*
- 8 Yong-Ping Ran (UM) *Repair-Based Scheduling*
- 9 Rens Kortmann (UM) *The Resolution of Visually Guided Behaviour*
- 10 Andreas Lincke (UT) *Electronic Business Negotiation: Some Experimental Studies on the Interaction between Medium, Innovation Context and Culture*
- 11 Simon Keizer (UT) *Reasoning under Uncertainty in Natural Language Dialogue using Bayesian Networks*
- 12 Roeland Ordelman (UT) *Dutch Speech Recognition in Multimedia Information Retrieval*
- 13 Jeroen Donkers (UM) *Nosce Hostem – Searching with Opponent Models*
- 14 Stijn Hoppenbrouwers (KUN) *Freezing Language: Conceptualisation Processes across ICT-Supported Organisations*
- 15 Mathijs de Weerd (TUD) *Plan Merging in Multi-Agent Systems*
- 16 Menzo Windhouwer (CWI) *Feature Grammar Systems - Incremental Maintenance of Indexes to Digital Media Warehouse*
- 17 David Jansen (UT) *Extensions of Statecharts with Probability, Time, and Stochastic Timing*
- 18 Levente Kocsis (UM) *Learning Search Decisions*
- 2004**
- 1 Virginia Dignum (UU) *A Model for Organizational Interaction: Based on Agents, Founded in Logic*
- 2 Lai Xu (UvT) *Monitoring Multi-party Contracts for E-business*
- 3 Perry Groot (VU) *A Theoretical and Empirical Analysis of Approximation in Symbolic Problem Solving*

- 2957 4 Chris van Aart (UvA) *Organizational Principles for* 3011
2958 *Multi-Agent Architectures* 3012
- 2959 5 Viara Popova (EUR) *Knowledge Discovery and* 3013
2960 *Monotonicity* 3014
- 2961 6 Bart-Jan Hommes (TUD) *The Evaluation of Busi-* 3015
2962 *ness Process Modeling Techniques* 3016
- 2963 7 Elise Boltjes (UM) *Voorbeeld_{IG} Onderwijs; Voor-* 3017
2964 *beeldgestuurd Onderwijs, een Opstap naar Abstrac-* 3018
2965 *Denken, vooral voor Meisjes* 3019
- 2966 8 Joop Verbeek (UM) *Politie en de Nieuwe Inter-* 3020
2967 *nationale Informatiemarkt, Grensregionale Politie* 3021
2968 *Gegevensuitwisseling en Digitale Expertise* 3022
- 2969 9 Martin Caminada (VU) *For the Sake of the Argu-* 3023
2970 *ment; Explorations into Argument-based Reasoning* 3024
- 2971 10 Suzanne Kabel (UvA) *Knowledge-rich Indexing of* 3025
2972 *Learning-objects* 3026
- 2973 11 Michel Klein (VU) *Change Management for Dis-* 3027
2974 *tributed Ontologies* 3028
- 2975 12 The Duy Bui (UT) *Creating Emotions and Facial Ex-* 3029
2976 *pressions for Embodied Agents* 3030
- 2977 13 Wojciech Jamroga (UT) *Using Multiple Models of* 3031
2978 *Reality: On Agents who Know how to Play* 3032
- 2979 14 Paul Harrenstein (UU) *Logic in Conflict. Logical Ex-* 3033
2980 *plorations in Strategic Equilibrium* 3034
- 2981 15 Arno Knobbe (UU) *Multi-Relational Data Mining* 3035
2982 16 Federico Divina (VU) *Hybrid Genetic Relational* 3036
2983 *Search for Inductive Learning*
- 2984 17 Mark Winands (UM) *Informed Search in Complex* 3037
2985 *Games*
- 2986 18 Vania Bessa Machado (UvA) *Supporting the Con-* 3038
2987 *struction of Qualitative Knowledge Models* 3039
- 2988 19 Thijs Westerveld (UT) *Using generative probabilis-* 3040
2989 *tic models for multimedia retrieval* 3041
- 2990 20 Madelon Evers (Nyenrode) *Learning from Design* 3042
2991 *facilitating multidisciplinary design teams* 3043
3044
- 2992 **2005** 3045
3046
3047
- 2993 1 Floor Verdenius (UvA) *Methodological Aspects of* 3048
2994 *Designing Induction-Based Applications* 3049
- 2995 2 Erik van der Werf (UM) *AI techniques for the game* 3050
2996 *of Go* 3051
- 2997 3 Franc Grootjen (RUN) *A Pragmatic Approach to the* 3052
2998 *Conceptualisation of Language* 3053
- 2999 4 Nirvana Meratnia (UT) *Towards Database Support* 3054
3000 *for Moving Object data* 3055
- 3001 5 Gabriel Infante-Lopez (UvA) *Two-Level Probabilis-* 3056
3002 *tic Grammars for Natural Language Parsing* 3057
- 3003 6 Pieter Spronck (UM) *Adaptive Game AI* 3058
- 3004 7 Flavius Frasinca (TU/e) *Hypermedia Presentation* 3059
3005 *Generation for Semantic Web Information Systems* 3060
- 3006 8 Richard Vdovjak (TU/e) *A Model-driven Approach* 3061
3007 *for Building Distributed Ontology-based Web Appli-* 3062
3008 *cations* 3063
- 3009 9 Jeen Broekstra (VU) *Storage, Querying and Infer-* 3064
3010 *encing for Semantic Web Languages* 3065
3066
- 10 Anders Bouwer (UvA) *Explaining Behaviour: Using*
Qualitative Simulation in Interactive Learning Envi-
ronments
- 11 Elth Ogston (VU) *Agent Based Matchmaking and*
Clustering - A Decentralized Approach to Search
- 12 Csaba Boer (EUR) *Distributed Simulation in Indus-*
try
- 13 Fred Hamburg (UL) *Een Computermodel voor het*
Ondersteunen van Euthanasiebeslissingen
- 14 Borys Omelayenko (VU) *Web-Service configuration*
on the Semantic Web; Exploring how semantics
meets pragmatics
- 15 Tibor Bosse (VU) *Analysis of the Dynamics of Cog-*
nitive Processes
- 16 Joris Graaumanns (UU) *Usability of XML Query Lan-*
guages
- 17 Boris Shishkov (TUD) *Software Specification Based*
on Re-usable Business Components
- 18 Danielle Sent (UU) *Test-selection strategies for prob-*
abilistic networks
- 19 Michel van Dartel (UM) *Situated Representation*
- 20 Cristina Coteanu (UL) *Cyber Consumer Law, State*
of the Art and Perspectives
- 21 Wijnand Derks (UT) *Improving Concurrency and*
Recovery in Database Systems by Exploiting Appli-
cation Semantics
- 2006**
- 1 Samuil Angelov (TU/e) *Foundations of B2B Elec-*
tronic Contracting
- 2 Cristina Chisalita (VU) *Contextual issues in the de-*
sign and use of information technology in organiza-
tions
- 3 Noor Christoph (UvA) *The role of metacognitive*
skills in learning to solve problems
- 4 Marta Sabou (VU) *Building Web Service Ontologies*
- 5 Cees Pierik (UU) *Validation Techniques for Object-*
Oriented Proof Outlines
- 6 Ziv Baida (VU) *Software-aided Service Bundling -*
Intelligent Methods & Tools for Graphical Service
Modeling
- 7 Marko Smiljanic (UT) *XML schema matching – bal-*
ancing efficiency and effectiveness by means of clus-
tering
- 8 Eelco Herder (UT) *Forward, Back and Home Again*
- Analyzing User Behavior on the Web
- 9 Mohamed Wahdan (UM) *Automatic Formulation of*
the Auditor's Opinion
- 10 Ronny Siebes (VU) *Semantic Routing in Peer-to-Peer*
Systems
- 11 Joeri van Ruth (UT) *Flattening Queries over Nested*
Data Types
- 12 Bert Bongers (VU) *Interactivation - Towards an e-*
cology of people, our technological environment, and
the arts
- 13 Henk-Jan Lebbink (UU) *Dialogue and Decision*
Games for Information Exchanging Agents

- 3067 14 Johan Hoorn (VU) *Software Requirements: Update, Upgrade, Redesign - towards a Theory of Requirements Change* 3121-3123
- 3068
- 3069
- 3070 15 Rainer Malik (UU) *CONAN: Text Mining in the Biomedical Domain* 3124-3125
- 3071
- 3072 16 Carsten Riggelsen (UU) *Approximation Methods for Efficient Learning of Bayesian Networks* 3126-3127
- 3073
- 3074 17 Stacey Nagata (UU) *User Assistance for Multitasking with Interruptions on a Mobile Device* 3128-3129
- 3075
- 3076 18 Valentin Zhizhkun (UvA) *Graph transformation for Natural Language Processing* 3130-3131
- 3077
- 3078 19 Birna van Riemsdijk (UU) *Cognitive Agent Programming: A Semantic Approach* 3132-3133
- 3079
- 3080 20 Marina Velikova (UvT) *Monotone models for prediction in data mining* 3134-3135
- 3081
- 3082 21 Bas van Gils (RUN) *Aptness on the Web* 3136-3137
- 3083
- 3084 22 Paul de Vrieze (RUN) *Fundamentals of Adaptive Personalisation* 3138-3139
- 3085
- 3086 23 Ion Juvina (UU) *Development of Cognitive Model for Navigating on the Web* 3140-3141
- 3087
- 3088 24 Laura Hollink (VU) *Semantic Annotation for Retrieval of Visual Resources* 3142-3143
- 3089
- 3090 25 Madalina Drugan (UU) *Conditional log-likelihood MDL and Evolutionary MCMC* 3144-3145
- 3091
- 3092 26 Vojkan Mihajlovic (UT) *Score Region Algebra: A Flexible Framework for Structured Information Retrieval* 3146-3147
- 3093
- 3094 27 Stefano Bocconi (CWI) *Vox Populi: generating video documentaries from semantically annotated media repositories* 3148-3149
- 3095
- 3096
- 3097 28 Borkur Sigurbjornsson (UvA) *Focused Information Access using XML Element Retrieval* 3150-3151
- 3098
- 3099 **2007** 3152-3154
- 3100 1 Kees Leune (UvT) *Access Control and Service-Oriented Architectures* 3155
- 3101
- 3102 2 Wouter Teepe (RUG) *Reconciling Information Exchange and Confidentiality: A Formal Approach* 3156-3157
- 3103
- 3104 3 Peter Mika (VU) *Social Networks and the Semantic Web* 3158-3159
- 3105
- 3106 4 Jurriaan van Diggelen (UU) *Achieving Semantic Interoperability in Multi-agent Systems: a dialogue-based approach* 3160-3161
- 3107
- 3108
- 3109 5 Bart Schermer (UL) *Software Agents, Surveillance and the Right to Privacy: a Legislative Framework for Agent-enabled Surveillance* 3162-3163
- 3110
- 3111
- 3112 6 Gilad Mishne (UvA) *Applied Text Analytics for Blogs* 3164-3165
- 3113
- 3114 7 Natasa Jovanovic' (UT) *To Whom It May Concern Addressee Identification in Face-to-Face Meetings* 3166-3167
- 3115
- 3116 8 Mark Hoogendoorn (VU) *Modeling of Change in Multi-Agent Organizations* 3168-3169
- 3117
- 3118 9 David Mobach (VU) *Agent-Based Mediated Service Negotiation* 3170-3171
- 3119
- 3120 10 Huib Aldewereld (UU) *Autonomy vs. Conformity an Institutional Perspective on Norms and Protocols* 3172-3173
- 3121
- 3122
- 3123
- 3124
- 3125
- 3126
- 3127
- 3128
- 3129
- 3130
- 3131
- 3132
- 3133
- 3134
- 3135
- 3136
- 3137
- 3138
- 3139
- 3140
- 3141
- 3142
- 3143
- 3144
- 3145
- 3146
- 3147
- 3148
- 3149
- 3150
- 3151
- 3152
- 3153
- 3154
- 3155
- 2008**
- 1 Katalin Boer-Sorbán (EUR) *Agent-Based Simulation of Financial Markets: A modular, continuous-time approach*
- 2 Alexei Sharpanskykh (VU) *On Computer-Aided Methods for Modeling and Analysis of Organizations*
- 3 Vera Hollink (UvA) *Optimizing hierarchical menus: a usage-based approach*
- 4 Ander de Keijzer (UT) *Management of Uncertain Data - towards unattended integration*
- 5 Bela Mutschler (UT) *Modeling and simulating causal dependencies on process-aware information systems from a cost perspective*
- 6 Arjen Hommersom (RUN) *On the Application of Formal Methods to Clinical Guidelines, an Artificial Intelligence Perspective*
- 7 Peter van Rosmalen (OU) *Supporting the tutor in the design and support of adaptive e-learning*
- 8 Janneke Bolt (UU) *Bayesian Networks: Aspects of Approximate Inference*
- 9 Christof van Nimwegen (UU) *The paradox of the guided user: assistance can be counter-effective*

- 3177 10 Wauter Bosma (UT) *Discourse oriented Summariza-* 3235
3178 *tion* 3236
- 3179 11 Vera Kartseva (VU) *Designing Controls for Network*
3180 *Organizations: a Value-Based Approach* 3237
- 3181 12 Jozsef Farkas (RUN) *A Semiotically oriented Cogni-*
3182 *tive Model of Knowledge Representation* 3238
- 3183 13 Caterina Carraciolo (UvA) *Topic Driven Access to* 3239
3184 *Scientific Handbooks* 3240
- 3185 14 Arthur van Bunningen (UT) *Context-Aware Query-* 3241
3186 *ing; Better Answers with Less Effort* 3242
- 3187 15 Martijn van Otterlo (UT) *The Logic of Adaptive Be-* 3243
3188 *havior: Knowledge Representation and Algorithms* 3244
3189 *for the Markov Decision Process Framework in First-* 3245
3190 *Order Domains* 3246
- 3191 16 Henriette van Vugt (VU) *Embodied Agents from a* 3247
3192 *User's Perspective* 3248
- 3193 17 Martin Op't Land (TUD) *Applying Architecture and* 3249
3194 *Ontology to the Splitting and Allying of Enterprises* 3250
- 3195 18 Guido de Croon (UM) *Adaptive Active Vision* 3251
- 3196 19 Henning Rode (UT) *From document to entity re-* 3252
3197 *trieval: improving precision and performance of fo-* 3253
3198 *cused text search* 3254
- 3199 20 Rex Arendsen (UvA) *Geen bericht, goed bericht. Een* 3255
3200 *onderzoek naar de effecten van de introductie van* 3256
3201 *elektronisch berichtenverkeer met een overheid op de* 3257
3202 *administratieve lasten van bedrijven* 3258
- 3203 21 Krisztian Balog (UvA) *People search in the enter-* 3259
3204 *prise* 3260
- 3205 22 Henk Koning (UU) *Communication of IT* 3261
3206 *architecture* 3262
- 3207 23 Stefan Visscher (UU) *Bayesian network models for* 3263
3208 *the management of ventilator-associated pneumonia* 3264
- 3209 24 Zharko Aleksovski (VU) *Using background knowl-* 3265
3210 *edge in ontology matching* 3266
- 3211 25 Geert Jonker (UU) *Efficient and Equitable exchange* 3267
3212 *in air traffic management plan repair using spender-* 3268
3213 *signed currency* 3269
- 3214 26 Marijn Huijbregts (UT) *Segmentation, diarization* 3270
3215 *and speech transcription: surprise data unraveled* 3271
- 3216 27 Hubert Vogten (OU) *Design and implementation* 3272
3217 *strategies for IMS learning design* 3273
- 3218 28 Ildiko Flesh (RUN) *On the use of independence re-* 3274
3219 *lations in Bayesian networks* 3275
- 3220 29 Dennis Reidsma (UT) *Annotations and subjective* 3276
3221 *machines- Of annotators, embodied agents, users,* 3277
3222 *and other humans* 3278
- 3223 30 Wouter van Atteveldt (VU) *Semantic network anal-* 3279
3224 *ysis: techniques for extracting, representing and* 3280
3225 *querying media content* 3281
- 3226 31 Loes Braun (UM) *Pro-active medical information re-* 3282
3227 *trieval* 3283
- 3228 32 Trung B. Hui (UT) *Toward affective dialogue man-* 3284
3229 *agement using partially observable markov decision* 3285
3230 *processes* 3286
- 3231 33 Frank Terpstra (UvA) *Scientific workflow design,* 3287
3232 *theoretical and practical issues* 3288
- 3233 34 Jeroen de Knijf (UU) *Studies in Frequent Tree Min-* 3289
3234 *ing* 3290
- 35 Benjamin Torben-Nielsen (UvT) *Dendritic mor-* 3291
phology: function shapes structure
- 2009**
- 1 Rasa Jurgelenaite (RUN) *Symmetric Causal Inde-*
pendence Models
- 2 Willem Robert van Hage (VU) *Evaluating Ontology-*
Alignment Techniques
- 3 Hans Stol (UvT) *A Framework for Evidence-based*
Policy Making Using IT
- 4 Josephine Nabukenya (RUN) *Improving the Qual-*
ity of Organisational Policy Making using Collabo-
ration Engineering
- 5 Sietse Overbeek (RUN) *Bridging Supply and De-*
mand for Knowledge Intensive Tasks - Based on
Knowledge, Cognition, and Quality
- 6 Muhammad Subianto (UU) *Understanding Classi-*
fication
- 7 Ronald Poppe (UT) *Discriminative Vision-Based Re-*
covery and Recognition of Human Motion
- 8 Volker Nannen (VU) *Evolutionary Agent-Based Pol-*
icy Analysis in Dynamic Environments
- 9 Benjamin Kanagwa (RUN) *Design, Discovery and*
Construction of Service-oriented Systems
- 10 Jan Wielemaker (UvA) *Logic programming for*
knowledge-intensive interactive applications
- 11 Alexander Boer (UvA) *Legal Theory, Sources of Law*
& the Semantic Web
- 12 Peter Massuthe (TU/e, Humboldt-Universität zu
Berlin) *Operating Guidelines for Services*
- 13 Steven de Jong (UM) *Fairness in Multi-Agent Sys-*
tems
- 14 Maksym Korotkiy (VU) *From ontology-enabled ser-*
vices to service-enabled ontologies (making onto-
logies work in e-science with ONTO-SOA)
- 15 Rinke Hoekstra (UvA) *Ontology Representation -*
Design Patterns and Ontologies that Make Sense
- 16 Fritz Reul (UvT) *New Architectures in Computer*
Chess
- 17 Laurens van der Maaten (UvT) *Feature Extraction*
from Visual Data
- 18 Fabian Groffen (CWI) *Armada, An Evolving*
Database System
- 19 Valentin Robu (CWI) *Modeling Preferences, Strate-*
gic Reasoning and Collaboration in Agent-Mediated
Electronic Markets
- 20 Bob van der Vecht (UU) *Adjustable Autonomy: Con-*
trolling Influences on Decision Making
- 21 Stijn Vanderlooy (UM) *Ranking and Reliable Clas-*
sification
- 22 Pavel Serdyukov (UT) *Search For Expertise: Going*
beyond direct evidence
- 23 Peter Hofgesang (VU) *Modelling Web Usage in a*
Changing Environment
- 24 Annerieke Heuvelink (VU) *Cognitive Models for*
Training Simulations
- 25 Alex van Ballegooij (CWI) *"RAM: Array Database*
Management through Relational Mapping"

- 3292 26 Fernando Koch (UU) *An Agent-Based Model for the* 3312 35 Wouter Koelewijn (UL) *Privacy en Politiegegevens;*
 3293 *Development of Intelligent Mobile Services* 3313 *Over geautomatiseerde normatieve informatie-*
 3294 27 Christian Glahn (OU) *Contextual Support of social* 3314 *uitwisseling*
 3295 *Engagement and Reflection on the Web* 3315 36 Marco Kalz (OUN) *Placement Support for Learners*
 3296 28 Sander Evers (UT) *Sensor Data Management with* 3316 *in Learning Networks*
 3297 *Probabilistic Models* 3317 37 Hendrik Drachsler (OUN) *Navigation Support for*
 3298 29 Stanislav Pokraev (UT) *Model-Driven Semantic In-* 3318 *Learners in Informal Learning Networks*
 3299 *tegration of Service-Oriented Applications* 3319 38 Riina Vuorikari (OU) *Tags and self-organisation: a*
 3300 30 Marcin Zukowski (CWI) *Balancing vectorized query* 3320 *metadata ecology for learning resources in a multi-*
 3301 *execution with bandwidth-optimized storage* 3321 *lingual context*
 3302 31 Sofiya Katrenko (UvA) *A Closer Look at Learning* 3322 39 Christian Stahl (TUE, Humboldt-Universität zu
 3303 *Relations from Text* 3323 Berlin) *Service Substitution – A Behavioral Ap-*
 3304 32 Rik Farenhorst and Remco de Boer (VU) *Architec-* 3324 *proach Based on Petri Nets*
 3305 *tural Knowledge Management: Supporting Archi-* 3325 40 Stephan Raaijmakers (UvT) *Multinomial Language*
 3306 *tects and Auditors* 3326 *Learning: Investigations into the Geometry of Lan-*
 3307 33 Khiet Truong (UT) *How Does Real Affect Affect Affect* 3327 *guage*
 3308 *Recognition In Speech?* 3328 41 Igor Berezhnyy (UvT) *Digital Analysis of Paintings*
 3309 34 Inge van de Weerd (UU) *Advancing in Software* 3329 42 Toine Bogers (UvT) *Recommender Systems for So-*
 3310 *Product Management: An Incremental Method En-* 3330 *cial Bookmarking*
 3311 *gineering Approach*